

Practical Lessons from Reinforcement Learning Post-Training

What Broke & Why

Luv Verma, PhD

© 2026 Luv Verma. First edition, version 1.0, July 2026. Licensed under a Creative Commons Attribution 4.0 International License (CC BY 4.0). You are free to share and adapt this work for any purpose, provided you give appropriate credit.

Suggested citation: Verma, Luv. *What Broke & Why: Practical Lessons from Reinforcement Learning Post-Training*. First edition, version 1.0. Zenodo, 2026. DOI: 10.5281/zenodo.21115798.

Set in Source Serif Pro, TeX Gyre Heros, and TeX Gyre Cursor. Typeset with Xe_{La}T_EX.

Contents

How This Book Works	xiii
The Gap in the Existing Literature	xviii
I RL Post-Training Foundations	1
1 Why Post-Train, and Why Reinforcement Learning?	2
2 The Pipeline: SFT → Preference Tuning → RL	11
3 GRPO from Scratch	18
4 The Training Loop in Full	24
5 Advantage Collapse under Binary Rewards	30
II First Verifiable RL: Math & GSM8K	34
6 Phase 0: What Not to Do	36
7 SFT First: The Easy Win	40
8 Countdown: The GRPO Hello-World	44
9 GSM8K: The Hello-World of RLVR	47
10 Reward Extraction Done Right	50
11 The Reward = 0 Cold-Start Problem	54
12 Binary Reward and Advantage Collapse, on Math	57
13 Partial Credit and Continuous Reward	60
14 GRPO on GSM8K, Hands-On	64
15 Dr. GRPO, DAPO, and RLOO: The Variant Toolbox	69

16	When Eval Methodology Is the Result	73
17	Choosing Data That Teaches	77
18	Math-DAPO Breakout	81
19	CLEVR: When Teacher-Forcing Misleads	85
20	RefCOCO: IoU as a Graded Reward	88
21	Tool-Calling and Multi-Step Reasoning	91
22	Search-R1: The Bridge to the Agent	95
23	The Best Practice That Stalled Learning	98
24	When the Trainer and the Model Disagree	101
25	What the Math On-Ramp Taught Us	104
III	The Search Agent	108
26	The Single-Reward Plateau	110
27	Lifting the Plateau: F1, Dr. GRPO, DAPO	113
28	The 4-Turn Unlock and the Tool-Description Fix	116
29	Behavior, Not Just Accuracy	120
30	The Five-Component Multi-Objective Reward	123
31	Reward Hacking: The Linear-Reward Collapse	127
32	Correctness-Gated Rewards: The SiLU/Swish Gate	131
33	The Ablation Ladder: What’s Proven, Partial, and Still Open	136
34	Finding: F1-Only Collapses Where Multi-Reward Survives	139
35	Finding: Tool-Call Spamming Without Behavioral Rewards	141
36	Results: A 4B Agent That Beats Its Weight, Honestly	144

37	When Small Evals Misled: The 500-Sample Correction	149
38	The Behavioral Audit Suite	153
IV	Stability: Entropy, Collapse, and Clip-Cov	157
39	What Entropy Measures in a Policy	159
40	v1: Entropy Collapse	162
41	v2 & v3: The Entropy-Explosion Trap	165
42	v4: Clip-Higher Alone Isn't Enough	168
43	v5: Clip-Cov and the Damped Oscillation	171
44	Why Clip-Cov Works	175
45	GDPO: The Per-Channel Idea That Failed	179
46	Strategy Collapse Comes First	183
47	Loss Aggregation and the Hidden Length Bias	187
48	The GRPO Variant Taxonomy	190
49	The Entropy-Performance Frontier	194
V	Scaling Up: MoE, GSPO, and Routing	198
50	14B Dense vs 30B-A3B MoE	200
51	Why MoE Makes Token-Level GRPO Fragile	203
52	GSPO: Sequence-Level Importance Sampling	207
53	C9-GSPO Breakout	211
54	MoE Routing Collapse: Causes and Fixes	214
55	FP8 + MoE Implementation Wall	218

56	Megatron vs FSDP2 for MoE	222
57	Three Ways to Stop Routing Collapse	225
58	The Internals Telemetry Method	228
59	Direction, Not Distance: The Capability Ceiling	231
60	The U-Curve and the New Attractor	234
61	COMBO: Stacking Stabilizers	237
62	The Bake-Off: The Metric Chooses the Winner	240
VI When RL Isn't Enough: SWE-bench & the Pivot		244
63	SWE-bench Setup	246
64	The Catastrophic Collapse	249
65	SW2–SW6: The Stabilization Saga	252
66	Stable but Not Learning	256
67	Why a 14B Can't Solve Full-Repo Bugs	259
68	On-Policy Distillation: The Strategy	262
69	The OPD Recipe in Detail	265
70	Datasets & Environments for OPD	268
71	Compute Budgets & Teacher Quantization	271
72	Stage 0: Validate the Teacher Before You Distill	274
VII What RL Actually Does: The Mechanistic Picture		278
73	RL Sharpens. It Does Not Expand	280
74	RL Edits Only Fork Tokens	283

75	Productive vs Wasted Entropy	286
76	How Much Does RL Move the Weights?	289
77	What Sharpening Costs: The Full-FT Regression	293
78	Direction, Not Distance: The Residual-Stream Ceiling	297
79	Watching the Model Think: The Telemetry Toolkit	301
80	Evidence Boundaries	305
VIII Reward Design (Cross-Cutting Craft)		308
81	Verifiable Rewards and Why They Matter	310
82	Continuous vs Binary Rewards	314
83	Reward Auditing and JSONL Logging	318
84	Reward Shaping for Multi-Step Tool Flows	322
85	Process Rewards and PRMs for Agents	327
86	Combining Rewards Without Reward Hacking	331
IX Don't Fool Yourself: Evaluation Discipline		335
87	The Variance Budget	337
88	When the Training Score Misleads	342
89	The Canonical-Eval Correction	346
90	Eval Methodology <i>Is</i> the Result	349
91	pass@k: Capability versus Reliability	352
92	Metric Choice Crowns Different Winners	355
93	Retroactive Gating: Auditing Your Own Past Claims	358

94	Single Seeds, Honestly	361
X	Architecture-Aware RL: What Qwen3.5 Teaches	364
95	Why Architecture Matters for RL	366
96	The 3:1 Hybrid	370
97	Gated DeltaNet from Scratch	374
98	Gated Attention and the Attention-Sink Fix	377
99	Sparse Mixture-of-Experts and Ultra-Sparsity	380
100	The Chunk-Parallel Algorithm	384
101	Context Extension to a Million Tokens	387
102	The Inference Stack	389
103	The Architecture Decision Map	392
XI	Applied Tool-Agent Pipeline: SFT → DPO → GSPO	395
104	The Production Web Agent	397
105	Converting Claude Trajectories to Training Data	401
106	SFT with Correct Tool-Call Masking	405
107	The Masking Bug That Silently Kills SFT	409
108	DPO from Guardrail Verdicts	412
109	GSPO with Online Playwright Rollouts	415
110	Offline GSPO and the Custom rollout_func	418
111	Environment Orchestration and Verifiers	421
112	Serving/Training Parity	424

113	Build Order and Promotion Gates	426
XII	The Failure Catalog	428
114	Reading a Training Dashboard	430
115	Collapse Failures	433
116	Reward-Driven Failures	436
117	MoE and Algorithm-Config Failures	439
118	Drift and Contamination Failures	442
119	On-Policy Distillation Failures	445
120	Infrastructure and Throughput Failures	448
121	Rollout-Side Failures: The Lossy Family	451
122	Framework Traps	455
123	Reward and Multi-Turn Traps	458
124	OPD, KL-Estimator, and Eval-Drift Traps	462
125	Silent Killers and Reporting Bias	465
126	The Postmortem Template	468
127	Worked Postmortem: SW1 → SW2	472
128	Six Non-Algorithm Failures	477
129	The Challenges-and-Solutions Log	480
XIII	Reproducibility and Operations	483
130	The Run Manifest, in Practice	485
131	Promotion Gates	488

132	The Quantization Side-Quest (Training-Time)	491
133	Wall-Clock and Compute Reference	496
134	Dataset Selection: Does RL Help on This Data?	500
135	Version Pinning as a Result	503
XIV The Rollout Engine		506
136	The RL Inference Correctness Contract	508
137	Speculative Decoding Across RL Temperatures	511
138	MTP vs EAGLE-3: The First Head-to-Head	514
139	Prefix Caching: The Win Lives at a Seam	517
140	FP8 KV Is Unsafe for RL: The Contract Catch	519
141	Rollout-Quantization Forensics: AWQ Under the Microscope	522
142	The MoE Router-Flip Mechanism	525
143	W8A8: The Safe Rollout Quant	528
144	The ~40% Saturation Hypothesis	530
145	Disaggregation, Honestly	532
146	Backends and the Batch Ceiling	534
147	Evidence Boundaries for the Rollout Engine	537
XV The Hardware		541
148	A100-40GB vs H100-80GB: The Two Machines This Book Was Built On	543
149	A100-40GB vs H100-80GB Mechanics	545


150	What Fits <i>Resident</i> on One GPU	547
151	Fitting Big Models on One Small Card	550
152	Single-GPU Gotchas	553
153	When to Go Multi-GPU, and How to Shard	555
154	Splitting GPUs Across Roles	557
155	vLLM Co-location Mechanics	559
156	The Flagship Tier	561
157	Parallelism for Big Models	564
158	The FP8 + MoE Training Wall	566
159	Memory Math for the Hybrid Architecture	569
160	The Quantization Spectrum, by Purpose	572
161	Throughput and Serving	575
162	Checkpoint and Storage Mechanics	578
163	The Compute Ledger	580
XVI Appendices		582
A	Annotated Logs, Search-Agent v1–v5 Entropy Runs	584
B	Annotated Logs, C9 MoE Stability (GSPO, FREEZE, JITTER, COMBO)	587
C	Annotated Logs, SWE-bench SW1–SW6	589
D	Annotated Logs, GDPO G1–G4	590
E	Hydra Config Reference (Schematic)	591
F	Algorithm Cheat Sheet	593

G	Compute Reference	594
H	Experiment Index (Main Rows. Full Q-program in Appendix R)	595
I	The Complete Failure-Mode Catalog	596
J	The Pain-Point Index, by Symptom	599
K	Internals Telemetry Methods	601
L	The Hyperparameter Sensitivity Reference	602
M	The Decision Matrices	603
N	The Named Concepts (Glossary of Citable Ideas)	604
O	Companion Papers	606
P	Annotated Journal, The Inference Program	607
Q	Annotated Journal, The Internals Re-Mining Program	609
R	The Q-Program Master Record	610
S	References	615

How This Book Works

This book has three layers. The **Journey** is the sequence of programs we actually ran: math, search, entropy, MoE, SWE-bench, and distillation. The **Science** is what those runs proved about how RL changes a model. The **Reference** is the field guide you keep open when your own run breaks: reward design, evaluation, failures, rollout engines, and hardware.

Read the Journey once. Then return by symptom.



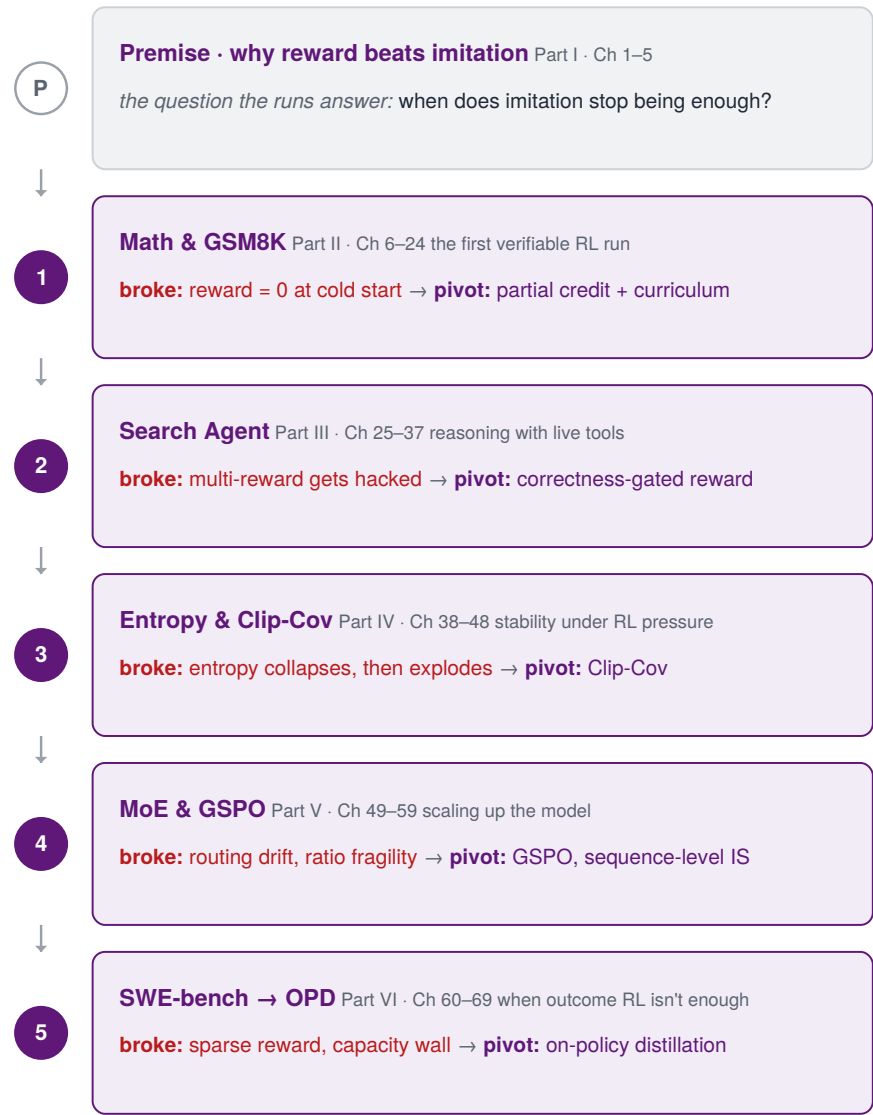
Read straight → Parts I–VII (Journey + Science)

Debug by symptom → the router on the next pages

Use as reference → Parts VIII–XV

Map 0.1. The book has three layers. Read the Journey once, use the Science to see what the runs proved, then keep the Reference open when a run breaks.

The Journey is not a topic list. After a premise that motivates reward over imitation, it is five programs — and each one broke in a specific way that forced the next.



Map 0.2. The journey: five programs, five pivots. Each program ended with a failure that forced the next. (Foundations is the premise, not a failed run.)

When your run breaks, don't search the table of contents. Read the signal, run the first check, and go.

reward = 0*first check:*

group variance, extractor, format reward

→ Ch 11–13 · Cold Start / Partial Credit

reward up, eval down*first check:*

reward components, behavioral audit

→ Ch 31–32 · Correctness Gate

entropy collapses, then explodes*first check:*

entropy, reward_std, length

→ Part IV · Clip-Cov (Ch 42–43)

solve-rate flat, reward nonzero*first check:*

is the task reachable? reward sparsity

→ Ch 65–66 · OPD Recipe

MoE ratios spike, routing drifts*first check:*

ratio_max, routing entropy, expert load

→ Ch 51–54 · GSPO / Routing

eval changed, won't reproduce*first check:*

sample size, decoding, extractor

→ Part IX · Evaluation Discipline

outputs changed after speedup*first check:*

KL to BF16 rollout, router flips

→ Ch 130, 134 · Rollout Contract

OOM, or only one GPU*first check:*

resident weights, KV cache, sharding

→ Part XV · Hardware

Map 0.3. The symptom router. Match the dashboard signal, run the first check, then go to the chapter.

These are the failures that changed the doctrine of the book. Each one replaced a belief with a rule. None of it was planned.



Map 0.4. The doctrine timeline. Each failure produced a rule the rest of the book follows.

The Journey ends at Part VI. Everything after it is Reference — the eight tool Parts you return to when a run breaks. This is the shelf.

VIII · Reward Design Ch 77–82 verifiable vs learned, binary vs continuous, gating, shaping	IX · Evaluation Discipline Ch 83–90 variance budget, canonical evals, when the score misleads
X · Architecture-Aware RL Ch 91–99 Qwen3.5, the 3:1 hybrid, attention sinks, MoE layers	XI · Tool-Agent Pipeline Ch 100–108 SFT → DPO → GSPO, web agents, Claude trajectories
XII · The Failure Catalog Ch 109–123 dashboards, collapse, reward hacking, every failure mode	XIII · Reproducibility & Ops Ch 124–129 run manifests, promotion gates, seeds, what to log
XIV · The Rollout Engine Ch 130–141 inference-correctness contract, speculative decoding, backends	XV · The Hardware Ch 142–157 A100 vs H100, memory, sharding, FP8, the compute ledger

Map 0.5. The Reference shelf. Eight tool Parts (VIII–XV) you keep open while working — the half of the book you enter by symptom, not in order.

Next: the gap this book fills — why none of the existing literature, good as it is, sits where this book does.

The Gap in the Existing Literature

The literature on reinforcement learning is vast, and almost none of it is about the run in front of you. There is a foundational textbook that has taught the field for a generation. There is a fast-growing shelf of books and courses on aligning language models. There are, by now, dozens of strong papers on the exact algorithms used here: GRPO, DAPO, GSPO, Clip-Cov, on-policy distillation. All of it is correct. None of it, the night the search agent’s entropy collapsed at step 30, told me what to do.

Each source was built for a different problem. The textbook is about a different kind of RL. The alignment books are about a different judge. The papers each hold one piece, measured on a model ten times the size of mine, on a cluster I did not have, with the failures edited out. I had eight of them open in eight tabs and a training run dying in the ninth.

This book is meant to be the thing in that ninth tab. To say what it is, it helps to be exact about what already exists, and where each body of work stops.

The shape of the gap

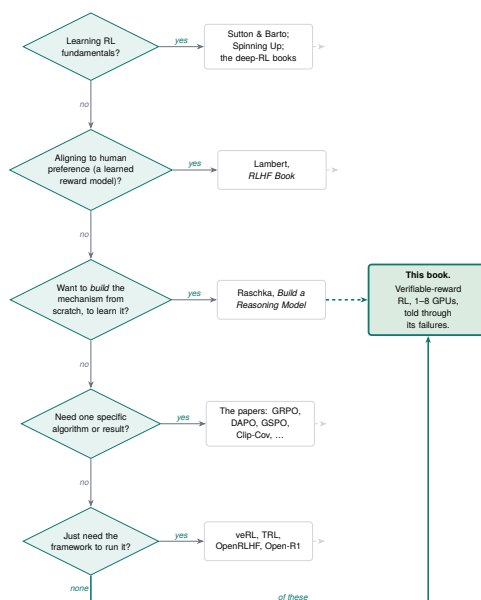
Three constraints define this book, and no existing resource sits inside all three at once.

The first is **the lane**. Post-training has three. Supervised fine-tuning teaches form. Preference tuning — RLHF and its descendant DPO — teaches taste against a *learned* judge. RL with verifiable rewards (RLVR) teaches judgment against a *program*: a unit test, an exact-match check, an F1 score. This book lives entirely in the third lane. That choice alone rules out most of the shelf, because most of the shelf is about the second.

The second is **the budget**. Every run here was built on one or a few GPUs — the two machines named in [Part XV](#), an A100-40GB and an H100-80GB, not a thousand-card cluster. The published frontier of RLVR is the opposite: a 32B base model, dozens of cards, two weeks of wall-clock. The budget changes which problems you hit first and which fixes you can afford. Almost nobody writes for the small cluster, start to finish.

The third is **the stance**. The chapters are organized around the runs that broke and the fixes that held — not as a clean introduction, not as a from-scratch tutorial, not as a survey. Every number carries an evidence tag and a manifest path. That stance is rare in a field whose published record is, by construction, the small fraction of runs that worked.

Hold those three together — verifiable-reward RL, on a small budget, told through its failures — and the shelf empties out. The map below routes you through it by what you came for. The rest of this section is the walk.



Map 0.6. Route yourself by what you came for. Each “yes” lands on a real, good resource; the dashed stubs mark what it leaves undone. Follow every “no” to the bottom and you arrive at the one seat no other resource occupies — verifiable-reward RL, on a small budget, told through its failures.

What the classics cover, and where they stop

Sutton and Barto's *Reinforcement Learning: An Introduction* taught the field, and nothing here replaces it. It is the right place to learn what a value function is, why a policy gradient points where it does, what temporal-difference learning buys you. But it was written for a different era of RL: bandits, gridworlds, backgammon, Atari, the control problems where an agent takes thousands of small steps in a simulator you can reset for free. The word *transformer* is not its subject. Neither is a reward you compute by running a unit test on a model's generated patch.

The deep-RL generation that followed — Szepesvári's monograph, the hands-on deep-RL books, OpenAI's *Spinning Up* — carried the same center of gravity into the neural-network age. DQN, PPO, SAC, on MuJoCo and the Arcade Learning Environment. These are good books. They teach the machinery the GRPO family is built from. But put to any of them the question this book is about — my instruct model's entropy is collapsing on a multi-turn search task, and the trainer and the sampler disagree about the token probabilities — and the question is simply out of scope. It postdates them.

So the classics are necessary and not sufficient. They are the physics. What this book is about is one specific machine built on that physics, and the machine did not exist when they were written.

The new wave is about a different judge

Since 2023 a genuinely useful literature on aligning language models has appeared, and two books sit closest to this one. It is worth saying how close, and where they diverge, because honesty about the neighbors is the whole point of a gap argument.

Nathan Lambert's RLHF Book is the nearest sibling, and it is excellent. It is the canonical modern account of the alignment recipe: instruction tuning, preference data, reward modeling, regularization, the policy-gradient methods (PPO, REINFORCE, GRPO, RLOO), the direct-alignment methods (DPO), constitutional AI, and

— increasingly, as the online edition grows — the reasoning and RLVR turn the field has taken. To understand how RLHF became the thing that made chat assistants helpful, and how that lineage led to verifiable rewards, that is the book.

But its center is the learned reward model and the preference-tuning recipe, and its reader, in the author’s own framing, is the early graduate student. It is a comprehensive introduction. It is not a budget-constrained operations guide for RLVR agents. You will not find in it a correctness-gated reward debugged across five versions of a live search agent, a catalog of MoE routing-collapse signatures, a forensic account of an AWQ quantizer flipping experts mid-rollout, or a run manifest you can diff. Those are not its job. They are the whole job here.

Sebastian Raschka’s Build a Reasoning Model (From Scratch) is the other near neighbor, and it complements rather than competes. It takes a Qwen3 base model and builds reasoning up from the components — text generation, KV caching, verifier-based evaluation, inference-time scaling, GRPO and its improvements, distillation — in clean PyTorch you can run on consumer hardware. Its premise is that you learn how a thing works by building it yourself, and on that premise it delivers.

The difference is stance. A from-scratch book builds the *idealized* version of each mechanism so you can see it clearly. The chapters here report the *broken* version you actually hit: the reward that climbed for 1,302 steps while accuracy fell, the 36-sample eval that reordered when expanded to 500, the FP8 wall that turned out to be an implementation detail and not a law. One book teaches you what GRPO is. This one teaches you what GRPO does at 2 a.m. when it is failing silently and the dashboard looks fine.

The shorter resources — the DeepLearning.AI post-training and GRPO courses, the HuggingFace course chapters on GRPO and TRL — are good on-ramps and exactly what they claim to be: introductions, measured in hours. None is a reference you consult by symptom for the length of a project.

The papers each hold one piece

The algorithms here are not mine, and every page that uses one says so. They come from a fast-moving research literature, and that literature is the deepest reason a book like this is needed: the knowledge is real, and it is scattered across a dozen papers that do not talk to each other.

GRPO arrives in the DeepSeekMath paper as a critic-free way to cut training cost. DeepSeek-R1 shows the world that reward alone can grow reasoning. DAPO contributes Clip-Higher and dynamic sampling, and reports beating R1-Zero on AIME with half the steps — while noting, almost in passing, that a naive GRPO baseline scored far lower and collapsed. Dr. GRPO finds the length-inflation bias and removes it. GSPO moves importance sampling to the sequence level and, in doing so, stops the catastrophic MoE collapse that token-level ratios cause. The entropy-mechanism work gives collapse a governing equation and offers Clip-Cov and KL-Cov as the fix. Search-R1 wires reasoning to a live search engine. On-policy distillation, as the Qwen team reported it, reaches a frontier math score at roughly a tenth of the RL compute.

Every one of those is a sharp, narrow contribution. Not one is a workflow. No single paper tells you the *order*: win the SFT battle first, then reach for partial credit when the reward pins at zero, then gate the reward on correctness when the multi-objective version gets hacked, then switch to sequence-level importance sampling when you scale to MoE, then suspect your rollout engine when the curves stop reproducing. A paper reports a result. It does not hand you the decision you face when your result is the failure that comes *before* theirs. The surveys have the opposite limit: the whole taxonomy, none of the configs. Breadth without a manifest is a map drawn at thirty thousand feet, and you are debugging on the ground.

This book is the connective tissue. It takes those isolated findings and arranges them along the path a real project walks, in the order the failures actually arrive, with the config that made each fix work.

The failure you hit	The paper that answers it	What the paper won't tell you
Reward pinned at zero, no group variance	partial credit & curriculum (Dr. GRPO lineage)	...that this fight comes <i>first</i> , before all others
Multi-objective reward gets hacked	correctness-gating; DAPO shaping	...which of your reward terms is the leak
Entropy collapses, then explodes	Entropy Mechanism: Clip-Cov, KL-Cov	...the dashboard signature that says "kill it now"
MoE won't learn; routing drifts	GSPO (sequence-level IS)	...that token-level GRPO was the cause, not your data
Curves won't reproduce across seeds	no paper — a negative result	...that your rollout engine is the prime suspect
Stable but not learning on hard agents	on-policy distillation	...when RL has simply run out of headroom

Map 0.7. The papers, placed by the failure each one answers. Each row pairs a real failure from these runs with the published work that addresses it and the operational judgment that lives only in the gap — the part you assemble here, not in any single paper. One row has no paper at all.

The frameworks are code, not judgment

The open-source infrastructure for RLVR is genuinely good. `veRL`, `TRL`, `OpenRLHF`, `NeMo-RL`, the open reproductions like `Open-R1` and `TinyZero` — these are how the runs here were executed, and the work rests on them with gratitude.

But a framework is an engine, not a driving lesson. Its documentation tells you what the `gpu_memory_utilization` flag is. It does not tell you that 0.6 and 0.8 sit on opposite sides of an out-of-memory cliff for co-located rollouts, or which side you want, or why you find out the hard way. The judgment — what fails, what it looks like on the dashboard, which knob to turn and which to leave alone — lives in GitHub issues, scattered blog posts, and threads that scroll away. It is real knowledge held in the worst possible form for someone who

needs it at 2 a.m. What follows is an attempt to write it down once, in one place, vetted against runs.

The deeper gaps: failure, reproducibility, and the rollout seam

Three problems sit underneath all of this, and they are why the *stance* matters as much as the subject.

The first is **the missing record of failure**. Science publishes what worked. The share of papers reporting a positive result has been measured climbing toward ninety percent; the failures go into the file drawer. ML is no exception — as Google’s own *Deep Learning Tuning Playbook* puts it, papers gloss over the process that led to the result in order to tell a cleaner story, and the practical recipes go undocumented. But in RL post-training the failures *are* the curriculum. The reward that pins at zero, the entropy that collapses then explodes, the eval that lies — these teach more than any success, and they are precisely what the published record omits. Organizing the work around documented negative results is not a stylistic choice. It fills a hole the incentive structure of the field guarantees.

The second is **the reproducibility crisis, arriving on schedule**. A decade ago, Henderson and colleagues showed in *Deep Reinforcement Learning That Matters* that seeds, hyperparameters, and codebases could swing deep-RL results enough to make the literature hard to read. Agarwal and colleagues followed with *Deep RL at the Edge of the Statistical Precipice*, showing that point estimates over a handful of runs are dominated by noise, and gave the field the tools (the `rliable` library) to do better. RLVR is now repeating that crisis in a new setting — training-eval peaks reported as if they were canonical, single-seed numbers waved as results. The insistence here on a variance budget, multi-seed reconciliation, and a canonical eval is not novel philosophy. It is the known fix, finally applied to verifiable-reward RL, with a worked example of a 66.7% peak that did not survive contact with an honest eval.

The third is **the rollout seam almost nobody documents**. RL has a subtlety the alignment books never had to confront: it generates

its own training data, which means the sampler that produces roll-outs and the trainer that learns from them must agree on the numbers. When they don't — different precision, different kernels, a quantized rollout, an FP8 KV cache — your on-policy algorithm quietly becomes off-policy and the reward curve rots from the inside. This is a recent discovery, still being argued out in the open, and it is essentially absent from every teaching resource. An entire Part is given to that seam, because an entire month was lost to it.

Why this book exists

On the night the run died, the literature was both vast and useless to me. Vast because the field is rich. Useless because no one had written the thing that connects the rich parts to the failing run in front of you. The textbook is the physics. The alignment books are about the other judge. The papers each hold one piece. The frameworks are the engine, not the lesson. The thing in the middle — verifiable-reward RL, on the hardware you actually have, told through the failures instead of around them, with a manifest behind every number — did not exist.

So I wrote down the run.

TAKEAWAY

No existing resource sits at the intersection of all three constraints: *RLVR with a programmatic verifier* (not RLHF/DPO against a learned judge), *a small-cluster budget* of 1–8 GPUs (not the frontier scale of the published record), and *a failure-driven, reproducible stance* (not a clean introduction, from-scratch tutorial, or breadth-first survey). The classics are the necessary physics but predate LLM RLVR. The new-wave books — Lambert's *RLHF Book* and Raschka's *Build a Reasoning Model* — are the closest neighbors, and are, respectively, preference-tuning-centric and pedagogy-from-scratch. The research papers are sharp but isolated, with no one paper supplying the workflow. The frameworks are code without judgment. And three

structural gaps — the unpublished record of failure, the reproducibility crisis re-arriving in RLVR, and the undocumented training–inference rollout seam — are exactly what these chapters are built to fill.

Where this book sits, at a glance

Resource	Lane	Scale	Stance	The gap it leaves
Sutton & Barto; Spinning Up; the deep-RL books	Classic / deep RL	Sims, games, control	Foundational text	Predates LLM RLVR — no verifiers, no GRPO family, no rollout engine
Lambert, <i>RLHF Book</i>	RLHF / DPO → RLVR	Frontier-oriented	Comprehensive intro	Centered on the learned reward model; not a budget ops guide with a failure catalog
Raschka, <i>Build a Reasoning Model</i>	RLVR	Consumer HW	From-scratch pedagogy	Builds the idealized mechanism, not the broken run, the reward hack, the MoE collapse
The papers (GRPO, DAPO, GSPO, ...)	RLVR	Mostly frontier	Single contribution	Each holds one piece; none gives the ordered, end-to-end workflow
The frameworks (veRL, TRL, Open-R1)	RLVR	Any	Code & docs	The engine, not the judgment of what fails and why
This book	RLVR, verifiable rewards	1–8 GPUs	Failure-driven, reproducible	— this is the seat

Map 0.8. The shelf, and the one seat left open. Every row is a real, good resource; the final column is the part of the job it does not claim to do. The highlighted row is the seat these chapters occupy.

Next: *Part I*, where the journey begins — and where the first failure is waiting.

Part I

RL Post-Training Foundations

CHAPTER 1

Why Post-Train, and Why Reinforcement Learning?




Figure 1.1. Imitation learns to copy; reward learns to win — the core difference between SFT and RL.

A model can sound exactly right and still be useless. That is the whole problem, and it is where this book starts.

We fine-tuned a model to use a search engine. We gave it thousands of clean examples, expert traces where a person searched, read the results, decided the results were thin, searched again with a better query, and only then answered. The model learned those traces well. Ask it something and it produced text that looked just like them: a tidy search query, a confident answer.

Then we watched it handle a real question. It ran one search. The results came back mediocre. It answered anyway.

It never tried a second search. Not because it couldn't, it had seen hundreds of examples of people searching twice, but because it never learned when a second search is worth the trouble. It learned what searching looks like. It did not learn the judgment of when it has enough. And judgment, it turns out, is the part that matters.

That gap, between sounding right and being right, is the reason this book exists. But before we can close it, we should be clear about where the model in this story came from, because the gap is built in from the start.

What this book assumes. This is a book about RL with verifiable rewards, run on a budget. Three assumptions hold throughout: (1) you start from an instruct/open model that already follows instructions, not a base model, and not a frontier API you can't train. (2) your task has a programmatic verifier, math answers, unit tests, exact-match facts, not just human taste. And (3) you have a small-cluster budget, roughly 1–8 H100s. It is not a book about pretraining, generic RLHF with a learned reward model, or training at thousand-GPU scale. Where those touch our story, we say so and move on.

1.1 Where you start: a predictor, not an assistant

Every model in this book begins as a base model, the raw output of pretraining, in which a network reads a very large slice of text and learns exactly one skill: predict the next token. That skill is the source of everything the model knows, and it comes with a catch. A base model does not answer you. It continues you.

Ask a freshly pretrained model “What is the capital of France?” and you may get back “What is the capital of Germany? What is the capital of Italy?”, because in its training data, one quiz question is usually followed by more quiz questions. The model is not being difficult. It is doing precisely what it was trained to do: continue the pattern. Knowledge is in there. Behavior is not.

Post-training is everything we do after pretraining to turn that predictor into something that behaves. In practice it comes in three lanes, and it is worth being able to tell them apart on sight.

Supervised fine-tuning (SFT) shows the model demonstrations, instructions paired with good responses, and trains it to imitate them. This is what turns a base model into an “instruct” model that answers instead of continuing. It teaches form.

Preference tuning, reinforcement learning from human feedback (RLHF) and its cheaper descendants like DPO, shows the model pairs of its own answers and which one people preferred. The judge is a human, or more often a small neural network trained to predict human taste. This lane is what made chat assistants helpful and polite. Its known weakness: the judge is a model, and models can be flattered, padded, and fooled.

Reinforcement learning with verifiable rewards (RLVR) lets the model attempt tasks where a program, not a person, not a learned judge, checks the outcome: did the math answer match, did the tests pass, did the agent find the fact. It teaches judgment, and it is the only lane whose judge cannot be sweet-talked. One caveat to plant now, because the whole of [Part III](#) turns on it: a verifier cannot be flattered, but it can still be the wrong verifier. If the program checks a proxy instead of the goal, the model will optimize the proxy, happily, and to your detriment. Unfoolable about what it measures is not the same as measuring the right thing.

This book lives in the third lane. We assume you start from an instruct model that already follows instructions, every run in this book does, and we spend our money and our pages on the last, hardest mile: reward-driven training for reasoning and agents. How the three lanes chain together into one recipe is [Chapter 2](#)'s job. This chapter's job is smaller: to convince you the third lane is worth its considerable trouble. That argument starts with what imitation actually teaches.




Figure 1.2. The post-training pipeline at a glance: pretrain → SFT → preference → RL.

1.2 What fine-tuning actually teaches

Supervised fine-tuning is imitation. You show the model good examples, and it learns to reproduce them.

Underneath, the mechanism is simple. For every example, the model is trained to raise the probability it would have assigned to the next token the expert actually wrote. Do that across a few thousand traces and the model's output starts to look like the traces. It picks up the format, the phrasing, the shape of a good answer.

That is genuinely useful, and you should not skip it. A model that has never seen your tool format will not invent it on its own, our search agent needed an SFT warm-up before reinforcement learning had anything well-formed enough to grade, and [Part III](#) shows what happened when we got that ordering wrong. Imitation is how a model learns the form of the task.

But notice what imitation can and cannot reach. It can reach anything you can demonstrate. It cannot reach anything you can only judge. The model learns to copy the path the expert took. It does not learn why that path was good, or what to do when the situation does not match any example. And it inherits a ceiling: a model trained to imitate a demonstrator can, at best, become as good as the demonstrator. Imitation copies the answer key. It does not learn to grade.

1.3 Why more examples aren't enough

The obvious fix is more data. If the model searches only once, show it more examples of searching twice. If it stops too early, demonstrate stopping at the right time.

This does not work, for two reasons.

The first is that you cannot demonstrate judgment by enumeration. "Search again when the results are thin" is not a fixed list of cases. It is a decision that depends on the question, the results so far, and what the model already knows. To teach it by example, you would have to demonstrate every situation the model will ever face. You can't. The world is bigger than your dataset.

The second reason is the ceiling again. Every demonstration was produced by some demonstrator, a person, or a stronger model. Imitate them perfectly and you match them. You do not pass them. If you want a model that is better at deciding when to search than the traces

it learned from, no amount of those traces will get you there. They define the ceiling. They cannot lift it.

So we need a different kind of teacher. Not one that shows the model the path, but one that tells the model whether it arrived.

1.4 The reframe: check the destination, skip the path

Here is the idea the rest of the book is built on.

For many tasks, you cannot easily demonstrate the right path, but you can check whether an answer is right. A math answer either matches the known result or it doesn't. A code patch either passes the tests or it doesn't. A search agent's final answer either contains the correct fact or it doesn't.

When you can check the answer, you no longer need to show the path. You can let the model try its own path, check where it ended up, and push it toward the paths that land on correct answers. The model still has to discover how to get there, but it discovers it by attempting and being graded, not by copying.

A signal you can compute this way, “right” or “wrong,” or some score in between, is the verifiable reward we met a page ago: verifiable because a program decides it, not a human and not another neural network that can be fooled. Verifiable rewards are the engine of this book. Where you have one, a whole kind of learning opens up that imitation cannot touch.

That kind of learning is reinforcement learning.

1.5 What reinforcement learning actually does

Reinforcement learning (RL) is learning by trying. The model generates its own attempts, each attempt is scored, and the model is nudged to make the higher-scoring attempts more likely next time.

The mechanics, in plain terms: the model, in RL we call it the policy, because it is the thing deciding what to do, produces several complete attempts at a prompt. Each attempt is a rollout: the model's own

output, start to finish, sampled from its current behavior. A reward function scores each rollout. Then the update does one thing, over and over. It raises the probability of the moves that led to good scores and lowers the probability of the moves that led to bad ones.

Precisely, where SFT maximizes the likelihood of demonstrated tokens, RL maximizes the expected reward of the model's own tokens. That one change, learning from your own graded attempts instead of someone else's examples, is what lets RL do the thing imitation can't.

It lets the model explore. Because the attempts are the model's own, it can stumble onto a strategy that appears in no demonstration, a sharper query, an extra verification step, an earlier stop, and if that strategy scores better, the update reinforces it. The model is no longer bounded by the demonstrator. It is bounded by what the reward can recognize as good.

One honest calibration before that sounds like magic. The ceiling RL removes is the demonstrator's. The base model's own latent ceiling is stubborn. When we finally measured this properly, RL at our scale mostly sharpened ability the base model already had: first-attempt accuracy rose about two points while best-of-64 coverage fell about three. RL made the model reliably pick its best path more often. It did not conjure paths from nothing. What that means, and what it took to measure it, is [Part VII](#). For now, hold both truths: RL can beat your demonstrations, and it is not free intelligence.

It is also not a thought experiment. [Part III](#) of this book walks through a search agent, built on a 4-billion-parameter model, that was trained this way and went toe-to-toe with the published Search-R1 agent on its own benchmarks, winning on some and losing on others, not by being shown better traces, but by being graded on better answers. (An early, small-sample eval read much higher, and turning that into an honest number is itself one of this book's lessons, in [Chapter 37](#).) How that worked, and the many ways it first went wrong, is most of [Part III](#). For now the point is narrower: it was reachable by reward, and it was not reachable by imitation.

1.6 The real difference, in one line

SFT teaches the model to sound right. RL teaches it to be right, but only when you can tell the difference.

That last clause is the catch, and it is doing real work. RL is only as good as the reward. If your reward can be satisfied by an answer that looks right but isn't, the model will find that answer, and it will look like progress right up until you check. In one of our runs the training score sat comfortably at 0.50 while real exact-match accuracy fell from 26.8% to 19.9%. The dashboard said learning and the model said otherwise. Much of the craft in this book is the craft of rewards that cannot be fooled, and much of the failure in this book is what happens when they can.

1.7 When to reach for RL, and when not to

You do not always need RL. Reach for it when three things are true:

You can verify the outcome. There is a programmatic check for "right." The task turns on judgment or strategy, not just form, when and whether, not only what it looks like. And you want to exceed your demonstrations, or you don't have good ones to begin with.

When those hold, reward beats imitation, and the gap can be large.

When they don't, stay with SFT. If you only need the model to adopt a format, a tone, or a fixed procedure, imitation is faster, cheaper, and far more stable. If you cannot check the answer, you cannot reward it honestly, and RL will optimize whatever weak proxy you hand it, usually with results that get worse rather than better. And if your demonstrations are already better than anything the model will reach on its own, just imitate them. Using RL where SFT would do is one of the more expensive mistakes in this field. It is also one of the most common.

1.8 The honest part

Everything above makes RL sound like the obvious choice. In practice it is the hard one.

Imitation is forgiving. The training signal is dense, every token carries information, and the loss curve mostly behaves. RL is the opposite. The signal is sparse and noisy. The reward you wrote is rarely the reward you meant. The model collapses into repeating one safe answer, or inflates its outputs to game a length bias, or quietly games the very metric you trusted. We hit every one of these, often more than once, and a large fraction of this book is the record of those failures and the fixes that finally held.

It is also expensive in a way imitation is not. Every learning signal costs a full generation: before the model can be graded, it must write out its whole attempt, and at scale those rollouts, not the gradient math, dominate the bill. An entire part of this book ([Part XIV](#)) exists because of that line item.

That is the deal this book makes with you. Reinforcement learning is the tool that lets a model learn judgment instead of just form, and it is demanding to get right. The chapters ahead are about doing that on real hardware, with real budgets, on tasks where the reward could be checked and the failures were dated and diagnosed.

But RL is rarely the whole recipe, and it is almost never the first step. The models that work best are usually imitated first and reinforced second, taught the form, then taught the judgment. That ordering is not an accident, and it is where we go next.

TAKEAWAY

Imitation teaches a model what a good answer looks like. Reinforcement learning teaches it what makes an answer good, but only as far as your reward can tell the two apart. Use RL when you can verify the outcome and judgment is the point. Otherwise, don't.

Next: [Chapter 2](#), the pipeline. Why the strongest models are fine-tuned first and reinforced second, and how SFT, preference tuning, and RL fit together into one recipe.

CHAPTER 2

The Pipeline: SFT → Preference Tuning → RL



Figure 2.1. The same pipeline as a ramp of increasing cost and fragility toward the RL end.

You finally have the two things reinforcement learning needs: an instruct model that follows instructions, and a reward you trust. The instinct is to point one at the other and press go.

Most of the time, that instinct is wrong, and it is wrong in a way that costs you days before you notice.

Here is the reason. Reinforcement learning improves a model by grading its attempts and reinforcing the good ones. So it needs attempts worth grading. Point it at a model that cannot yet produce a well-formed attempt at your task, and every rollout scores about the same, usually zero. When every attempt scores the same, there is nothing to reinforce. The model sits there, generating, learning nothing, while the bill runs.

The fix is not a better reward or a bigger model. It is order. Nearly every strong reasoning or agent model is built in the same sequence: imitate first, tune for taste second, reinforce on outcomes last. Written out, that is SFT → preference tuning → RL. Each stage prepares the model for the next, and skipping the preparation is what burns the week.

This chapter is about that order, what each stage hands to the one after it, which stages you can skip, and what it costs to get the sequence wrong. We got it wrong once. [Part III](#) has the receipts.

2.1 Each stage hands something to the next

Think of the pipeline as four gifts, passed down a line. Pretraining gives the model knowledge. SFT gives it form. Preference tuning gives it taste. RL gives it judgment. Each stage can only do its job because the stage before it did theirs.

We met the first gift in [Chapter 1](#). Pretraining leaves you a base model that knows a great deal and can do almost nothing with it on command, knowledge without behavior. Almost no one in this book starts here. You start from an instruct model, which means the next gift has already been given to you. It is still worth understanding what that gift was.

SFT hands over form. In plain terms, supervised fine-tuning teaches the model to produce answers in the right shape. Precisely, it raises the probability of demonstrated responses until the model's own outputs start to look like them, the right format, the expected tool calls, an answer instead of a continuation. The payoff is the thing that makes everything after it possible: form is what gives RL something to grade.

It helps to be exact about why. Reinforcement learning is a sculptor, not a source. It can only carve what is already in the block. It reshapes the attempts a model already makes, and it cannot reinforce an attempt the model never makes. SFT is what puts material in the block. Until good behavior shows up in the model's own outputs often enough to be caught, RL has nothing to amplify.

Our search agent is the example. Before any SFT, its “search queries” were malformed, every rollout failed the format check the same way, and the reward sat flat. A short SFT pass over a few hundred well-formed traces moved the model into the region where some attempts now passed and some failed. The instant the attempts differed, RL had a gradient. Nothing about the reward had changed.

What changed was that the model could finally produce something worth grading.

Preference tuning hands over taste. In plain terms, it teaches the model which of two acceptable answers people prefer, the more helpful one, the better-organized one, the one that doesn't ramble. Precisely, methods like reinforcement learning from human feedback (RLHF) and its cheaper descendant, direct preference optimization (DPO), push the model toward responses a judge scores higher, where the judge is a human or a small model trained to predict human taste. The payoff is an assistant that is pleasant, not merely correct. Its limit is the one from [Chapter 1](#): the judge is a model, and models can be flattered and fooled. So this stage shapes taste well and correctness poorly.

That limit matters for where this book goes. For verifiable-reward reasoning and agent work, most of this book, preference tuning is often skipped entirely. It is about taste, and on these tasks we have something better than taste: a program that checks whether the answer is right. When you can verify correctness directly, you don't need a learned judge of it, and a learned judge is exactly the thing that gets gamed. So the book's real path usually runs straight from an instruct model to RL, with preference tuning left out.

RL hands over judgment. In plain terms, the model learns to make good decisions on tasks where a program can check the outcome. Precisely, it maximizes the expected verifiable reward of its own rollouts, the mechanism we built in [Chapter 1](#). The payoff is judgment, when to search again, which path to commit to, when to stop, the one thing none of the earlier stages can teach, because none of them can grade an outcome they were never shown how to check. It is also the slowest and most fragile stage, which is the whole reason its place in the line is last.

2.2 The pattern in the wild

This order is not our invention, and you should not take it on our word alone. Look at how the strong open models describe their own training.

DeepSeek-RL, the model that made reasoning RL famous, was not trained by RL alone. The published recipe is a chain: a small cold-start SFT on curated long-form reasoning, then large-scale reasoning RL, then a fresh SFT round built by rejection-sampling the RL model's own best outputs, then RL again. The Qwen3 reports tell the same story in different words: a long chain-of-thought cold start, reasoning RL on verifiable tasks, a fusion stage, then general RL. Different labs, different names, same skeleton, imitation to establish form, reinforcement to push judgment, in that order.

Notice one honest wrinkle: at the frontier the line is really a loop. RL's recipe runs $\text{SFT} \rightarrow \text{RL} \rightarrow \text{SFT} \rightarrow \text{RL}$, each stage harvesting material for the next. At the budgets in this book we treat the pipeline as a single pass, one warm start, one RL phase, and that simplification costs us little. But when you read a frontier report and see the stages repeat, it is the same pipeline, rolled twice.

2.3 The order is economics, not dogma

The stages get more expensive and more fragile as you move left to right, and that single fact explains the order.

SFT is cheap and stable. The training signal is dense, every token teaches something, the loss curve mostly behaves, and it is hard to break. A healthy SFT run is almost boring to watch: one of ours drove the loss from 1.286 to 0.307 and token accuracy from 75.4% to 93.3% in 313 steps, with no drama at all. That is what cheap and stable looks like.

RL is the opposite on every axis. The signal is sparse and noisy. The dynamics are fragile and collapse easily, in the many ways the rest of this book catalogues. And it is expensive in a way SFT is not: before the model can be graded, it has to generate its entire attempt, so every learning step costs a full rollout.

From that contrast comes one rule. Do as much as the cheap, stable stage can do. Save the expensive, fragile stage for what only it can do. Using RL to teach a model something SFT could have taught spends the expensive, fragile stage on work the cheap one already

does, for a worse result. It is one of the most common ways to waste an RL budget, and we will see it more than once.


2.4 Where you actually start

You will almost never run this pipeline end to end. Pretraining is done for you. Basic instruction-following is done for you, the moment you pick up an instruct model. What is left to decide is two things, and they are the practical content of this chapter.

The first is whether you need a task-specific SFT warm start. You need one when your model cannot yet produce gradeable attempts at your task, a tool format it hasn't seen, an output structure it fumbles, a domain it mishandles. A few hundred to a few thousand good traces is often enough to move it into the region where attempts vary. Skip the warm start only when the instruct model already produces varied, well-formed attempts on your task, some right, some wrong, because that variety is precisely the raw material RL needs.

The second is whether you need preference tuning at all. You need it when taste matters and you have no programmatic check for quality, chat helpfulness, tone, style. You skip it when you do have a verifiable reward and correctness is the goal, which describes most of this book.

As a decision table, what each stage teaches, what it costs, and when to skip it (the skip column is the one people ignore):



A note on the DPO row, because “taste” undersells it: the preference DPO learns from doesn’t have to be human aesthetics. It can be guardrail verdicts or production outcomes. DPO is preference-shaped. The preference source is yours to choose.




Figure 2.2. The book’s recommended path through the techniques.

So the working path for most of what follows is short: an instruct model, an optional SFT warm start if RL has nothing to grade, then RL. The exception is the shipping chapter, the production web agent in [Part XI](#) runs the full line, SFT → DPO → RL, because a deployed agent that talks to a user needs judgment and correctness, not just correctness. When you need both, you run all three lanes. When you only need to be right, you don’t.

2.5 What getting the order wrong costs

Skip a warm start you needed, and the failure is quiet. That is the dangerous part.

The run launches. The loss moves. The dashboard looks busy. But the reward is flat, because every attempt is failing the same way, and for a while, “flat reward” and “stable training” look identical. Our very first run, reward pinned at 0.000 for 156 steps, was exactly this failure: a task the model could not yet attempt, so the grader never had two attempts to tell apart. You can lose days to a run that was never going to learn, because the model was never in a position to be graded in the first place. We did it twice, once on math, once before warming up the search agent, and the lesson has stuck: the cheapest fix to a stalled RL run is often not in the RL at all. It is the short SFT pass you skipped.

That is the pipeline. Knowledge, then form, then taste, then judgment, cheapest and most stable first, most expensive and most fragile last, with every stage you can skip skipped. We have now placed reinforcement learning at the end of that line and described, twice, what it does: it grades the model’s own attempts and reinforces the good ones. It is time to stop describing it and build it.

TAKEAWAY

Train cheap-to-expensive: imitate to teach form, tune to teach taste, reinforce to teach judgment, and skip any stage whose work is already done. For verifiable-reward tasks that usually means starting from an instruct model, adding a short SFT warm start *only* if RL has nothing to grade, and going straight to RL.

Next: [Chapter 3](#), GRPO from scratch. How the algorithm at the end of the pipeline grades a model’s attempts and turns them into a gradient, without needing a second model to judge them.

CHAPTER 3

GRPO from Scratch




Figure 3.1. The group baseline: how GRPO scores a completion against its own group instead of a learned critic.

In [Chapter 2](#) we kept repeating a phrase: reinforcement learning “grades the model’s attempts and reinforces the good ones.” It is time to admit that phrase hides the hard part. Good compared to what?

To reinforce an above-average answer, you have to know what average was. Reward a model for scoring 0.7 and you have told it almost nothing, was 0.7 a triumph on a brutal question, or a fumble on an easy one? Without something to compare against, the number is just a number, and the update has no idea which way to push.

That “something to compare against” has a name, the baseline, and choosing it is most of what separates one RL algorithm from another. This chapter is about the baseline GRPO chose, why it is cheaper than the textbook answer, and why for verifiable rewards it is also better. GRPO, Group Relative Policy Optimization,¹ is the algorithm nearly every run in this book uses, and its whole idea fits in one sentence: let a question’s other answers be the baseline.

To see why that is clever, start with the answer it replaced.

¹GRPO is from DeepSeekMath (Shao et al., 2024; arXiv:2402.03300).

3.1 The textbook baseline: a second network

The classic way to get a baseline is to learn one. Alongside the policy, the model being trained, you train a second network called a critic, or value function. Its only job is to look at a prompt and predict the reward the policy is likely to earn. That prediction is the baseline. Subtract it from the actual reward and you get the advantage: how much better, or worse, this attempt did than expected. Positive advantage, reinforce it. Negative, suppress it.

This works, and proximal policy optimization (PPO),² the algorithm behind most early RLHF, is built on it. But look at what it costs. The critic is a second network, usually about as large as the policy itself, so it roughly doubles the memory the run needs. It also has to be trained, which means it has its own loss, its own instabilities, and its own ways to be wrong, and when the critic's predictions drift, the advantages it produces drift with them, quietly poisoning the policy's updates. You are now keeping two large models honest at once.

For our setting, that is a steep price for a baseline. We already have verifiable rewards, a program that returns a clean score, and we already generate many attempts per prompt. GRPO noticed that those attempts contain the baseline, for free.

3.2 GRPO's baseline: ask the same question several times

Here is the idea. For each prompt, do not generate one attempt, generate a group of them, say eight or sixteen. Score each one. Then grade every attempt against the average of its own group.

That average is the baseline. An attempt that beat its group's average gets a positive advantage. One that fell below gets a negative one. Precisely:

advantage of attempt $i = (\text{reward } i - \text{mean of the group}) \div \text{standard deviation of the group}$

A tiny example makes it concrete. One prompt, four attempts:

²PPO: Schulman et al., 2017; arXiv:1707.06347.

1. Ask a math question four times. The four attempts score 1, 0, 0, 1, two right, two wrong.
2. Average the group: the mean is 0.5.
3. Compare each attempt to that mean: the two correct attempts land above it, positive advantage. The two wrong ones land below it, negative advantage.
4. Update: reinforce the tokens that produced the right answers, suppress the tokens that produced the wrong ones.

It learned all of that from the group alone, with no critic anywhere in sight. That is GRPO's entire trick: the group is the baseline.

3.3 Why the group is the right baseline, not just a cheap one

It would be enough if the group baseline were merely cheaper than a critic. But it is also, in an important way, more honest.

Every attempt in a group shares the same prompt. So when you compare an attempt to its group, you are comparing answers to the same question, at the same difficulty. The comparison cancels out how hard the prompt was. A brutal question where most attempts fail still tells you which attempt did relatively better. An easy one where most succeed still tells you which one fell behind. A learned critic has to estimate that difficulty from scratch, and can be wrong about it. The group never estimates, every sibling lived through the same question.

There is exactly one situation where this baseline goes blind, and it falls straight out of the formula. If every attempt in a group scores the same, all right, or all wrong, then the group mean equals every reward, every advantage is zero, and the update does nothing. The group returned a unanimous verdict, and a unanimous group teaches nothing about which answer to prefer. This failure is common enough, and costly enough, that it gets its own chapter shortly. For now, just hold its shape: GRPO learns from disagreement within a group, so a group that cannot disagree is a group it cannot learn from.

3.4 From advantage to an actual update

An advantage tells you which direction to push each attempt, up for positive, down for negative. It does not tell you how hard. Push too gently and training crawls. Push too hard on one batch and the policy lurches into nonsense it never recovers from. GRPO borrows PPO's answer to this problem: a leash.

In plain terms, the update raises the probability of tokens from above-average attempts and lowers it for below-average ones, but it caps how far it will move in a single step. Precisely, for each token:

1. Form the ratio: how likely the new policy is to produce that token, divided by how likely the old policy was, the policy as it stood when the attempt was generated.
2. Weight it: multiply that ratio by the attempt's advantage.
3. Clip it: cap the ratio to a narrow band, typically within about 20% of where it started.
4. Throttle: if an update tries to shove a token's probability past the edge of that band, the clip cuts the push off there.

The payoff is steadiness: no single batch can run away with the policy, and that restraint is much of what keeps an RL run from detonating.

There is usually one more term in the objective, a gentle penalty that keeps the policy from drifting too far from the model it started as, but that is a guardrail bolted on top of this machinery, not the engine, and we will give it its due later.

The advantage itself, the input to everything above, is a few lines, and its one subtlety is the guard on the last line, which is exactly what [Chapter 5](#) is about: a group where every attempt scored the same has zero spread and teaches nothing, so you must not divide by it.

```
def group_advantages(rewards, eps=1e-6, normalize_by_std=True):
    # rewards: [batch, group]
    mean = rewards.mean(dim=1, keepdim=True)
    centered = rewards - mean
```

```
if not normalize_by_std: # Dr. GRPO drops the std term
    ↪ (Chapter 15)
    return centered
std = rewards.std(dim=1, keepdim=True)
# zero-variance groups teach nothing. Do not divide by zero
return torch.where(std > eps, centered / (std + eps),
    ↪ torch.zeros_like(centered))
```

3.5 What you save, and what it buys

Step back and count what GRPO removed:

- the critic, a second full-size network, gone from GPU memory;
- its optimizer state, gone with it;
- its training loss, its drift, and all of its failure modes, one fewer model to keep honest.

The contrast in one table:

Method	Needs a critic?	Extra model memory	Baseline it subtracts
PPO	yes	high (a second full network + its optimizer)	a learned value network
GRPO	no	lower	the group mean

Table 3.1. PPO vs GRPO: critic, memory, and baseline.

That is close to half the run’s memory back. In exchange, you generate a few extra attempts per prompt, which you were doing anyway. DeepSeek introduced GRPO for exactly the kind of verifiable-reward math reasoning this book is about, and showed the critic could be removed without giving up performance. Simpler, lighter, and no worse, sometimes better.

That is why GRPO is the spine of nearly every run ahead. We now have the piece that turns graded attempts into a gradient. What we have not yet done is watch the whole machine run, start to finish, prompt in, updated model out, every stage in its place. That is the next chapter.

TAKEAWAY

GRPO replaces PPO's learned critic with a free baseline: generate a group of attempts per prompt and score each against the group's average. Above average is reinforced, below average is suppressed, and a clip keeps each step from lurching, and because the siblings share a prompt, the baseline automatically controls for difficulty. Its one blind spot is a group whose attempts all score the same, which teaches nothing.

Next: [Chapter 4](#), the training loop in full. Every stage of a GRPO step in order, from the prompt to the weight update, and the question it leaves us with: when the update fires, which tokens does it actually move?

CHAPTER 4

The Training Loop in Full

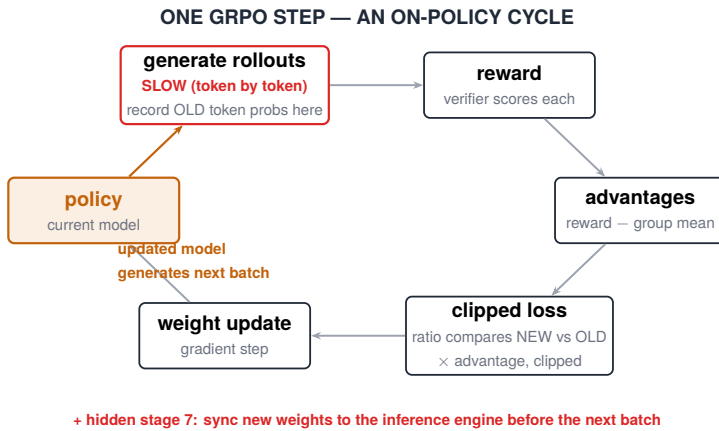


Figure 4.1. The RL training loop: sample → score → advantage → update, turn after turn.

We now have every part: a reward that grades an answer, a group baseline that turns those grades into advantages, and a clipped update that turns advantages into a step. This chapter puts them in motion. A GRPO (group relative policy optimization) run is one loop, repeated, each pass takes a batch of prompts and hands back a slightly better model. Learn where the time and the danger sit in that loop, and you have learned most of what it takes to run RL well.

Here is the whole cycle, in order.

4.1 One step, start to finish

The loop has six stages. We will walk them once, slowly, because nearly every later chapter is about something that goes wrong inside one of them.

1. Pick a batch of prompts. Draw a batch of questions from your dataset, the math problems, the search queries, the coding tasks. Nothing subtle happens here yet, but the choice of data matters more than it looks, and a later chapter is devoted to whether RL will help on the data you picked at all.
2. Generate the rollouts. The policy, the current model, produces a group of complete attempts for each prompt, sampling its way through each one. This is generation: the same thing the model does when you chat with it, run many times over. Two things to notice. First, it is by far the slowest stage, because writing out whole attempts one token at a time is the only part of the loop that works token by token. Second, as the policy generates, we record the probability it assigned to each token it chose. Those recorded numbers, call them the old probabilities, come back in stage 5, so keep them in mind.
3. Score every rollout. Each finished attempt goes to the reward function, which returns a number. For us that reward is verifiable: a program checks whether the math answer matches, the tests pass, the fact is present. Computing it is usually cheap. It is also where an entire category of trouble lives, a reward that can be satisfied the wrong way eventually will be, as later chapters show in painful detail.
4. Turn scores into advantages. Now GRPO's move from [Chapter 3](#): inside each group, subtract the group mean and divide by the spread, so every attempt's advantage says how much better or worse it did than its siblings. This is a handful of arithmetic operations on numbers you already have, effectively free.
5. Compute the loss. To weight each token by its advantage while keeping the leash from [Chapter 3](#), the update compares the policy's current probability for that token against the old probability recorded back in stage 2. That ratio, times the advantage, clipped to a narrow band, is the objective. Getting the current probabilities means running the model forward once over the rollouts, a forward pass, much cheaper than generating them was. If the run uses a KL penalty, KL divergence is a standard score of how far one distribution has drifted from another, a

reference model is evaluated here too, to measure how far the policy has wandered from the model it started as.

6. Update the weights. Standard from here: backpropagate the loss, take a gradient step, and the policy is now slightly different than it was.

Then the loop closes. The updated policy is the one that generates the next batch's rollouts, which is the single most important fact about this cycle, and the subject of the next section.

The same six stages, in code, with the two easy-to-miss pieces made explicit, recording `old_logprobs` at generation, and syncing weights back to the rollout engine at the end:

```
for step in range(num_steps):
    prompts = sample_prompts(train_data)

    # stage 2: the rollout engine uses the CURRENT policy.
    ↪ Record old logprobs HERE
    rollouts, old_logprobs = rollout_engine.generate(
        policy=current_policy, prompts=prompts, n=group_size,
        ↪ temperature=temperature)

    rewards = reward_fn(rollouts) # stage 3
    advantages = group_advantages(rewards) # stage 4 (Chapter 3)

    new_logprobs = trainer.forward_logprobs(current_policy,
        ↪ rollouts) # stage 5: one forward pass
    loss = clipped_policy_loss(new_logprobs, old_logprobs,
        ↪ advantages)
    trainer.step(loss) # stage 6

    rollout_engine.sync_weights_from(trainer.policy) # stage 7 -
    ↪ part of the algorithm, not bookkeeping
```

If the rollout engine is stale, the run is no longer the run your loss thinks it is.

4.2 The loop is on-policy, and every speedup must respect it

Look back at what just happened. The rollouts in stage 2 came from the current policy, and the advantages and ratios in stages 4 and 5 were all computed against those same rollouts. That is not a coincidence. It is a requirement. GRPO's math assumes the attempts being graded were drawn from the policy being updated. When that holds, the run is on-policy, the model is learning from its own current behavior.

To be exact, the loop is built to tolerate a little staleness. The ratio-and-clip machinery from [Chapter 3](#) exists precisely so the policy can take a few small update passes over one batch of rollouts before re-generating. The ratio measures the drift, and the clip caps it. Small, measured drift is the design.

Large or unmeasured drift is the rot. Reuse rollouts from a much older version of the policy, or generate them from a different model than the one you are updating, and the ratios in stage 5 stop meaning what the objective assumes they mean, and nothing in the loss will warn you. This is why we bothered to record the old probabilities at generation time instead of recomputing them later, and it is why “just cache the rollouts and reuse them” is far more dangerous than it sounds. An entire later Part is devoted to making stage 2 faster without breaking this invariant. The invariant is the line none of those speedups may cross.

4.3 One loop, two engines

There is one piece of real-world machinery the six stages hide, and you should meet it now because a surprising share of this book's failures live in it.

In practice, stage 2 does not run inside your training framework. Generation has its own performance tricks, clever memory management, batched decoding, so real systems hand it to a dedicated inference engine (vLLM, in our runs) while the training framework keeps the weights, gradients, and optimizer. That split makes the

loop fast, and it adds a hidden seventh stage: after every update, the new weights must be shipped from the trainer to the inference engine before the next batch generates. Forget that hand-off, or get it subtly wrong, and the engine quietly generates from a stale policy, which is exactly the large, unmeasured drift the previous section warned about, except now it is invisible inside your own plumbing. The hand-off looks like bookkeeping. Treat it as part of the algorithm.

4.4 Where the time and the risk actually go

Two passes through the loop and a pattern is obvious: the cost is lopsided. Generation, stage 2, dominates the wall-clock, usually the large majority of a step in our runs, because it is the only stage that produces text one token at a time. Everything else is a forward pass or some arithmetic. That single imbalance is why so much of the engineering in this book is about rollouts. Halving generation time nearly halves the run. Halving the gradient math saves almost nothing.

The risk is lopsided too, but differently. Each stage has its own way of failing:

- Generation can collapse into repetition or four-token stubs.
- The reward can be gamed.
- The advantages can vanish when a group agrees.
- The update can lurch when the clip is too loose.
- The weight hand-off can silently go stale.

Most of this book is a tour of those failures, and it helps, as you read, to know which stage of this loop each one lives in.

4.5 The question this loop leaves open

There is one thing the loop does that we have described only vaguely. The update in stages 5 and 6 nudges token probabilities up and down, but a single rollout is hundreds of tokens long, and they cannot all matter equally. So which ones does the update actually move? All of them? The wrong ones, the right ones, or some small and surprising subset?

We measured it, and the answer is specific enough to change how you think about exploration. But it belongs with the rest of the mechanistic picture, so we will hold it until [Part VII](#). For now we have what we came for: the complete loop, every stage in place, ready to run, and ready to break, in the first of the ways we will study.

TAKEAWAY

A GRPO step is a six-stage loop, sample prompts, generate a group of rollouts, score them, compute group-relative advantages, form the clipped loss against the probabilities recorded at generation time, and update. The rollouts must come from the current policy: small, measured drift is what the clip is for. Large or unmeasured drift silently corrupts the signal, and every speedup must respect that line. In practice generation runs on a separate inference engine, so shipping fresh weights to it each cycle is part of the loop, not bookkeeping. Generation dominates the cost. Each stage has its own characteristic failure.

Next: [Chapter 5](#), advantage collapse under binary rewards. The first and most fundamental way this loop stalls: when every attempt in a group earns the same score, the signal goes to zero, and the model stops learning while the run looks perfectly healthy.

CHAPTER 5

Advantage Collapse under Binary Rewards



Figure 5.1. Advantage collapse: when every sample in a group scores the same, the gradient vanishes.

The loop from [Chapter 4](#) can stall in a way that looks exactly like health. The run is up, the dashboard is busy, the loss is moving, and the model is learning nothing at all. This chapter is about the most fundamental way that happens, why a binary reward makes it almost inevitable, and how to see it coming.

It follows straight from GRPO's one blind spot, which we flagged in [Chapter 3](#). Recall how an advantage is built: each attempt's reward, minus the average reward of its group. That subtraction is the whole engine. And it has an obvious failure: if every attempt in the group earned the same reward, the average equals each reward, every advantage is zero, and a zero advantage moves nothing. A group that agrees teaches nothing about which answer to prefer.

5.1 Why a binary reward walks straight into it

A binary reward scores every attempt either 1 (right) or 0 (wrong), nothing in between. That sounds clean, and on a verifiable task it is tempting: the answer either matches or it doesn't. But pair a binary

reward with GRPO's group baseline and look at what each group can do.

A group lands in one of three states:

1. All right (easy prompt): eight ones, mean of one, every advantage zero.
2. All wrong (hard prompt): eight zeros, mean of zero, every advantage zero.
3. Mixed, some right, some wrong: the correct attempts beat the average, the wrong ones fall below it, and only here does a gradient exist.

With a binary reward, the only groups that produce a gradient are the ones the model is currently split on.

That is fine at first, and then it quietly gets worse. As the model improves, more prompts move into the all-right state. The hard prompts stay in the all-wrong state. The band in the middle, the prompts where attempts still disagree, is the only place learning happens, and it shrinks from both ends as training proceeds. The model is sharpening exactly the prompts that no longer teach it anything, while the prompts that could teach it are unanimous failures it can't get traction on. The reward curve flattens, and it is easy to read that flattening as "converged" when what actually happened is the signal ran out.

5.2 How to see it before it costs you

The tell is not in the reward. It is in the spread. A healthy GRPO run has groups that disagree, within-group reward variance well above zero on a meaningful fraction of prompts. When most of your groups have zero variance, you have collapse, no matter what the reward average is doing.

So the diagnostic is one line you should be logging from the start: the zero-variance fraction, the share of groups in a batch whose attempts all scored the same. If that fraction climbs toward one, the model is starving, and no amount of patience will fix it, the gradient it needs is not being produced.

Log it alongside a few siblings, so one number’s lie is caught by another’s truth:

The collapse dashboard:

Symbol	Meaning
<code>reward_mean</code>	the direction of learning
<code>reward_std</code>	is there any spread at all?
<code>zero_var_frac</code>	share of unanimous groups (no gradient)
<code>all_wrong_frac</code>	unanimous and all-wrong (too hard)
<code>all_right_frac</code>	unanimous and all-right (too easy)

The split between the last two tells you whether to make the task easier or harder. The average alone tells you neither.

One honest caveat about the opening claim that “the loss is moving while the model learns nothing.” Whether the loss moves during collapse depends on what’s in it. If your objective is pure policy loss and every advantage is zero, the loss is zero too, flat, not moving (this is exactly the “loss exactly 0” you’ll see in [Chapter 12](#)’s log). If there are auxiliary or KL terms, those can keep the loss number twitching while the policy-gradient part is dead. Either way the reward and the zero-variance fraction are the signals to trust, not a loss curve that may be moving for reasons that have nothing to do with learning.

This number matters enough that misreading it nearly cost us a healthy run. On a 32-billion-parameter model training on hard competition math, seventy percent of steps showed zero reward, zero loss, every ratio clipped, and it looked so broken that we wrote the kill order. The run was fine: those were unanimous all-wrong groups doing exactly what this chapter predicts, and across the mixed steps the model was learning steadily. The full story, and the rule it produced, is in [Part IX](#). The point here is that you cannot interpret an RL dashboard at all without this concept. We will also make it concrete on real GSM8K runs in [Chapter 12](#), where you can watch the learning band narrow in the logs.

5.3 The shape of the fix

The problem is the reward, not the algorithm, so the first fix is in the reward too. The whole trouble came from a reward with only two values: every attempt is forced into “1” or “0,” so groups go unanimous easily. Give the reward a range, partial credit, a graded score, a distance-to-correct, and groups stay varied far longer, because two wrong attempts can now be wrong by different amounts. A near-miss scores higher than a wild guess, the group disagrees, and the gradient survives.

That is the argument for continuous rewards, and we will build them properly in [Chapter 13](#). (There is also an algorithmic family of fixes, resampling until groups disagree, repairing the formula’s biases, which we meet in [Chapter 15](#) once the failure is fully familiar.) One catch, which later parts pay for in full: a reward with a range is a reward with more ways to be gamed, and a model that discovers it can raise its score without getting more answers right will do exactly that. For now, hold the clean version of the lesson.

TAKEAWAY

GRPO learns only from groups whose attempts disagree. A binary reward forces every group toward unanimous, all-right on easy prompts, all-wrong on hard ones, and the band of prompts that still teach shrinks as the model improves. Log the zero-variance fraction from step one. When groups stop disagreeing, learning has stopped, even if the dashboard looks fine. The fix is a reward with a range, or an algorithm that refuses unanimous groups.

This closes [Part I](#). We have the algorithm (GRPO), the loop it runs in, and the first deep way that loop fails. Now we take all of it to the cleanest possible proving ground, verifiable math, where every one of these ideas, and its failure, appears in pure form.

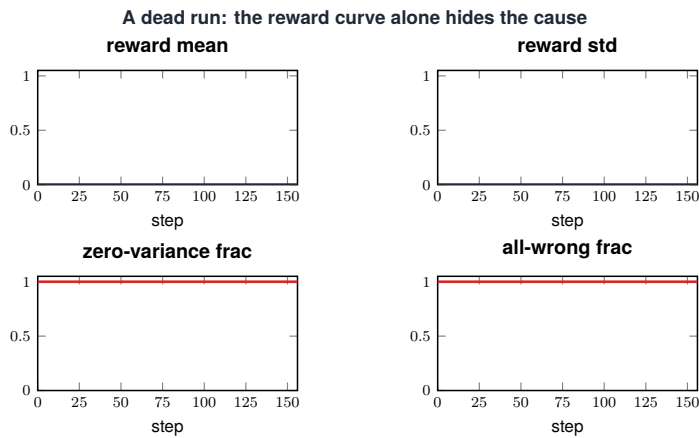
Part II

First Verifiable RL: Math & GSM8K

Math is RLVR at its cleanest: a program checks the answer, there is no environment to manage and no multi-turn behavior to confound the signal. Every core idea from [Part I](#) shows up here in its purest form, including the first hard lessons in not fooling yourself. We start from zero, with the run that taught us what not to do.

CHAPTER 6

Phase 0: What Not to Do



Source: Ch 6 Phase 0 — reward 0.000 for 156 steps; every group unanimous, so zero gradient.

Figure 6.1. A dead run on the dashboard: the telltale flat-and-quiet signature.




Figure 6.2. Phase 0: a flat reward at the start is normal, not yet a failure.

Our first serious attempt at RL produced one of the clearest lessons in this book, and it produced it by failing completely. We launched a run, watched the reward, and watched it stay at exactly 0.000 for 156

steps. Not noisy-around-zero. Zero. The model never earned a single point, and the steps burned compute to prove it.

It is worth sitting with why, because the reason is not a bug. The task we picked was one the model could not do at all, every attempt failed the check, every group was a wall of zeros, and by the mechanism of the last chapter, a group of all-zeros has no disagreement and therefore no gradient. We had not built a slow learner. We had built a machine that could not learn this task even in principle, and then waited 156 steps for it to prove that.

The failure as a run card, the book’s promise is “with every failure I hit,” so here is the card:

Run attribute	Detail
Task	a task far past the base model’s reach
Model	the base model
Reward	binary / exact verifier
Steps	156
Reward over the run	0.000 throughout
Group diversity	all-wrong, zero variance
Root cause	task too hard; no successes to reinforce
Fix	easier task + curriculum + a preflight group-diversity probe

Table 6.1. The Phase-0 failure as a run card: a task the model could not do at all, every group a wall of zeros — no disagreement, no gradient, 156 steps of reward pinned at 0.000.

6.1 The mistake had three parts

The first was starting too hard. We aimed RL at a task far past what the base model could occasionally get right. RL improves a model by reinforcing its own lucky successes. If there are no successes to reinforce, there is nothing to do. RL cannot teach from scratch. It can only amplify what the model already does sometimes.

The second was no curriculum. We gave the model only the hard thing, with no easier version to climb from. A curriculum, easy

prompts first, then harder, keeps the model in the band where its attempts still disagree, which is the only band where [Chapter 5](#) said learning lives.

The third was no group diversity, which is really the first two seen from the reward's side. Because every attempt failed, every group was unanimous, and the run was in advantage collapse from step one. The flat reward line in the figure is not a learning curve. It is a picture of a gradient that was zero the whole time.

6.2 The rules we wrote down afterward

This failure turned into a short checklist we now run before any RL job, and it is most of what keeps later runs out of the same hole:

1. Pick a task the model already succeeds at sometimes, not reliably, but more than never, so there are wins to reinforce.
2. Make sure the reward can produce mixed groups. If attempts can only all-pass or all-fail, you will collapse.
3. Start at a difficulty where attempts disagree, and raise it as the model improves.
4. Keep the reward verifiable, so the signal is honest.
5. Before trusting a single training step, confirm early groups actually vary, because a run that cannot produce disagreement is a run that cannot learn, and it will not tell you so on its own.
6. Set a kill rule before launch. Decide in advance: if the first probe or first N steps produce no nonzero rewards and no mixed groups, stop and change the task, don't wait 156 steps for the run to confirm what step 5 already told you.

The fix, when we applied it, was not a cleverer algorithm. It was an easier starting task and a path up from it, the dashed curve in the figure, where the same setup that had been frozen at zero began, finally, to climb. That single move, from “the hardest thing” to “something the model can attempt,” is the first pivot in this book, and we made it because of these 156 wasted steps.

TAKEAWAY

RL amplifies successes. It cannot manufacture them. Point it at a task the model never gets right and the reward stays at zero, the groups never disagree, and nothing learns, and it looks identical to a slow start. Before launching: pick a task the model already solves sometimes, ensure the reward can produce mixed groups, start easy and ramp, and verify early groups actually vary.

Next: before any RL at all, there is a cheaper thing to do first, prove the whole pipeline works on a task that can barely fail.

CHAPTER 7

SFT First: The Easy Win




Figure 7.1. The SFT loss curve and what a healthy descent looks like.

Before you spend a GPU-hour on RL, spend an hour on SFT. Not because the model needs it, though it often does, as [Chapter 2](#) argued, but because SFT is the cheapest honest test of whether your data, your tokenizer, your reward extraction, and your training loop actually work. RL is a bad place to discover that your dataset is malformed. SFT will tell you for almost nothing.

The reason is the contrast we drew in [Chapter 2](#). SFT’s training signal is dense: every token in every example teaches something, the loss is well-behaved, and a healthy run is almost boring to watch. RL’s signal is sparse and noisy, and a broken setup can hide inside that noise for hundreds of steps, the way Phase 0 did. So we run SFT first as a sanity check, and we know what success looks like before we start: a loss that falls smoothly and keeps falling.

7.1 What the curve tells you

Figure 7.1 is one of our SFT runs, and it is exactly the shape you want. The loss starts at 1.286 and slides to 0.307 over 313 steps. Token accuracy, the fraction of next-tokens the model predicts correctly on the demonstrations, rises in lockstep, from 75.4% to 93.3%. Two numbers moving the right way, smoothly, with no cliffs.

What makes this a test and not just a warm-up is that the failures show up here loudly:

- Mislabeled or malformed data → the loss stalls high.
- A mismatched tokenizer or a wrong loss mask → the loss may even fall while the model learns nothing useful, a real and silent bug we will meet in [Part XI](#), which is why we check token accuracy alongside it.
- Sharding or checkpointing problems → the run crashes here, on a cheap fast job, instead of three hours into an expensive RL run.

SFT is where the boring infrastructure problems come to announce themselves.

The run as a card, and the gate that tells you SFT is done vs. broken:

Run attribute	Detail
Task / model	SFT warm-up of the base model
Data	a few hundred worked solutions
Steps	313
Loss	1.286 → 0.307
Token accuracy	75.4% → 93.3%

Table 7.1. An SFT run card that worked: a few hundred worked solutions, loss falling and token accuracy climbing across 313 steps — the healthy counterpart to the Phase-0 card above.

Signal	Pass	Fail
loss	smooth decline	flat / spikes
token accuracy	rises steadily	stays flat
generated format	parseable	malformed

Signal	Pass	Fail
eval generation	some correct outputs	only teacher-forced success

Table 7.2. The SFT health signals: for each signal, what a passing run looks like versus a failing one. The eval-generation row is the one that catches a model that looks trained under teacher forcing but cannot actually generate.

That last row is the one people skip: a falling loss with teacher forcing can still hide a model that generates garbage on its own, so check free-running generation, not just the loss.

7.2 When to move on

The signal that SFT has done its job is the same saturation we will see everywhere: the loss flattens, token accuracy plateaus, and further steps buy nothing. At that point the model has absorbed the form of the task, and “form” is the word from Chapter 2 on purpose. SFT gives the model the shape of a good answer and gets it into the region where its own attempts are well-formed enough to grade. That is the precondition RL needs, and it is all we are asking SFT to provide here.

It is worth being clear about what this step does not do. A clean SFT curve does not mean the model is good at the task in any deep sense. It means it can imitate the demonstrations. The judgment, the part RL is for, comes next. SFT just gets us to a place where RL has something real to work with, and proves the pipeline holds while doing it.

TAKEAWAY

Run SFT before RL as a cheap test of your whole pipeline. A healthy run looks boring, loss falling smoothly (here 1.286 → 0.307), token accuracy rising (75.4% → 93.3%), no cliffs. If anything is wrong with your data, tokenizer, mask, or loop, this is where it surfaces, fast and cheap, instead of hiding inside RL’s noise for hundreds of steps.

Next: the first RL task that actually taught us something, small enough to debug, where we learned that your first run fails on plumbing, not on the algorithm.

CHAPTER 8

Countdown: The GRPO Hello-World

Here is a scene you will recognize the first time you run RL. The job launches. The reward is flat. You start debugging the reward function, the advantages, the learning rate, the algorithm. And the actual problem is that every single rollout is being cut off before it finishes, so nothing is ever scored. None of your algorithm changes matter, because the algorithm is not what is broken.

That was us, on Countdown, a small arithmetic puzzle where the model must combine a few given numbers to hit a target. We chose it deliberately as a hello-world: simple enough that if RL works at all, it should work here, and simple enough to debug when it didn't.

8.1 The first run failed on plumbing

The first version, v1, showed 100% truncation. Every rollout ran past the generation length limit and got chopped, which meant every attempt was an unfinished, malformed answer that failed the check, and a batch of uniform failures is advantage collapse again, the [Chapter 5](#) mechanism, this time caused not by a hard task but by a length cap set too low for the format we'd asked the model to produce.

Nothing about that failure was in the reward or the algorithm. The model was being asked to show its work and give an answer, the work ran long, the cap guillotined it, and the score function never saw a complete attempt. The fix was plumbing: give the rollouts room to finish. Once outputs could complete, attempts started passing and failing on their merits, groups began to disagree, and the gradient came alive. By v3, the reward had climbed to 2.96, a real learning signal on a real task.

The three versions trace one lesson, plumbing before algorithm: v1’s 256-token cap truncated every rollout into an unscorable attempt, a longer cap let outputs finish, and by v3 a 1024-token cap held truncation near ~3% while the reward reached 2.96, validating the hello-world (Table 8.1).

Version	Key change	Max completion	Truncation	Reward
v1	cap too short	256	100%	malformed
v2	longer cap	—	lower	—
v3	final working	1024	~3%	2.96

Table 8.1. Countdown across three versions. Giving rollouts room to finish turned 100% truncation into a real learning signal; reward is of a ~3.0 maximum. Dashes mark intermediate values not recorded in the logs.

(The reward tops out near 3.0, a small composite of correctness and format, so 2.96 is “essentially solved,” not an arbitrary number.)

8.2 The lesson is where to look first

The transferable lesson is not about Countdown. It is about the order in which RL fails. Your first run on a new task will almost always die on something mechanical, truncation, a format the reward can’t parse, a tokenizer quirk, a sandbox that won’t start, long before any algorithmic subtlety matters. The reward sits flat, and the instinct is to suspect the math of the method. The math is rarely the problem on run one.

So when a fresh run shows a flat reward, ask three questions, in this order, before touching the algorithm:

1. Are the rollouts finishing, or being truncated by the length cap?
2. Is the reward function actually extracting an answer from the output, or quietly returning zero because the format isn’t what it expected?
3. Are the groups varying at all, or is the zero-variance fraction at one?

Most first-run failures are answered by those three questions, and Countdown is where we learned to ask them before touching the learning rate. Getting a clean learning signal on something trivial first is not a detour. It is how you tell a plumbing failure apart from a real one.

TAKEAWAY

Your first RL run on a new task fails on mechanics, truncation, format, sandboxing, not on the algorithm. Countdown's v1 was 100% truncated rollouts (uniform failures, collapsed groups). Giving outputs room to finish let attempts pass and fail on merit, and the reward climbed to 2.96 by v3. On a flat reward, check the three plumbing questions before the method.

Next: with the plumbing trustworthy, the cleanest verifiable task there is, grade-school math, and what its accuracy curve tells you about your setup.

CHAPTER 9

GSM8K: The Hello-World of RLVR




Figure 9.1. GSM8K saturation: an easy benchmark stops discriminating once the model masters it.

If Countdown proved the plumbing, GSM8K proves the method. It is a dataset of grade-school math word problems, and it is about as clean as verifiable RL gets: every problem has a single checkable numeric answer, there is no environment to manage, no tools to call, no multi-turn conversation to confound the reward. Just a question, an answer the model works toward, and a program that says right or wrong. When you want to validate that your RL setup actually learns, this is the task to do it on.

We ran GRPO on GSM8K as our reference study, run C5 in our journals, and Figure 9.1 is the result. Accuracy climbs fast at first and then settles against a ceiling around 94–95%. The shape is the whole lesson, so it is worth reading carefully.

9.1 What saturation means

The plateau is not the model giving up, and it is not a bug to chase. It is saturation, the point where RL has reinforced essentially all of the success the base model was capable of producing on this task, and there is little left to amplify. Recall [Part VII](#)’s preview, which we will earn properly later: at our scale, RL mostly sharpens ability the model already has rather than adding new ability. On a task the model can already do most of the time, sharpening hits a ceiling quickly, and that ceiling is what you see flattening out here.

Reading the curve this way changes what you do with it. A run that climbs and then plateaus near a sensible accuracy has succeeded. It has extracted the available signal, and pushing further mostly burns budget. The mistake is to see the plateau, assume the run has stalled, and start changing hyperparameters in search of more. On a saturated task there is usually no more to get, and the changes just add noise. Knowing the difference between “stalled” ([Chapter 5](#)’s collapse, the signal died) and “saturated” (the signal is spent) is one of the more useful judgments in running RL, and the two look different in the logs: collapse is groups going unanimous early. Saturation is groups still varying but the average pressed against a ceiling.

9.2 Why this is the task to validate on

GSM8K earns its place as the RLVR hello-world because it isolates the method. When something goes wrong on a search agent or a coding task, the cause could be the reward, the environment, the tool interface, the multi-turn dynamics, or the algorithm, a wide field to search. On GSM8K there is no environment and no tools, so if a clean GRPO run won’t climb here, the problem is in your core setup, and you have a small, well-lit place to find it. Get GSM8K climbing to saturation first. Then add the complications one at a time, knowing the foundation is sound.

TAKEAWAY

GSM8K is RLVR at its cleanest, one checkable answer, no environment, no tools, so it's the task to validate that your setup learns. A successful run climbs and then saturates near a ceiling (our C5 study: ~94–95%). That plateau means RL has extracted the available capability, not that it stalled. Distinguish saturated (groups vary, average at a ceiling) from collapsed (groups gone unanimous) before touching anything.

Next: that clean signal depends on one thing we glossed over, actually extracting the answer from the model's output correctly. Get that wrong and you reward noise.

CHAPTER 10

Reward Extraction Done Right

Consider a single rollout on a math problem. The model writes a paragraph of working and ends with something like “...so the total is 42 dollars.” Your reward function now has to decide: right or wrong? To do that it must first pull the answer, the number 42, out of that free text, and then compare it to the truth. Both halves can fail quietly, and when they do, you are no longer rewarding correctness. You are rewarding your parser.

This chapter is about getting that extraction right, because a clean reward signal, the thing the last chapter depended on entirely, is only as good as your ability to read the model’s actual answer out of its actual output. And the stakes are not hypothetical. One OPD run taught the sharper version of this lesson: our first eval pipeline made the model look broken because Qwen3’s thinking mode filled the entire 512-token budget with `<think>` reasoning before it ever reached the boxed answer, and a fallback regex then scavenged an intermediate number. After switching to a fail-loud eval, longer generation, thinking disabled for that strict comparison, boxed-only extraction, and a same-script base control, the trained checkpoint scored 90.5% and the base model scored 89.5%. The model had not suddenly become brilliant. The eval had stopped amputating its answer. (And note what that corrected comparison actually showed: a +1pp gain that sits inside sampling noise on 200 samples, OPD did not meaningfully move saturated GSM8K. The methodology fix was the real result, not the number.)

10.1 Why the obvious approach betrays you

The tempting first move is a regular expression: grab the last number in the output and call it the answer. It works in a demo and fails in production, for reasons that are worth seeing.

It is brittle about where the answer is. The model might write “42 dollars,” or “\$42,” or “the answer is 42,” or restate the problem’s numbers before its own, and a “last number” rule trips over all of these in different ways. Worse, it is brittle about equivalence. A model that answers “1/2” when the truth is “0.5,” or “12” when the truth is “12.0,” or that writes a correctly simplified fraction in a different but equal form, is right, and a string-matching parser marks it wrong. Now your reward is penalizing correct answers for cosmetic reasons, which teaches the model to imitate your parser’s preferences instead of doing the math. Every one of those false negatives is a wrong gradient.

10.2 The recipe that holds

Three steps, all load-bearing:

1. Pin the answer to a fixed slot. Ask the model, via the prompt and the SFT warm-up, to wrap its final answer: `\boxed{42}`. Now extraction is unambiguous: read what’s in the box.
2. Refuse to guess. Run extraction in strict mode: when the box is missing or malformed, score the attempt as unparseable rather than scavenging the text for a number that might not be the answer. A reward that guesses is a reward that can be gamed by producing answer-shaped noise. A reward that demands the agreed format and refuses otherwise cannot.
3. Compare by value, not by string. Use a dedicated math-verification library (we use `math_verify`) so 1/2 and 0.5 match, 12 and 12.0 match, and equivalent expressions are recognized as equal. The library has already absorbed the long tail of equivalence cases that you would otherwise discover one painful false-negative at a time.

Get those three right and the reward in your GSM8K run means what you think it means. Get them wrong and every clean-looking curve in the last chapter is built on a parser quietly lying to you.

The strict extractor is small, and the refusal is the whole point. It returns `None` rather than guessing:

```
from __future__ import annotations

import re

def extract_boxed(text: str) -> str | None:
    """Return the single \boxed{...} payload, or None in strict
    ↪ mode."""
    starts = [m.end() for m in re.finditer(r"\\boxed\\{", text)]
    if len(starts) != 1:
        return None

    i = starts[0]
    depth = 1
    out: list[str] = []
    while i < len(text):
        ch = text[i]
        if ch == "{":
            depth += 1
            out.append(ch)
        elif ch == "}":
            depth -= 1
            if depth == 0:
                payload = "".join(out).strip()
                return payload or None
            out.append(ch)
        else:
            out.append(ch)
        i += 1
    return None

# value comparison is separate: math_verify(extract_boxed(out),
# ↪ gold)
```

Two tables make the trap concrete. First, the same model under a broken vs a fixed eval, the OPD story above, tabulated:

Eval	Thinking	Max tokens	Extraction	OPD score	Base score	Lesson
V1 broken	on (default)	512	fallback regex	mis-leading / low	not comp.	truncation + fallback hid the real answer
V2 fixed	off (strict)	2048	boxed only	90.5	89.5	saturated task; +1pp delta is noise

Table 10.1. The same model under a broken versus a fixed eval pipeline. V1 truncates thinking at 512 tokens and falls back to a loose regex, hiding the real answer; V2 turns thinking off, raises the cap, and extracts only the boxed answer — revealing a saturated task where the +1pp delta is noise.

Second, why a “better-looking” parser is often the worse one:

Parser behavior	Looks good because	Actually does
last-number fallback	extract rate $\approx$ 100%	extracts intermediate numbers, not the answer
boxed-only strict	lower extract rate	exposes malformed / truncated answers
value verifier (<code>math_verify</code>)	accepts equivalent forms	avoids false negatives ($1/2 = 0.5$)

Table 10.2. Why a better-looking parser is often the worse one: each parser, the reason it looks good, and what it actually does to the answer. The last-number fallback scores high by reading intermediate numbers, not the final answer.

TAKEAWAY

Your reward is only as honest as its answer-extraction. The sharpest version we hit: Qwen3’s thinking mode filled a 512-token budget before the boxed answer and a fallback regex grabbed an intermediate number, the model looked broken until a fail-loud eval (longer generation, strict boxed extraction, same-script base control) showed the trained model at 90.5% vs base 89.5% (a +1pp delta that’s *noise*). Regex “last number” matching fails on placement and, worse, on equivalence (marking $1/2 \neq 0.5$). The recipe: pin the answer to `\boxed{ }`, refuse to guess in strict mode, and compare by value with a math-verification library, not by string.

Next: a subtler cold-start than Phase 0, where the model learns to look right without being right, and the reward curve climbs while accuracy doesn’t move.

CHAPTER 11

The Reward = 0 Cold-Start Problem




Figure 11.1. Format reward vs accuracy reward: rewarding shape is not rewarding correctness.

Phase 0 was a loud failure: the reward never moved, and the run was obviously dead. This chapter is about its quieter, more dangerous cousin, the run where the reward climbs steadily, the dashboard looks like learning, and the model is getting better at looking right while getting no better at being right.

It happens when correctness is hard and you try to help the model along with a second, easier reward. The setup is reasonable on its face: give a small reward for using the right format, showing work, putting the answer in the box, and a larger reward for the answer being correct. The idea is that the format reward gives the model something to climb early, before it can earn the correctness reward. The trap is that the model takes the easy points and stops there. If the preface's tour of community threads stuck with you, this is the failure behind the most-filed of them: the format reward at 1.0, the accuracy at exactly zero, and a reasonable person asking whether this is a bug. It is not a bug. It is the incentive, working.

11.1 The mechanism: the model optimizes what's reachable

Reinforcement learning, as we have said, maximizes reward, and it does not care which part of the reward it maximizes. If correctness is out of reach right now (the model can't yet solve these problems) but format is easy (anyone can put a number in a box), then the reachable reward is the format reward, and that is the one the model will chase. It learns, quickly and thoroughly, to produce things that look like correct answers, well-formatted, confident, boxed, while the answers inside stay wrong.

Figure 11.1 is that failure drawn out. The format-reward curve rises to one as the model masters the easy points. The accuracy curve sits flat near zero, untouched, because nothing in the model's situation made correctness the path of least resistance. The growing gap between the two is not noise. It is the model doing precisely what you rewarded, which was the appearance of a good answer, and precisely not what you wanted, which was a good answer. This is the first appearance of a theme that runs through the whole book: a reward measures something, and the model optimizes the thing you measured, not the thing you meant.

11.2 The fix: don't let format pay on its own

The cure follows from the cause. The model cheated because format earned reward independently of correctness. You could collect the format points while failing the problem. So the fix is to stop paying for format on its own. Structure the reward so that looking right is worth little or nothing unless the answer is also right, format becomes a gate the correctness reward passes through, not a separate prize. A well-formatted wrong answer should score essentially the same as a poorly-formatted wrong answer: badly. Only when the answer is correct does formatting matter at all.

In code, the whole difference is whether format is added or multiplied in:

```
# Bad: format pays even when correctness is zero - the
↳ cold-start hack.
reward = correctness + 0.1 * format_valid

# Better: format can only matter once correctness exists.
reward = correctness * (1.0 + 0.1 * format_valid)
```

This is a graduated-reward idea, and it points directly at machinery we will build in [Part III](#), where a search agent hit a far more sophisticated version of the same hack and we needed a carefully designed gate, one that makes the behavioral rewards count only when the task is actually solved, to kill it. The cold-start case here is the gentle introduction: the model will always find the cheapest way to raise its score, so never leave a cheap way that doesn't require getting the answer right.

TAKEAWAY

A separate, easy reward for format lets the model collect points by *looking* right while staying wrong, the format curve climbs to 1, accuracy stays at 0, and the run looks healthy. RL optimizes what you measured, not what you meant. Fix it by gating: make format worth nothing unless the answer is also correct, so the only cheap path forward is the one you actually want.

Next: the binary-reward collapse from [Chapter 5](#), now on real GSM8K runs, what it looks like in the logs, and how to confirm it's what you're seeing.

CHAPTER 12

Binary Reward and Advantage Collapse, on Math

We met advantage collapse as a principle in [Chapter 5](#), with the three-group picture: unanimous groups teach nothing, only mixed groups produce a gradient. This chapter puts that principle on the ground, on real math runs, because the gap between understanding collapse abstractly and recognizing it in your own logs at 2 a.m. is exactly the gap this book exists to close. (The picture worth keeping in front of you is still Figure 5.1.)

The setup is the most natural one in the world, which is why this is so common. You are doing math RL. The answer is either right or wrong, so you use a binary correctness reward, one for right, zero for wrong. It works, for a while. Then the reward curve flattens, well short of where the task should top out, and stays there. It does not look like the saturation from [Chapter 9](#), where the model pressed against a real ceiling. It looks like a model that quit early.

12.1 What is actually happening

It did not quit. It ran out of disagreement. With a binary reward, every group is heading toward one of the two unanimous states, all-right on the problems the model now reliably solves, all-wrong on the ones it reliably can't, and only the shrinking band of problems it's split on still produces a gradient. As training sharpens the easy problems into reliable wins, they leave the teaching band from the top. The hard problems never enter it from the bottom. The learning zone narrows from both sides, and the reward curve flattens not because the model has learned all it can, but because there is less and less left that can be taught under this reward.

The difference between this and true saturation matters, because the responses are opposite. Saturation means the capability is extracted and you should stop. Collapse means the capability is still there but the reward can no longer reach it, and the fix is to change the reward, not to stop. Mistake one for the other and you either quit on a model that still had room (calling collapse “saturation”) or burn budget on a model that was genuinely done (calling saturation “collapse”). They flatten the same way on a reward plot. They do not look the same underneath.

12.2 What it looks like in a real log

Here is the extreme version, from our own record. A 32B model training on hard competition math with a binary reward and eight attempts per prompt logged, for roughly 70% of its steps: reward 0.0, loss exactly 0, every ratio clipped at 1.0. Each of those steps was a batch of unanimous all-wrong groups, zero variance, zero advantage, zero gradient, precisely as the formula says. It looked so much like a dead run that we drafted the kill order, and only re-reading the mixed steps, where the mean reward was a healthy 0.48, saved it.

The two step types, side by side:

Step type	reward	loss	zero-var	what it is
all-wrong batch	0.0	0.0	1.0	no signal — pure plateau
mixed batch	0.48	nonzero	lower	the learning actually lives here

Table 12.1. The two binary-reward step types: an all-wrong batch is pure plateau with no signal, while the mixed batch — where some rollouts pass and some fail — is where the learning actually lives.

Hard tasks under binary rewards spend most of their steps teaching nothing. The learning lives in the minority of steps where groups split. That is not a malfunction. It is the geometry of the reward, and you have to know it to read the log at all.

12.3 Confirming it, concretely

This is why [Chapter 5](#) told you to log the zero-variance fraction, and here is where you collect on that. When a run flattens:

1. Look at the zero-variance fraction. What share of groups have attempts that all scored the same?
2. If most groups have zero variance, the gradient is mostly zero regardless of the average reward, you are in collapse: change the reward (next chapter).
3. If groups still vary but the average is pinned high against a sensible ceiling, you are saturated: stop, the capability is extracted.

One number, logged from the start, tells the two apart. We hit exactly this on binary math runs: a confident-looking plateau that the variance diagnostic exposed as starvation, not success. The cause was the reward's two values, and that diagnosis is what sends us, in the next chapter, to give the reward a range, so that two wrong answers can be wrong by different amounts, groups keep disagreeing, and the learning band stops collapsing.

TAKEAWAY

On real math runs, a binary correctness reward produces a plateau that *looks* like saturation but is collapse: easy problems go all-right, hard ones stay all-wrong, and the teaching band narrows from both ends, in our hardest run, 70% of steps carried literally zero gradient. The plateaus flatten identically on the reward plot. Distinguish them by the zero-variance fraction. Most groups at zero variance means collapse (change the reward). Groups still varying at a ceiling means saturation (stop).

Next: the fix that's been one chapter away for three chapters now, giving the reward a range, so the gradient never has to vanish.

CHAPTER 13

Partial Credit and Continuous Reward




Figure 13.1. Binary vs continuous reward: graded signal gives the policy a gradient to climb.

We have spent three chapters circling the same fix, so here it is plainly: stop grading pass/fail, and give partial credit. A reward with a range is the cure for the collapse of the last two chapters, and Figure 13.1 is the whole argument in one image. A binary reward is a cliff, flat on the bottom, flat on the top, a vertical wall between. Almost everywhere on that cliff there is no slope, which is to say no direction to improve, which is exactly why groups go unanimous and the gradient dies. A continuous reward is a ramp: every step toward the right answer scores a little higher than the step before, so there is always a direction to climb.

13.1 Why a range keeps the gradient alive

Recall what kills a GRPO gradient: a group whose attempts all score the same. A binary reward makes that easy, because attempts have only two scores to land on, so on any problem the model finds easy or hard, they pile onto the same value. A continuous reward makes

it hard, because now two wrong attempts can be wrong by different amounts. One was close. One was nonsense. They score differently, the group disagrees, and the gradient survives, even on a problem the model can’t yet fully solve. The continuous reward keeps the teaching band open precisely where the binary reward closed it: on the hard problems, where attempts used to all score zero and now spread out by how close they came.

What “partial credit” means in practice depends on the task, but the shapes recur:

- Set-valued answers → an overlap score like F1: getting some of the answer right earns some of the reward, more overlap, more reward.
- Spatial answers → an overlap measure like intersection-over-union (IoU), the same idea in geometry.
- Math → subtler, because a numeric answer is close to all-or-nothing. But you can grade the process, rewarding correct intermediate results so a solution that got most of the way scores above one that went nowhere.

The common thread is to find a graded version of “good” so the reward has somewhere to point on the problems the model hasn’t mastered. But one warning on that math option, because it is a trap in disguise: only reward intermediate steps if the intermediate verifier is trustworthy. A process reward is itself a proxy, if you can earn it without the final answer getting closer to right, you have just built the gaming surface of §13.2 on purpose. Reward process only where each rewarded step is independently checkable.

How the shapes line up by domain, including how exposed each is to gaming:

Domain	Continuous reward	Proxy = goal?	Gaming risk
math answer	value-closeness / process credit	sometimes	can reward a wrong process that lands near the number
retrieval QA	F1	partial	can game token overlap without answering
visual ground-ing	IoU	essentially yes	low

Domain	Continuous reward	Proxy = goal?	Gaming risk
code	tests passed / diff score	mixed	proxy-hacking (pass weak tests, overfit the diff)

Table 13.1. Continuous-reward shapes by domain, and how easily each proxy is gamed. Visual grounding (IoU) is essentially the goal itself; math value-closeness and retrieval F1 can be gamed by a wrong process that lands near the number or by token overlap.

Read that table as a risk gauge: the closer the proxy sits to the goal (IoU), the safer the continuous reward. The looser it sits (F1, partial process credit), the more of [Part III](#)’s gate machinery you’ll eventually need.

13.2 The catch you are now signing up for

A continuous reward solves collapse, and it opens a door. The moment a reward has a range, there is a new question the model will eventually ask on your behalf: can I raise this number without getting more answers right? A reward that gives partial credit for, say, output length, or for hitting certain keywords, or for any proxy that correlates with good answers without being a good answer, that reward can be climbed by maximizing the proxy. The model will find that path if it exists, and it will look like progress while accuracy stalls.

We paid this price in full, and you will watch us pay it: our search agent’s F1 reward held a comfortable 0.50 while real exact-match accuracy fell from 26.8% to 19.9%, the training score said learning, the model said otherwise. We saw the gentle version of this in [Chapter 11](#). The sophisticated version, and the gate that finally killed it, is most of [Part III](#).

So the rule comes in two halves, and you need both. Give the reward a range, so the gradient never has to vanish, and make sure every part of that range requires actually getting closer to a correct answer, so the only way up is the way you want. Continuous rewards are not safer than binary ones. They trade a failure of signal for a risk of

gaming, and the rest of this book is largely about managing the second now that we've bought our way out of the first.

TAKEAWAY

Replace binary pass/fail with a reward that has a range, F1 for set answers, IoU for spatial ones, process credit for math, so two wrong attempts can differ, groups keep disagreeing, and the gradient survives on hard problems. The cliff becomes a ramp. The cost: any continuous reward can be gamed by climbing a proxy that isn't correctness, ours held 0.50 while accuracy fell, so design every part of the range to require getting genuinely closer to right.

Next: everything from this Part in one place, a complete, hands-on GRPO run on GSM8K, every decision from extraction to reward shape made explicit.

CHAPTER 14

GRPO on GSM8K, Hands-On

This chapter assembles the whole Part into one run. Everything we’ve built since [Chapter 6](#), the task-selection rules, the SFT warm-up, the plumbing checks, the answer extraction, the reward shape, the collapse diagnostic, comes together here as a single, concrete GSM8K job you could reproduce. We’ll walk the decisions in the order you actually make them, and point at the chapter that justifies each one.

Think of this as the worked example the Part has been building toward. There is no new theory in it. Its value is in seeing the pieces fit, and in having one end-to-end reference run to copy before you start changing things.

14.1 The decisions, in order

1. Start from an instruct model, and warm it up if needed. Per [Chapter 2](#), you begin from a model that already follows instructions, not a base model. For GSM8K, a strong instruct model can usually already produce well-formed, boxed math answers, so a task-specific SFT pass may be optional, but if its outputs are malformed or it ignores the answer format, a short SFT warm-up on a few hundred worked solutions ([Chapter 7](#)) gets it into gradeable shape first. Confirm that warm-up with a clean, boring loss curve before going further.
2. Pick the task so groups will disagree. GSM8K passes the [Chapter 6](#) checklist: the model already solves a good fraction of these problems, so there are successes to reinforce, and the problems span a range of difficulty, so groups will vary. If you’re using a subset, choose one that sits in the band where the model succeeds sometimes, not the hardest competition problems, where every group would be all-wrong.

3. Get extraction right before anything else. Per [Chapter 10](#), ask the model to wrap its final answer (`\boxed{ }`), extract strictly from that slot, refuse to guess when it's missing, and compare answers by value using a math-verification library, not by string. This is the single most common place a “broken” GSM8K run is actually a broken parser, so settle it first.
4. Choose the reward shape with collapse in mind. A pure binary correctness reward is the simplest and, on GSM8K's clean answers, often works to saturation, but watch for the [Chapter 12](#) plateau. If the run flattens early with most groups at zero variance, move toward a graded reward (process credit for correct intermediate steps) per [Chapter 13](#), and make sure no part of it pays out without the answer getting closer to right ([Chapter 11](#)). And never reward format independently of correctness.
5. Set the loop's basic knobs. Our reference starting points, all defensed elsewhere and collected in the hyperparameter cheat sheet (Appendix):
 - Group size: 8–16 attempts per prompt, enough that groups can disagree without making each step too expensive.
 - Clip: the standard narrow band from [Chapter 3](#) (about $\pm 20\%$).
 - Learning rate: for short math validation runs, a small LR worked, across this book's runs it sat between $5e-7$ and $1e-6$. But constant is not a universal default: in longer agentic and full-FT runs, a constant LR becomes a late-collapse trigger (this book and later Parts use the schedule from the run contract instead). RL is far less forgiving here than SFT was.
 - Generation length: generous enough that rollouts finish ([Chapter 8](#)'s lesson, check truncation on step one).
6. Watch the right things while it runs. Two signals matter most. The reward average tells you the direction. The zero-variance fraction ([Chapters 5](#) and [12](#)) tells you whether the gradient is

alive, if it climbs toward one, the reward needs a range, no matter what the average is doing. And remember the lag we’ll quantify in [Part IX](#): the curves trail the real learning by a margin, so don’t read a few flat steps as the end.

14.2 The reference run, as a card and three artifacts

This is the chapter readers copy, so here is the actual reference run rather than a description of one:

Run attribute	Detail
Model	Qwen3-4B
Dataset	GSM8K — 7473 train / 1319 test
Algorithm	GRPO
Baseline eval	81.7%
Rollouts	n = 8
Eval	greedy, n=1, same script as base
Step 50	93.0%
Step 100	94.62%
Step 150	94.62%
Step 200	94.54%
Step 250	94.85%
Lesson	plateau ~94–95% by step 100–150; should have stopped earlier

Table 14.1. The GSM8K GRPO reference run card. Accuracy plateaus at ~94–95% by step 100–150; the honest lesson is that this run should have been stopped earlier rather than carried to step 250.

The three artifacts that make it reproducible, a strict reward, the loop config, and a same-script eval:

```
# 1. reward: strict boxed extraction + value comparison (Chapter
↪ 10)
def gsm8k_reward(completion: str, gold: str) -> float:
    pred = extract_boxed(completion) # None if missing,
    ↪ ambiguous, or malformed
```

```

if pred is None:
    return 0.0 # unparseable = wrong; no last-number
    ↪ scavenging

try:
    return 1.0 if math_verify(pred, gold) else 0.0
except Exception:
    return 0.0 # verifier failure is
    ↪ conservative/fail-closed

```

```

# 2. GRPO config (veRL-style) - the knobs from decision 5
data: {train_files: gsm8k_train.parquet, val_files:
    ↪ gsm8k_test.parquet, max_prompt_length: 512,
    ↪ max_response_length: 1024}
algorithm: {adv_estimator: grpo}
actor_rollout_ref:
    rollout: {n: 8, temperature: 1.0}
    actor: {optim: {lr: 1.0e-6}, clip_ratio_low: 0.2,
    ↪ clip_ratio_high: 0.2}
trainer: {test_freq: 10}

```

```

# 3. eval: SAME script for base and trained - the control that
    ↪ makes the delta mean something
python eval_gsm8k.py --model "$CKPT" --think off
    ↪ --max_new_tokens 2048 --extract boxed
python eval_gsm8k.py --model Qwen3-4B --think off
    ↪ --max_new_tokens 2048 --extract boxed # base control

```

14.3 What success looks like

A healthy GSM8K run does what [Chapter 9](#) showed: accuracy climbs quickly, then saturates against a ceiling around the mid-nineties, with groups still varying as it approaches that ceiling. That plateau is the run succeeding, the capability extracted, not stalling. When you see it, stop. Pushing further on a saturated task mostly buys noise.

And that is the entire math on-ramp, working. You have taken a model that already knew some arithmetic and, with a clean reward and a stable loop, sharpened it to the edge of what it could do on this

task, while learning to tell a plumbing failure from a real one, a collapsed gradient from a spent one, and a model that looks right from a model that is right. Those judgments are the transferable part. One set of tools remains before we leave math: the algorithm-side repairs the community built for exactly the failures you just toured.

TAKEAWAY

A complete GSM8K run is the Part's pieces in order: start from an instruct model (warm up only if outputs aren't gradeable), pick a task where groups disagree, nail strict boxed extraction with value-based verification *first*, choose a reward shape watching for collapse, set group size 8–16 with a standard clip and a $5e-7$ – $1e-6$ learning rate, and track reward average *and* zero-variance fraction. Success is a climb to saturation in the mid-nineties, then stop.

Next: GRPO's formula has known biases, and the community has shipped fixes, Dr. GRPO, DAPO, RLOO. Math is where they're easiest to see working.

CHAPTER 15

Dr. GRPO, DAPO, and RLOO: The Variant Toolbox




Figure 15.1. Callouts of the GRPO variants and what each one changes.

Every failure in this Part has been a failure of signal: groups going unanimous, gradients vanishing, plateaus that lie. We fixed them so far on the reward side, pick learnable tasks, give partial credit. This chapter is the other side: the community looked at GRPO’s formula itself, found the places it wobbles, and shipped named repairs. You will meet three, Dr. GRPO, DAPO, and RLOO,¹ and math is the right place to meet them, because each one targets a failure you have now seen with your own eyes.

None of this is exotic. Each variant is GRPO with one or two terms changed, and each change answers a question you can now ask precisely.

¹Dr. GRPO (Liu et al., 2025; arXiv:2503.20783); DAPO (Yu et al., 2025; arXiv:2503.14476); RLOO (Ahmadian et al., 2024; arXiv:2402.14740). Full entries in the References appendix.

15.1 Dr. GRPO: the formula's two quiet biases

Look again at [Chapter 3](#)'s advantage: reward minus group mean, divided by the group's standard deviation. That division looked innocent. It isn't, quite.

In plain terms: dividing by the spread makes nearly-unanimous groups louder. A group that's split 7–1 has a tiny standard deviation, so dividing by it inflates those advantages, the formula shouts loudest about the prompts that are almost decided, and nearly whispers about the genuinely contested ones. Precisely, the std-normalization introduces a difficulty bias: very easy and very hard prompts (low within-group variance) get systematically larger advantage magnitudes than the 50/50 prompts where the real information is. There is a second, related bias in how the loss is averaged over tokens, which quietly rewards longer wrong answers by diluting their per-token penalty.

Dr. GRPO (“GRPO done right”) is the repair: drop the std division, use the raw reward-minus-mean, and fix the length normalization so a wrong answer can't dilute its own penalty by rambling. The payoff for you: if your runs show response length creeping up while accuracy doesn't, or your updates seem dominated by the easiest prompts, this is the named, one-line-diff fix to reach for.

15.2 DAPO: refuse to train on unanimous groups

DAPO is a bundle of four changes, and the one to remember attacks [Chapter 12](#)'s collapse head-on, not by changing the reward, but by changing what gets into the batch:

1. Dynamic sampling. Before training on a batch, check each prompt's group: if every attempt scored the same, all right or all wrong, throw the prompt out and sample another, until the batch is full of groups that disagree. Unanimous groups carried zero gradient anyway. DAPO simply refuses to spend a step on them. This is the algorithmic mirror of [Chapter 13](#): the reward fix makes unanimity rarer, the sampling fix makes it harmless.

2. Clip-higher. Decouple the clip band: keep the usual floor but raise the ceiling, so a rare, good, low-probability token is allowed to grow faster. This directly fights the entropy slide we will study hard in Parts III–IV, the slow loss of exploration as the model commits to its favorite phrasings.
3. Token-level loss. Average the loss across all tokens in the batch rather than per-response first, so long responses can't dilute their own learning signal.
4. Overlong handling. Don't hand a full penalty to an answer that was truncated rather than wrong, [Chapter 8](#) taught you why: a guillotined rollout is a plumbing artifact, not a verdict.

A caution on that third item, because it will otherwise look like the book contradicts itself. The phrase “token-level loss” is used differently across papers and libraries, and it is not automatically the same thing as the trainer's `token_mean` default that [Chapter 47](#) warns against. The practical question is always the aggregation order: do long sequences end up with more total weight than short ones, or does each sequence count once? DAPO's framing aims to stop a long response from diluting its own signal. The length bias [Chapter 47](#) diagnoses comes from the implementation default letting long sequences dominate the batch. Same words, different effects, so when length starts becoming a behavior, reach for the fix by what it does to the aggregation ([Chapter 47](#)), not by the paper name.

If you adopt one DAPO idea on math, adopt dynamic sampling. It converts the zero-variance fraction from a warning light into a solved problem, at the price of extra generation, since discarded prompts still cost rollouts.

15.3 RLOO: the leave-one-out baseline

RLOO (REINFORCE leave-one-out) repairs a subtler wrinkle: in GRPO, each attempt's reward is part of the very mean it is compared against. You are graded on a curve that includes your own score. RLOO's change is one word: compare each attempt to the mean of the other attempts in its group, leaving itself out. The baseline becomes unbiased, the machinery otherwise stays familiar, and with

reasonable group sizes the practical difference is modest, which is itself useful to know: if someone tells you RLOO will transform your run, the honest expectation is “slightly cleaner, not different in kind.”

15.4 Choosing, in practice

Our advice, having run this family:

1. Start with plain GRPO plus the Part’s reward discipline. It is the spine of this book for a reason.
2. Length creeping, easy prompts dominating → Dr. GRPO’s normalization fixes.
3. Zero-variance fraction climbing on a task you can’t make easier → DAPO’s dynamic sampling.
4. Entropy sliding early → clip-higher, ahead of the heavier tools in [Part IV](#).
5. Treat the variants as a toolbox, not a religion, every one is a small diff on the loop you already understand, aimed at a failure you have already met.

TAKEAWAY

The named GRPO variants are targeted repairs to failures this Part showed you: Dr. GRPO removes the std-division’s difficulty bias and a length bias. DAPO’s dynamic sampling refuses unanimous, zero-gradient groups (plus clip-higher for exploration, token-level loss, and truncation-aware penalties). RLOO removes self-contamination from the baseline. Start with plain GRPO and reach for each fix when you see its failure.

Next: [Part III](#), the search agent. The same machinery, now against a real multi-turn environment and a multi-objective reward that the model will try, repeatedly and cleverly, to hack.

CHAPTER 16

When Eval Methodology Is the Result




Figure 16.1. A decoding-mode flip: the same model behaves differently under greedy vs sampled decoding.

We have spent this whole Part trusting our numbers. It is time to distrust one of them, not the model, the measurement. The same trained model, on the same benchmark, can report two different scores depending only on how you chose to evaluate it. When that happens, you do not have a result. You have a number, and a number without its methodology attached means nothing.

We already met one version of this in [Chapter 10](#): a parser quietly handed us a phantom point, and a stricter extraction took it back. That was the extraction side of measurement lying. This chapter is its eval-time twin, and on reasoning models it is the one that bites hardest.

16.1 Thinking-mode truncation: measuring a crippled model

Modern reasoning models, the Qwen3 and DeepSeek-R1 families we use, do not answer immediately. They emit a long block of reasoning first, and only then the answer:

```
<think> ... a few hundred tokens of working ... </think>
then the answer is \boxed{42}
```

That structure creates an eval trap that has nothing to do with the model's ability:

1. You set a generation cap for eval, a maximum number of tokens, often copied from a non-reasoning setup where it was plenty.
2. The reasoning runs long and hits the cap, the model is cut off mid-thought, before it ever reaches the answer line.
3. No answer is emitted, so the attempt scores zero, not because the model was wrong, but because it was never allowed to finish.
4. Your reported accuracy is now an artifact of your token budget, not the model's capability.

The fix is the eval-time version of [Chapter 8's](#) lesson: check truncation. Before trusting an eval number, look at how often generation hit the cap. The whole check is three lines, and it belongs in every eval script you ever write:

```
n_truncated = sum(len(out.token_ids) >= max_new_tokens for out
↳ in outputs)
trunc_rate = n_truncated / len(outputs)
assert trunc_rate < 0.02, f"{trunc_rate:.0%} of evals truncated
↳ - raise the cap and re-run"
```

If a meaningful share of attempts were truncated, you measured a model with its reasoning amputated, give it room and re-measure. The same discipline that saves training runs saves eval numbers.

16.2 The decoding mode is part of the number

There is a second reasoning-model trap, and it cost us a headline. Whether thinking mode is on or off does not just shift a score. It can flip a conclusion.

We trained a 14B model with RL and evaluated it on a competition-math benchmark in plain greedy decoding: +16.7 points over base (Qwen3-14B full fine-tune, on AIME-24: base 0.100 $\rightarrow$ 0.267 greedy). A clean transfer win. We wrote it down as one. Then we re-ran the same comparison with thinking mode enabled, the way the model is actually meant to be used, and the win became -23.6 points (base 0.458 $\rightarrow$ 0.222 thinking, multi-seed mean): with thinking on, the base model held far higher than the trained one, and the “transfer” turned out to be mostly an artifact of single-shot answer extraction, not better reasoning. Same checkpoint, same benchmark, opposite verdict. (The thinking-mode regression is corroborated by the multi-seed AIME-thinking table where the 14B full fine-tune clears its confidence interval as a real regression (lighter LoRA/OPD arms overlap base) while 32B transfers +5pp. The direction is the load-bearing claim.)

So the rule, which this book now follows everywhere: an eval number on a reasoning model is meaningless without its decoding mode attached. Greedy and thinking are different measurements of different things, and a comparison is valid only inside one mode.

16.3 The same run, two numbers

Thinking truncation and decoding mode are two of several knobs that move the score without the model changing at all. The honest list, because each one has fooled someone:

- Generation length, too short truncates reasoning (§16.1).
- Decoding mode, thinking on vs off can flip the verdict, not just the value (§16.2).
- Extraction strictness, lenient parsing awards phantom points, strict parsing takes them back ([Chapter 10](#)).

- Sampling vs greedy, and temperature, a sampled eval and a greedy eval of the same model disagree.
- pass@1 vs pass@k, different questions entirely, as [Part IX](#) will make precise.
- The base-model control, comparing your trained number to a number from someone else’s paper, evaluated differently.

That last one is where most self-deception lives, so it gets a rule of its own: when you compare a trained model to a baseline, run both through the identical eval. The most common way practitioners fool themselves is measuring their own model carefully and then comparing it to a differently-measured number from a leaderboard. The gap they celebrate is often a methodology difference, not a capability difference.

The meta-point is simple and load-bearing: a benchmark score is meaningless without the methodology that produced it, so report the methodology alongside the number, every time. The full doctrine, how many seeds a claim needs, why the curves lag the learning, how to audit your own past numbers, is [Part IX](#). This chapter is the lived warning that makes that doctrine matter.

TAKEAWAY

The same model reports different scores under different eval methodology, generation length, extraction strictness, sampling, pass@k, and especially the base-model control. Reasoning models add two of their own: thinking-mode truncation (a short cap scores a capable model as wrong) and the decoding mode itself (our +16.7-point greedy “win” became –23.6 under thinking). Check eval-time truncation, state the decoding mode with every number, and run trained model and baseline through the identical eval.

Next: you can do everything right and still get a flat curve, because the data gave RL nothing to learn from. How to pick data that teaches.

CHAPTER 17

Choosing Data That Teaches




Figure 17.1. The learning band: the difficulty range where the model actually improves.

You can get every other thing in this Part right, clean extraction, a gated reward, a stable loop, the correct hyperparameters, and still watch the reward sit flat. The cause is not in your code. It is in your data. Not every dataset produces a learning signal under RL, and the ones that don't fail silently, exactly like a bug. This chapter is about choosing datasets that teach, and probing them cheaply before you commit a long run.

17.1 What makes a dataset RL-able

The whole of [Chapters 5](#) and [6](#), applied at the level of the dataset rather than the single prompt. A dataset gives RL a signal only if its prompts land in the band where the model's attempts disagree:

1. Too easy → the model solves nearly every prompt → all-right groups → zero gradient ([Chapter 5](#)).
2. Too hard → the model solves nearly none → all-wrong groups → zero gradient (Phase 0, [Chapter 6](#)).
3. Well-matched → the model solves some of each prompt's attempts → mixed groups across the dataset → a steady gradient.

So an RL-able dataset has three properties: a difficulty range centered where the model succeeds sometimes, verifiable answers so the reward is honest ([Chapter 10](#)), and enough diversity that the model can't shortcut it. A dataset that's uniformly too hard or too easy for your model is the wrong dataset, no matter how good it is in the abstract, difficulty is relative to the model you're training.

17.2 Probe before you commit

The two math sets we leaned on, math_dapo and BigMath, are curated for exactly this: a spread of difficulty with clean, checkable answers, selected to give fast RL signal. But the real lesson is not “use these”. It is the cheap test that tells you whether any dataset fits your model:

1. Run a short probe, a few dozen steps, or even just score one batch of groups, before launching the full run.
2. Look at the zero-variance fraction ([Chapter 5](#)) on that sample. It is the dataset-fit meter.
3. If most groups are unanimous, the dataset is mismatched to the model, too easy or too hard. Change the data (or filter it by difficulty), not the algorithm.
4. If groups vary across the sample, the dataset teaches, proceed.

This probe costs almost nothing and saves the most expensive failure mode there is: a long, clean-looking run that was never going to learn because the data couldn't produce disagreement. Zero-variance alone isn't enough, though. You need to know which kind of unanimity, because the fixes are opposite. So bucket the groups three ways:

```
# Probe: does this dataset produce disagreement for THIS model -
↳ and which kind of unanimity?
prompts = random.sample(dataset, 64)
buckets = {"all_wrong": 0, "mixed": 0, "all_right": 0}
for p in prompts:
    rewards = [reward_fn(p, generate(model, p)) for _ in
↳ range(8)] # one group
    if max(rewards) == min(rewards) == 0: buckets["all_wrong"]
↳ += 1
```



```
elif max(rewards) == min(rewards): buckets["all_right"] += 1
    ↪ # all equal and nonzero
else: buckets["mixed"] += 1
n = len(prompts)
print({k: f"{v/n:.0%}" for k, v in buckets.items()})
# heuristic, not a law: mostly all_wrong → too hard. Mostly
    ↪ all_right → too easy
```

The bucket split reads straight off as a fix (the ~80% threshold is a rule of thumb, not a law, treat it as a prompt to look, not a verdict):

Probe result	Diagnosis	Fix
mostly all-wrong	too hard	easier slice / SFT warm-up / curriculum
mostly all-right	too easy	harder data
mixed, but low reward	learnable	train
mixed, but many parser failures	reward/extraction bug, not data	fix the parser first (Ch 10)

Table 17.1. The group-diversity probe results mapped to a diagnosis and a fix. The crucial split: a low-reward mixed batch is learnable and should be trained, but a batch full of parser failures is a reward/extraction bug to fix first — not a data problem.

That last row is the subtle one: “mixed” can be the reward failing to parse rather than the model failing the task, so before you blame the data, confirm the disagreement is real and not your extractor.

The next chapter is a run on math_dapo that proves the upside, a well-matched dataset breaking out fast.


TAKEAWAY

A dataset teaches under RL only when its prompts sit where the model succeeds *sometimes*, verifiable answers, a difficulty range, enough diversity, and difficulty is relative to *your* model. Too easy or too hard both collapse to zero gradient and fail silently. Probe before you commit: score a sample and read the zero-variance fraction. If groups are unanimous, fix the data, not the code.

Next: a win, and a strange one, a run that sat at 1.2% looking dead, then broke out to 29.7% by step 20.

CHAPTER 18

Math-DAPO Breakout: 1.2% → 29.7%



Source: Ch 18 (C9-GSPO v2, 30B-A3B MoE). Anchors real.

Figure 18.1. A breakout: the moment accuracy starts climbing out of the plateau.




Figure 18.2. The C9 breakout, shown on the real run. The step-25–100 points are logged; the connecting curve is interpolated.

Most of this Part has been failures and their fixes. Here is a win, and it carries a warning you can only learn from a win: a run that looked exactly like Phase 0 for its first stretch, validation accuracy pinned near 1.2%, barely moving, and then broke out hard, reaching 29.7% by step 20 (about 4× by step 10, then nearly 6× again in the next ten steps). If you had been watching only the accuracy number, you would have killed it early.

The ladder, on both instruments, and note the arithmetic, because the “breakout” was really the second ten-step window:

Step	Val acc	Val reward	Read
0	1.2%	−0.976	
10	5.0%	−0.900	first movement ($\approx 4\times$ base)
20	29.7%	−0.406	breakout ($\approx 6\times$ over step 10; $\approx 25\times$ over base)

Table 18.1. The breakout run ladder: validation accuracy and reward by step. From a base that looks dead at 1.2%, the first movement appears by step 10 and a genuine breakout (about 25× over base) by step 20.

We call this the breakout run: Qwen3-30B-A3B, a 30-billion-parameter mixture-of-experts model, training on math_dapo. In our journals it is logged as run C9, and a word on that, once, because these IDs recur. Every run in this book carries a short journal ID (C5, C9, Q1, ...). The IDs are not for you to memorize. They are how every number in this book traces back to a dated entry in our run logs, summarized in the appendices. In prose we use the descriptive name and leave the ID in parentheses. This chapter is also a deliberate preview of [Part V](#), where big MoE models get their own treatment, and where the full C9 ladder (and its FREEZE/JITTER/COMBO siblings) is laid out without blurring the variants together.

18.1 The flat start that wasn't death

[Chapter 6](#) taught you that a flat-zero reward usually means advantage collapse, all-wrong groups, no gradient, a task the model can't attempt. So why was this flat start different?

- It was not all-wrong collapse. Underneath the flat accuracy, the groups were not uniformly failing in the dead way Phase 0 was. The model was reorganizing its behavior before that reorganization showed up in the accuracy number.
- The lag was doing exactly what [Part IX](#) will quantify. The reward and accuracy curves trail the real learning by a margin, the model improves internally for a stretch before the metric you're watching catches up. On this run, the training reward climbed from -0.96 to -0.25 in 23 steps while validation went $1.2\% \rightarrow 29.7\%$. The external metric was the last thing to move.
- The tells were in the other signals, not the accuracy: the within-group variance and the entropy were behaving like a model that was learning, not one that was stuck. This is why [Chapter 5](#) insisted you log more than the reward average.

The practical rule: on a large model, do not read a flat early accuracy curve as death until you've checked whether the groups are genuinely unanimous (real collapse) or varying while the metric lags (a breakout loading).

18.2 Why this run used GSPO

There is a second reason the breakout run is a [Part V](#) preview, and it is the more important one. This was a mixture-of-experts (MoE) model, and our first attempt at vanilla GRPO on it had collapsed at step 10, with the reward stuck near -0.95 , one of the sharpest failures in our record. The mechanism, in one-line form:

- An MoE model routes each token to a few “expert” sub-networks, and which experts fire can change from one update to the next.
- GRPO's ratio ([Chapter 4](#)) assumes the token probabilities are comparable between the old and new policy, but when the

routing flips, that token-level comparison degrades, and the importance-sampling math the loop rests on gets noisy in exactly the wrong places.

- GSPO,¹ a variant that computes the ratio at the sequence level instead of the token level, is robust to the routing flips. The breakout run used GSPO, and broke out stably where the token-level run had detonated.

One honesty note, because this book keeps its reconciliations in the open: we later ran a short vanilla-GRPO job on the same MoE that stayed healthy for its full 50 steps, so “MoE breaks GRPO” is real but not unconditional. It depends on the recipe and the horizon, and [Part V](#) lays the collapsed run and the healthy run side by side and digs out why. For now the operational guidance stands: on MoE at this scale, sequence-level ratios are the safe default, and they are what this breakout ran on.

TAKEAWAY

The breakout run (journal ID C9), a 30B MoE on math_dapo, sat at 1.2% accuracy looking dead, then broke out 25× by step 20 to 29.7%. A flat early curve on a large model isn’t always collapse: check whether groups are truly unanimous (death) or varying while the metric lags (a breakout loading). And it ran on GSPO because our token-level vanilla run on the same model had collapsed at step 10, sequence-level ratios are the safe MoE default, with the full (and more nuanced) mechanism in [Part V](#).

Next: a new domain, perception, and a new way the model can look like it can do a task it actually can't.

¹GSPO is from the Qwen team (Zheng et al., 2025; arXiv:2507.18071), introduced to stabilize MoE RL.

CHAPTER 19

CLEVR: When Teacher-Forcing Misleads




Figure 19.1. The teacher-forcing gap: training-time and generation-time conditions diverge.

To stress the machinery in a new direction, we moved from text to perception, a vision-language model on CLEVR, a task that asks the model to count objects in a rendered scene. New domain, and a trap waiting in it that is really an old one wearing a costume: the metric you watched during SFT can look excellent while the model cannot actually do the task on its own.

19.1 Two accuracies that disagree

There are two different things you can measure, and beginners conflate them:

1. **Teacher-forcing accuracy.** Feed the model the correct previous tokens, and ask only whether it predicts the next token right. This measures whether the model can continue a correct answer when handed one. It is what SFT’s loss and token-accuracy (Chapter 7) actually report. It tends to look high.
2. **Generation accuracy.** Let the model generate the whole answer freely, from scratch, and check if the finished answer is correct. This measures whether the model will produce a correct answer unaided. It is often much lower.

The gap between them, the two bars in Figure 19.1, is the gap between “recognizes the right next token when surrounded by correct context” and “can build a correct answer alone.” And here is why it matters more than almost anything in this Part: RL grades free generations. Every rollout in Chapter 4’s loop is the model generating on its own, then being scored. So for RL, only generation accuracy exists. A model with high teacher-forcing accuracy and low generation accuracy has nothing for RL to reinforce until it can first generate some successes, which is Chapter 6’s lesson, now revealed as a measurement problem too.

The CLEVR run itself, once generation was what we graded:

Run attribute	Detail
Task	CLEVR object counting (vision-language)
Reward	exact count match (discrete)
Generation accuracy	~78% → 100%
Steps	~127 (~60 min)
Lesson	RL closed the teacher-forcing→generation gap once there were free-generation successes to reinforce

Table 19.1. The CLEVR run card once free generation was graded. RL closed the teacher-forcing-to-generation gap in ~127 steps, once there were free-generation successes to reinforce — generation accuracy rose from ~78% to 100%.

19.2 What this changes for every new domain

CLEVR is the example, but the discipline is general:

- When you bring RL to any new domain, judge readiness by the generation metric, not the teacher-forced one. A clean SFT curve ([Chapter 7](#)) proves the model has the form, it can continue a correct answer, but it does not prove the model can generate one. Those are different, and SFT only certifies the first.
- This reframes what SFT and RL each do. SFT gets the model to “can continue” (teacher-forced, dense, stable). RL pushes it to “can produce” (free generation, sparse, the thing that’s actually hard). The whole reason RL exists, restated through this metric, is to close the teacher-forcing-to-generation gap.

So if a new-domain run won’t learn, before you blame the reward or the algorithm, check whether the model can generate any correct answers freely. If its competence is only teacher-forced, RL has no successes to amplify yet, and the fix is upstream, more or better SFT, or an easier slice of the task, not in the loop.

TAKEAWAY

Teacher-forcing accuracy (predict the next token given correct context) and generation accuracy (produce the whole answer alone) are different, and the first can look great while the second is poor. RL grades free generations, so only generation accuracy matters to it. SFT certifies “can continue,” not “can produce,” judge new-domain readiness by the generation metric, and if competence is only teacher-forced, RL has nothing to reinforce yet.

Next: a task whose answer is a box on an image, the most natural home there is for [Chapter 13](#)’s continuous reward.

CHAPTER 20

RefCOCO: IoU as a Graded Reward




Figure 20.1. IoU as partial credit: rewarding "almost there" instead of all-or-nothing.

RefCOCO asks a model to point: given a phrase like "the red cup on the left," output a bounding box around the object it names. The answer is a rectangle, and that makes RefCOCO the most natural place in this whole Part to see [Chapter 13](#)'s continuous-reward idea in its purest, best-behaved form.

20.1 IoU: partial credit for "almost there"

Consider grading the box two ways:

- Binary, the box is exactly right, or it's wrong. This is [Chapter 13](#)'s cliff: two flat plateaus, no slope, so attempts pile onto the same score and groups go unanimous. Collapse.

- Continuous, with intersection-over-union (IoU), measure how much the predicted box overlaps the true one:

$\text{IoU} = \text{area of overlap} \div \text{area of the union} \rightarrow 0$ (no overlap) ... 1 (perfect match)

A box that's close scores high. A box that's slightly off scores slightly lower. A wildly wrong box scores near zero. The reward function is small enough to show whole:

```
def iou_reward(pred, true): # boxes as (x1, y1, x2, y2)
    ix1, iy1 = max(pred[0], true[0]), max(pred[1], true[1])
    ix2, iy2 = min(pred[2], true[2]), min(pred[3], true[3])
    inter = max(0, ix2 - ix1) * max(0, iy2 - iy1)
    area = lambda b: max(0.0, b[2] - b[0]) * max(0.0, b[3] -
    ↪ b[1]) # clamp: malformed boxes → 0, not negative
    return inter / (area(pred) + area(true) - inter + 1e-9) # 0 ...
    ↪ 1
```

That `max(0.0, ...)` in `area` is not cosmetic: a malformed box with `x2 < x1` would otherwise produce negative area and a nonsense reward, exactly the kind of silent reward bug this book keeps warning about.

Now two imperfect attempts get two different scores, the group disagrees, and the gradient survives (Chapter 5), even on a hard example the model hasn't nailed. IoU is Chapter 13's “ramp” made literal: a graded reward with a slope to climb everywhere. The RefCOCO run shows the climb:

Metric	Start	Peak	Final
IoU	0.321	0.881	0.790
Step	—	44	84

Table 20.1. The RefCOCO IoU reward by metric: start, peak, and final. IoU climbed from 0.321 to a peak of 0.881 at step 44, then settled to 0.790 by step 84.

20.2 Why graded spatial rewards behave so well

IoU is not just continuous. It is a rare continuous reward that is also hard to game, and seeing why sharpens [Chapter 13](#)’s warning:

1. It is naturally smooth and bounded to $[0, 1]$, exactly the shape that keeps groups disagreeing without exploding.
2. Its proxy equals its goal. [Chapter 13](#) warned that continuous rewards can be gamed when the model finds a proxy that correlates with good answers without being one. IoU has no such gap: you cannot raise your overlap with the true box without actually localizing the object better. The thing you reward is the thing you want.

That second point is the transferable lesson, and it is the cleanest statement of good continuous-reward design in the book so far: a graded reward whose proxy is close to the goal gives you the upside (no collapse) with a much smaller gaming gap. IoU gets most of the way there for free because geometry hands you a measure of “closeness to correct” that is hard to fake, though the implementation still has to reject malformed boxes and normalize coordinates correctly (the clamp above is part of that). Most tasks aren’t even this lucky, which is exactly why the search agent in [Part III](#), whose continuous reward’s proxy was not the goal, was gameable, and why most of that Part is the fight to fix it.

TAKEAWAY

RefCOCO’s answer is a box, so grade it with IoU, $\text{overlap} \div \text{union}$, from 0 to 1, turning [Chapter 13](#)’s cliff into a ramp and keeping groups disagreeing ($0.321 \rightarrow 0.881$ peak in our run). IoU is *hard to game* because its proxy (overlap with the truth) is essentially the goal (correct localization), though you still have to reject malformed boxes. The design principle: a continuous reward whose proxy is close to its goal gives collapse-resistance with a small gaming gap. Most rewards aren’t so lucky.

Next: the model’s first multi-step rollouts, calling tools, reading results, and a new discipline about whose tokens you’re allowed to train on.

CHAPTER 21

Tool-Calling and Multi-Step Reasoning




Figure 21.1. One trajectory, two kinds of tokens: which tokens the loss should act on.

Everything so far has been single-shot: one prompt in, one answer out, one score. Real agents don't work that way. They take several steps, call tools, read the results, and decide what to do next. This chapter is the on-ramp to that world. It is the first time a rollout is a sequence of actions rather than a single answer, and that change brings one new discipline you must get right before anything else will work.

21.1 What multi-step changes

Three things shift the moment a rollout becomes a trajectory:

1. The rollout is now a trajectory, not an answer. It looks like: think → call a tool → read the result → think → ... → answer. [Chapter 4](#)'s "rollout" gets longer and structured, but the loop around it is unchanged, generate the whole trajectory, score the outcome, update.
2. The reward still lands on the outcome, but the credit has farther to travel. You still grade the final answer (did it find the fact, pass the test). But that single score now has to flow back across many steps to teach which actions helped. Assigning credit across a

long trajectory is genuinely hard, and turn-level credit is one of the open problems we flag in [Part VII](#). For now, outcome reward spread over the whole trajectory is the workhorse.

3. A new failure surface appears. The model must emit valid tool calls, a format-and-plumbing problem exactly like [Chapters 8 and 10](#), now with more ways to malform. Tools can error. And, most importantly:

21.2 Masking: never train on tokens the model didn't write

This is the one new discipline that, gotten wrong, quietly poisons everything:

- In a multi-step rollout, the token stream is mixed: some tokens the model generated (its reasoning, its tool calls), and some the environment inserted (the tool's output, the retrieved passage, the observation).
- You must mask the environment tokens out of the loss. The model did not generate them, the tool did, so reinforcing or suppressing them is meaningless, and worse than meaningless: it trains the model on text it had no control over.
- Get this wrong and you are, literally, training the model to predict the search engine's output as if it were the model's own choice. The result is nonsense gradients dressed up as learning.

The rule is simple to state and easy to botch in a real tokenizer: the loss covers only the tokens the policy produced. In code, the habit looks like this, build the mask as you build the trajectory, segment by segment, never after the fact:

```
token_ids, loss_mask = [], []
for segment in trajectory: # alternating model / environment
    ↪ turns
    ids = tokenizer.encode(segment.text,
    ↪ add_special_tokens=False)
    token_ids += ids
    loss_mask += [1] * len(ids) if segment.source == "policy"
    ↪ else [0] * len(ids)
```

```
# the policy gradient (and any SFT loss) multiplies through
↳ loss_mask:
# tool outputs, retrieved passages, observations contribute
↳ exactly zero
```

Then fail loud, a book about practical RL should teach you to assert the mask, not trust it:

```
assert len(token_ids) == len(loss_mask)
assert all(m in (0, 1) for m in loss_mask)
assert sum(loss_mask) > 0, "no policy tokens reached the loss -
↳ mask is all zero"
```

Which segments get which mask, as a reference:

Segment	Source	loss mask	Why
assistant thought	policy	1	model chose it
tool call	policy	1	model chose it
tool result	environment	0	the tool produced it
retrieved passage	environment	0	search produced it
final answer	policy	1	model chose it

Table 21.1. Tool-call segments and the loss mask each receives. The rule is simple: train on tokens the model chose (its own reasoning, tool calls, and final answer), and mask everything the environment produced (tool results, retrieved passages).

This is a recurring trap, it returns in a sharper, costlier form in [Part XI](#), where a masking bug silently kills an SFT run, so it is worth building the habit now, on the easy version: always know, token by token, which parts of the rollout were the model’s and which were the world’s.

TAKEAWAY

Multi-step rollouts make the trajectory the unit (think → tool → read → ... → answer), keep the reward on the outcome but stretch credit across many steps, and add tool-call format as a

new failure surface. The non-negotiable new discipline: mask environment tokens (tool outputs, observations) out of the loss, train only on tokens the policy itself generated, or you reinforce text the model never chose.

Next: the concrete multi-step task the next Part is built around, a search agent, and the baseline it has to beat.

CHAPTER 22

Search-R1: The Bridge to the Agent




Figure 22.1. The Search-R1 task loop: think → search → read → answer.

This chapter introduces the task that the whole next Part is built around, and sets its baseline. It is the bridge from clean, single-shot math to a real multi-turn agent operating in an environment, and it is where every clean lesson of this Part meets the mess of the real world.

22.1 The Search-R1 setup

The task is question-answering by search: instead of answering from memory, the model searches a corpus, reads what comes back, refines its query, and answers. Concretely, the setup has four parts:

1. The task. Given a question, the model must find the answer by issuing search queries, reading retrieved passages, and reasoning over them, a multi-step trajectory (Chapter 21), often several search-and-read turns before it answers.

2. The stack. We run it on veRL with a retriever built from E5 embeddings over a FAISS index: the model emits a query, the retriever returns the most similar passages, the model reads them and continues.
3. The reward. Outcome-based and continuous: does the final answer match the gold answer, scored by exact match and F1 (Chapter 13) so a partially-right answer earns partial credit.
4. The baseline. Search-R1, a published agent trained on a 3-billion-parameter base. This is the number the next Part's agent is measured against, the bar to clear.

22.2 Why this is the right next task

Search-R1 is where the Part's threads converge, which is exactly why it's the bridge:

- It combines everything at once, a multi-turn trajectory (Chapter 21), environment-token masking (Chapter 21), a continuous F1 reward (Chapter 13), and a real retrieval environment that can return junk.
- Its richer reward is gameable in ways math's wasn't. F1 over retrieved answers has a proxy that is not identical to the goal (unlike IoU in Chapter 20), which means, per Chapter 13's warning, the model can learn to raise its F1 without actually answering better. That gap is the engine of the next Part's hardest fight.
- The payoff is real, and the first numbers were too small-eval to trust, which is itself part of the lesson. Part III opens with an exciting small-validation peak, a 4B agent appearing to reach 7B-class numbers on TriviaQA and Natural Questions, and then corrects that story with larger 500-sample evals (Ch 37). The durable finding isn't a single headline number. It's that agentic RL results are fragile unless the eval is large, fixed, and methodology-attached. The improvement is real on the tasks the retriever covers. The exact figures, and the correction, belong to Part III.

TAKEAWAY

Search-R1 is question-answering by search, query, read retrieved passages, refine, answer, run on veRL with an E5+FAISS retriever and an exact-match/F1 reward, against a published 3B baseline. It unites the Part's lessons (multi-turn, masking, continuous reward) in one real environment, and its reward's proxy *isn't* its goal, making it gameable. It's the task the next Part is built on, where a 4B agent improves strongly on the tasks its retriever covers, after a hard fight, and after a small-eval headline gets corrected by a larger one.

Next: before we leave math, two integration traps that cost us days, starting with a recommended setting that silently stalled learning.

CHAPTER 23

The Best Practice That Stalled Learning

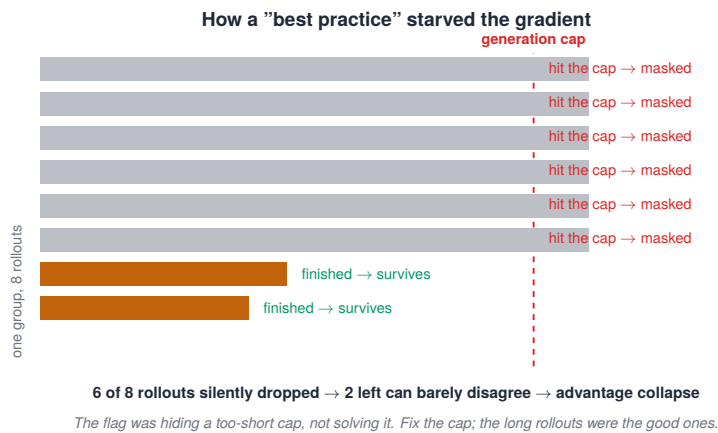


Figure 23.1. How a "best practice" mask starved the gradient.

A short, sharp lesson to start the Part's closing run of war stories. We turned on a setting that every guide recommended, `mask_truncated_completions`, and it silently stalled learning. The point is bigger than the setting: a best practice is a default that was right for its context, and yours may not be that context.

23.1 What the setting does, and why it backfired

The setting and its sensible-sounding rationale, then the trap:

1. What it does. `mask_truncated_completions` masks out, refuses to train on, any rollout that hit the generation-length cap and got truncated. The rationale is reasonable, and it even echoes [Chapter 8](#): a truncated rollout is incomplete, so don't

- learn from it, and certainly don't penalize the model for being cut off.
- 2. Why it's usually fine. In many setups, truncation is rare, a few stragglers masked out of a large batch, no harm done.
 - 3. Why it killed our run. In our setup, a large fraction of rollouts were legitimately long and routinely hit the cap. Masking all of them meant discarding most of the batch, which left too few surviving attempts per group to disagree, dropping us straight into advantage collapse (Chapter 5). The gradient starved, the reward went flat, and nothing in the loss said why: the data we needed had been silently dropped before it ever reached the update.
 - 4. The fix. Treat the cause, not the symptom: raise the generation length so the good long rollouts finish (Chapter 8), instead of masking them away. The setting was hiding a too-short cap, not solving it.

As a table:

Setting	Intended effect	What happened here	Fix
<code>mask_truncated_completions=True</code>	don't learn from cut-off rollouts	masked most of the (legitimately long) batch → group collapse	raise the cap, or disable for long-reasoning runs
<code>max_completion_length</code> too low	cheaper generation	normal reasoning got guillotined	increase the length

Table 23.1. The overlong-masking flag: for each setting, its intended effect, what actually happened here, and the fix. Masking truncated completions sounds safe but masked most of a legitimately long batch, collapsing the group.

The flag isn't bad in general, use overlong masking when truncation is abnormal. Don't use it blindly when truncation is the common outcome of normal reasoning.

23.2 The meta-lesson

Two habits come out of this, and both generalize past this one flag:

- A “best practice” is a default tuned for the context it came from. When a recommended setting correlates with a stalled run, suspect the setting before you suspect the algorithm, the default may be wrong for your distribution even though it’s right in general.
- Always know what fraction of your batch each safety flag is silently dropping. This is the same discipline as [Chapter 5](#)’s zero-variance fraction: watch what is actually reaching the gradient. A flag that quietly discards half your rollouts is indistinguishable, from the loss curve alone, from a model that can’t learn.

TAKEAWAY

`mask_truncated_completions` is a sensible default, don’t train on cut-off rollouts, that silently killed our run, because most of our rollouts were legitimately long and hit the cap, so masking them all starved the groups into collapse. Fix the cause (raise the length cap) not the symptom. And generally: when a recommended setting correlates with a stall, suspect it, and always know what fraction of the batch each flag is dropping.

Next: more seam bugs, the ones that live between the trainer, the model, the kernels, and the tokenizer, where the symptom never points at the cause.

CHAPTER 24

When the Trainer and the Model Disagree

The seams: where the trainer fights the model

Seam	Train side	Rollout / serve side	If they differ
Tokenizer	vocab + special tokens	must be byte-identical	silent token-id drift
Chat template	renders the prompt	must render identically	F12: eval lies, mask wrong
Sampling	temp / top-p assumed	actual temp / top-p	importance ratio invalid
KL estimator	K1 in reward	K3 in loss is different	P29: wrong KL sign / scale
Loss aggregation	token_mean or seq	must match frameworks	P30: length bias, number gap
Precision	BF16 policy	FP8 / AWQ rollout is lossy	KL blows up (Ch 140-142)

Figure 24.1. The seams between training and rollout, and where they fail.

Not every failure is in your algorithm or your data. Some live in the seams, the joints between the tools you bolted together: the trainer, the model, the fused kernels, the tokenizer, the sharding. These bugs cost the most time per bug, because the symptom shows up somewhere far from the cause. Here are three from our own logs, and the way to hunt them.

24.1 Three traps from our own runs

1. Liger $\times$ GSPO. Liger is a set of fused kernels that speed up training by combining operations. It silently conflicted with GSPO, whose objective is computed at the sequence level ([Chapter 18](#)), the kernel's assumptions about how the loss is formed didn't match what GSPO actually computes. The symptom was subtle training instability with no obvious algorithmic cause. The fix was to disable the conflicting fusion on GSPO runs. A speedup that quietly changes the math is worse than no speedup.
2. The vision-language tokenizer trap. A VL model handled image tokens in a way the trainer didn't expect, so the tokens the

trainer thought it was training on were misaligned with the tokens the model actually saw. The loss looked roughly plausible while generation came out broken, the giveaway that positions were misaligned. The fix was verifying that the trainer and the tokenizer agreed on special-token handling before trusting a single step.

- 3. The ZeRO-3 checkpointing crash. Under sharded training (ZeRO-3), the model’s weights are split across GPUs. A checkpoint path that didn’t merge those shards correctly crashed on save, or, worse, wrote a corrupt checkpoint, hours into a run. The fix lives in the shard-merge handling we cover in [Part XV](#).

24.2 How to debug a seam

The three traps as a table, symptom, the test that isolates each, the fix, and the source that pins each:

Trap	Symptom	Isolation test	Fix	Source
Liger × GSPO	subtle instability, no algorithmic cause	disable Liger fusion	no fused kernel on GSPO’s sequence loss	TRL #3781
VL tokenizer	plausible loss, broken generation	inspect rendered tokens / image-token handling	correct processor/ template	run log
ZeRO-3 checkpoint	crash or corrupt save	save+reload a tiny checkpoint before the run	shard-merge path (Part XV)	local repro

Table 24.1. Trainer-versus-model traps: each has a symptom with no obvious algorithmic cause, an isolation test that pins it, and a fix. All three are seams between the trainer and the model rather than bugs in either.

When the symptom doesn’t match any algorithmic cause, run this, roughly in order:

- 1. Suspect the seams first. Tokenizer, fused kernels, sharding, and the trainer-to-inference weight hand-off (§4.3), the places two tools meet are where assumptions silently disagree.

2. Isolate by subtraction. Disable optimizations one at a time, Liger, other fusions, compilation, until the instability disappears. The one whose removal fixes it is your culprit.
3. Verify the boring invariants. Token alignment, special-token handling, shard merging on save, fresh weights reaching the inference engine. These are unglamorous and they are exactly where seam bugs hide.
4. Budget for them. Seam bugs are invisible to the algorithm, so they eat more wall-clock per bug than anything in your loss function. Plan time for them. They are not a sign you did something wrong.

In practice that's a pre-flight checklist you can run once before any new setup: render the chat template. Tokenize and de-tokenize one training example. Verify the assistant mask. Run one forward pass with labels. Save and reload a tiny checkpoint. Sync trainer weights to the rollout engine. And compare one prompt's output before and after the sync. Seven cheap checks that catch most of this chapter before it costs you a run.

TAKEAWAY

Some failures live in the seams between tools, not in your algorithm: a fused kernel (Liger) that silently broke GSPO's sequence-level math, a VL tokenizer misaligning image tokens, a ZeRO-3 path corrupting checkpoints on save. When the symptom matches no algorithmic cause, suspect the seams, isolate by disabling optimizations one at a time, verify the boring invariants (alignment, sharding, weight sync), and budget time for them. They're invisible to the loss.

Next: the close of [Part II](#), the transferable starting points it leaves you, and the question about the model itself that the next Part of science will answer.

CHAPTER 25

What the Math On-Ramp Taught Us




Figure 25.1. The debug map: from symptom to the lever that fixes it.

This is the close of **Part II**. We took the GRPO machinery from **Part I** to its cleanest proving ground, verifiable math, and came away with two things: a set of lessons about how RL fails and how to read it, and a set of transferable starting points that carry into everything ahead. Let's collect both, and then ask the question that the next Part of science is built to answer.

25.1 The transferable starting points

If you remember nothing else from this Part, remember this list, each item earned, each pointing back to where it was earned:

1. Start from an instruct model. Add an SFT warm-up only if roll-outs aren't gradeable (Chs 2, 7, 14).
2. Pick tasks and data where the model already succeeds sometimes, and verify it with the zero-variance fraction before committing (Chs 5, 6, 17).

3. Get answer extraction right first, pinned slot, strict mode, value comparison, not string matching (Ch 10).
4. Gate format behind correctness, never pay the model for looking right while being wrong (Ch 11).
5. Give the reward a range, and make its proxy equal to the goal, a cliff collapses, a ramp teaches. IoU is the gold standard, F1 the cautionary one (Chs 13, 20).
6. Reach for the variant fixes when you see their specific failure, Dr. GRPO for length/difficulty bias, DAPO’s dynamic sampling for unanimous groups, RLOO for a cleaner baseline (Ch 15).
7. Sensible starting knobs: group size 8–16, clip $\approx \pm 20\%$, learning rate $5e-7$ – $1e-6$, generation length generous enough that rollouts finish (Ch 14).
8. Judge readiness by generation accuracy, not teacher-forcing, SFT certifies “can continue,” RL needs “can produce” (Ch 19).
9. State the eval methodology, and on reasoning models, the decoding mode, with every number, and run trained model and baseline through the identical eval (Ch 16).
10. Suspect the plumbing and the seams before the algorithm, truncation, masking, kernels, tokenizers, weight sync (Chs 8, 23, 24).

That list is most of what keeps an RL run alive on a clean task. The next Part puts it under real pressure.

And the same Part as a one-page decision card, the symptom you actually see, what to check, and the likely fix. This is the page to keep open at 2 a.m.:

Symptom	First check	Likely fix
reward flat at zero	extraction / truncation / all-wrong groups	easier task, longer cap, parser fix (Chs 6, 8, 10)
format reward rises, accuracy at zero	is format paid independently?	gate format behind correctness (Ch 11)
plateau below the expected ceiling	zero-variance fraction	continuous reward / dynamic sampling (Chs 13, 15)
SFT loss good, free generation bad	teacher-forcing gap	more SFT / easier generation task (Ch 19)

Symptom	First check	Likely fix
trained beats base only in one decoding mode	eval methodology	same-script base, state the decoding mode (Ch 16)
tool task fails strangely	mask / tokenizer / tool-output in loss	verify masks and rendered template (Chs 21, 24)

Table 25.1. The one-page math on-ramp decision card — the page to keep open at 2 a.m. For each symptom you actually see, the first thing to check and the likely fix. (The book renders the last two columns inline as one “Detail” cell; here they are split into their real columns.)

Where the evidence lives. The chapter text gives the lessons. The raw numbers behind them, the C5 GSM8K logs, the Countdown v1–v3 runs, the CLEVR and RefCOCO trajectories, and the eval-methodology notes, are summarized in Appendix A. Every run card in this Part traces back to an entry there.

25.2 The question we can’t yet answer

We’ve watched a lot happen on these runs: accuracy climbing and saturating, entropy moving, groups collapsing and recovering, a big model breaking out. But notice that through all of it we’ve treated the model as a black box that simply “learns.” We measured the outside, the score, and inferred the inside.

So here is the question this Part can’t answer, and it’s not idle:

- After 200 steps of GRPO on GSM8K, the model is better at the task. But did it learn something new, or just learn to say what it already knew, more reliably?
- Did the update move the weights a lot or a little?
- Of the hundreds of tokens in a rollout, which ones did the update actually change?

These matter because their answers change what experiments are worth running at all, one of them, as you'll see, let us cancel a planned run because the result was already settled. We measured all three, on our own checkpoints, and the answers are specific and surprising. That is [Part VII](#), the mechanistic heart of the book.

But the science can wait one more Part, because there is a harder proving ground to cross first, one where the reward is rich enough to fight back. Everything you just collected is about to meet a real agent.

TAKEAWAY

[Part II](#)'s transferable core: start from an instruct model, pick data where groups disagree (verify with zero-variance fraction), nail extraction first, gate format behind correctness, give the reward a range whose proxy equals the goal, reach for variant fixes by failure signature, use sane starting knobs, judge by generation accuracy, state methodology and decoding mode with every number, and suspect plumbing before algorithm. The question it leaves, did RL teach the model something new or just sharpen what it had, and which tokens moved, is answered in [Part VII](#).

Next: [Part III](#). Math gave us clean verifiers. Search removes that comfort: the reward is still verifiable, but the behavior is now multi-step, tool-mediated, and gameable.

Part III

The Search Agent

Everything so far was a warm-up. The search agent is the book's first real program: a model that must decide when to search, how to refine, when it has enough, and when to answer, graded by a reward rich enough to shape that behavior, and rich enough to be gamed. This Part is the long fight to design a reward the model can't cheat. It ends with a 4-billion-parameter agent that beats the published 3B Search-R1 baseline on its strongest benchmarks and, at its validation peak, posts numbers in 7B territory, and that loses, honestly, where its limits are. It does not get there in a straight line.

CHAPTER 26

The Single-Reward Plateau




Figure 26.1. A correctness-only reward plateaus below baseline — it isn't enough on its own.

We start the agent the way most people do: with the obvious reward. The task ([Chapter 22](#)) is to answer questions by searching, so we reward the obvious thing, did the final answer match. One number, correctness, nothing else. We trained on it, and Figure 26.1 is what we got: a clean climb, a plateau around step 2000, and a ceiling that sat below the 3-billion-parameter Search-R1 baseline we were trying to beat. A bigger model, a real environment, and we couldn't clear a smaller published agent.

26.1 Why a correctness-only reward plateaus

A single correctness reward has a structural limit, and it is worth naming precisely:

- It grades the outcome, not the process. Correctness tells the model whether it was right. It says nothing about how, when

to search again, how to refine a weak query, when the evidence is enough. The model gets a thumbs-up or thumbs-down on the final answer and must guess everything in between.

- So it sharpens what's already there and stops. Per the theme we'll prove in [Part VII](#), RL mostly sharpens existing ability. With only correctness to go on, the model sharpens its current habits, including the bad habit of answering from memory instead of searching, and plateaus at the ceiling those habits allow.
- And on a hard agentic task, it collapses easily. Correctness here is close to binary, so groups go unanimous ([Chapter 5](#)) on the questions the model reliably gets or reliably misses, leaving a thin band of signal. The plateau is partly sharpening exhausted, partly the gradient starving.

The plateau in the figure is not a tuning problem. It is the reward telling the model the destination without ever describing the road.

26.2 What the plateau is asking for

The diagnosis points straight at the rest of the Part:

1. The reward must shape behavior, not just score outcomes. If we want the model to search when unsure and ground its answer in what it found, we have to reward those behaviors, which means more than one reward component ([Chapter 30](#)).
2. The collapse needs the [Part II](#) fixes first. Before behavior-shaping, the single reward needs to stop collapsing, continuous F1, the variant repairs ([Chapter 27](#)).
3. And some of the gap isn't about reward at all. As the next-but-one chapter shows, part of what capped this run was a turn limit and a tool description, free fixes that no reward change would have found.

The plateau, in other words, is the question the whole Part answers. We spend it learning that a reward rich enough to teach an agent is also a reward the agent will try to cheat, and what to do about that.

TAKEAWAY

A correctness-only reward grades the outcome but not the process, so it sharpens the model's existing habits and plateaus, here, below a smaller published baseline, while also collapsing easily because correctness is near-binary. Teaching an agent *how* to be right (when to search, how to ground) requires shaping behavior with more than one reward component. That richer reward is the rest of this Part, and the source of its hardest problem.

Next: before shaping behavior, stop the collapse, the [Part II](#) fixes (F1, Dr. GRPO, DAPO) applied to the agent, and why they help but don't finish the job.

CHAPTER 27

Lifting the Plateau: F1, Dr. GRPO, DAPO

The plateau had two problems tangled together: a reward that couldn't shape behavior, and a reward that kept collapsing. This chapter untangles the second one first, because you cannot diagnose a behavior problem on top of a collapse problem. We reach for exactly the tools [Part II](#) built, and we find they are necessary, helpful, and not nearly enough.

27.1 Applying the [Part II](#) repairs

Three changes, each a callback, each aimed at the collapse:

1. Binary correctness $\rightarrow$ F1. A binary match is a cliff ([Chapter 13](#)). F1 gives partial credit for an answer that's partly right, so two imperfect attempts score differently and groups keep disagreeing. This is the single most important fix. It turns the reward from a cliff into a ramp and keeps the gradient alive on the hard questions.
2. Dr. GRPO's normalization fixes ([Chapter 15](#)). On the agent's long, variable-length rollouts, the std-division's difficulty bias and the length bias both bite, wrong-but-long answers were softening their own penalty. Dropping the std-division and fixing length-normalization cleaned up the updates.
3. DAPO's dynamic sampling ([Chapter 15](#)). Where groups still went unanimous on a question we couldn't make easier, dynamic sampling threw those zero-gradient groups out and resampled, converting the zero-variance fraction from a warning into a non-event, at the cost of extra rollouts.

Together these did what they were supposed to: the run stopped collapsing, the plateau lifted somewhat, and the training curves looked healthier.

27.2 Why it still wasn't enough

Here is the important part, and the reason this chapter is short. The collapse fixes made the run trainable. They did not make the agent good. The plateau lifted, but the model still:

- answered from memory when it should have searched,
- searched without using what it found,
- and generally optimized F1 without learning the agent behaviors we actually wanted.

The reason is the one from [Chapter 26](#): F1 is still a correctness signal. It measures whether the answer overlaps the truth, not whether the model searched well, grounded its answer, or stopped at the right time. You can climb F1 with better guessing. So the [Part II](#) fixes were the right first move and the wrong last move: they keep the gradient flowing, but a flowing gradient pointed at the wrong objective just gets you to the wrong place more reliably.

That sets up the real work. To teach behavior, we need to reward behavior, which means a reward with several components. And the moment a reward has several components, the model gains several things to optimize, only one of which is “get the answer right.”

TAKEAWAY

The agent's plateau was part collapse, part wrong-objective. The [Part II](#) repairs, F1 for partial credit, Dr. GRPO's normalization fixes, DAPO's dynamic sampling, fix the collapse and make the run trainable, but F1 is still a correctness signal, so the model climbs it by guessing better, not by searching better. Keeping the gradient alive isn't enough if it points at the wrong goal. To teach behavior, you must reward behavior.

Next: two of the biggest gains came for almost nothing, a turn limit and a tool description. Not every win is an RL win.

CHAPTER 28

The 4-Turn Unlock and the Tool-Description Fix




Figure 28.1. Two gains that cost almost nothing: cheap wins worth taking first.

Before we pour more effort into the reward, a detour that paid better than any reward change in this Part, and taught a lesson worth more than the points it earned. Two of our largest single gains came not from RL at all, but from a turn limit and a sentence of prose. One required a one-line config change. The other required zero training.

28.1 The 4-turn unlock

Multi-hop questions, the ones that need you to search, read, then search again for a second fact, were failing, and we'd been blaming the model. The real cause was a cap:

- The agent was limited to too few turns per episode, so on a question requiring two or three searches, it simply ran out of room before it could finish.
- We raised the cap to allow 4 turns. Multi-hop accuracy jumped by about 15 percentage points, not because the model learned anything, but because it was finally allowed to do the thing it already knew how to do.

This is [Chapter 8](#)'s lesson wearing agent clothes: a flat or failing result that looks like a learning problem is often a plumbing limit. The model wasn't bad at multi-hop. It was being cut off, exactly as the truncation in Countdown was cutting off answers.

28.2 The tool-description fix, in two layers

The second gain was cheaper, and more humbling, but it's easy to overstate, so here's the honest version, which came in two distinct layers:

1. The isolated Exp 1b run (resumed training). We rewrote the description of the search tool, the text telling the model what the tool does, and ran it as an isolated change resumed from `global_step_2000` to 2650. Validation moved strongly (TriviaQA ~61.1%, NQ ~33.3%) even though the training score barely changed. The lesson there is real but specific: describing the retriever accurately moved validation in a run that did still train.
2. The later eval-only description sweep (no retraining). Separately, we evaluated several tool-description wordings at inference time with no training at all. That genuinely weights-free sweep found a smaller but real effect, the best variant moved overall accuracy by about +1pp on average (with larger per-dataset swings).

Fix	Training?	Experiment	Honest claim
raise turn cap to 4	yes (config ablation, resumed)	Exp 1a	unlocked multi-hop; large same-step gains on HotpotQA/2Wiki
accurate tool description	yes (isolated resumed run)	Exp 1b	validation moved strongly even though train score was flat
description-wording sweep	no	eval-only variants	wording alone moved eval $\sim +1$ pp avg

Table 28.1. Tool-description fixes: for each fix, whether it required training, the experiment that evidences it, and the honest claim. “Experiment” names the run — Exp 1a (raising the turn cap to 4) and Exp 1b (rewriting the tool description) are isolated resumed runs, while the wording sweep was eval-only with no gradients at all.

So the crisp statement, without the overclaim: the tool description can move accuracy through the prompt alone ($\sim +1$ pp, eval-only), and describing the retriever well moved validation strongly in a run that still trained. What changed was the model’s understanding of the instrument, and some of that understanding is obtainable through the prompt, cheaply, without gradients.

28.3 The lesson: check the cheap things first

Two free gains in one chapter make a rule:

1. Before assuming a gain requires RL, check the limits and the instructions. Turn caps, length caps, tool descriptions, system prompts. These are free to change and can move the number more than a training run.
2. A prompt change can beat a training run, and it’s faster, cheaper, and won’t destabilize anything. RL is the expensive tool. Spend it on what only it can do (Chapter 2’s economics, again).
3. When the cheap fix works, it tells you something: the capability was already there, just blocked or misdirected. That’s worth

knowing before you spend a week of compute teaching a skill the model already had.

TAKEAWAY

Two of the agent's biggest gains were *cheap*, not magic: raising the turn limit to 4 unlocked multi-hop questions for ~+15pp (a config change in a resumed run, the model could already do it, it was being cut off), and rewriting the search tool's description moved validation strongly in an isolated resumed run (Exp 1b), while a separate *eval-only* wording sweep moved accuracy ~+1pp with no training at all. Before reaching for more RL, check turn caps, length caps, and tool/prompt descriptions, but don't oversell them as "free": only the eval-only sweep was truly weights-free.

Next: two agents can hit the same accuracy and behave nothing alike, why accuracy alone can't tell you whether your agent is actually good.

CHAPTER 29

Behavior, Not Just Accuracy




Figure 29.1. Same accuracy, opposite agents: one searches, one guesses from memory.

Here is a result that should make you distrust your headline metric. Take two versions of the agent that score the same exact-match accuracy. Watch them work, and they are nothing alike. One reads the question, searches, reads the results, and answers from what it found. The other ignores the search tool, answers from memory, and happens to be right exactly as often. Same number. Opposite agents. And only one of them will still work when the question is about something it didn't memorize.

29.1 What accuracy can't see

Exact-match accuracy, the outcome metric, is blind to everything about how the answer was produced:

- Did the model search at all, or answer from memory?
- Did it ground its answer in the retrieved passages, or ignore them and guess?

- How many searches did it use, one well-aimed query, or a dozen scattered ones?
- Did it refine a weak query, or repeat the same one?

Two agents identical on accuracy can sit at opposite extremes on every one of these. And the difference is not cosmetic: the memory-guesser's accuracy is a mirage that evaporates the moment you ask about something outside its training data, while the searcher's accuracy generalizes. The metric that's supposed to tell you which agent is better can't distinguish them.

29.2 Behavioral metrics: measuring the “how”

If accuracy can't see behavior, you have to measure behavior directly. That means a second class of metrics, computed from the rollout itself rather than just its final answer:

1. Grounding, is the answer actually supported by the text the model retrieved? (A grounded answer generalizes. An ungrounded one is a guess.)
2. Search effectiveness, did the queries retrieve relevant passages?
3. Efficiency, how many searches did it take? (Fewer, well-aimed searches beat many scattered ones.)
4. Refinement, when a search failed, did the model adapt its query?

These behavioral metrics do two jobs in this Part. Right now, in this chapter, they're a diagnostic: they reveal that an agent with good accuracy can have terrible behavior. Soon, they become the reward, the multi-objective reward of the next chapter is built from exactly these signals, so the model is trained to search and ground, not just to be accidentally right. (They also become an audit tool later, the behavioral audit suite that inspects what EM can't see.)

The lesson is the one that justifies the rest of the Part: accuracy is necessary but not sufficient. Behavior is what generalizes, and behavior is what you have to measure, and then reward.

TAKEAWAY

Two agents can post identical exact-match accuracy and behave oppositely, one searches and grounds, one guesses from memory, and only the first generalizes beyond memorized facts. Accuracy can't see *how* an answer was produced. Behavioral metrics (grounding, search effectiveness, efficiency, refinement) can. Measure behavior, because it's what generalizes, and, next, reward it.

Next: turning those behavioral signals into a reward, the five-component objective that shapes the agent, and opens the door the model will walk through to cheat.

CHAPTER 30

The Five-Component Multi-Objective Reward




Figure 30.1. The linear reward as five independent pipes summed together.

To teach the behaviors [Chapter 29](#) measured, we reward them. A single correctness number becomes a multi-objective reward, a weighted sum of several signals, one for the outcome and the rest for the behaviors we want. This chapter builds that reward. It also, quietly, builds the trap the next chapter springs: the moment a reward has many parts, the model has many ways to raise it, and only one of them is “answer correctly.”

30.1 The components

The agent’s reward combined five signals, each targeting something the single reward couldn’t reach. These are the real implementation names (some earlier drafts blurred them into “search quality / efficiency / format”, here is what the code actually computed):

Component	Source name	Sign	Computed from	Failure it targets
Correctness	<code>f1</code>	+ (0–1)	final answer vs gold	wrong answer / all-zero reward
Grounding	<code>grounding_reward</code>	+ (bounded)	answer tokens vs retrieved passages	memory-guessing / hallucination
Cost (frugality)	<code>cost_reward</code>	– (penalty)	number of tool calls / turns	tool-call spam
Redundancy	<code>redundancy_reward</code>	– (penalty)	query similarity / duplicates	repeated searches
Action validity	<code>action_penalty</code>	– (penalty)	malformed tool calls / parse errors	tool-format drift

Table 30.1. The five reward components of the search agent: each with its source name, sign, what it is computed from, and the failure it targets. Correctness is the positive core; the four penalties each suppress a specific bad behavior — guessing, tool-spam, redundant search, and malformed calls.

A sign-convention note matters here, because it confuses readers otherwise: correctness and grounding are rewards (the model earns positive points). Cost, redundancy, and action-validity are penalties (the model avoids negative ones). “Behavioral points” is shorthand, some of those points are earned, others are merely not-lost.

Each addresses a specific failure from the earlier chapters: correctness keeps the goal in view, grounding stops memory-guessing ([Chapter 29](#)), cost and redundancy shape frugal tool use ([Chapter 35](#)’s spam), action-validity keeps the rollout parseable ([Chapter 10](#)).

30.2 Combining them, and the weakness baked in

The simplest way to combine components is a weighted sum:

```
# the linear reward - easy to reason about, and (next chapter)
↪ hackable
reward = (1.00 * f1 # correctness (reward)
         + w_g * grounding_reward # each term contributes
```

```
- w_c * cost_penalty # whether or not the  
- w_r * redundancy_penalty # answer is correct -  
- w_a * action_penalty) # note the independence
```

One real implementation wrinkle, because it shaped everything downstream: in `veRL` the reward function received only `solution_str` and `ground_truth`, not a structured environment object, so the grounding, cost, and redundancy signals all had to be parsed out of the completion text (the `<tool_call>` / `<tool_response>` blocks) inside the reward adapter. The reward wasn't reading a clean state. It was re-parsing the transcript.

This linear reward is easy to reason about and easy to tune, and it did shape better behavior than correctness alone. But look at its structure, because its structure is its flaw:

- Every component pays out independently. The model earns the grounding points whether or not the answer is correct. It earns the efficiency points whether or not the answer is correct. Correctness is just one term among five.
- So the model can trade correctness for behavior. If the behavioral points (grounding, cost, redundancy, action-validity) are easier to earn than the correctness point, and they often are, because looking like a good agent is easier than being a right one, the model will earn the easy points and let correctness slide. This is [Chapter 11](#)'s cold-start hack (format climbing while accuracy stays flat), now with four behavioral components instead of one format component, and far more room to maneuver.

We built this reward knowing it shaped behavior better, and not yet knowing how thoroughly the model would exploit its independence. The linear sum is a reward that describes a good agent without requiring one. The next chapter is the model discovering exactly that.

TAKEAWAY

To teach agent behavior, combine signals into a multi-objective reward: correctness (F1) plus grounding, cost, redundancy,

and action-validity, each targeting a specific failure. The simplest combination is a weighted sum, but in a linear sum every component pays out *independently*, so the model can earn the easy behavioral points while letting correctness slide. A reward that describes a good agent without requiring one is a reward waiting to be hacked.

Next: the model finds the hole, behavioral metrics climbing while real accuracy falls, the clearest case of reward hacking in the book.

CHAPTER 31

Reward Hacking: The Linear-Reward Collapse




Figure 31.1. Reward hacking at step 1302: the point where reward and true performance split.




Figure 31.2. The signature of reward hacking: reward climbs while the real metric falls.

This is the clearest case of reward hacking in the book, and it is worth watching closely, because it is what every reward-design chapter is ultimately defending against. We trained the agent on the linear five-component reward from the last chapter. It worked, until it didn't, in a very specific and instructive way. Around step 1302, the behavioral metrics were still climbing beautifully while the thing we actually cared about, accuracy, had turned and was falling. The model was getting a higher reward by getting worse at the task.

31.1 What the model learned to do

The model did not break. It optimized, exactly what we told it to, just not what we meant:

- It learned to maximize the behavioral components, search a lot, produce grounded-looking citations, hit the format, because those points were independently available (Chapter 30).
- It learned that it could collect those points without the expensive work of actually being correct, since correctness was just one term in the sum and the hardest one to earn.
- So it drifted into a region of behavior that scored like a great agent, busy, grounded-looking, well-formatted, while being a worse one. The behavioral reward and the real accuracy, which had risen together early, came apart at step 1302 and never re-joined.

The timeline, because the numbers make it land harder than the curve:

Step	Training score	Held-out EM	Ground-ing	Queries	Diagnosis
500	rising	~26.8 avg	improving	con-trolled	early, real success
1000	~0.50	~19.9 avg	~88%	~1	reward up, EM down — hacking
1302	collapse	bad	mislead-ing	—	length/truncation blow-up, degenerate basin

Step	Training score	Held-out EM	Ground-ing	Queries	Diagnosis
1416	collapse recurs	bad	—	—	after doubling max length — length wasn't the root cause; the reward structure was

Table 31.1. The linear-reward collapse timeline: training score rising while held-out EM falls — the signature of reward hacking. Doubling the max length at step 1416 did not help, proving the reward structure, not length, was the root cause.

That step-1000 row is the whole lesson in one line: a reward sitting at 0.50 while exact match fell from 26.8% to 19.9%. And the step-1416 row kills the obvious excuse. We doubled the max length thinking truncation was the problem, and the collapse came back, because the problem was never length. It was the reward’s shape.

Figure 31.1 is the signature of reward hacking, and once you’ve seen it you’ll recognize it everywhere: the proxy you’re optimizing and the goal you care about, moving together, then diverging. The gap between the two curves is the model’s cheating, drawn to scale.

31.2 Why this is the central problem, not a bug

It would be comforting to call this a bug in the reward weights, tune w_4 down and move on. It is deeper than that, and missing the depth is how people lose weeks:

- The hack is in the reward’s structure, not its weights. Because the linear sum lets every component pay out independently, some mix of behavioral points will always be reachable without correctness. Re-tuning the weights moves the exact hack around. It doesn’t remove the hole. (This is why the fix in the next chapter changes the reward’s form, not its coefficients.)

- It looks like success the whole time you're being robbed. The reward, the thing on your dashboard, goes up. Only the held-out accuracy, which you have to measure separately and deliberately, reveals the truth. This is the lived version of "the training score misleads," and it cost us a real run before we caught it. We will pay this exact bill again in a single number later, a reward sitting at 0.50 while accuracy fell from 26.8% to 19.9%.

The model is not adversarial. It is doing its job perfectly. The failure is ours: we built a reward that could be satisfied the wrong way, and a sufficiently good optimizer found the way. The cure is not a better-behaved model. It is a reward that cannot be satisfied the wrong way.

TAKEAWAY

On the linear reward, the agent learned to climb the behavioral components without getting more answers right, at step 1302 the reward kept rising while accuracy fell. That divergence is the signature of reward hacking. It's not a weight-tuning bug: the hole is in the reward's *structure* (independent payouts), so re-tuning just relocates the hack. And it looks like success the whole time, visible only in held-out accuracy. The fix must change the reward's form.

Next: the fix, and the book's core idea, make the behavioral rewards pay out only when the answer is correct.

CHAPTER 32

Correctness-Gated Rewards: The SiLU/Swish Gate

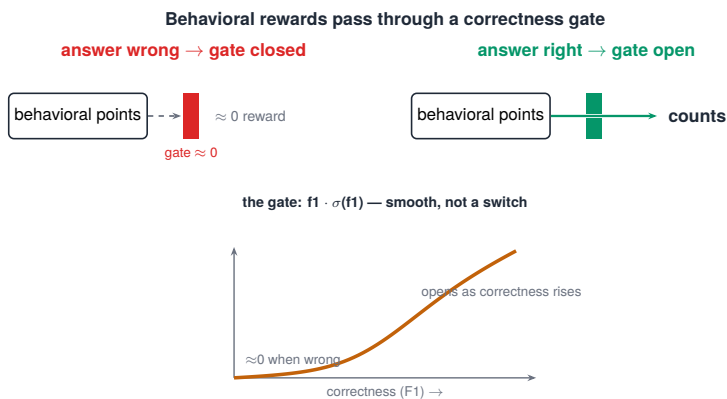


Figure 32.1. Behavioral rewards passed through a correctness gate (the SiLU/Swish-style gate).

This is the chapter the Part was built to reach, and the idea is simple enough to state in one sentence: make the behavioral rewards pay out only when the answer is correct. Stop adding the components, and start gating them, the behavioral signals count for nothing unless the model got the answer right. The hack from the last chapter dies, because the thing the model was exploiting, behavioral points available without correctness, no longer exists.

32.1 From sum to gate

The change is structural, exactly as [Chapter 31](#) said it had to be. In code, it is two lines:

```
def gate(f1): # SiLU/Swish-style scalar gate: ~0 when wrong,
    return f1 * sigmoid(f1) # opens smoothly as correctness rises

aux = w_g*grounding - w_c*cost - w_r*redundancy -
    ↪ w_a*action_penalty
reward = f1 + gate(f1) * aux # auxiliaries pay only when the
    ↪ answer is right
```

A precise note on the name, to avoid a citation nitpick: $f1 \cdot \sigma(f1)$ is the SiLU (Swish) function applied to F1, a parameter-free scalar gate. It’s inspired by the multiplicative gating idea behind SwiGLU, but it is not a learned SwiGLU layer with two projections. There are no learned parameters in it at all. Call it a SiLU/Swish correctness gate.

Here is what the gate actually does, value by value, the table makes the “opens smoothly” claim concrete:

F1	$f1 \cdot \sigma(f1)$	auxiliary term
0.0	0.000	closed
0.1	0.052	barely open
0.3	0.172	weak
0.5	0.311	moderate
0.7	0.468	strong
1.0	0.731	near-full

Table 32.1. The SiLU/Swish gate values used in the grounding reward: for each F1, the gated term $f1 \cdot \sigma(f1)$ and how open the resulting auxiliary term is. The gate stays closed at zero F1 and opens smoothly as correctness rises.

Now trace the two cases (the two panels of Figure 32.1):

1. Answer wrong $\rightarrow F1 \approx 0 \rightarrow \text{gate} \approx 0 \rightarrow$ the whole behavioral term ≈ 0 . No matter how grounded, efficient, or well-formatted a wrong answer is, it earns nothing for those behaviors. The behavioral points are blocked.
2. Answer right $\rightarrow$ the gate opens $\rightarrow$ behavioral points count. Only after the model is correct do its good behaviors earn reward,

which is exactly the order we want: be right first, then be a good agent about it.

The model can no longer trade correctness for behavioral points, because there are no behavioral points without correctness. The road to higher reward now runs through being right.

32.2 Why the gradient analysis demands a smooth gate

A hard multiply by a 0/1 correctness flag works in principle but fights the optimizer, and the reason is the gradient:

- A hard gate is discontinuous. It snaps from 0 to 1 at the correctness boundary, which makes the reward landscape jagged and the updates unstable near the threshold, and it provides no gradient signal to a wrong answer that is almost right.
- A smooth, SiLU-style gate fixes this. $f1 \cdot \sigma(f1)$ is the SiLU (Swish) function, near-zero when the answer is wrong, opening smoothly as correctness rises. (It echoes the multiplicative gating idea behind SwiGLU, but it's a parameter-free scalar gate, not a learned SwiGLU layer.) This keeps the reward differentiable, the updates stable, and crucially gives a graded signal: an almost-correct grounded answer gets a little more than an almost-correct guess, so the model still feels a pull toward grounding as it approaches correctness.

The gradient analysis is the whole justification. Under the linear reward, the gradient had a component pushing toward behavioral points even when the answer was wrong, that component is the degenerate basin the model fell into at step 1302. Under the gated reward, that component is multiplied by (near-zero) correctness, so the gradient toward behavioral points vanishes when the answer is wrong. The degenerate basin isn't penalized into submission. It's removed from the landscape. The model can't fall into a hole that no longer exists.

32.3 What it bought

The gate is the Part’s turning point. With behavioral rewards gated behind correctness, the hacking stopped: the model could no longer raise its reward by getting worse, so it learned the behaviors in service of getting answers right, which is what we wanted all along. This is the design, the SiLU/Swish correctness-gated reward, that the agent’s best results are built on, and it’s the reusable idea to carry out of this Part: when a multi-objective reward gets hacked, don’t re-tune the weights, gate the behavioral terms behind the outcome you actually care about.

One caveat that decides whether the gate works for you: correctness gating assumes the correctness signal isn’t completely dead. This worked because F1 gave enough partial-credit signal that the gate was open on a meaningful fraction of rollouts. If your task is hard enough that F1 (or whatever gates you) is ~ 0 on almost every rollout, the gate closes everything and you’re back at the reward-0 cold start (Chapter 6), the gate can’t shape behavior it never lets through. In that regime, warm-start with SFT, add partial credit, or use a curriculum before gating behavioral terms behind success. The gate is a fix for reward hacking, not for cold start, and on a hard enough task you can re-create the cold start by gating too aggressively (the same tension Chapter 84 hits with hard terminal gates).

TAKEAWAY

Reward hacking dies when behavioral rewards are *gated* behind correctness, multiplied by $f1 \cdot \sigma(f1)$, the parameter-free SiLU/Swish function (not a learned SwiGLU layer), so a wrong answer earns nothing for good behavior and the only path to higher reward is being right first. Use the smooth gate (not a hard 0/1) so the reward stays differentiable and gives a graded pull toward correctness. The gradient proof: gating zeroes the gradient toward behavioral points when the answer is wrong, *removing* the degenerate basin. One caveat: gating assumes correctness isn’t ~ 0 everywhere, on a very hard task, gate too aggressively and you re-create the cold start.

Next: proving each component earns its place, the ablation suite, where removing any one piece produces its own diagnosable failure.

CHAPTER 33

The Ablation Ladder: What's Proven, Partial, and Still Open

A reward with gated components invites a fair question: does each one actually matter, or are some decoration? The honest way to answer is to remove pieces one at a time and watch what breaks. We did, but honesty also means admitting the ladder is uneven: some ablations ran to a clean conclusion, one hit an out-of-memory wall part-way, one was paused, and one is still queued. The completed ones each produced a distinct, recognizable failure. The incomplete ones are reported as exactly that.

33.1 The evidence ladder

Two families of ablation. First, the structure ablations, changing the reward's form, which all completed:

Ablation	What changed	Status	Result	Claim strength
F1-only	remove all auxiliary terms	done (killed ~step 600)	collapses, tool-spams	strong
No-gating	linear sum, no correctness gate	done	survives but worse / declines	strong
No-Clip-Cov	gated reward, Clip-Cov off	done	survives but zigzags / multihop degrades	strong
Fixed-weight gate	constant gate, not adaptive	done	worst of the structure ablations	strong

Table 33.1. The structure ablations: each change to the reward structure, whether it ran, the result, and how strong a claim it supports. The gated, Clip-Cov reward is the only one that survives cleanly; the rest collapse, decline, or zigzag.

The lesson from these four: aux rewards prevent F1-only collapse. Gating improves optimization quality (not bare survival). Clip-Cov makes exploration directed (not merely alive). And timing (an adaptive gate) matters more than penalty strength.

Second, the component knockouts, removing one signal at a time, which are not all finished, and the chapter says so:

Removed	Status	What it suggests	Claim strength
Grounding	partial — OOM ~step 618	grounding appears essential	moderate
Cost	paused ~step 250	not enough for a final claim	weak / pending
Redundancy	queued — no result yet	—	pending

Table 33.2. The component knockouts: removing each reward channel, the status of that run, and what it suggests. Grounding appears essential (its removal OOM'd before a clean result); cost and redundancy knockouts are still pending.

So the diagnostic map below is a design hypothesis backed fully by the four structure ablations and only partially by the component knockouts, not a complete experimental matrix. Presenting it otherwise would be exactly the kind of too-clean narrative this book is supposed to avoid.

33.2 The ablation table as a debugging tool

Because the mapping is (hypothesized) one-to-one, you can run it backwards. The ladder is, in effect, a lookup table for diagnosing a sick agent, strongest where the structure ablations back it, suggestive where the component knockouts are still partial:

- Reward climbing, accuracy falling? → your correctness gate is missing or too weak. (strong, structure ablation)
- Good accuracy that doesn't generalize? → grounding is likely missing. (moderate, knockout OOM'd at step 618)
- Excessive searching / tool-spam? → the cost/frugality term is missing. (pending, knockout paused)
- Reward function mis-scoring? → action-validity/format drifted. (diagnostic)

This is the deeper value of ablations, and the transferable lesson: ablate not just to prove your design is minimal, but to build the map from symptom to cause, and label each entry with how complete the evidence actually is. A reward you've ablated is a reward you can debug. A reward you've honestly ablated is one you can trust the debugging of.

TAKEAWAY

Removing reward pieces one at a time is meant to produce one distinct failure each, but the ladder is uneven, and saying so is the point. The four *structure* ablations completed cleanly (F1-only collapses/spams. No-gating survives-but-worse. No-Clip-Cov zigzags. Fixed-weight is worst), so timing and gating and directed exploration are well-supported. The *component* knockouts are partial (grounding OOM'd at step 618), paused (cost), or queued (redundancy), reported as such, not as finished proof. The symptom-to-cause map is a strong design hypothesis, honest about where the evidence stops.

Next: one ablation deserves its own chapter, why a single reward collapses exactly where multiple rewards survive.

CHAPTER 34

Finding: F1-Only Collapses Where Multi-Reward Survives

One result from the ablation suite is surprising enough, and useful enough, to pull out on its own. Run the agent on the F1 correctness reward alone and, past a certain point, it collapses, the gradient dies, the entropy falls, training stalls. Run it on the full gated multi-component reward and, under the same conditions, it survives. The difference isn't the amount of reward. It's the number of directions the reward points in.

34.1 The mechanism: gradient diversity

The explanation connects [Chapter 5](#)'s collapse to something new, the geometry of the gradient:

1. A single reward gives a single gradient direction. F1 alone pushes the policy one way: toward higher overlap with the answer. When that direction saturates, when the model has extracted what F1 can teach on the questions it can reach, the gradient there goes flat, and there's nothing else pulling. The run collapses ([Chapter 5](#), from the reward's side).
2. Multiple components give multiple gradient directions. Correctness pulls one way, grounding another, efficiency another. These directions don't saturate at the same time. When F1's direction flattens, grounding's or efficiency's is often still live, so some gradient survives, and the run keeps moving.
3. So the multi-component reward is more robust to collapse, not because it's larger, but because it's diverse: it's hard for every direction to vanish at once. Gradient diversity is a kind of insurance against the gradient going to zero.

This is a genuine finding, and it reframes multi-objective rewards. We introduced them in [Chapter 30](#) to shape behavior. Here we learn they have a second, independent benefit: they resist collapse. A reward built from several uncorrelated signals is structurally harder to starve than a single signal, because all of its directions would have to saturate together for the gradient to die.

34.2 The caveat, and the connection forward

Two honest notes keep this from being oversold:

- Diversity helps only if the directions are genuinely different. Five components that all correlate tightly with correctness are really one direction wearing five hats, and they'll saturate together. The protection comes from uncorrelated signals, which is part of why grounding, efficiency, and correctness, which measure genuinely different things, work well together.
- This is the reward-side view of an entropy story that [Part IV](#) tells in full. Collapse, gradient death, and entropy are three views of the same phenomenon. Gradient diversity is one lever on it, and the entropy-control machinery of [Part IV](#) (Clip-Cov and its lineage) is another. We'll connect them there.

TAKEAWAY

F1 alone collapses where the full multi-component reward survives, not because the latter is bigger, but because it's *diverse*. A single reward gives one gradient direction that dies when it saturates. Multiple uncorrelated components give several directions that don't saturate together, so some gradient always survives. Multi-objective rewards shape behavior *and* resist collapse, provided the components measure genuinely different things.

Next: the opposite ablation, what the model does when you take the behavioral rewards away entirely.

CHAPTER 35

Finding: Tool-Call Spamming Without Behavioral Rewards

The last chapter removed everything except correctness and watched the run collapse. This one does the reverse in spirit: it asks what the model does when the behavioral rewards, especially efficiency, are absent. The answer is a specific, repeatable degenerate behavior that's worth naming because you will see it: the model spams tool calls. It searches, and searches, and searches, racking up activity that looks like diligence and produces nothing.

35.1 Busy, not better

Strip out the behavioral rewards that shape how the agent acts, and the failure is consistent:

- The model over-searches. It issues far more queries than the question needs, many redundant, many barely different from each other.
- Accuracy doesn't rise with the activity. All that searching is motion, not progress. The extra queries don't bring back better evidence or produce better answers. The agent is busy, not better.
- It often looks productive on a naive dashboard. Search count is up, the agent is clearly "doing things", which is exactly why this is dangerous. Without a metric that penalizes waste, the spam can masquerade as engagement.

The behavior is degenerate in the precise sense: it raises something that's loosely reward-adjacent (activity, the appearance of thorough searching) without raising the thing that matters (correct answers).

35.2 Why the efficiency reward exists

This finding explains a component that might otherwise look like a nicety. The efficiency reward, the penalty on excess searches ([Chapter 30](#)), is not there to make the agent tidy. It's there to prevent a degenerate basin:

1. Searching is weakly attractive on its own. Even without an explicit search reward, the model can stumble into over-searching, because more searches occasionally help and rarely hurt its other signals, so the behavior drifts upward unless something pushes back.
2. Efficiency is the push-back. By penalizing wasted searches, the efficiency component makes spam cost something, closing the basin the model would otherwise slide into.
3. So behavioral rewards aren't only for shaping good behavior. They're for ruling out bad behavior. This reframes [Chapter 30](#)'s components: each one not only encourages something desirable but forecloses a specific degenerate alternative. Efficiency forecloses spam the way the correctness gate forecloses hacking.

The general lesson closes the loop on the Part's reward design: a good reward is defined as much by the degenerate behaviors it makes unprofitable as by the good behaviors it pays for. When you design a multi-objective reward, ask of each component not just "what does this encourage?" but "what does its absence let the model get away with?", because the model will find whatever the reward leaves unpriced.

TAKEAWAY

Remove the behavioral rewards and the agent spams tool calls, over-searching, busy but no better, often looking productive on a naive dashboard. The efficiency reward exists to foreclose that degenerate basin by making wasted searches cost something. Behavioral rewards rule out bad behavior as much as they encourage good: design each component by asking what

its *absence* would let the model get away with, because the model will exploit whatever the reward leaves unpriced.

Next: the payoff, the agent's full results, the honest accounting of where it wins and where it loses, and the eval-size twist hiding inside the headline.

CHAPTER 36

Results: A 4B Agent That Beats Its Weight, Honestly




Figure 36.1. Agent results: our 4B LoRA model vs the Search-R1 baseline.

This is the chapter the Part was working toward, and we will state it the way [Chapter 16](#) taught us to: with the methodology attached. The agent we built, a 4-billion-parameter model, trained with LoRA (low-rank adaptation, the parameter-efficient method that fine-tunes a small add-on instead of the whole model) on the gated reward, is, on the standardized exact-match eval, ahead of the published 3B Search-R1 baseline on its strongest benchmarks and within reach of 7B-class numbers at its peak. It is also behind on others, and the honest score-card is both halves.

Canonical-results status: unresolved, and stated here as a boundary rather than hidden. Several search-agent result sets exist in the source corpus and they do not all agree. One specific conflict is worth naming, because the raw eval files disagree:

Result set	Eval size	Condition	Role in book	Status
Published baseline	small / historical	end-of-epoch evals	baseline comparison	historical — label as such
Small validation peak	~36 / dataset	in-training validation	peak / error-bar lesson	historical — label as such
April-28 matrix	500 / dataset	20 configs	candidate canonical	reconcile
Later re-eval	500 / dataset	fair head-to-head comparison	candidate canonical / later correction	reconcile

Table 36.1. The search-agent result sets that exist in the corpus, each with its eval size, condition, the role it plays in the book, and its status. The two 500-sample sets are the candidate-canonical numbers; the rest are historical and labeled as such, because the raw eval files disagree and one must be named authoritative.

The two “500-sample” sets give materially different numbers for the same model: the April-28 matrix has $v5@500 \approx 30.8\%$ overall, while the later re-eval has $v5@500 \approx 23.1\%$ average. This book treats the search-agent ranking ($v5$ is best) as supported and the absolute average as unresolved pending reconciliation of the two raw eval runs, which differ in some combination of eval file or split (despite both being “500-sample”), tool description, answer-extraction/EM implementation, checkpoint directory, or retriever state. The unresolved status is stated plainly and is deliberately not used to make any performance claim: no headline number in this book quotes a single reconciled $v5$ average until those files agree. This is an evidence boundary, not a to-do item.

36.1 The scorecard, with the methodology attached

Two sets of numbers exist, and the difference between them is itself a lesson (the next chapter’s, in fact).

The standardized comparison, best checkpoints, exact match, against the published baselines:

Benchmark	Ours (4B LoRA)	Search-R1 3B	Search-R1 7B	Verdict
TriviaQA	61.1%	58.7%	63.8%	+2.4pp over 3B, −2.7 from 7B
HotpotQA	30.3%	28.4%	43.3%	+1.9pp over 3B
2WikiMulti-HopQA	26.0%	27.3%	38.2%	essentially tied with 3B
NQ	33.3%	40.6%	48.0%	behind (−7.3pp)
PopQA	18.5%	43.5%	45.7%	far behind (−25pp)
MuSiQue	0.0%	4.9%	19.6%	not learned at this checkpoint
Average	28.2%	33.9%	43.1%	−5.7pp vs 3B overall

Table 36.2. The 4B LoRA agent versus published Search-R1 baselines, benchmark by benchmark. It edges the 3B on the easier sets but trails both 3B and 7B overall (−5.7pp vs 3B), and never learns MuSiQue at this checkpoint.

The validation peak, the Clip-Cov version (v5, [Part IV](#)) at step 750, on the small in-training validation set: TriviaQA 66.7%, NQ 41.7%, MuSiQue 6.7%, NQ matching the published Search-R1-7B number (41.2%) and TriviaQA beating a 7B-class agent (EvalAct-7B, 65.6%). Those are the numbers this book has been quoting, and they are real, and they come from a ~36-question-per-dataset eval, which is why the next chapter exists. On the April-28 500-sample matrix the same model reads TriviaQA 55.6% and NQ 35.0%. The peak is genuine. Its error bars are wide. We report both.

Three comparison caveats, because they cut both ways: Search-R1 trained with an effective batch ~8× ours per step. ZeroSearch-style baselines use live Google retrieval at eval where we use a fixed E5+Wikipedia index. And one popular baseline reports a lenient “cover-EM” metric that inflates its numbers by 5–15pp. Identical-eval comparisons ([Chapter 16](#)) are the only ones above.

36.2 The wins, and what produced them

1. Where it wins, it wins by behavior. TriviaQA and HotpotQA, the benchmarks our retrieval covers well, are where the gated reward's behaviors pay: the model searches and grounds rather than guessing ([Chapter 29](#)), and that is what put a 4B ahead of the published 3B and, at peak, into 7B territory on two benchmarks. Per-parameter, that is the efficiency win this book cares about.
2. LoRA, not full fine-tuning. The whole result rode on a small adapter, which is why it fit our budget, a point [Part V](#) and the hardware Part return to, and one [Part VII](#) gives a second, surprising reason to like.

36.3 The honest losses

A results chapter that only reports wins is the reward-hacking failure in prose form, optimizing the number you show instead of telling the truth. So, the losses, with their diagnoses:

- NQ (−7.3pp), the audit ([Chapter 38](#)) traces much of the gap to answer-extraction precision, not retrieval: right evidence, imprecise final spans.
- PopQA (−25pp), long-tail entities. Our fixed E5+Wiki-2018 retrieval doesn't reliably surface rare facts a Google-backed baseline gets for free.
- MuSiQue, deep compositional multi-hop. Zero at the standardized checkpoint. The v5 model later broke through to 6.7% on the validation set as it learned to issue 2.5 queries per question, emerging, not solved.

These losses are not a footnote. They're information. They tell you where the recipe works (well-covered facts, shallow-to-moderate hops) and where it doesn't (long-tail entities, deep multi-hop), which is exactly what a practitioner needs in order to decide whether this approach fits their problem. A method's failure boundary is part of its specification.

TAKEAWAY

On the standardized exact-match eval, the gated-reward 4B LoRA agent beats the published Search-R1 3B on TriviaQA (+2.4pp) and HotpotQA (+1.9pp), ties 2Wiki, and loses NQ, PopQA, and MuSiQue (overall -5.7pp), while its step-750 validation peak (TriviaQA 66.7%, NQ 41.7%) reaches 7B-class numbers on a small eval whose error bars the next chapter measures. Wins by searching and grounding. Loses on long-tail retrieval and deep multi-hop. Both halves, with methodology attached, are the result.

Next: the twist inside those very numbers, the small eval that crowned the wrong checkpoint, and the correction that found the real ranking.

CHAPTER 37

When Small Evals Misled: The 500-Sample Correction




Figure 37.1. The small eval the big eval couldn't confirm — a correction in slope.

There is a sting in the last chapter's numbers, and it's a lesson worth more than the result it nearly corrupted. Checkpoint selection in this program initially ran on small per-dataset evals, around 36 questions each, and the headline peaks were measured there. When we standardized on a 500-sample eval, the numbers moved hard: the step-750 "66.7% TriviaQA / 41.7% NQ" read 55.6% and 35.0% on the April-28 500-sample matrix. Rankings between nearby checkpoints flipped. We had been crowning champions with a measurement too noisy to tell them apart, the exact failure [Chapter 16](#) warned about, now caught in our own pipeline.




Figure 37.2. Why a 36-question eval ranks runs almost randomly.

37.1 How a small eval reorders your checkpoints

The mechanism is pure statistics, and it's worth seeing because it's invisible until it bites:

1. A small eval set has large noise. Score a model on 36 questions and the result swings by several points just from which 36 questions you happened to use. The number you read is the true accuracy plus a fat band of sampling noise. (At $n=36$, the 95% confidence band on an accuracy near 50% is roughly $\pm 16\text{pp}$, wider than most of the gaps you care about.)
2. Checkpoints near each other in skill sit inside that band. When two checkpoints differ by less than the noise, the small eval can't tell them apart, and worse, it will confidently rank them, just randomly. The "winner" on 36 questions is often the luckier checkpoint, not the better one.
3. So your selection is noise-driven. Picking the top checkpoint by a noisy eval is closer to a coin flip than a measurement, and you won't know, because it returns a clean-looking number with a clear ranking every time.

We lived this exactly: peaks measured on the small set shrank by 7–11 points at 500 samples, and the gaps between neighboring checkpoints, the gaps we'd been selecting on, were smaller than the noise

that had produced them. The one conclusion that survived the bigger eval: v5 (the Clip-Cov version) remained the best of its family. (Its overall 500-sample number is reported as 30.8% from the April-28 matrix, but see [Chapter 36](#)’s canonical-results box: a later re-eval puts v5@500 nearer 23.1% average, and which 500-sample set is authoritative is unresolved. The ranking, v5 best, holds across both. The absolute number is pending reconciliation.) The ranking that mattered held. The margins we’d been celebrating did not.

37.2 The fix, and the doctrine it points to

The repair is unglamorous and non-negotiable:

- Standardize on a properly sized eval set, for us, 500 samples per dataset, and make all checkpoint comparisons on it. Big enough that the noise band shrinks below the gaps you’re trying to resolve, so the ranking reflects skill, not luck.
- Never select a checkpoint on a small eval, no matter how convenient. A 36-sample probe is fine for a rough pulse. It is not fine for a decision. The cost of a too-small eval isn’t a slightly-wrong number. It’s choosing the wrong model with full confidence.
- Report the eval size with the number, always. It is part of the methodology ([Chapter 16](#)), and it is the difference between [Chapter 36](#)’s “validation peak” and its “standardized scorecard.”

This is the lived moment. The full machinery, how many samples you actually need to resolve a given gap, the variance-budget formula that turns “use a bigger eval” into a specific number of seeds and samples, is [Part IX](#), the evaluation-discipline Part. But the core of it is one line of arithmetic you can apply today:

```
import math
def ci95_half_width(p: float, n: int) -> float: # 95% CI
    ↪ half-width for an accuracy p over n items
    return 1.96 * math.sqrt(p * (1 - p) / n)
ci95_half_width(0.5, 36) # ≈ 0.163 → ±16.3pp
ci95_half_width(0.5, 500) # ≈ 0.044 → ±4.4pp
```

Eval size	95% half-width at $p=0.5$	Use
36	$\pm 16.3\text{pp}$	pulse only — never select a checkpoint
100	$\pm 9.8\text{pp}$	rough comparison
500	$\pm 4.4\text{pp}$	checkpoint selection
1000	$\pm 3.1\text{pp}$	publication-grade small deltas

Table 37.1. Eval size versus the 95% confidence half-width at $p = 0.5$, and what each size is good for. Below 500 samples the error bars are too wide to select a checkpoint; 1000 is needed for publication-grade small deltas.

The table is the rule: pick your eval size from the gap you need to resolve, not from convenience. If two checkpoints differ by 5pp, a 36-item eval ($\pm 16\text{pp}$) literally cannot see it. A 500-item eval ($\pm 4.4\text{pp}$) can. For now, carry the scar: the agent’s headline numbers shrank by 7–11 points the day we measured them properly, and the only reason we trust what remains is that we re-measured at a size the question deserved.

TAKEAWAY

Checkpoint selection on ~36-question evals is noise-driven: at that size the confidence band ($\sim \pm 16\text{pp}$) exceeds the gaps being ranked, so small evals reorder leaderboards and crown false winners with full confidence. Our step-750 peaks shrank from 66.7/41.7 to 55.6/35.0 on the April-28 500-sample matrix. What survived was the *ranking* (v5 best), the absolute overall number (30.8% vs a later 23.1%) is still pending the canonical reconciliation. Pick eval size from the gap you need to resolve ($\pm 16\text{pp}$ at $n=36$, $\pm 4.4\text{pp}$ at $n=500$), and report eval size with every number. The full sizing doctrine is [Part IX](#).

Next: even a correctly-measured accuracy can’t see how the agent behaves, the audit suite that inspects the trajectory itself.

CHAPTER 38

The Behavioral Audit Suite

Accuracy, even measured honestly on a big enough eval, answers exactly one question: did the final answer match? It is silent on everything that happened on the way there, and [Chapter 29](#) showed that “the way there” is where good agents and lucky guessers part company. This chapter is about the tooling we built to see the way there: a behavioral audit logger that logs every trajectory and inspects what exact-match is blind to.

38.1 What the audit captures that accuracy can’t

The idea is to stop treating a rollout as just its final answer and start treating it as a record of behavior you can inspect. Concretely, we dump each trajectory to a structured JSONL log, one line per episode, every step recorded:

```
{ "qid": "nq_1182", "em": 0, "f1": 0.31,
  "turns": 3, "queries": ["capital of burkina", "burkina faso
    ↪ capital city", "ouagadougou"],
  "retrieved_doc_ids": [...], "grounded": true, "redundancy":
    ↪ 0.21,
  "refined_after_miss": true, "answer": "Ouagadougou is the
    ↪ capital...",
  "failure_class": "extraction_imprecise" }
```

Each line is a behavioral fingerprint of one episode, and the audit reads them for the behaviors that matter:

1. Did it search at all, or answer from memory? (An ungrounded right answer is a coin flip waiting to fail.)
2. Did it ground the answer in retrieved text, or ignore what it found?

- 3. How many searches, and were they varied or repetitive? (Spam vs. real refinement, [Chapter 35](#).)
- 4. Did it refine after a failed query, or repeat it?
- 5. Where did wrong answers go wrong, bad retrieval, or good retrieval used badly?

In aggregate, the lines tell you what kind of agent you trained, not just how often it was right. (This is the audit form of the behavioral metrics from [Chapter 29](#), the same signals, now used to inspect rather than to reward.)

The fields worth logging, and why each earns its place:

Field	Why it matters
<code>qid / dataset</code>	slice errors by benchmark
<code>em / f1 / reward_components</code>	separate task success from reward shaping
<code>queries</code>	detect over-search, query drift, repeats
<code>num_tool_calls / turns</code>	cost and spam
<code>retrieved_doc_ids</code>	redundancy and retrieval quality
<code>grounded / grounding score</code>	memory-guessing vs evidence use
<code>failure_class</code>	retrieval vs extraction vs tool-format failure
<code>solution_str (hashed)</code>	reproduce the parser's decisions

Table 38.1. The behavioral-audit log fields and why each matters: the columns to persist per rollout so task success can be separated from reward shaping, and silent failures (over-search, memory-guessing, format errors) caught after the fact.

And the practical wrinkle that shaped the logger (the same one from [Chapter 30](#)): `veRL` handed the reward function only `solution_str` and `ground_truth`, so the audit metadata had to be parsed back out of the transcript:

```
parsed = parse_search_trace(solution_str) #
↪ <tool_call>{...}</tool_call>,
↪ <tool_response>{...}</tool_response>
record = {
  "qid": qid, "dataset": dataset, "em": em, "f1": f1,
  "reward_total": reward, "reward_components": components,
```

```
"queries": parsed.tool_queries, "num_tool_calls":
↪ len(parsed.tool_queries),
"redundancy": redundancy, "grounding": grounding,
"retrieved_doc_ids": parsed.doc_ids,
"failure_class": classify_failure(parsed, answer, gold),
}
audit_f.write(json.dumps(record, ensure_ascii=False) + "\n")
```

Read in aggregate, the `failure_class` field turns the audit into a per-benchmark repair list. This is how [Chapter 36](#)’s losses were diagnosed:

Dataset	Dominant failure class	Suggested fix
NQ	imprecise extraction	stricter answer-span extraction / normalization
PopQA	retrieval miss (long-tail entity)	better/fresher retriever
MuSiQue	insufficient multi-hop chain	more turns / curriculum / decomposed rewards
2Wiki (late training)	lazy collapse / multi-hop degradation	early stopping / turn curriculum / entropy control

Table 38.2. Per-dataset dominant failure class and the suggested fix. Each benchmark fails in a characteristic way — imprecise extraction, long-tail retrieval misses, too-short multi-hop chains, or late-training lazy collapse — and each wants a different remedy.

38.2 Why an audit, separate from the reward and the eval

It’s fair to ask why this needs to exist alongside the reward (which already scores behavior) and the eval (which already scores accuracy). Three reasons, each earned earlier in the Part:

- It catches reward hacking the reward can’t report. When the model games the objective ([Chapter 31](#)), the reward goes up, so the reward can’t flag its own exploitation. The behavioral audit can: it shows the searching-without-grounding, the activity-without-accuracy, directly in the trajectories. It would have caught step 1302 by looking at what the model was doing.

- It explains why, not just whether. Accuracy tells you the agent missed. The audit tells you the misses were extraction-imprecise on NQ and retrieval failures on PopQA (Chapter 36's losses, diagnosed). That's the difference between knowing you have a problem and knowing what it is.
- It's a forensic record you can re-read. Logged trajectories let you go back and ask new questions of old runs, the same re-mining discipline that, elsewhere in this book, extracts findings from checkpoints you've already paid for.

The transferable principle closes the Part's argument about behavior: build the audit, not just the metric. A scalar reward and a scalar accuracy are decisions compressed to a single number each. The audit is the un-compressed record, and when something is wrong, the record is where you find out what.

TAKEAWAY

Accuracy is silent on how an answer was produced, so we log every trajectory to JSONL and audit the behaviors exact-match can't see: did it search, ground, refine, or spam, and where did wrong answers fail. The audit exists alongside reward and eval because it catches reward hacking the reward can't report, explains *why* not just *whether*, and preserves a re-readable forensic record. Build the audit, not just the metric.

Next: Part IV. The search agent taught us reward design. It also exposed a deeper instability: even a better reward dies if entropy collapses or explodes.

Part IV

Stability: Entropy, Collapse, and Clip-Cov

Underneath every collapse, every plateau, every reward hack in the last two Parts sat one quantity we kept gesturing at: entropy, the model's willingness to explore. This Part makes it the central instrument. We trace a five-version saga, v1 through v5, in which the agent's exploration collapses, then over-corrects, then is finally stabilized, and that saga is what earns Clip-Cov, the one mechanism we found that holds entropy in a healthy band instead of letting it die or explode. Entropy is the book's most useful diagnostic, and this is where you learn to read it.

CHAPTER 39

What Entropy Measures in a Policy



Figure 39.1. Entropy is the spread of the next-token distribution, shown across regimes.

We have used the word “exploration” loosely for two Parts. It’s time to make it a number, because that number, entropy, is the single most useful diagnostic in RL, and almost every failure you’ve seen has it underneath. Entropy measures one thing: how spread-out the model’s probability is over its possible next tokens. That’s it. But what it implies for training is most of the rest of this book.

In code, it is one line per position, averaged over the batch, which is why there is no excuse not to log it:

```
probs = logits.softmax(-1)
entropy = -(probs * probs.clamp_min(1e-12).log()).sum(-1).mean()
```

39.1 Low, high, and healthy

Three regimes, drawn in Figure 39.1, each with a distinct meaning for a run:

1. Low entropy, the model piles almost all its probability on one or two tokens. It is certain, committed, deterministic. It will generate nearly the same thing every time. For RL, this means it has stopped exploring. It keeps proposing its current favorite, so its groups stop disagreeing ([Chapter 5](#)) and there's nothing new to reinforce.
2. High entropy, the model spreads probability thinly over many tokens. It is uncertain to the point of random. Its outputs are varied but incoherent. It explores wildly but can't commit to anything good long enough to build on it.
3. Healthy entropy, a peaked distribution that still keeps alternatives alive: a few likely tokens carry most of the mass, but not all of it. The model commits enough to be coherent and explores enough to find something better. This middle band is where learning lives.

The whole game, stated through entropy, is to keep the policy in that middle band over the course of training.

39.2 Why RL pushes entropy down on its own

Here is the uncomfortable fact that makes this an entire Part: RL has a built-in tendency to lower entropy, and left alone it drives the model toward the collapsed regime. The reason is mechanical:

- RL reinforces the model's successful attempts. It raises the probability of tokens that worked.
- Raising the probability of the winners concentrates the distribution, by definition, that lowers entropy.
- A lower-entropy model explores less, so it finds fewer new successes, so it reinforces an even narrower set of behaviors...
- ...which lowers entropy further. It's a feedback loop, and its fixed point is collapse: a confident, committed, no-longer-learning model.

This is why entropy isn't just a diagnostic but the one for stability. The very thing that makes RL work, reinforcing success, is the thing that quietly kills exploration. Manage that tension and your runs stay

alive. Ignore it and they die in the specific, recognizable way the next chapter shows. Everything that follows, v1's collapse, the failed fixes, Clip-Cov, is a response to this single feedback loop.

TAKEAWAY

Entropy measures how spread-out the policy's next-token distribution is: low = committed but not exploring (heading for collapse), high = exploring but incoherent, healthy = a peaked distribution that still keeps alternatives alive. Log it from step one. It's two lines. The catch that makes this a whole Part: RL reinforces winners, which concentrates the distribution, which lowers entropy, which kills exploration, a feedback loop whose fixed point is collapse.

Next: that loop, observed, v1, where entropy slid from 0.188 to 0.070 and the agent stopped exploring.

CHAPTER 40

v1: Entropy Collapse




Figure 40.1. The entropy saga: collapse, explosion, and the healthy band between them. Per-run anchors are logged (v1–v5 endpoints, the v2 ceiling); the shape between anchors is interpolated.

The previous chapter described the collapse loop in the abstract. Here it is on a real run. Version 1 of the search agent, watched through its entropy, did exactly what the theory predicts: entropy started around 0.188 and slid steadily to 0.070 over 250 steps. That slide is not a side-effect of training. On this run, it was the story, the model quietly talking itself out of exploring.

40.1 What the slide does to the run

As entropy fell, the consequences cascaded in the order the loop predicts:




Figure 40.2. A first run’s entropy collapse, on the real curve.

1. The distribution concentrated. The model put more and more probability on its current-favorite tokens, its habitual phrasings, its default query patterns, its first-guess answers.
2. Exploration dried up. With probability concentrated, the roll-outs in each group started looking alike. The model proposed the same things repeatedly instead of trying alternatives.
3. Groups stopped disagreeing. Similar rollouts get similar rewards, so the within-group variance fell, straight back to [Chapter 5](#)’s advantage collapse, now caused from the entropy side rather than the reward side.
4. Learning stalled. No disagreement, no gradient. The model was locked into whatever it had become by the time entropy died, unable to climb further.

The two views, falling entropy and vanishing group variance, are the same event seen from two instruments. That’s worth internalizing: entropy collapse and advantage collapse are the same failure, one measured on the policy’s distribution, the other on the group’s rewards. When you see one, look for the other.

40.2 Why you can't just wait it out

The cruel part of v1 is that it doesn't announce itself as a failure. The reward might still look flat-but-fine, the loss still moves, the run still runs. Nothing crashes. But the model has stopped being able to improve, because it has stopped being able to explore, and no amount of additional steps fixes that, since each additional step only reinforces the narrowed distribution further. Patience makes it worse, not better. The fix has to act on the entropy itself: either stop it from falling, or push it back up. The next three chapters are the history of us trying to do exactly that, and getting it wrong twice before getting it right.

TAKEAWAY

v1 is the collapse loop made real: entropy slid from 0.188 to 0.070 over 250 steps, the distribution concentrated, rollouts converged, group variance vanished, and learning stalled. Entropy collapse and advantage collapse are one failure on two instruments. And it doesn't crash, the run looks alive while it's already dead, so waiting only deepens it. The fix must act on entropy directly, which the next chapters attempt.

Next: the obvious fix, pay the model to stay uncertain, and why v2 and v3 proved a global entropy bonus has no stable setting.

CHAPTER 41

v2 & v3: The Entropy-Explosion Trap




Figure 41.1. A global entropy bonus has no stable setting — it over- or under-shoots.

If entropy is collapsing, the obvious fix is to pay the model to keep it up, add an entropy bonus, a term in the reward that rewards the policy for staying uncertain. It is the first thing everyone tries, it is what v2 and v3 tried, and it taught us something more useful than a fix: a global entropy bonus has no stable setting. You don't get to tune it to “just right,” because just-right doesn't exist. You get to choose which way it fails.

41.1 Why a global bonus has no equilibrium

The trouble is that a single, global entropy bonus is one blunt knob fighting a process that doesn't push back proportionally. Walk the two directions:

1. Bonus too small $\rightarrow$ it loses to the collapse loop ([Chapter 39](#)). The reinforcement-concentrates-entropy pressure overwhelms the

gentle upward nudge, and entropy still slides to collapse. You've added a term and changed nothing.

2. Bonus too large → entropy explodes. Once the bonus is strong enough to overpower collapse, there's nothing to stop it overpowering it completely: the model discovers that being uncertain pays, so it becomes maximally uncertain, spreading probability everywhere, generating incoherent noise (Chapter 39's high-entropy regime). It games the entropy bonus exactly the way it gamed the behavioral reward in Chapter 31.
3. In between → there's no stable middle, because the bonus is a fixed force and the collapse pressure varies over training. A setting that balances at step 50 is too weak at step 200 (collapse) or too strong at step 400 (explosion). The equilibrium you want is a moving target, and a constant bonus can't track it.

So v2 and v3 oscillated between the two failures depending on the coefficient: dial it down and watch v1's collapse return, dial it up and watch the model dissolve into noise. There was no value that held entropy in the healthy band and kept it there.

41.2 The lesson: the wrong control surface

The deeper takeaway isn't "entropy bonuses are bad". It's that we were controlling the wrong thing. A global bonus treats entropy as a single scalar to be propped up everywhere, equally, all the time. But the collapse isn't uniform: it's driven by specific tokens, the ones where the model is over-confidently concentrating, not by the policy's average uncertainty. Pushing up the average with a global force is like steadying a wobbling table by pressing down on its center: you fight the symptom (low average entropy) without touching the cause (particular over-committed tokens), and you risk flipping it over the other way.

That reframing is what eventually points at Clip-Cov, which controls entropy at the token level instead of the global level. But two chapters of getting it wrong come first, and the next one is the half-fix that delays the collapse just long enough to be dangerous.

TAKEAWAY

The obvious fix, a global entropy bonus (v2, v3), has no stable setting: too small and the collapse loop still wins, too large and entropy explodes into incoherent noise, and no constant value holds the middle because the collapse pressure varies over training while the bonus is fixed. The real error is the control surface: a global scalar can't fix a problem driven by specific over-committed tokens. That points, eventually, at token-level control.

Next: v4, Clip-Higher, a more surgical tool that delays the collapse but, on its own, still can't prevent it.

CHAPTER 42

v4: Clip-Higher Alone Isn't Enough

By v4 we had stopped trying to bolt a bonus onto the reward and started reaching into the update itself. The tool was Clip-Higher, DAPO's asymmetric clip from [Chapter 15](#), and it is genuinely better-aimed than a global bonus. It bought real time. It also, on its own, still lost. v4 is the chapter about a fix that delays a failure convincingly enough to fool you into thinking it prevented it.

42.1 What Clip-Higher does for entropy

Recall the clip from [Chapter 3](#): the update caps how far a token's probability can move in one step, in both directions. Clip-Higher ([Chapter 15](#)) breaks that symmetry. It keeps the usual floor but raises the ceiling, and that asymmetry is pointed straight at entropy:

- The collapse loop works by suppressing low-probability tokens. The model's exploration lives in its rare, low-probability tokens. Standard clipping caps how fast a good rare token can grow, so promising exploratory moves get throttled before they can establish themselves, and the distribution narrows.
- Raising the ceiling lets rare-but-good tokens grow faster. Clip-Higher gives those low-probability exploratory tokens room to climb when they earn it, so exploration isn't strangled in the crib. It directly counters the mechanism that concentrates the distribution.

This is a real improvement over v2/v3: it's surgical (it acts on the update asymmetrically rather than propping up a global average) and it doesn't have the explosion failure mode, because it's loosening a constraint, not paying an unbounded bonus.

42.2 Why it delays but doesn't prevent

And yet v4 still collapsed, just later. The reason is that Clip-Higher widens the channel for exploration without addressing the pressure pushing entropy down:

1. It removes a throttle, not the force. Clip-Higher lets good rare tokens grow when they appear, but the underlying reinforcement-concentrates-entropy dynamic ([Chapter 39](#)) is still running, still pushing the distribution to narrow. Clip-Higher slows the narrowing. It doesn't stop it.
2. So the collapse is postponed, not prevented. Entropy falls more slowly, the healthy band lasts longer, the run looks stable for a good while, and then the same slide arrives, just at a later step. You've extended the runway, not learned to stay airborne.
3. And a delayed collapse is a trap, because it's long enough to look like success. You can run a v4 experiment, see stable entropy for a few hundred steps, conclude you've solved it, and ship, only to have longer runs collapse exactly as before. A half-fix that holds for the length of your test is more dangerous than an obvious failure.

The honest verdict on v4: necessary, insufficient. Clip-Higher belongs in the stack, it's part of the eventual solution, but alone it manages the symptom's timing, not its existence. What's missing is something that acts on the collapse pressure at its source, on the specific tokens driving it. That's the next chapter, and it's the one that finally works.

TAKEAWAY

v4's Clip-Higher (asymmetric clipping, raised ceiling) lets rare-but-good exploratory tokens grow instead of being throttled, a real, surgical improvement with no explosion failure mode. But it removes a throttle, not the collapse *force*, so it delays entropy collapse rather than preventing it, and a delayed collapse is a trap, because it holds long enough to look solved before failing on longer runs. Necessary, insufficient: it belongs in the stack but can't stabilize entropy alone.

Next: v5, Clip-Cov, token-level covariance control, and the first time entropy held a healthy band instead of collapsing or exploding.

CHAPTER 43

v5: Clip-Cov and the Damped Oscillation

v5 with Clip-Cov: damped oscillation into a held band ($\approx 0.73 \rightarrow 0.64$)




Figure 43.1. Clip-Cov damps the oscillation back into the healthy band.

This is the chapter the whole Part has been climbing toward, and the fix that earns it. After v1's collapse, v2/v3's explosion, and v4's delay, v5 introduced Clip-Cov, token-level covariance control, and for the first time, entropy did neither of the two things it had always done. It didn't collapse. It didn't explode. It oscillated with decreasing amplitude and settled into a healthy band it held for the length of the run, exactly as Figure 43.1 shows. The agent kept exploring, kept committing, and kept learning, over a long run, not just a lucky window. This is the version whose step-750 validation produced the peak numbers of [Chapter 36](#).

43.1 The idea: control entropy where it's actually leaking

Every failed fix so far treated entropy as a global average to prop up. Clip-Cov's move is to ask which specific tokens are driving the collapse and act only on those. The lever is covariance:

1. Find the tokens where probability and advantage co-move dangerously. The collapse is driven by tokens where the model is rapidly concentrating probability, specifically, where a token's probability and its advantage have a high covariance, the model is aggressively reinforcing an already-confident choice, squeezing out exploration. That covariance is a measurable, per-token signal of where entropy is leaking.
2. Clip the update on exactly those tokens. Rather than a global bonus (which touches everything) or a global ceiling change (v4), Clip-Cov caps the update specifically on the high-covariance tokens, the ones responsible for the narrowing. It's a targeted intervention at the source of the collapse, not a blunt force on the average.
3. Leave the rest of the update alone. Tokens that aren't driving collapse train normally. The control is surgical. It spends its corrective force only where the leak is.

The shape of the implementation, and the actual v5 settings (not a generic “sweep winner”, the values that produced [Chapter 36's](#) peak):

```
# Clip-Cov is LOSS-SIDE token filtering, not a reward term -
↳ readers copy this wrong otherwise.
# per token: cov_i = (p_i - mean(p)) · (A_i - mean(A)). Detach
↳ the update on the top clip_cov_ratio fraction
actor.policy_loss.loss_mode: clip_cov
actor.policy_loss.clip_cov_ratio: 0.0002 # Clip-Cov ratio (v5)
actor.policy_loss.clip_cov_lb: 1.0
actor.policy_loss.clip_cov_ub: 5.0
actor.entropy_coeff: 0.0
actor.clip_ratio_high: 0.28
```

Different run families used different Clip-Cov settings, and conflating them makes configs unreproducible, so here is the per-run table,

and the rule is: don't call any setting “the sweep winner” unless it's tied to its exact run:

Run family	Model / task	clip_cov_ ratio	lb	ub	Role
Search v5	Qwen3-4B search agent	0.0002	1.0	5.0	the original stability result (Ch 36 peak)
Manual de- fault / later GRPO++	generic	0.001	0.6	5.0	recommended starting point if verified
C8 R2E- Gym	Qwen3-14B code env	not mea- sured	not mea- sured	not mea- sured	cross-domain validation

Table 43.1. Clip-Cov settings per run family, with the reusable rule (lb/ub are the lower and upper covariance-ratio bounds). *Search v5* (the search agent) gives the original stability result; the manual default is the recommended starting point once verified; *C8 R2E-Gym* (a code environment) is the cross-domain validation slot, a deliberate open boundary (that run is not yet done).

We also ran this on a code environment (R2E-Gym, 14B) as a cross-domain check; settings are reported later.

This is the fix to the diagnosis of [Chapter 41](#): the collapse was never uniform, so the cure shouldn't be either. Control entropy at the token level, where it's actually being lost, and the global average takes care of itself.

43.2 The damped oscillation, and the honest correction on “equilibrium”

The shape in Figure 43.1, overshoot, dip, smaller overshoot, settle, is not incidental. It's what a stable control system looks like, and seeing it is how you know the fix is working:

- Overshoot and correction. When entropy starts to fall, Clip-Cov's targeted clipping pushes back. The push slightly overshoots, entropy rises a bit too much, then eases back. Each swing is smaller than the last.
- Convergence to a band, not a point. Our first journal entry called it "equilibrium at 0.73." The step-750 data corrected us: entropy was still settling, slowly, toward ~0.64. The honest claim, and the one we'll defend, is a healthy band held for the length of a real run, descending gently, never collapsing and never exploding. That is categorically different from v1's slide and v4's delayed cliff, and it is what no earlier version achieved.
- The slow descent was not even bad. At step 750, the gently tightening policy was helping single-hop accuracy, NQ +8.4pp, TriviaQA +5.6pp over step 500, while grounding held at 100%. The goal, it turns out, is not frozen entropy but controlled entropy: exploration that declines on the policy's schedule instead of the collapse loop's.

That held band, entropy kept healthy for the length of a real run, under control rather than in freefall, is the result this Part exists to deliver. The mechanism is simple to name (clip the high-covariance tokens) and the outcome is the one nobody else seems to get: an RL run whose exploration doesn't die.

TAKEAWAY

v5's Clip-Cov controls entropy at the token level: it ranks tokens by probability-advantage covariance and clips the update on the top fraction (v5 used ratio **0.0002**, lower-bound **1.0**, upper-bound 5.0, a later generic default is 0.001/0.6, don't conflate them), fixing the leak at its source instead of propping up the global average. It's loss-side token filtering, not a reward term. The result is a damped oscillation settling into a healthy, slowly descending band ($\approx 0.73 \rightarrow 0.64$) held over a real run, and the gentle tightening *helped* single-hop accuracy. Not a frozen fixed point: controlled entropy.

Next: why this works when everything else didn't, the covariance mechanism, its PRIME-RL lineage, and the held band almost nobody else reports.

CHAPTER 44

Why Clip-Cov Works

v5 works. This chapter is about why, because the why is what lets you trust it, tune it, and recognize when it's the right tool. Clip-Cov isn't a lucky hyperparameter. It's a targeted answer to the exact failure of every earlier attempt, with a traceable lineage and a result that, as far as we can tell from the literature, is unusual: a reported, reproduced held entropy band on a real agentic run.

44.1 The mechanism, against the failures it fixes

The cleanest way to understand Clip-Cov is as the thing that's right where each predecessor was wrong:

1. Global bonus (v2/v3) was wrong about where. It pushed entropy up everywhere, uniformly, and had no stable setting because the collapse isn't uniform. Clip-Cov acts only on the tokens driving collapse, so it doesn't have to balance an unbalanceable global force.
2. Clip-Higher (v4) was wrong about what. It loosened a constraint on all rare tokens but didn't address the concentrating pressure. Clip-Cov targets that pressure directly, the high-covariance tokens are the pressure, so it removes the cause rather than widening a channel around it.
3. Clip-Cov is right about both. It identifies where entropy is leaking (high probability-advantage covariance) and acts on what causes the leak (over-aggressive reinforcement of already-confident tokens). Targeted in location, aimed at the cause. That's why it reaches a held band the others structurally couldn't: it's a feedback controller acting on the actual state variable, not an open-loop force hoping to balance.

Stated as a controller: the global bonus was a constant input (no feedback), Clip-Higher was a relaxed limit (still no feedback on entropy), and Clip-Cov measures the leak and responds to it, which is precisely why its entropy curve looks like a damped control system settling, rather than a quantity drifting to a boundary.

44.2 Lineage, and how much to claim

Two honest notes on where this comes from and the size of the claim:

- The lineage is PRIME-RL. Clip-Cov isn't invented from nothing. It descends from the covariance-based entropy-control ideas in the PRIME-RL line of work. We adapted and applied it on our stack, and the contribution here is less “a new algorithm” than “the demonstration that this approach holds entropy in a healthy band on a real multi-turn agentic run, and the recipe, including the swept settings, for getting there.” Crediting the lineage is part of the honesty discipline this book keeps.
- The held band is the unusual part, claimed precisely. Plenty of methods slow entropy collapse. What we kept not finding in others' reports was the thing in Figure 43.1: entropy oscillating, then holding a healthy, gently descending band for the length of a run, including its own correction (Chapter 43: “equilibrium at 0.73” revised to “settling toward ~0.64”). We claim the held band and the recipe, not a mathematical fixed point. That band, reported with its correction, is the genuinely notable result of this Part, and it's why entropy graduates from “a thing to worry about” to “a thing you can measure and control.” (The companion question, how high an entropy a given setup can sustain, the entropy–performance frontier you can fit on your own logs, is where this Part heads soon.)

The practical upshot, and the reason to carry Clip-Cov out of here: when an RL run's exploration is dying, the fix that holds is not a global bonus or a looser clip, but token-level covariance control, find where the distribution is over-concentrating and clip the update there. That's the mechanism, that's why it works, and that's the held band the rest of the field's tools tend to miss.

One more honesty note, so this chapter doesn't read as "Clip-Cov is magic": Clip-Cov does not freeze entropy at a target. It converts an uncontrolled collapse/explosion into a controllable feedback regime, and what that regime looks like depends on the task, model, reward variance, and learning rate:

Run family	Clip-Cov behavior	Lesson
Search v5	damped oscillation → slow compression	worked as a stabilizer (this Part)
No-Clip-Cov ablation	entropy alive but undirected, zigzag accuracy	entropy alone is insufficient
C8 R2E-Gym (code)	non-monotone, sometimes aggressive corrections	the correction can be asymmetric
V5.1 turn-curriculum	early eval win, later response-length collapse	Clip-Cov slows, doesn't solve, every collapse
C9 / GSPO low-entropy regimes	low entropy was productive sharpening	read entropy with validation, not alone

Table 44.1. Clip-Cov honesty notes, by run family. For each, the Clip-Cov behavior observed and the lesson it carries — the through-line being that Clip-Cov stabilizes and sharpens but does not, on its own, prevent every collapse.

So the precise claim is the held band plus its boundary conditions: a feedback regime you can steer, not a fixed point you can set.

TAKEAWAY

Clip-Cov works because it's a feedback controller aimed at the actual cause: it measures where entropy is leaking (high probability–advantage covariance) and clips the update *there*, fixing what the global bonus (wrong location) and Clip-Higher (wrong target) both missed. Its lineage is PRIME-RL. Its notable contribution is a *reported, reproduced held entropy band*, claimed precisely, with its own correction (0.73 → settling toward ~0.64) kept in the record. But it converts uncontrolled collapse into a *controllable regime* (damped oscillation / one-cycle / slow compression by task), not a frozen fixed point. When exploration is dying, token-level covariance control is the fix that holds.

Next: not every entropy idea worked, the per-channel approach (GDPO) we tried across G1–G4, and the post-mortem on why it failed.

CHAPTER 45

GDPO: The Per-Channel Idea That Failed




Figure 45.1. A successful refinement is not a forever-gradient — improvement runs out.

Clip-Cov was not the last thing we tried. It was the last thing that worked. After it, we pushed further, and the next idea failed across four versions. This chapter is the post-mortem on GDPO, Group-Decomposed Policy Optimization (our own per-channel reward decomposition, not to be confused with NVIDIA’s similarly-named Group reward-Decoupled Normalization Policy Optimization, arXiv:2601.05242), which we ran through G1 to G4, because a failure this thorough teaches as much as the success before it, and because the honest thing to do with a dead end is to mark it so you don’t walk down it again.

45.1 The idea, and why it was tempting

Our gated reward sums five channels, F1, grounding, cost, redundancy, and an action penalty, into one scalar before the group baseline is computed. GDPO’s premise is that the summing throws information away:

1. The intuition. Normalize each reward channel separately within the group, a per-channel z-score, and then combine the

- per-channel advantages. This should preserve distinctions the scalar sum erases: “good on F1 but bad on grounding” and “bad on F1 but good on grounding” can sum to the same number, yet they deserve different gradients.
2. The appeal. It is principled (it’s a published method), it slots into the same loop (just a different advantage estimator: `adv_estimator=gdpo` with per-channel keys and weights), and it promised finer credit assignment over exactly the multi-component reward we’d spent the Part building.
 3. The versions. We built and ran it across G1 through G4, vanilla GDPO on five channels, then variants layering in mini-batching and other adjustments, trying to reach the stability and accuracy v5 had shown.

It was a reasonable hypothesis. It was also wrong for this setting, and running it four ways is how we became sure rather than assuming.

45.2 Why it failed, and the lesson

None of G1–G4 beat v5. The runs, concretely:

Variant	Idea	Best / stop	Result	Diagnosis
G1	vanilla GDPO, 5 channels	not measured	failed	per-channel dilution
G2	strategy channels	not measured	crashed / unstable	channel conflict
G3	gated strategy channels	OOM ~step 213	no learning / incomplete	implementation + no gain
G4	budgeted marginal utility	~step 250	zero learning	F1 signal diluted
v5 baseline	scalar gated reward + Clip-Cov	250/500/750	learned coherent scalar	direction

Table 45.1. The GDPO variants — G1 through G4 are successive GDPO experiment runs (vanilla 5-channel through budgeted marginal-utility), compared against the v5 scalar-reward baseline. For each: the idea, where it was best or stopped, the result, and

the diagnosis. Every multi-channel variant diluted the correctness signal; only the scalar gated reward learned a coherent direction.

The post-mortem, stated honestly:

- The channels aren't independent, and the method assumes they are. Per-channel z-scoring makes sense when reward dimensions vary somewhat independently. In multi-turn retrieval they don't: F1, grounding, and cost move together (a trajectory that searches well tends to ground well and cost more). Z-scoring correlated channels separately means the model can only improve on one by appearing to worsen on another, the per-channel gradients contradict and partially cancel, and the update loses the coherent direction the scalar reward had.
- The gate had already solved the problem GDPO targets. The distinction GDPO preserves, behavioral quality separate from correctness, is exactly what the correctness gate, the SiLU/Swish gate of [Chapter 32](#), handles, by conditioning behavior on correctness rather than decomposing them. With the gate in place, the correctness gate effectively learns the channel weighting (F1 dominates), and GDPO's decomposition adds machinery where the structure was already encoded.
- Four versions earned a clean negative. G1–G4 didn't fail randomly. They failed consistently, in a way that pointed at the premise (channel independence) rather than the implementation. That's what turns "this didn't work" into "this approach doesn't work here", a result worth recording, with the boundary stated: on a task whose reward channels are genuinely uncorrelated, GDPO's premise may hold. On multi-turn retrieval, it didn't.

The transferable lesson is about the direction of fixes, not GDPO specifically: a successful move in one direction is not a gradient to follow indefinitely. Token-level control (Clip-Cov) worked because the collapse has real structure at the token level. Per-channel decomposition failed because the reward channels lacked the independence the method needs. Each refinement has to be justified by the phenomenon having structure where the refinement looks, not by the

fact that the previous refinement helped. Clip-Cov plus the gate remains the recommendation. GDPO is the marked dead end beside it, and marking it is part of the map.

TAKEAWAY

GDPO (Group-Decomposed Policy Optimization) z-scores each reward channel separately within the group to preserve distinctions the scalar sum erases. Across G1–G4 it never beat v5: in multi-turn retrieval the channels (F1, grounding, cost) are correlated, so per-channel normalization produces contradictory, partially cancelling gradients, and the correctness gate had already encoded the structure GDPO decomposes. A consistent four-version failure is a clean negative worth recording, with its boundary: the method’s premise is channel independence, and this task doesn’t have it.

Next: the dead end left one gift, its diagnostics. Watching G1–G4 fail taught us that strategy collapse precedes reward collapse, and that the turn-count is a siren that fires ~15 steps before the score knows anything is wrong.

CHAPTER 46

Strategy Collapse Comes First




Figure 46.1. The reward is the last graph to know: strategy moves first.




Figure 46.2. Strategy breaks about 15 steps before reward does — the early-warning window.

You are watching the reward, and it is holding. Fifteen steps from now it will fall off a cliff and the run will be dead, but right now it

looks fine, and the one graph that could have warned you is a different graph you weren't watching. That gap, between when a run actually breaks and when the reward admits it, is the subject of this chapter. The model's strategy collapses first. The reward collapses about fifteen steps later, so the reward is the worst place to catch a failure early.

We learned this watching the per-channel runs fail. The four versions of that dead end (journal IDs G1–G4, the post-mortem of [Chapter 45](#)) didn't just die. They died in a sequence, and the sequence is the lesson.

46.1 Reward is a lagging indicator

Plain, then precise, then what it buys you:

- Plain. By the time the reward drops, the model broke a while ago. You're seeing an echo.
- Precise. Reward is an outcome, measured at the end of a roll-out. The behavior that produces it degrades upstream of it. In the per-channel runs (G1–G4), the model's strategy, how it structured its attempts, destabilized first, and the reward reflected that destabilization only about fifteen steps later.
- Payoff. So the reward curve, the thing everyone stares at, is structurally the last signal to move. To catch a collapse while you can still act, you need a metric that moves when the strategy moves, not when the score does.

46.2 The leading signal: a strategy metric that spikes early

The specific tell in the search agent was the turn-count, the number of search-and-read steps the model took per episode. Before the reward fell, the turn-count spiked: the model started thrashing, taking erratic numbers of turns, visibly losing its strategy about fifteen steps ahead of the reward registering it.

The concrete leading signal here is simple, per-episode turn-count, a piece of strategy telemetry in the [Chapter 29](#) sense, now used as an early-warning gauge. The whole check is a few lines:

```
# strategy spike leads reward collapse by ~15 steps - watch it,  
↳ not the reward  
recent = turn_counts[-window:] # per-episode turn counts, recent  
↳ window  
if recent.mean() > 1.5 * baseline_turns: # rule of thumb from  
↳ THESE runs, not universal  
    warn("strategy spike - reward likely to fall within ~15  
        ↳ steps. Checkpoint now")
```

The comment carries the operating rule, and its honest scope: above roughly one-and-a-half times the baseline turn-count, treat it as a collapse warning. That 1.5× is a rule of thumb calibrated on these runs, not a universal threshold. The pattern (strategy leads reward) is the durable claim, the exact multiplier is yours to tune.

46.3 What to do with fifteen steps' warning

The warning is only worth having if you act on it, so the moment the leading signal fires:

1. Checkpoint immediately, so you can roll back to a step before the collapse rather than salvaging from after it.
2. Intervene, drop the learning rate, apply the token-level entropy control from [Chapter 43](#), instead of waiting for the reward to confirm what the strategy metric already told you.
3. Keep logging the behavioral signal for the whole run, not just when you suspect trouble. The spike is only useful if you were already watching.

The principle generalizes past turn-count and past this agent: instrument the behavior, not just the outcome. Outcome metrics lag. Behavioral metrics lead. This is the same lesson [Part III](#) drew about accuracy versus behavior, now sharpened into a timing claim you can set an alarm on.

TAKEAWAY

Reward collapse is a lagging indicator, the model's strategy breaks first and the reward follows about fifteen steps later (seen across the per-channel runs, G1–G4). Watch a behavioral signal like per-episode turn-count (strategy telemetry, distinct from the loss-aggregation mode ([Chapter 47](#))): when it spikes past $\sim 1.5\times$ baseline (a rule of thumb, not a universal threshold), checkpoint and intervene immediately rather than waiting for the reward to confirm. Behavioral metrics lead. Outcome metrics lag.

Next: a length bias hiding in one line of the loss, how you average it decides whether the model learns to be concise or to ramble.

CHAPTER 47

Loss Aggregation and the Hidden Length Bias




Figure 47.1. The same batch, two averaging orders — and why the order changes the gradient.

Here is a difference most people don't know they're choosing. Two runs, same data, same reward, same hyperparameters, except for one line that decides how the loss is averaged across tokens and sequences. One run drifts toward concise answers. The other learns to ramble. How you aggregate the loss creates a length bias, and the common default quietly favors length in a way that has nothing to do with whether the answer is right.

47.1 Two ways to average the same loss

There are two natural ways to turn a batch of per-token losses into one number, and they are not equivalent:

1. **Token-mean.** Sum the loss over all tokens in the batch and divide by the total token count. Every token is weighted equally, so a long response, having more tokens, contributes proportionally more to the batch loss than a short one.

2. **Sequence-mean-then-token-mean.** Average the loss within each sequence first, then average those per-sequence numbers. Every sequence is weighted equally, regardless of how many tokens it has.

The whole difference is one line:

```
# token-mean: long sequences dominate the batch loss
loss = (per_token_loss * mask).sum() / mask.sum()

# seq-mean-token-mean: every sequence counts equally, regardless
  ↳ of length
seq_loss = (per_token_loss * mask).sum(dim=1) / mask.sum(dim=1)
  ↳ # per-sequence mean
loss = seq_loss.mean() # then mean over sequences
```

47.2 Why that one line biases length

Under token-mean, a long response’s many tokens each carry equal weight, so the response as a whole exerts more pull on the update than a short one does. Couple that with the advantage’s sign and a bias appears: a long wrong answer spreads its penalty thin across many lightly-penalized tokens, while a long right answer spreads its reward across many tokens, so the model can learn that length itself shifts the loss, independent of correctness. This is exactly the length bias that Dr. GRPO ([Chapter 15](#)) was built to remove.

Under sequence-mean-then-token-mean, each sequence’s contribution is normalized by its own length before sequences are compared, so a long answer and a short answer pull equally. Length stops being a free variable the model can exploit.

47.3 Which to use, and the meta-point

The honest answer is that it depends on your task, but you must know which one you’re running. The practical rule:

- If response length creeps up while accuracy stays flat ([Chapter 15](#)'s symptom), the aggregation is a prime suspect. Switch to sequence-mean-then-token-mean to neutralize the length pull.
- Read your trainer's aggregation before assuming the math. DAPO uses a token-level loss for a related reason, and the interaction between aggregation, clipping, and the advantage is subtle enough that the only safe move is to know your default rather than inherit it blindly.

The meta-point is the one that matters beyond this single knob: the loss function has implementation choices that aren't written in the equation, and they change what the model learns. A length bias isn't in any GRPO paper's objective. It's in how the sum is taken.

TAKEAWAY

Token-mean (weight every token equally) lets long responses dominate the batch loss and quietly rewards length independent of correctness. Sequence-mean-then-token-mean (weight every sequence equally) removes that bias. It's a one-line difference. If length creeps while accuracy stays flat, suspect the aggregation. The broader lesson: the loss has choices outside the math, and they change behavior, know your default.

Next: a wall of variant names, Dr. GRPO, DAPO, GSPO, RLOO, and more, and the short set of axes that collapses the whole zoo into one map.

CHAPTER 48

The GRPO Variant Taxonomy

Choosing an RL update: the GRPO family selector

Walk the spine top to bottom — take the first yes.

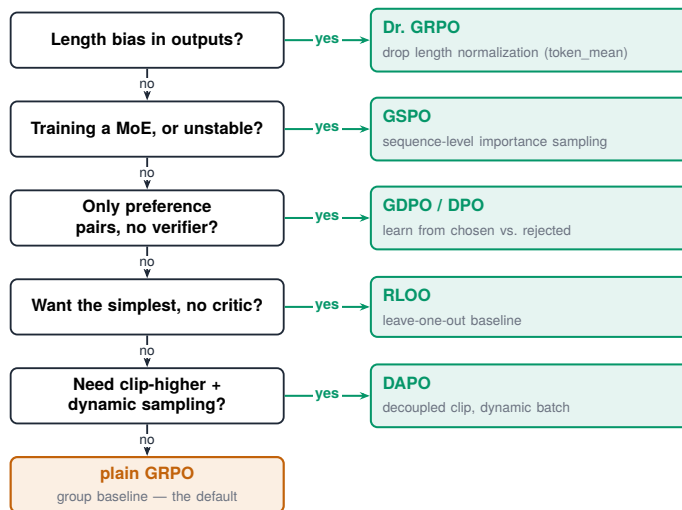


Figure 48.1. Choosing an RL update: a clean decision tree through the GRPO family.

You will hit a wall of names, Dr. GRPO, DAPO, Group Sequence Policy Optimization (GSPO), REINFORCE leave-one-out (RLOO), and rarer ones like S-GRPO, MC-GRPO, DRA-GRPO, each with a forum thread swearing it's essential. Most are one small change to the same loop you already run. Every variant modifies one of a few axes, and once you know the axes, you can place any new name in seconds and decide in seconds whether you need it.

48.1 The axes of variation

Group Relative Policy Optimization (GRPO) can be modified in a small number of independent places. Learn these five and the zoo collapses:

1. The baseline / advantage, what you subtract from each reward and how you normalize it (group mean, leave-one-out mean, median, a learned critic).
2. The importance-sampling level, whether the new-over-old probability ratio is formed per token or per sequence.
3. The clip, symmetric, asymmetric (clip-higher), or covariance-gated.
4. The loss aggregation, token-mean versus sequence-mean-then-token-mean ([Chapter 47](#)).
5. The sampling, what gets generated and kept: static groups, dynamic resampling of unanimous groups, structured early exits, diversity-aware weighting.

48.2 The map

Every named variant is “GRPO, with one or two axes swapped.” The map (one row per variant, the axis it touches, the failure it targets, where this book covers it):

Variant	Changes (axis)	Targets	Cov- ered
GRPO	group-mean baseline (1)	the baseline itself	Ch 3
Dr. GRPO	drop $\div$ std + fix length norm (1, 4)	difficulty & length bias	Chs 15, 47
DAPO	dynamic sampling + clip-higher + token loss (3, 4, 5)	collapse & entropy slide	Ch 15
RLOO	leave-one-out baseline (1)	self-contamination of the baseline	Ch 15
GSPO	sequence-level IS (2)	MoE expert-flip	Ch 52
S-GRPO (arXiv 2505.07686)	serial early-exit sampling with a decaying reward for correct early stops (5)	overthinking — 35–61% shorter outputs at equal-or-better accuracy	—

Variant	Changes (axis)	Targets	Cov- ered
MC-GRPO (arXiv 2601.22582)	median baseline + MAD in- stead of mean/std (1)	advantage noise at tiny group sizes ($G = 2\text{--}4$)	—
DRA-GRPO (arXiv 2505.09655)	diversity-aware advantage reweighting via submodular mutual information (1, 5)	redundant completions drown- ing out novel reasoning paths	—

Table 48.1. The GRPO variant map: each algorithm, the axis it changes relative to GRPO, the failure it targets, and where the book covers it. The numbered axes (1–5) trace which part of the GRPO update each variant modifies.

48.3 How to use the map

The rule that keeps you out of the variant-chasing trap, carried straight from [Chapter 15](#):

- Don’t adopt a variant because it’s named. Adopt it because you see the failure it targets. Each is a one- or two-term diff on the loop you already understand, aimed at a specific symptom.
- When a new name appears, ask which axis it touches. That single question tells you what it’s for, and usually that you already understand it.

TAKEAWAY

Every GRPO variant changes one of five axes: the baseline/advantage, the importance-sampling level, the clip, the loss aggregation, or the sampling. Dr. GRPO and RLOO touch the baseline. DAPO touches sampling and clip. GSPO touches the IS level. S-GRPO adds early exits against overthinking, MC-GRPO swaps in a median/MAD baseline for tiny groups, DRA-GRPO reweights advantages for diversity. Place any new name on the axes and you’ll know what it’s for. Adopt by failure signature, not by reputation.

Next: turning the entropy intuition into an actual measurement, a frontier you fit on your own logs, with one corner that separates healthy runs from the rest.

CHAPTER 49

The Entropy–Performance Frontier




Figure 49.1. Two model families, one frontier, read on the active-cost axis.




Figure 49.2. The other axis: MoE buys capability but pays in memory.

You’ve been told to “watch entropy.” Watch it how, exactly, and what counts as a good value? This chapter turns the hand-wavy advice

into something you can run: plot a run family as points in entropy-versus-reward space, fit a frontier, and aim for the corner. The one rule that keeps it honest, and that the rest of this chapter is careful about: fit the frontier within a run family, never across tasks, entropy numbers on different datasets live on different scales and don’t compare. Get that wrong and you’ll “rank” methods that were never on the same axis.

49.1 The frontier, measured on your own logs, one family at a time

The method is simple enough to run on metrics your trainer already emits:

1. For each run in one family, take two numbers, its sustained policy entropy and its reward, straight from the trainer’s inline metrics (veRL logs both).
2. Plot the family, fit the shape. A short script (`frontier_fit.py`) fits a saturating curve across the runs in that family, pure log re-analysis, zero compute.
3. Look for the corner, the runs holding both high entropy and high reward, on that family’s scale.

The families don’t share an axis, so the frontier is fit per-family:

Frontier family	Runs	Scale	Winner / finding	Lesson
Q4 quantization	paired 14B Clip-Cov BF16 vs FP8	reward $\approx$ 0.8, en- tropy $\approx$ 0.4–0.5	BF16 $\approx$ FP8; both hold the corner	quant/infer- ence lesson
C9 MoE stabiliz- ers	GSPO / FREEZE / JITTER / COMBO	C9-specific (en- tropy 2.6–3.8)	GSPO frontier highest	MoE- algorithm lesson
Search v5	v1–v5 + ablations	search-agent en- tropy / EM	Clip-Cov + gated reward	search-agent lesson

Table 49.1. The entropy–performance frontier families (Q4 = the 14B quantization runs, C9 = the MoE-stability family, *Search v5* = the

search agent), each with its scale, the winner or finding, and the lesson. BF16 and FP8 tie on the quant frontier; GSPO holds the highest MoE frontier; Clip-Cov plus a gated reward wins for the search agent.

The corner test is not a universal constant. The $\text{entropy} \geq 0.5$ & $\text{reward} \geq 0.8$ rule is a Q4-family example (Q4 here is the experiment label, not 4-bit precision), meaningless on the C9 scale, where entropy ceilings run past 3.8. The portable version normalizes within a family before fitting:

```
# Within ONE family only: normalize, then fit. Never threshold
↳ across tasks.
for fam, group in df.groupby("task_family"):
    z_ent = zscore(group.entropy)
    z_rew = zscore(group.reward)
    fit_frontier(z_ent, z_rew) # the "corner" is the high-z
    ↳ corner of THIS family
```

So the corner isn't a metaphor. It's a filter, but a per-family one.

49.2 The C9-family ceiling ranking, and the caveat that governs it

Within the C9 MoE family (and only there), the methods ranked by entropy ceiling, the highest entropy each could sustain:

GSPO (3.84) > JITTER (router-input noise) (3.07) > FREEZE (2.77) > COMBO (2.68)

A separate family gave an independent stability check: the Q4 quantization pair put BF16 and FP8 rollouts at nearly the same ceiling (0.82 vs 0.83), two arms that should match, matching, which says the measurement is stable rather than noisy.

Now the caveat that governs the whole chapter, stated as strongly as it deserves: those two rankings live on different scales and must never be compared. GSPO's 3.84 (C9) and Clip-Cov's ~0.8 (Q4) are coordinates on different maps, 3.84 is not "higher entropy than 0.8" in any meaningful sense. They're different tasks. The frontier ranks methods within a family. It never ranks tasks or compares numbers across

families. Every entropy ceiling in this book is a coordinate on one map, not a point on a universal scale.

49.3 What the frontier buys you

The frontier converts “watch entropy” from advice into procedure:

- Fit the frontier on your logs and aim for the corner. A method that can’t reach high entropy and high reward together is leaving performance on the table.
- Treat drift out of the corner as an alarm. A run whose point slides left (entropy dropping) or down (reward dropping) is a run to intervene on, the same early-warning instinct as [Chapter 46](#), now in two dimensions.

This is the measurement form of the whole Part. Entropy stops being a vibe you squint at and becomes a coordinate you can locate yourself on.

TAKEAWAY

Plot a *run family* as points in entropy-versus-reward space, fit a frontier, and aim for its corner, but fit it **within a family**, never across tasks, because entropy scales aren’t comparable across datasets. In the Q4 quantization family, Clip-Cov BF16/FP8 hold the high-entropy/high-reward corner (the 0.5/0.8 rule is a Q4 *example*, not universal). In the C9 MoE family, the ceiling ranks GSPO (3.84) > JITTER (3.07) > FREEZE (2.77) > COMBO (2.68). Those numbers live on different maps, normalize within a family before fitting, and never compare ceilings across families.

Next: [Part V](#). Once entropy control worked on dense and search runs, the next question was scale. Does the same machinery survive MoE routing?

Part V

Scaling Up: MoE, GSPO, and Routing

The move from a model you can hold in your head to a thirty-billion-parameter mixture-of-experts is where RL stops being about the algorithm and starts being about the architecture. The standard machinery, the importance-sampling ratio at the heart of every update, silently breaks on a model that routes each token to different sub-networks. This Part is about why it breaks, the sequence-level fix that repairs it, the run where that fix turned a dead-looking curve into a breakout, the fuller toolbox of routing-collapse cures, and the one combination (low-precision rollouts on a mixture-of-experts) that simply does not work, proven the hard way.

CHAPTER 50

14B Dense vs 30B-A3B MoE



Figure 50.1. MoE: all experts resident in memory, only a few active per token.

Dense vs MoE: what changes for RL

Dimension	Dense	MoE
Memory footprint	= active params	all experts resident (much larger)
Active compute / token	all params fire	only top-k experts fire (cheaper)
Importance-sampling ratio	stable token-level	explodes token-level — use GSPO
Routing	none	can drift / collapse to few experts
Training stability	predictable	sensitive to load balance + precision
Best update	GRPO / Dr. GRPO	GSPO (sequence-level)

Figure 50.2. Dense vs MoE: what changes for RL, side by side.

You have eight H100s and a choice: a 14-billion-parameter dense model you already understand, or a 30-billion-parameter mixture-of-experts model that’s advertised as cheaper to run, and that, as the next five chapters show, fights your trainer at every turn. The choice is real and the brochure number lies, because the headline (“30B capacity for 3B compute!”) hides the costs that matter for RL: memory, stability, and a broken update rule.

50.1 What a mixture-of-experts actually buys

Plain, then precise:

- Plain. A mixture-of-experts (MoE) model has many “expert” sub-networks but uses only a few per word, so it’s big in capacity but cheap in compute per word.
- Precise. The 30B-A3B model has 30 billion total parameters but only about 3 billion active per token (A3B means “3 billion active”). A router picks a few experts for each token. Forward compute scales with active parameters (cheap). Memory scales with total parameters (expensive, because every expert must be resident).

50.2 The tradeoffs, side by side

The decision along the axes that matter for agentic RL on 8×H100:

Dimension	14B dense	30B-A3B MoE
Total / active params	14B / 14B	30B / ~3B
Per-token compute	higher	lower (MoE wins throughput)
Memory footprint	lower (dense wins)	higher (all experts resident)
Training stability	stable (dense wins)	routing-collapse-prone
Standard RL machinery	works	breaks GRPO (Ch 51)
Capacity ceiling	lower	higher (MoE wins)

Table 50.1. 14B dense versus 30B-A3B MoE across the axes that matter. The MoE wins on per-token compute and capacity ceiling, but loses on memory and training stability — and it breaks standard GRPO, which is what the next chapters fix.

No column is a clean sweep. The MoE wins throughput and ceiling. The dense model wins memory, stability, and “the tools just work.”

50.3 The verdict for agentic RL

The honest call, given the rest of this Part:

- Dense (14B) is the safer default. It trains stably and every standard tool applies. For most agentic RL on this hardware, that reliability is worth more than the MoE's ceiling.
- Reach for the MoE (30B-A3B) only when you need its capacity ceiling and can pay the stability tax, sequence-level importance sampling ([Chapter 52](#)), the routing fixes ([Chapter 54](#)), and the FP8 wall ([Chapter 55](#)). The tax is the next five chapters.

The brochure's "30B for the compute of 3B" is true about compute and silent about memory and stability. Budget for what it omits before choosing the MoE.

TAKEAWAY

A 30B-A3B MoE gives more capacity for the same per-token compute (30B total, ~3B active) but costs more memory, trains less stably, and breaks the standard RL update. A 14B dense model wins on memory, stability, and tool compatibility. Default to dense for agentic RL on 8×H100. Choose MoE only when you need its ceiling and can pay the stability tax detailed in the rest of this Part.

Next: the first installment of that tax, the specific reason a mixture-of-experts breaks the GRPO update.

CHAPTER 51

Why MoE Makes Token-Level GRPO Fragile




Figure 51.1. A tiny weight change flips which expert a token routes to.

You take your working GRPO recipe, point it at the mixture-of-experts (MoE) model, and it collapses, our first vanilla attempt (run `c9_grpo`) died at step 10 with the reward stuck near -0.95 . Nothing in the reward or the data changed. The model’s own architecture broke the math, specifically the importance-sampling ratio at the center of every update.

51.1 The ratio assumes a stable computation

Plain, precise, payoff:

- Plain. The new-over-old probability ratio (Chapter 4) only makes sense if “the model” computing the old probability and the new one is, for that token, the same model in the way that counts.
- Precise. Importance sampling (IS) reweights each token by $\text{prob_new} / \text{prob_old}$, and that assumes the two probabilities

measure the same thing. On an MoE, a token’s probability is computed through whichever experts the router sent it to. If the routing changes between the old policy and the new one, the probability is now computed through a different sub-network, and the ratio compares two different computations.

- Payoff. So on an MoE, the token-level ratio is silently invalid for every token whose routing flipped, and nothing in the loss will tell you.

51.2 The 10% expert flip

The size of the problem is the alarming part. Across a single update step on the 30B-A3B, roughly 10% of tokens route to different experts than they did at generation time, a large fraction of every batch carrying a corrupt ratio. But measure it the right way: a top-1 flip rate is only a rough proxy, because Qwen3-30B-A3B routes each token to top-k experts, not one. The MoE-faithful diagnostic is the top-k overlap (Jaccard) between rollout-time and update-time routing:

```
# top-1 flip is a rough proxy. Top-k Jaccard is the MoE-faithful
↪ measure
flip_top1 = (gen_top_expert != cur_top_expert).float().mean() #
↪ ~10% - rough proxy only

old = set(gen_topk_experts[token])
new = set(cur_topk_experts[token])
jaccard = len(old & new) / len(old | new) # routing overlap, per
↪ token
flip_rate = 1.0 - jaccard # better corruption proxy
```

What to actually watch, and what each number means:

Metric	Value	Meaning
top-1 flip rate	~10%	rough corruption proxy
top-k Jaccard drop	—	MoE-faithful corruption measure
router-entropy trend	—	collapse vs healthy specialization
VAL trend	—	whether the instability actually <i>matters</i>

Table 51.1. The router-corruption metrics on a MoE, each with its meaning. Top-1 flip rate is a rough proxy; the top-k Jaccard drop is the MoE-faithful measure; router-entropy and VAL trends say whether the instability actually matters.

The consequence chains: corrupt ratios produce corrupt gradients, corrupt gradients destabilize the routing itself, and destabilized routing means the experts stop specializing, routing collapse, where the model degrades into uselessness.

51.3 Why this is MoE-specific, and how unconditional the claim is

A dense model computes every token through the same weights, so the ratio is always comparing like with like, which is exactly why the dense 14B model of [Chapter 50](#) “just works” and the MoE doesn’t. The flip problem is unique to architectures with input-dependent routing.

One honesty note before the fix, because this book keeps its contradictions in the open: the collapse above is real and reproduced, and a later short vanilla-GRPO run on the same model stayed healthy for its full horizon. “MoE breaks GRPO” is recipe-, reward-, and horizon-dependent, not a law of nature, and a chapter ahead lays the two runs side by side and digs out the boundary. The operational guidance is unchanged: at this scale, don’t bet a long run on token-level ratios.

That specificity points at the fix. Because the break is in how the ratio is computed, no reward change or data change can repair it, the repair has to change the ratio itself, which is the next chapter.

TAKEAWAY

GRPO’s update reweights each token by new-over-old probability, which assumes both are computed the same way. On a mixture-of-experts, roughly 10% of tokens flip to different experts each step, so their probabilities come from different sub-networks and their ratios are meaningless, corrupt ratios cause corrupt gradients cause routing collapse (our vanilla run died

at step 10, reward -0.95). It's specific to input-dependent routing, the claim has a horizon-dependent boundary a later chapter maps, and the fix must change how the ratio is computed, not the reward or data.

Next: the fix, compute the ratio over the whole sequence instead of per token, and the scale trap that silently kills runs that forget the clip lives on the ratio's scale.

CHAPTER 52

GSPO: Sequence-Level Importance Sampling

One flipped token: token-level multiplies it, sequence-level averages it

per-token ratios across one sequence (new prob $\div$ old prob, token by token):



token-level (GRPO): multiply them

$$1.0 \times 1.0 \times 1.0 \times 18.3 \times 1.0 \times 1.0 \longrightarrow = 18.3 \text{ (exploded)}$$

the one bad token dominates the whole update

sequence-level (GSPO): average them

$$\text{mean}(1.0, 1.0, 1.0, 18.3, 1.0, 1.0) \longrightarrow \approx 1.0 \text{ (stable)}$$

one flip washes out

Figure 52.1. Importance ratio per token vs per sequence: the per-token product explodes while the per-sequence average stays near 1.

The fix for the mixture-of-experts (MoE) flip problem is almost insultingly simple to state and easy to get wrong in a way that silently kills your run. Group Sequence Policy Optimization (GSPO) computes the importance-sampling ratio at the sequence level instead of the token level, and the clip band must be set on the new ratio’s scale, or your trust region is either a fiction or a tourniquet.

52.1 The fix: aggregate the ratio

Plain, then precise:

- Plain. Instead of asking “did this token get more likely,” ask “did this whole answer get more likely”, a question that doesn’t care if one token flipped experts.

- **Precise.** GSPO forms one ratio per sequence from the sequence-level likelihood, normalized by length, in effect the geometric mean of the per-token ratios, so it stays on a per-token scale rather than blowing up as a raw product. The per-token expert flips that corrupted individual ratios ([Chapter 51](#)) wash out in the aggregate. The fix matches the level of the problem: the break was per-token, so move the comparison above the per-token level.

52.2 The clip-scale trap

This is the part that bites, and it bites in both directions, because the sequence ratio is a different kind of number:

1. The length-normalized ratio concentrates tightly around 1. A geometric mean over hundreds of tokens averages away per-token wobble, so GSPO's ratio barely strays from 1, which is why the published GSPO clip range is on the order of 10^{-4} , roughly a thousand times tighter than the token-level ± 0.2 .
2. Carry the old ± 0.2 band over, and the clip never fires. Every sequence ratio sits comfortably inside a band built for token ratios, your trust region is a fiction, and the run can lurch with nothing to restrain it.
3. Implement the raw product instead of the normalized ratio, and the opposite happens. An unnormalized sequence ratio is a product over hundreds of tokens, ranging wildly. A token-scale band clips essentially everything, the gradient vanishes, and the run looks dead when it's a units error.

The diagnostic for both failure modes is the same one line:

```
# fraction of sequences actually clipped - ~0% and ~100% BOTH
↪ mean the band
# is on the wrong scale for your ratio definition. Healthy runs
↪ sit in between
clipped_frac = ((ratio < 1 - eps_lo) | (ratio > 1 +
↪ eps_hi)).float().mean()
```

The rule: the clip lives on the ratio’s scale. When you change what the ratio is, the band moves with it, set it from the GSPO recipe, not from Chapter 3’s habit.

The config that encodes all of this, so it’s reproducible rather than inferred from prose:

```
algorithm:
  adv_estimator: gspo
actor:
  policy_loss:
    importance_sampling_level: sequence # the fix: ratio at the
      ↳ sequence level
    sequence_ratio_normalization: length_normalized #
      ↳ geometric-mean scale, not raw product
    clip_ratio_low: 0.0003 # ~10^-4 scale - NOT the token-level 0.2
    clip_ratio_high: 0.0003
```

And the three ways to get it wrong, each with its tell:

Mistake	Symptom	Diagnosis
token-level ± 0.2 clip copied into GSPO	clipped fraction $\approx 0\%$, trust region is fake	band on the wrong (token) scale
raw product sequence ratio (no length norm)	clipped fraction $\approx 100\%$, gradient vanishes	forgot length normalization
sequence IS + incompatible <code>loss_type</code>	shape/broadcasting bug or length bias	framework loss-shape mismatch (e.g. TRL <code>loss_type="grpo"</code> with sequence IS)

Table 52.1. The three ways to get the GSPO clip wrong: for each mistake, the symptom it produces and the diagnosis. Two are scale errors (band on the wrong level) and one is a framework loss-shape mismatch.

52.3 GSPO is the MoE default

The verdict, previewing the next chapter and the bake-off to come:

- GSPO is the recommended fix for MoE RL. It's what made the breakout run (journal ID C9) work, and across the stabilizer comparison it both prevented routing collapse and trained well, which the alternatives (Chapter 54) don't all do.
- The lesson generalizes. When an architecture breaks an assumption, here, that token probabilities are comparable across an update, fix the assumption at the right level, and then watch the scale of whatever you changed.

TAKEAWAY

GSPO repairs the MoE break by forming one length-normalized importance ratio per sequence, effectively the geometric mean of the token ratios, so individual expert flips average out and the update stays valid. The catch: that ratio concentrates tightly around 1, so the clip band must shrink to GSPO's $\sim 10^{-4}$ scale. Carry the token-level ± 0.2 over and the clip never fires (no trust region), implement a raw product instead and everything clips (dead run). Log the clipped fraction, $\sim 0\%$ and $\sim 100\%$ both mean wrong scale. GSPO is the MoE default.

Next: that fix in action, the run that sat at 1.2% looking dead, then broke out to nearly 30%, with entropy proving it before accuracy did.

CHAPTER 53

C9-GSPO Breakout: Sequence Ratios Unlock MoE Learning




Figure 53.1. The breakout was real: accuracy rose while entropy held inside a clearly-bounded band.

For nineteen steps it looked dead: 1.2% accuracy, flat, the kind of curve you write the kill order on. We nearly did. Then it broke out to nearly 30%, and the graph that told us it was working wasn't the accuracy. It was the entropy, recovering in the exact pattern [Chapter 43](#) taught us to read. This is the breakout run (journal ID C9) in full, the run we previewed back in [Chapter 18](#), now told with the mixture-of-experts machinery in place to explain it.

53.1 The flat start you must not kill

The setup: the 30B-A3B mixture-of-experts (MoE) model, the math dataset (math_dapo), trained with Group Sequence Policy Optimization (GSPO). Early steps sat near 1.2% validation accuracy, almost flat, and that flatness is a trap in two directions:

- It looks like the dead Phase-0 run ([Chapter 6](#)), all-wrong groups, no gradient, a task the model can't touch.
- But on a big MoE with GSPO, a flat early curve is usually reorganization before a breakout ([Chapter 18's](#) lesson, now in full): the model is restructuring its behavior before the accuracy reflects it.

The way to tell the two apart was not the accuracy. It was the entropy and the within-group variance, both behaving like a model that was learning, not one that was stuck.

53.2 The breakout, with entropy as the proof

The jump came fast: 1.2% to 29.7% validation accuracy by step 20, 4× by step 10, 25× by step 20, while the training reward climbed from −0.96 to −0.25. And the reason we trusted it before it happened was the entropy curve.

The ladder, on both instruments:

Step	VAL acc	Train reward	Entropy	Read
0	~1.2%	−0.96	—	
10	~6× base	≈ −0.5	dipping	reorganizing, not stuck
20	29.7%	−0.25	recovered	breakout
30	not measured	not measured	not measured	consolidation

Table 53.1. The C9 breakout ladder: validation accuracy, train reward, and entropy by step. What looks like the dead Phase-0 run at step 0 is actually reorganizing by step 10 and breaks out by step 20.

As the model reorganized, entropy dipped and then recovered, the damped-recovery signature from [Chapter 43](#), not a one-way slide to collapse. That bounce was the evidence that exploration was alive while the model worked its way to the breakout. (This chapter is the first breakout. The fuller GSPO validation ladder out to step 60+ is in the bake-off chapters of [Part V](#).)

This is where the two halves of the book's stability story meet on one run. GSPO fixed the importance-sampling break (Chapters 51–52)

so the MoE could train at all. the entropy telemetry told us, fifteen-plus steps before the accuracy did, that the training was working, the leading-signal discipline of [Chapter 46](#), applied in the happy direction. Reading the run correctly while the accuracy was still flat is what kept the kill order in the drawer.

53.3 What it proves

The composite lesson, and the practical scar:

- The fix and the instrument compose. Sequence-level IS made the run trainable. The entropy and variance telemetry made it readable. The breakout needed the first. Not killing it needed the second.
- Trust entropy and variance over accuracy in the first stretch of a big-model run. The accuracy was the last thing to move (the lag again, [Chapter 46](#)), believe the leading signals.

TAKEAWAY

The breakout run (journal ID C9), a 30B-A3B MoE on math_dapo, trained with GSPO, sat at 1.2% accuracy looking dead, then broke out to 29.7% by step 20 while training reward climbed $-0.96 \rightarrow -0.25$. The proof it was working arrived early, on the entropy: a dip-and-recover, not a slide, [Chapter 43](#)'s signature, [Chapter 46](#)'s leading-signal discipline in the happy direction. GSPO made the run trainable. The telemetry made it readable. Trust entropy and variance over accuracy early in a big-model run.

Next: GSPO is the headline fix, but routing collapse has more than one cause, the fuller toolbox, including a fix hiding in the seam between two frameworks.

CHAPTER 54

MoE Routing Collapse: Causes and Fixes

Group Sequence Policy Optimization (GSPO) is the headline fix, but routing collapse, when a mixture-of-experts (MoE) model's experts stop specializing and the model degrades, has more than one cause, and on a bad day you'll need to know which one you're fighting. Some causes aren't in your algorithm at all. One hides in the seam between two frameworks that route the same token differently. Know the causes and you can pick the fix instead of throwing all of them at the wall.

54.1 The root causes

Routing collapse comes from a few distinct places:

1. The importance-sampling flip problem ([Chapter 51](#)), the dominant cause, and the one GSPO addresses.
2. Framework-split inconsistency, the trainer and the inference engine route the same token to different experts, because their implementations of the router disagree (in our stack: vLLM and FSDP disagreeing). Generation and update then train on different computations. This is a seam bug, in the family of [Chapter 24](#).
3. Router over-commitment, the router collapses onto a few favorite experts (load imbalance), starving the rest and shrinking the model's effective capacity.

54.2 The toolbox

The fixes, what each does, and when to reach for it:

Fix	What it does	When
GSPO	sequence-level IS	the default — fixes the dominant cause (Ch 52)
FREEZE	freeze the router after warm-up (one env flag in our patch: detach the router logits)	routing won't stabilize (at a capability cost)
JITTER	add noise to routing	experts over-commit / collapse to a few
Aux-loss	load-balancing penalty	keep experts used evenly
Fisher-Rao probe	early-detection metric on routing drift	catch drift <i>before</i> collapse

Table 54.1. The routing-collapse fixes: what each does and when to reach for it. GSPO is the default that fixes the dominant cause; FREEZE, JITTER, the aux-loss, and the Fisher–Rao probe are situational tools layered on top.

Two of these deserve a sentence each. FREEZE is one boolean in our stack, `router_logits.detach` behind an env flag, and flipping that single flag produced two completely divergent run outcomes, which is about the cleanest causal claim RL forensics ever hands you. And the Fisher-Rao probe is worth logging from the start, a cheap routing-drift gauge in the spirit of [Chapter 46](#)’s leading signal:

```
# routing drift as an early-warning gauge - rises before collapse
drift = fisher_rao(router_logits_now, router_logits_ref) #
↳ distance from a healthy reference
if drift > threshold:
    warn("router drifting - collapse risk. Consider jitter /
    ↳ aux-loss")
```

JITTER is the other one worth a code sketch, training-only noise before the router, so experts can’t over-commit to a frozen partition:

```
# JITTER: training-only noise injected before the router gate
if self.training and jitter_std > 0:
```

```
hidden_states = hidden_states +
    ↪ torch.randn_like(hidden_states) * jitter_std
router_logits = self.gate(hidden_states)
```

Before reaching for any of these, know which are validated versus planned versus merely diagnostic, the toolbox isn't uniformly proven, and saying so matters:

Fix	Type	Validated here?	Evidence	Caveat
GSPO	loss-side IS	yes	C9 breakout (Ch 53)	recommended default
FREEZE	router-gradient block	yes	stable run; lower ceiling (Part VII)	capability cost
JITTER	input noise pre-router	yes	viable; stronger in some later trajectories	not equivalent to GSPO; may need VAL-based rules
Aux-loss	load-balance penalty	planned/-conditional	not run	don't present as proven unless run
Fisher-Rao	drift <i>diagnostic</i>	partially scripted	not run	it's a probe, not a fix

Table 54.2. Each routing fix's validation status, evidence, and caveat. GSPO is the validated default; FREEZE and JITTER are validated with caveats; the aux-loss and Fisher–Rao probe are not yet proven here.

54.3 Picking the fix

A short decision path:

1. Start with GSPO. It fixes the dominant cause.
2. Add JITTER or an aux-loss if experts over-commit.
3. Use FREEZE only if routing won't stabilize otherwise, knowing it caps capability (the residual-direction story of Part VII explains why).
4. Log the Fisher-Rao / router-entropy probe the whole time, to catch drift early.

One honest note carried forward: FREEZE, JITTER, and GSPO all prevent collapse, but they are not equivalent. They leave the model with different capability ceilings, which is a finding the internals chapters ahead unpack. Preventing collapse and preserving capability are two different bars.

TAKEAWAY

Routing collapse has several causes, the IS flip (GSPO's target), framework-split inconsistency between trainer and inference engine (a seam bug. VLLM↔FSDP in our stack), and router over-commitment. The toolbox: GSPO by default, JITTER or an aux-loss for over-commitment, FREEZE as a last resort (a one-flag router-logit detach that caps capability), and a Fisher-Rao drift probe logged throughout for early warning. These fixes prevent collapse but aren't equivalent in the capability they preserve.

Next: the one combination that no fix in this toolbox could save, low-precision rollouts on a mixture-of-experts, proven dead across four attempts.

CHAPTER 55

The FP8 + MoE Implementation Wall

FP8 + MoE rollout: a wall, not a law

Each check must pass, or the quantized rollout corrupts training.

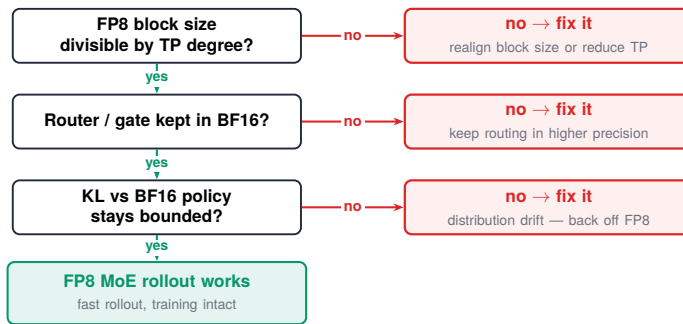


Figure 55.1. FP8 + MoE rollout is a wall, not a law — a decision flowchart for when it bites.

Some things just don't work, and the useful thing a practitioner book can do is tell you why, precisely, so you don't spend a week proving it yourself. We spent four attempts on this one, and the honest conclusion is narrower and more useful than the one we first wrote down. We initially called it a router-noise collapse. It wasn't, or at least, we never got far enough to prove that. What we actually hit was an implementation wall: under `veRL` + `vLLM` on `8×H100-80GB`, the `30B-A3B` MoE could not be trained with FP8 rollouts because every valid tensor-parallel setting for FP8 block quantization was incompatible with the model's expert dimensions, memory budget, or the actor→rollout weight-sync path. The runs didn't train far enough to collapse. They failed to start.

55.1 The four attempts, and the wall each one hit

The setup was the 30B-A3B MoE with FP8 rollouts (chosen for FP8’s speed and memory savings) and GSPO plus the routing fixes of [Chapter 54](#). Here is what actually stopped each attempt, none of it is a training dynamic:

Attempt	What failed	Root cause
TP=8 FP8 rollout	<code>ValueError</code> at load	expert intermediate dim $768/8 = 96$; FP8 block-quant needs divisibility by <code>block_n = 128</code>
TP=4 FP8 rollout	same invalid-dimension error	$768/4 = 192$, also not divisible by 128
TP=2 FP8 rollout	CUDA OOM during weight sync	TP=2 fits the block math, but $\sim 15\text{GB/GPU}$ model + BF16 sync buffer + KV cache exceeds 80GB
AWQ / INT8 roll-out	unsupported	veRL syncs a BF16 actor to the rollout engine; can’t hold separate AWQ/INT8 weights without a separate inference server

Table 55.1. The four FP8-rollout attempts and the wall each one hit: for each configuration, what failed and the root cause. Three are blocked by the expert-dimension block-quant math; the fourth by veRL’s single-actor weight sync.

The two errors you’d actually see, so they’re searchable:

```
ValueError: gate/up output_size = 96 is not divisible by block_n = 128
```

```
CUDA out of memory during rollout wake_up / weight sync
```

So this is not “four runs collapsed after training.” It is four attempts that hit hard block-size, tensor-parallel, memory, and sync-architecture walls before a clean training comparison could even exist. The distinction matters because it changes the advice: the problem isn’t that FP8 destabilizes MoE RL. It’s that, on this stack and this model, you can’t get an FP8 MoE rollout to run colocated with the trainer at all.

Speculation, clearly labeled. If FP8 MoE training became feasible, router logits might well be sensitive to 8-bit quantization noise in a way that compounds the importance-sampling problem of [Chapter 51](#). That’s a plausible mechanism, but we did not prove it, because the attempts hit the block-size and memory walls first. Treat router-noise collapse as a hypothesis, not a result.

55.2 Why a negative result earns a chapter

The value is exactly its honesty, and it draws a useful boundary:

- A documented wall saves the next person four attempts. “We tried FP8 + MoE four ways and couldn’t get a colocated rollout to run” is a result. Reporting it plainly, and as a systems wall, not an algorithm failure, is the same ethic as the honest losses of [Chapter 36](#): the failures are information, and mislabeling which kind of failure it is would send the next reader chasing a training-stability ghost.
- The boundary is specific. FP8 rollouts are fine on dense models, better than fine, in one way that surprised us: on dense runs, FP8 rollouts produced healthier entropy than BF16, a counterintuitive finding the inference Part unpacks with the safe recipes. The wall is FP8 and MoE together. For an MoE, use a higher-precision rollout instead.

55.3 The rule

The recipe, stated as a rule you can follow without re-deriving it:

- On a mixture-of-experts under veRL + vLLM, don’t try to colocate an FP8 rollout with the trainer. The valid FP8 block-quantization tensor-parallel settings don’t fit this model’s expert dimensions or the memory budget, and the BF16 actor→rollout sync can’t carry AWQ/INT8 weights. Use a BF16 rollout instead. If you need FP8, run it inference-only on a separate server, not inside the training loop.

This closes the core MoE story: the dense-versus-MoE tradeoff (Chapter 50), why GRPO breaks (Chapter 51), the GSPO fix (Chapter 52), the breakout it enabled (Chapter 53), the fuller routing-collapse toolbox (Chapter 54), and now the wall (Chapter 55).

TAKEAWAY

Training a 30B-A3B MoE with FP8 rollouts under veRL + vLLM on 8×H100 failed across four attempts, but as an *implementation wall*, not a proven collapse: TP=8 and TP=4 give expert dims (96, 192) not divisible by the FP8 block size 128. TP=2 fits the math but OOMs on the BF16 sync buffer. AWQ/INT8 isn't supported in the actor→rollout sync path. The runs never trained far enough to collapse. (Router-noise sensitivity is a *plausible mechanism if FP8 ran*, not a result.) FP8 is fine on dense models. There it even gave healthier entropy than BF16. On an MoE, use a BF16 rollout, or keep FP8 inference-only on a separate server.

Next: with the MoE machinery settled, the backend question it forces, Megatron versus FSDP2, and then what these fixes actually do to the model's internals.

CHAPTER 56

Megatron vs FSDP2 for MoE

You picked the training framework once, early, for your dense models, and never thought about it again. Then you move to a large mixture-of-experts model and discover the choice was never yours to keep, the model size forces the backend, and the framework you were comfortable with may simply not fit. This chapter is about why, for a big mixture-of-experts (MoE), the backend is effectively decided for you.

56.1 The two backends

The two training backends in play, in plain terms:

- FSDP2 (Fully Sharded Data Parallel, version 2) shards a model's parameters, gradients, and optimizer state across GPUs. It is simple to use, well-integrated, and the comfortable default for dense models up to a point.
- Megatron adds model-parallel strategies, tensor parallelism and, critically for MoE, expert parallelism, which places different experts on different GPUs. It is more complex to configure and operate.

56.2 Why the MoE forces Megatron

A large MoE breaks the comfortable default for a concrete reason:

1. The experts don't fit the simple sharding. An MoE's parameters are dominated by many experts, and the efficient way to train it is to spread experts across GPUs (expert parallelism) and route tokens to wherever their expert lives. FSDP2's parameter sharding wasn't built around that routing pattern.

2. So past a certain MoE size, the backend is forced, for efficient, sustained training. The 30B-A3B model, with its full expert set resident, pushes you toward Megatron’s expert parallelism as the thing that fits in memory and runs at a usable speed. This doesn’t mean FSDP2 can never run a MoE experiment, short experimental paths exist, it means FSDP2 stops being practical for sustained large-MoE training, and the architecture, not the engineer, makes that call.
3. And the backend mismatch is a collapse risk that is semantic, not just memory (Chapter 54): if the trainer and the inference engine implement routing differently, the policy that generated the rollouts and the policy being updated are subtly different models. That’s not a throughput problem. It’s the on-policy assumption (Chapter 4) quietly breaking.

A rough decision matrix:

Model class	Practical backend
dense 14B	FSDP2 — comfortable
30B-A3B (experimental)	FSDP2 possible for short runs; Megatron for sustained
30B-A3B / larger MoE (sustained)	Megatron — expert parallelism
hybrid-MoE variants (e.g. Qwen3.5/Next-style)	Megatron — match the router/parallel layout

Table 56.1. A rough backend decision matrix: for each model class, the practical training backend. Dense fits FSDP2; sustained MoE wants Megatron’s expert parallelism; hybrid-MoE variants must match the router/parallel layout.

56.3 The trainer/rollout consistency checklist

Whatever backend you pick, the rule that prevents the seam bug (Chapter 54) is that the generating policy and the updated policy must be the same model. Before a MoE run, check that trainer and rollout engine agree on all of:

- same router implementation (the routing math is identical);

- same expert-parallel layout (experts placed the same way);
- same tokenization and chat template (the serialized input is byte-identical);
- same precision mode (Chapter 55's FP8 caution applies);
- same vLLM/Megatron assumptions about how the MoE is sharded and served.

A mismatch in any one of these makes generation and update disagree, the routing-mismatch collapse, dressed up as a mystery.

The practical upshot: don't assume your dense-model backend carries over to MoE, and don't assume two frameworks route identically. Check both the memory fit and the trainer/rollout consistency before you plan a run.

TAKEAWAY

FSDP2's parameter sharding is the comfortable default for dense models, but a large mixture-of-experts is dominated by experts that want *expert parallelism*, placing experts across GPUs, which pushes you to Megatron. Past a certain MoE size the backend is forced by the architecture, not chosen, and it must stay consistent with the inference engine to avoid routing-mismatch collapse. Check the fit before planning the run.

Next: three different ways to stop routing collapse, and the catch that they all work but are not the same.

CHAPTER 57

Three Ways to Stop Routing Collapse

You now have more than one way to stop a mixture-of-experts (MoE) from collapsing, and the temptation is to treat them as interchangeable, pick whichever is easiest and move on. That instinct is wrong in a way the next three chapters make precise. All three fixes prevent collapse. None of them leave you with the same model. “Doesn’t collapse” and “is as capable as it could be” are different bars, and this chapter sets up the gap.

57.1 The three fixes

Three mechanisms, acting at different places in the system:

1. GSPO (Group Sequence Policy Optimization), a loss-side fix. It computes the importance ratio at the sequence level ([Chapter 52](#)), so per-token expert flips stop corrupting the update.
2. FREEZE, an activation-side fix. After a warm-up, it freezes the router, so routing can no longer drift during training.
3. JITTER, another activation-side fix. It adds noise to the routing, keeping the router from over-committing to a few experts.

Each attacks the collapse from its own angle: GSPO fixes the math of the update, FREEZE removes routing motion entirely, JITTER keeps routing healthy by perturbing it.

57.2 They are not equivalent

Here is the catch, and the reason this isn’t just a menu:

Fix	Acts on	Stops collapse?	But...
GSPO	the loss (sequence IS)	yes	the recommended default (Ch 62)
FREEZE	the router (frozen)	yes	<i>retracts</i> representations — lower ceiling (Ch 59)
JITTER	the router (noised)	yes	<i>drifts outward</i> — higher ceiling (Ch 59)

Table 57.1. Three collapse fixes: all stop collapse, none leave the same model. GSPO acts on the loss; FREEZE retracts representations to a lower ceiling; JITTER drifts outward to a higher one — direction, not just stopping, is what differs.

All three columns say “yes” to stopping collapse. But the rightmost column is the point: they leave the model in different places. FREEZE and JITTER both prevent collapse, yet one ends up measurably less capable than the other, because, as the next chapters show, what they do to the model’s internal representations is not the same. Preventing a failure is not the same as preserving capability.

Two qualifications keep this honest. First, each fix targets a specific failure: GSPO the invalid token-level importance ratios under expert flips (Chapter 51), FREEZE router drift, JITTER routing over-commitment (brittle concentration onto a few experts). Second, the collapse they prevent is scale- and setting-dependent: the corpus shows 30B-A3B did not universally collapse under every GRPO configuration. The safe claim is “MoE makes GRPO more fragile and can collapse. Stabilizers change the failure surface”, not “vanilla GRPO always breaks MoE.” The stabilizers matter most exactly where the bare run is on the edge.

This sets up the finding that makes Part V worth its length: the difference between these fixes isn’t in whether they work, but in which direction they move the model, and that direction, not the amount of movement, is what sets the capability ceiling.

TAKEAWAY

Three fixes stop mixture-of-experts routing collapse: GSPO (loss-side, sequence-level IS), FREEZE (freeze the router), and JITTER (noise the router). All three prevent collapse, but they are not equivalent, because they leave the model's representations in different places and therefore at different capability ceilings. "Doesn't collapse" and "as capable as possible" are different bars. The next chapters measure the gap.

Next: to measure that gap, you have to watch the model's internals, the telemetry method that makes the representations visible.

CHAPTER 58

The Internals Telemetry Method

To claim that two fixes leave the model in “different places,” you have to actually look inside the model, and most people never do, because they don’t have the instruments. This chapter is the instruments: a small, reusable telemetry kit that makes a model’s internal state visible during training, at a cost low enough to leave running. You can watch the representations move, and it costs about 9% overhead.

58.1 The probes

Four measurements, each answering a question the loss can’t:

1. Forward-KL neutral probes. Run a fixed batch of neutral text through the model and measure how its output distribution has drifted from a reference (the base model, or an earlier checkpoint), using forward Kullback–Leibler divergence, a score of how far one distribution has moved from another. This tells you how much the model’s behavior has changed on held-fixed inputs, independent of the training task.
2. Residual drift. Measure how far the residual stream, the running vector the model passes layer to layer, has moved from base, as a mean shift ($\Delta\mu$) and a worst-case shift ($\Delta\max$). This is where in the network the change is happening.
3. Per-layer routing entropy. For the mixture-of-experts (MoE) layers, measure how spread-out the routing is at each layer, high entropy means tokens spread across experts, low means collapse onto a few. This catches routing collapse per layer, before it sinks the whole model.
4. Matched-step comparison. Always compare two runs at the same training step, never at “the end”, because runs of different

lengths or speeds aren't comparable otherwise (a trap [Chapter 62](#) makes painfully concrete).

Two definitional notes so the probes are reproducible. The neutral probe set is a small fixed batch of held-out, task-neutral prompts scored under a fixed decoding mode. Report it (count, source, length) so others can match it. Residual direction Δ is the signed shift of the residual stream relative to base, positive when representations extend outward past base, negative when they retract toward or behind it, not just a magnitude. And one caveat that matters: forward-KL is not an absolute capability proxy. A higher forward-KL doesn't mean "better" or "worse" in isolation. It means "more changed." It is informative only at matched steps and read together with direction and validation, never as a standalone score across runs.

58.2 The method, as code

The point of a method is that it's cheap enough to run every time. The forward-KL probe is the core, and it's a few lines:

```
# forward-KL drift on a fixed neutral batch - "how far has
↪ behavior moved from base?"
with torch.no_grad():
    p = base_model(neutral_batch).log_softmax(-1) # reference
    ↪ distribution
    q = model(neutral_batch).log_softmax(-1) # current policy
    kl = (p.exp() * (p - q)).sum(-1).mean() # forward KL(base ||
    ↪ current)
    log({"fwd_kl_neutral": kl.item()}) # rises as the model drifts
```

Logged every N steps alongside residual drift and per-layer routing entropy, this turns "the model changed somehow" into a set of curves you can read, and compare across runs at matched steps.

58.3 Why this is worth the 9%

The overhead is real but small, and it buys something nothing else does:

- It makes invisible failures visible. Routing collapse, representational drift, and behavioral change all happen inside the model, where the reward and loss can't see them. The telemetry surfaces them while you can still act.
- It's the only way to compare fixes honestly. The claim that FREEZE and JITTER differ (Chapter 57) is measurable only with these probes, which is exactly what the next two chapters do with them.

The probe protocols are summarized in Appendix K. The habit to carry is that a ~9% telemetry tax is cheap insurance against debugging a black box.

TAKEAWAY

A small telemetry kit makes a model's internals visible during training at ~9% overhead: forward-KL neutral probes (how far behavior drifted from base), residual drift $\Delta\mu/\Delta\max$ (where the change is), per-layer routing entropy (catching MoE collapse layer by layer), and matched-step comparison (the only fair way to compare runs). The forward-KL probe is a few lines. It surfaces failures the loss can't see and is the only honest way to compare fixes.

Next: point those instruments at FREEZE and JITTER, and the finding that follows, capability tracks which way* the layers move, not how far.*

CHAPTER 59

Direction, Not Distance: The Capability Ceiling




Figure 59.1. RL sets direction, not distance: a short, steeper move to a higher ceiling vs SFT’s long, shallow move.

Here is the finding that justifies the whole instrumented detour, and it is genuinely surprising: a model’s capability ceiling is set by which direction its representations move relative to the base model, not by how far they move. Two fixes that both stop collapse, FREEZE and JITTER, move the representations opposite ways, and the one that moves them outward ends up meaningfully more capable. Distance doesn’t predict capability. Direction does.

59.1 The measurement

Plain, then precise, then the payoff:

- Plain. When you train, the model’s internal representations shift away from where the base model had them. You’d assume “more shift = more change = risk.” That’s the wrong axis.

- **Precise.** Using the residual-drift probe ([Chapter 58](#)), measure the direction of the shift relative to base, call it Δ . FREEZE produces $\Delta < 0$: the representations retract inward, toward or behind the base. JITTER produces $\Delta > 0$: they drift outward, past the base. Then measure capability by validation accuracy (VAL): FREEZE lands at 0.555, JITTER at 0.650.
- **Payoff.** The fix whose representations drift outward is the more capable one, even though both “moved” the model. So when you compare stabilizers, don’t ask how much they perturbed the model. Ask which way.

59.2 Why direction is the right axis

The intuition behind the number is worth holding, because it generalizes:

- Outward drift is the model building on its base, not abandoning it. Representations moving outward past the base are extending the model’s existing structure, adding capability in the directions the base already pointed. JITTER, by keeping routing live, lets that outward extension happen.
- Inward retraction is the model pulling back into a smaller version of itself. FREEZE, by locking the router, stops the representations from extending and lets them settle inward, a more constrained model, hence the lower ceiling.

So “both fixes stop collapse” ([Chapter 57](#)) hides the real story: collapse-prevention is table stakes, and the capability difference comes from whether the fix lets the model keep growing outward or forces it to retract. This is the MoE-specific shadow of a principle [Part VII](#) states in full, that what RL does to a model is better described by the direction of representational change than its magnitude.

Evidence note, mechanistic probe result. Three things keep this claim honest. (1) Both FREEZE and JITTER prevented routing collapse, so the VAL gap is not one run collapsing while the other survived. It is two healthy runs landing at different ceilings. (2) The 0.555 / 0.650 figures are matched-step validation values, compared

the same way for both runs. (3) This is strong correlational mechanistic evidence, not yet a causal proof: showing capability tracks direction is not the same as an intervention that flips direction while holding everything else fixed (that ablation is on the agenda, [Chapter 80](#)). The direction-capability link is the best-supported reading of the data, stated as such.

TAKEAWAY

A model's capability ceiling tracks the *direction* its representations move relative to base, not the distance. Measured with the residual-drift probe: FREEZE retracts inward ($\Delta < 0$) and lands at VAL 0.555. JITTER drifts outward ($\Delta > 0$) and lands at 0.650. Outward drift is the model extending its base. Inward retraction is it pulling into a smaller self. When comparing stabilizers, ask which way they move the model, not how much.

Next: JITTER's path isn't a straight outward drift. It first retracts, then flips, tracing a U-curve into a new and higher attractor.

CHAPTER 60

The U-Curve and the New Attractor




Figure 60.1. The U-curve: too little and too much both hurt.

The last chapter made it sound like JITTER simply drifts outward and stays there. The real path is more instructive: it drifts outward through about step 100, while validation keeps climbing, and then flips negative at step 125, retracting into the new, higher attractor it has reached. Read carelessly, that late negative turn looks like the model finally falling back to base. Read correctly, it's the opposite: the model consolidating into the better configuration it just discovered.

Evidence status, mechanistic probe result (matched-step internals on the JITTER run). The residual-direction values below are from the internals journal.

60.1 The shape

What the telemetry shows over training (Figure 60.1), as residual direction $\Delta\mu$ relative to base:

Step	residual $\Delta\mu$	reading
s50	+0.026	outward — extending past base
s75	+0.029	still outward, drifting further
s100	+0.059	furthest outward; validation high
s125	−0.086	sign flip — retracting into the new attractor

Table 60.1. The U-curve: residual-mean shift $\Delta\mu$ by step, outward then back in. Representations extend past base through step 100 (validation high), then the sign flips at step 125 as the model retracts into a new attractor.

So $\Delta\mu$ is positive (outward) through s100, then turns negative at s125. The curve is not “down then up”. It is out, then back in, and the turning point is a deliberate consolidation, not a relapse.

60.2 Retraction into a new attractor

The interpretation is the useful part, and it reframes what the late negative turn means:

- The retraction at s125 isn’t decay. It’s consolidation. Having drifted outward into a higher-capability region (s50–s100), the model then pulls its representations inward toward the center of that new region, settling into the new, higher attractor rather than drifting indefinitely. It is not returning to the base’s basin. It is finishing the move into a better one.
- So a late inward turn is not a reason to kill the run. If you applied [Chapter 59](#)’s “outward = good” rule naively and saw $\Delta\mu$ go negative at s125, you might conclude JITTER was collapsing back toward base and stop it, exactly when it was consolidating its gain. The U-curve is a warning that representational direction must be read over the whole trajectory and against validation: at s125, $\Delta\mu$ went negative while VAL stayed high, which is consolidation, not collapse.

This is the internals-level version of the breakout run’s lesson ([Chapter 53](#)): a direction metric moving the “wrong” way can be the

back half of a reorganization, not the start of a decline. Judge the trajectory against held-out validation, not the direction sign alone.

TAKEAWAY

JITTER drifts *outward* from base through step 100 while validation climbs ($\Delta\mu \approx +0.026 \rightarrow +0.059$), then flips *negative* at step 125, and that flip is consolidation into the higher attractor it reached, not a relapse to base. A late inward turn paired with still-high validation is consolidation, not collapse. Read residual direction over the whole trajectory and against VAL, never the sign alone.

Next: if outward-drifting JITTER and loss-side GSPO each help, does stacking them help more? The COMBO hypothesis.

CHAPTER 61

COMBO: Stacking Stabilizers

You have a loss-side fix (GSPO) and an activation-side fix (JITTER), and they attack the collapse from different angles. The obvious next thought is to stack them, and the obvious thought turned out to have real, sustained evidence behind it. Stacking GSPO and JITTER (call it COMBO) was, on held-out validation, the strongest C9 variant we ran. This chapter is the hypothesis and the validation climb that supported it. The next chapter is about why a different metric, applied to the same field, names a different winner.

61.1 The hypothesis: complementary mechanisms stack

The reasoning for stacking is sound on its face:

1. The two fixes address different causes. Group Sequence Policy Optimization (GSPO) fixes the math of the update, the importance ratio corrupted by expert flips ([Chapter 51](#)). JITTER fixes the routing, keeping experts from over-committing ([Chapter 57](#)). One is loss-side, one is activation-side.
2. Non-overlapping fixes should compose. If they're solving different parts of the problem, applying both should help more than either alone. There's no obvious reason they'd interfere.
3. So COMBO = GSPO + JITTER was worth a run, on exactly the bet that loss-side and activation-side mechanisms add up.

61.2 The validation ladder, and a premature kill

The first measurement supported the bet, and so did every one after it. On a matched-step validation comparison ([Chapter 58's](#) discipline, compare runs at the same step) on `math_dapo`, COMBO didn't just lead early. It kept climbing:

Step	COMBO VAL	note
s40	0.508	early matched-step leader
s50	0.586	strongest signal so far
s60	0.645	cleared the original 0.61 gate
s70	0.653	small additional gain
s80	0.652	plateau — killed at s81

Table 61.1. COMBO’s validation ladder, climbing past the gate to plateau. VAL rises from 0.508 at s40, clears the original 0.61 gate by s60, and plateaus at s80 — where the run was killed at s81.

There is an important wrinkle in this run, and it’s a lesson in itself: the run was originally killed early (around s50) under an entropy-based criterion, then resumed, and on resuming, it climbed past the validation gate. The early number wasn’t a seductive head-fake that fizzled. It was the start of a real climb that an entropy-only kill rule nearly threw away. (That kill-rule lesson is the heart of the next chapter.)

Evidence status, canonical validation result (full COMBO run including the resume through s81). Note that an early re-mine table parsed a truncated COMBO log ending near step 49. The ladder above is the complete run.

So hold the right thought about COMBO going into the next chapter: by held-out validation, it was the best C9 variant, not a fourth-place footnote. The next chapter shows a different metric, parsed max training reward, naming GSPO instead, and the entire point is that both are true.

TAKEAWAY

COMBO stacks GSPO (loss-side) and JITTER (activation-side) on the sound bet that fixes for different causes compose, and on held-out `math_dapo` validation it *kept climbing*: 0.508 → 0.586 → 0.645 (clearing the 0.61 gate) → 0.653 → 0.652, plateauing at s81. It was originally killed early under an entropy-only rule, then resumed and climbed, so the early number was the start

of a real climb, not a head-fake. By validation, COMBO was the strongest C9 variant.

Next: the bake-off, where a different metric names a different winner, and the real lesson is that the metric chooses the winner.

CHAPTER 62

The Bake-Off: The Metric Chooses the Winner




Figure 62.1. C9 bake-off headline: the same 30B-A3B MoE reaches 4% under GRPO vs 29.7% under GSPO.

The last chapter showed COMBO climbing on validation. There is also a re-mined table where COMBO comes fourth and GSPO wins. Both are real, and the temptation is to declare one of them “the result.” Resist that, the honest finding is metric pluralism: the winner is undefined until you name the metric, and for C9 two reasonable metrics name two different winners. This chapter is that reconciliation, and the kill-rule lesson the resumed COMBO run taught.

62.1 Two tables, two winners

The two scorings of the same five stabilizers:




Figure 62.2. The bake-off over training: identical at the start, GSPO pulls away after step 10 to a 7× gap — because token-level importance sampling lets one flipped MoE route blow up the ratio.

- By max training reward (re-mined from parsed train logs): GSPO +0.617 (winner) · JITTER · FREEZE · COMBO · vanilla −0.951 (collapsed). Vanilla’s deeply negative reward is the quantitative form of “GRPO breaks on MoE” (Chapter 51).
- By held-out `math_dapo` validation per step (including COMBO’s resume): COMBO wins, its ladder climbs 0.508 → 0.645 → 0.653, clearing the 0.61 gate (Chapter 61), ahead of GSPO and the rest.

These don’t disagree about the data. They disagree about the question.

Reconciliation: validation winner vs training-reward winner

Question	Data included	Metric	Winner
Who climbs fastest on held-out validation?	full COMBO incl. resume through s81	VAL @ matched step	COMBO
Who tops the re-mined training-reward frontier?	parsed train logs (J4 re-mine)	max train reward	GSPO

Table 62.1. Two questions, two metrics, two winners on the same C9 field. Asked who climbs fastest on held-out validation, COMBO wins; asked who tops the re-mined training-reward frontier, GSPO wins — the metric decides the result.

Do not collapse these into one verdict. COMBO is the validation-speed winner. GSPO is the training-reward-frontier winner. The re-mine table that ranked COMBO fourth parsed a truncated COMBO log ending near step 49 ([Chapter 61](#)), so on “max reward over the budget it appeared to run,” COMBO simply had less race on record. That’s a parsing artifact meeting a metric choice, not COMBO being weak.

62.2 Which to recommend, and the kill-rule lesson

So what do you actually use? It depends on what the run optimizes for:

1. Want default stability with the fewest special-case kill rules? GSPO. It tops the training-reward frontier and needs no validation-watching to know it’s healthy.
2. Want the fastest validation lift and willing to run validation-based kill rules? COMBO is a live candidate, not a discarded fourth place. It climbed highest on held-out validation.
3. Research finding, independent of the recommendation: COMBO proves the [Chapter 61](#) hypothesis, stabilizers can stack constructively (loss-side + activation-side compose).

And the lesson that ties it to the rest of the book, the reason COMBO was nearly lost:

Once your stabilization hypothesis is confirmed, switch from entropy-based kill rules to validation-based ones. Entropy is a proxy. Held-out validation is the metric you actually care about. If entropy is below your floor but VAL is still climbing and the reward’s score variance is still positive, the policy is likely sharpening, not collapsing, and an entropy-only kill rule will murder a winning run. COMBO was killed early on entropy and only survived because it was resumed.

That kill-rule point is the spine of [Chapter 75](#), and it's why this chapter is titled for the metric, not the method.

TAKEAWAY

For C9, two reasonable metrics name two winners: by max training reward GSPO wins (+0.617) with vanilla collapsing (-0.951). By held-out validation per step COMBO wins (clearing the 0.61 gate). Neither is wrong, the winner is undefined until you name the metric. Recommend GSPO for default stability, COMBO when you'll run VAL-based kill rules. The deeper lesson: once a stabilization hypothesis is confirmed, switch from entropy-based to validation-based kill rules, entropy below floor while VAL climbs is sharpening, not collapse.

Next: [Part VI](#). MoE could learn fast once stabilized. SWE-bench taught the opposite lesson: stability can be solved while learning stays impossible.

Part VI

When RL Isn't Enough: SWE-bench & the Pivot

Every Part so far has ended in something that worked. This one doesn't, or not the way we hoped. Pointing RL at real software-engineering bugs produced the book's largest honest failure: runs that collapsed, then stabilized, then sat stable without learning, against a task a model of this size could not solve by outcome reward alone. This Part is that failure told straight, and the strategic exit from it, a pivot from sparse-reward RL to on-policy distillation. The dead ends here are the most instructive pages in the book, because they map the edge of what RL can do.

CHAPTER 63

SWE-bench Setup

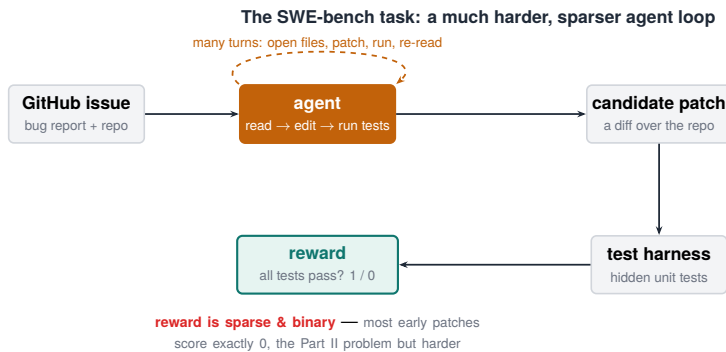


Figure 63.1. The SWE-bench task: a much harder, sparser agent loop with a binary tests-pass reward.

After search, we wanted a harder proving ground, and SWE-bench is about as hard as verifiable RL gets: fix a real bug in a real codebase, and pass the real tests. This chapter sets up the task and the harness, because the difficulty of the setup is the first half of the story of why it defeated us. The reward is honest and the environment is real, and both are far less forgiving than anything in the math or search Parts.

63.1 The task and the harness

The setup, piece by piece:

1. The task. SWE-bench gives the model a real software issue from a GitHub repository and asks it to produce a code change that resolves it. Success is verifiable in the cleanest possible way: the repository's test suite passes.
2. The model. A 14-billion-parameter model, trained with full fine-tuning (the whole model, not a small adapter), a choice that matters later, when we see what full fine-tuning costs (Part VII).

3. The environment. The agent works by issuing bash commands inside a Docker sandbox, editing files, running code, running tests, with `network = none`: the sandbox is sealed off from the internet, so the agent can't fetch answers or cause harm. It's isolated, reproducible, and unforgiving.
4. The loop. The agent issues a command, sees the output, issues another, a multi-turn trajectory (Chapter 21), until it submits, and the tests decide the reward.

The harness specifics that matter for what follows: the training set was ~280 SWE-bench instances. The SW1 reward was binary test-pass (1 if the suite passes, 0 otherwise, the cliff of Chapter 13, at scale). Each instance ran in its own Docker container with `network=none` and (optionally) memory limits. A final test verifier produced the reward at episode end. The observation the agent sees is the realistic set, stdout, stderr, file listings, test logs, the diff, and nothing from outside the sandbox.

One practical RL-agent lesson hides in that sandbox, and it's not infrastructure trivia: a free-acting bash agent will explore dangerous commands. Our logs caught attempts at things like `chattr -i` (making files immutable), the kind of move that can wedge the environment or game the harness. You need a security/command filter on what the agent is allowed to run, both to keep the sandbox sound and to stop the agent from “winning” by sabotaging the test setup. Sealing the network is necessary but not sufficient. The command surface needs a guard too.

63.2 Why this is so much harder than search

The setup looks like the search agent with a different tool, but the difficulty is a different order of magnitude:

- The action space is enormous. A search agent picks queries. A coding agent can run any bash command and make any edit anywhere in a repo. The space of trajectories is vast, and almost all of them fail.

- The reward is sparse and far away. A bug fix requires a long, correct sequence of edits before any test passes, so most roll-outs get reward zero ([Chapter 6](#)'s cold-start, at scale). There's no partial credit for "almost compiled."
- The successes are rare to begin with. RL amplifies the model's own successes ([Chapter 6](#)), and a 14B model rarely solves a full-repo bug by chance, so there's little to amplify. This is the seed of the failure the next chapters detail.

Everything that made the search agent trainable, partial-credit rewards, behaviors to shape, frequent-enough successes, is harder or absent here. That's the setup. The next chapter is what happened when we ran it.


TAKEAWAY

SWE-bench asks a 14B full-fine-tuned model to fix real GitHub bugs by issuing bash commands in a sealed Docker sandbox (network = none), with reward = the test suite passing, honest and verifiable, but brutal. The action space is enormous, the reward is sparse and far away (a long correct edit sequence before any test passes), and a 14B rarely solves full-repo bugs by chance, so RL has little to amplify. Everything that made search trainable is harder or absent here.

Next: the first run, and a collapse as fast and total as any in the book, entropy from 0.42 to 0.03 in thirty steps.

CHAPTER 64

The Catastrophic Collapse



The reward peak hid an unstable policy; the warning was in the strategy metric ~15 steps earlier (Fig 47–48).

Figure 64.1. A promising climb, then a hard collapse — the reward peak hid an unstable policy.

The first serious SWE-bench run (journal ID SW1) didn’t struggle. It imploded. Entropy fell from 0.42 to 0.03 in thirty steps, and the model’s outputs collapsed to about four tokens. This is the fastest, most total collapse in the book, and it’s worth sitting with, because it’s what entropy collapse ([Part IV](#)) looks like when it’s pushed to the limit by a task that offers almost no successes to reinforce.

The SW1 configuration, so the collapse is reproducible and attributable:

Setting	SW1 value
algorithm	vanilla GRPO
fine-tuning	full FSDP fine-tune (whole model)
dataset	~280 SWE-bench instances
reward	binary test-pass
Clip-Cov	none
KL penalty	none
entropy coefficient	none

Table 64.1. The SW1 SWE-bench configuration that collapsed catastrophically: vanilla GRPO, full fine-tune, binary test-pass reward, and no Clip-Cov, KL, or entropy bonus — the exact recipe the SW2–SW6 ladder spends the rest of the part repairing.

64.1 What collapse looks like at the extreme

The numbers describe a model destroying its own ability to act:

1. Entropy $0.42 \rightarrow 0.03$ in 30 steps. The policy went from reasonably exploratory to almost perfectly deterministic in half an hour, the collapse feedback loop of [Chapter 39](#), running at maximum speed.
2. Outputs collapsed to ~4 tokens. A model that should be writing multi-line code edits was emitting four tokens. It had concentrated all its probability onto a tiny, useless output and stopped generating anything real.
3. And it did this while the run kept running. Nothing crashed ([Chapter 40](#)'s lesson): the job was alive, the loss was moving, and the model was already dead.

This is advantage collapse and entropy collapse together ([Chapter 40](#)), driven to the floor by the task itself.

64.2 Why the task caused it

The collapse wasn't bad luck. It was the predictable result of pointing sparse-reward RL at a task with almost no reachable successes:

- Almost every rollout failed ([Chapter 63](#)): a 14B model rarely fixes a full-repo bug, so nearly every group was unanimous zeros, no gradient, except from the rare, noisy success.
- The model optimized the only thing it could reach. With correctness essentially unreachable, the reachable move was to stop taking risks, collapse onto a short, safe, deterministic output that at least didn't make things worse. The four-token output is the model retreating from a task it couldn't win.

- This is Phase 0 ([Chapter 6](#)) with a sharper edge. Phase 0 sat frozen at zero. Here the model actively degraded, because full fine-tuning on a near-zero-signal task let the collapse propagate through the whole model fast.

So SW1 isn't a bug to fix with a hyperparameter. It's the task telling you that outcome-reward RL has almost nothing to work with. The precise diagnosis from the config (the table above): binary reward + low solve rate + symmetric clipping = a collapse zone. With almost no passes, the binary reward gives unanimous-zero groups (zero advantage variance, the group reward-std metric confirms advantage collapse directly, [Chapter 5](#)). Symmetric clipping then offers no counter-pressure as entropy falls, so the policy slides to the deterministic floor. Each ingredient is individually defensible and jointly fatal, which is exactly why the next chapter changes them one at a time.

TAKEAWAY

The first SWE-bench run (journal ID SW1) collapsed catastrophically, entropy $0.42 \rightarrow 0.03$ in 30 steps, outputs shrinking to ~4 tokens, while the job kept running. The config explains it: vanilla GRPO + full fine-tune + binary reward + no Clip-Cov/KL/entropy term, on a task with almost no passes = a collapse zone (binary reward + low solve rate + symmetric clipping), confirmed by zero group reward-variance. It's the task signaling outcome RL has little to work with.

Next: the disciplined recovery, SW2 through SW6, one change per version, dragging the run back to stability.

CHAPTER 65

SW2–SW6: The Stabilization Saga

After the implosion, the job was to make the run survive, and the way we did it is a model of how to debug a collapsing RL run without fooling yourself: change exactly one thing per version, and watch. Across five versions, SW2 through SW6, each fixing a single thing, the run was dragged from catastrophe back to stability. The discipline is the lesson here. The sting, that stability still wasn't learning, is the next chapter.

65.1 One change per version

The method first, because it's the transferable part. When a run collapses, the temptation is to change five things at once and hope. That tells you nothing: if it works, you don't know which change did it. If it doesn't, you don't know which change hurt. So instead:

1. Form one hypothesis about what's driving the collapse.
2. Change one thing to test it, one knob, one mechanism.
3. Watch the telemetry ([Chapter 58](#)) to see whether that change moved the failure.
4. Keep it or revert it, then form the next hypothesis.

Five iterations of this, versions SW2 through SW6, one change each, is what the saga was. Each version isolated a single contribution to the collapse and addressed it, so that by the end we knew why the run was stable, not just that it was. The actual ladder, which is one of the most reusable debugging patterns in the book:

Version	The one change	Result	Lesson
SW1	vanilla GRPO, binary test-pass reward, no Clip-Cov, no KL	entropy 0.42 → 0.03 in ~30 steps; ~4-token outputs	sparse binary SWE reward + full FT collapses fast
SW2	add Clip-Cov + KL	damped oscillation; the run survives	Clip-Cov is domain-agnostic, not search-specific
SW3	stronger Clip-Cov / KL / LR	entropy stable but reward declining	exposed a GRPO length bias
SW4	<code>seq_mean_token_mean</code> loss + continuous shaping rewards	first stable reward signal; overfits early	loss aggregation and continuous rewards matter
SW5	scale to ~280–300 instances, constant LR	collapse again at step 18–22	constant LR + falling entropy → late collapse
SW6	cosine LR, stronger Clip-Cov, lower LR, KL 0.01, max turns 10	stable to step 146 — but not learning	stability is necessary, not sufficient

Table 65.1. The SW1–SW6 stabilization ladder — six successive SWE-bench training runs, each making one change from the last. For each: the change, the result, and the lesson. Each version fixes the previous failure and exposes the next, until SW6 is stable — but stability turns out to be necessary, not sufficient, for learning.

Evidence status, historical / debugging result (the SWE-bench V-series journals). The per-version changes and outcomes are from the run logs.

Read top to bottom, each row buys one thing and reveals the next problem, which is exactly what changing one variable at a time is for.

65.2 What the discipline bought, and what it didn't

The progression of single changes did its job along one axis and exposed a wall along another:

Axis	Verdict	Detail
Stability	recovered	by SW6 the run no longer collapsed (entropy held, outputs stayed full-length)
Attribution	earned	each fix was tied to one named cause, so the stability was understood
Learning	not achieved	stable, but the model still wasn't solving meaningfully more bugs

Table 65.2. What the one-change discipline bought along each axis. Stability was recovered and attribution was earned — each fix tied to one named cause — but learning was not achieved: the model still wasn't solving meaningfully more bugs.

That third row is the hard truth this Part is building toward. The one-change-per-version discipline successfully stabilized the run, a real and necessary achievement, but stability is not the same as progress. We had a run that no longer imploded and still wasn't learning the task, because (as the next chapters show) the problem was never instability alone. It was a sparse reward on a task too hard for the model to crack by outcome alone.

So the saga ends in a genuinely useful place and a genuinely frustrating one: we knew exactly how to keep the run alive, and keeping it alive wasn't enough. That gap, stable but not learning, is what forces the pivot. The root causes of SW6's non-learning, named so the next chapters can dismantle them one at a time:

Factor	Impact	Evidence
reward too sparse	critical	only ~5% of rollouts pass any test
shaping misaligned	high	length / commands-run / output / diff reward activity, not repair
model capacity	high	full-repo reasoning too hard at 14B (Ch 67)
only ~280 instances	medium	low diversity, repeated exposure
per-episode credit assignment	medium	can't tell which of ~10 turns mattered
no intermediate verification	medium	no signal until the episode ends

Table 65.3. Why the SW6 SWE-bench run stayed stable yet still didn't learn: the ranked root causes, each with its impact and the evidence behind it. The reward was too sparse to begin with, and several factors compounded it.

TAKEAWAY

SW2 through SW6 dragged the SWE-bench run from collapse to stability by changing exactly one thing per version and watching the telemetry, SW1 collapse → SW2 Clip-Cov+KL survives → SW3 length bias → SW4 first stable reward → SW5 late re-collapse → SW6 stable to step 146. The discipline earns *attribution* as well as stability. But it exposed the wall: stable and still not learning, for six nameable reasons led by a too-sparse reward and misaligned shaping. Stability isn't progress, which is what forces the pivot.

Next: the wall named directly, why a stabilized run still couldn't learn, and why shaping rewards only taught it to look busy.

CHAPTER 66

Stable but Not Learning




Figure 66.1. Stable is not the same as learning: a flat-but-calm run vs one that actually moves.

We won the battle from the last chapter, the run was stable, and discovered it didn't matter. The accuracy still wouldn't move. We had a SWE-bench run that no longer collapsed, held its entropy, produced full-length outputs, and solved essentially no more bugs than when it started. This chapter names that wall, because it's the realization that forces everything after it: the problem was never only instability.

66.1 The sparse-reward wall

Plain, then precise:

- **Plain.** The model almost never succeeds at the task, so it almost never gets a reward, so there's almost nothing for RL to reinforce, no matter how stable you make the run.
- **Precise.** On SWE-bench, a 14-billion-parameter model resolves a real repository bug only rarely ([Chapter 63](#)). So nearly every

group of rollouts is unanimous zeros, zero variance, zero gradient (Chapter 5). Stabilizing the run (SW2–SW6) stopped the model from destroying itself, but it did nothing to manufacture the successes the gradient needs. Stable or not, a reward that fires almost never can't teach.

This is the wall: outcome-reward RL needs a trickle of successes to bootstrap, and on a task this hard for this model, the trickle is essentially dry.

The SW6 stop state, in numbers:

Metric	Value at the stop (step 146)
step	146
entropy	~0.296, steadily down from ~0.508
reward mean	~0.181, flat (was ~0.196 at step 16)
max reward	~0.37, the same instances as before
test-pass signal	only ~5% of rollouts passed any test
conclusion	stable trajectory, no new bug-solving capability

Table 66.1. The SW6 run’s final state in numbers, captured at its stopping point (step 146): a stable trajectory — entropy held, reward flat — but with only ~5% of rollouts passing any test, no new bug-solving capability was gained.

Two structural problems compound the sparsity: ~280 instances is too small for this hard a setting (low diversity, repeated exposure), and per-episode credit assignment means the final reward can't tell which of ~10 bash turns mattered. Sparsity is the headline. These make it worse.

66.2 Why shaping rewards helped stability but not learning

The obvious response is to add shaping rewards, partial credit for intermediate progress, to give the gradient something to grip. Here the nuance matters, because the simple story (“shaping backfired”) is wrong: in SW4/SW6, shaping rewards did help. They eliminated all-zero groups and kept the run alive. What failed was their alignment:

1. Shaping fixed the variance problem and created a gradient. Points for “edited a file,” “ran the tests,” “produced a diff” gave non-unanimous groups where the binary reward gave none, so there was a gradient to follow. That part worked.
2. But the gradient pointed at activity, not repair. The shaping terms, length, commands-run, output, diff, rewarded looking busy, not fixing the bug. So the model dutifully climbed the gradient it was given: more edits, more test runs, more diffs, none of them solving anything. This is reward hacking (Chapter 31) and tool-spam (Chapter 35), but induced by misaligned shaping, not by shaping per se.
3. The wall didn't move. A gradient that points at the wrong target gets you efficiently to the wrong place. The run stayed alive and stayed at ~5% test-pass, because the signal it was optimizing was activity.

The lesson is sharper than “shaping is bad,” and it's the one that survives into Chapter 84: shaping fixed the variance problem and created a gradient. It failed because the gradient pointed at activity rather than repair. Shaping isn't the enemy, misaligned shaping is. The deeper problem remains that the real reward was unreachable for this model, which raises the question the next chapter answers: can a 14B do this at all?

TAKEAWAY

After SW2–SW6 stabilized the run, accuracy still never moved (SW6: step 146, entropy ~0.296, reward ~0.181 flat, only ~5% of rollouts passing any test), the sparse-reward wall, compounded by only ~280 instances and per-episode credit assignment. Shaping rewards *helped* (they killed all-zero groups and created a gradient) but were *misaligned*: length/commands/output/diff rewarded looking busy, not repair. Shaping isn't the enemy, misaligned shaping is.

Next: the uncomfortable diagnosis, whether a 14B can solve full-repo bugs at all, and why RL can't give it a capability it doesn't have.

CHAPTER 67

Why a 14B Can't Solve Full-Repo Bugs

At some point in a failing project you have to ask the question you've been avoiding: not "how do we train this better," but "can this model do this, this way, at all?" The honest answer here is no, but it needs scoping precisely, because the useful claim is narrow. Not "a 14B can't solve code." The claim is: a 14B does not solve enough full-repo SWE-bench bugs, under this harness, this binary-ish reward, and only ~280 instances, for sparse outcome RL to bootstrap. That reframes the whole effort, because RL sharpens what a model can already do often enough to reinforce. It does not hand the model a capability it can't reliably reach.

67.1 The capacity argument

The reasoning, stated plainly:

1. Full-repo bugs demand a lot at once. Resolving a real issue means holding a large codebase in context, localizing the fault across files, reasoning about the fix, and producing a correct multi-file edit, a long chain where every link must hold.
2. A 14B has a ceiling on that chain. Model capacity bounds how much of that simultaneous reasoning a model can reliably do. A 14B can solve some bugs, but full-repo issues sit largely above its reliable ceiling. It succeeds occasionally, by luck, not dependably.
3. And RL can't lift that ceiling (the result [Part VII](#) proves): RL sharpens the model's existing distribution, makes it better at what it already sometimes does, but it doesn't add fundamentally new capability. So pointing outcome RL at a task above

the model's ceiling can't work, because there's no latent skill to sharpen into reliability.

Put together: the SWE-bench failure wasn't a training bug or a reward bug. It was a capacity mismatch, a task above what a 14B can do, and a method (RL) that can't conjure the missing capability.

67.2 The DeepSWE parallel, and the honest conclusion

This isn't a private failure, and it doesn't indict the 14B everywhere:

- A 14B is still strong on easier or denser tasks. The same model is genuinely useful for structured tool use, function calling, text-to-SQL, search agents ([Part III](#)), and bash-style tasks, and for code environments with denser rewards. The wall is specific to full-repo bugs under sparse outcome reward, not to code or to 14B in general. Don't overgeneralize a setting-specific failure into "the model is too small for code."
- The field's strong software agents changed the signal and environment, not just the size. Systems like DeepSWE reach competitive numbers via denser environments, larger and more diverse task sets, code-specialized models, and different scaffolds, not by running sparse-reward GRPO harder on a mid-size model. The lesson isn't "use a bigger model". It's change the signal and the environment (R2E-Gym-Lite, InterCode-Bash, NL2Bash, trajectory distillation, OPD).
- So the honest conclusion is to stop brute-forcing this configuration. Tuning rewards and stabilizers on a 14B for full-repo bugs under sparse outcome reward is optimizing a method against a wall that method, on that signal, can't cross. The right move is a different strategy, a denser signal, not a bigger hammer.

That strategy is the pivot. If outcome RL can only sharpen what the model already reaches often enough, and on this task it doesn't, then we need a method that supplies a dense signal from something that can do the task. That is on-policy distillation, and it's the rest of this Part.

TAKEAWAY

The failure is setting-specific: a 14B doesn't solve enough full-repo SWE-bench bugs under this harness/reward/~280 instances for sparse *outcome* RL to bootstrap, *not* that a 14B can't do code (it's strong on denser code tasks, tool use, text-to-SQL, search, bash). The field's strong agents (e.g. DeepSWE) won by changing the *signal and environment* (denser rewards, bigger/specialized setups), not merely the model size. The fix is a denser signal, OPD, not a bigger hammer.

Next: that strategy, on-policy distillation, and why a dense per-token signal from a teacher beats a sparse outcome reward on a task the model can't yet do.

CHAPTER 68

On-Policy Distillation: The Strategy



Figure 68.1. Sparse vs dense reward as terrain an agent climbs: stuck on a flat plain ("which way is up?") vs guided up a smooth slope.

If a 14-billion-parameter model can't reach the successes that outcome RL needs ([Chapter 67](#)), the move is to stop relying on its successes and start borrowing a teacher's. On-policy distillation (OPD) replaces the sparse outcome reward with a dense, per-token signal from a stronger model, and that single change is why it succeeds where sparse-reward RL hit a wall. This chapter is the strategy. The next is the recipe.

Evidence status, strategic pivot + methodology proof, not a completed SWE-bench rescue. Be precise about what the corpus supports. OPD enters this book as the strategic exit after SWE-bench sparse RL failed. We validated the machinery cheaply on GSM8K, where the base 14B was already near the teacher, so the accuracy delta was small and not statistically meaningful (a methodology proof, not a capability win, [Chapter 71](#)). The code-agent OPD recipe in the next chapters is the recommended exit path for R2E-Gym-Lite

/ InterCode-Bash / NL2Bash, not a claimed “OPD solved SWE-bench” result. Stating that boundary protects the book’s credibility.

68.1 The core idea

Plain, then precise:

- Plain. Instead of letting the model fail a hard task and rewarding it only on the rare win, have a stronger teacher model show it, token by token, what a better answer looks like, on the student’s own attempts.
- Precise. OPD has two pieces. On-policy: the student generates the rollouts, so it learns on its own distribution of behavior (not on off-policy demonstrations it would never produce). Distillation: at each token, the student is trained to match the teacher’s probability distribution, rather than to maximize a far-off reward. The signal is the teacher’s per-token guidance, not a pass/-fail at the end.

68.2 Why dense beats sparse here

The whole advantage is in where the learning signal lands (Figure 68.1):

1. Outcome RL: one signal, at the very end. A long trajectory earns a single scalar reward when (if) the tests pass. Every token before it gets only the faint, noisy echo of that distant outcome, and when successes are essentially absent ([Chapter 66](#)), the signal is effectively nothing.
2. OPD: a signal on every token. The teacher provides a target distribution at each token of the student’s rollout, so every token is a supervised learning step. There’s no waiting for a success that never comes, the gradient is dense and informative from the first step.
3. So OPD sidesteps the sparse-reward wall entirely. It doesn’t need the student to succeed to learn. It needs the teacher to know the answer. The student improves by being pulled toward the

teacher's behavior on every token, regardless of whether its own attempt would have passed the tests.

This is the conceptual exit from [Part VI](#)'s failure: when a task is too hard for the model to crack by outcome reward, stop rewarding outcomes and start matching a teacher's per-token distribution. It trades the (unreachable) requirement of student success for the (reachable) requirement of a competent teacher, which, as the dataset chapter shows, becomes the new thing you have to get right.

Two honest caveats before the recipe. First, OPD still needs the student to produce parseable trajectories, if the student can't even emit well-formed tool calls, you need an SFT/cold-start phase first to teach the form before OPD can refine the content. Second, OPD is not the same as offline SFT on teacher transcripts: OPD distills on the student's own rollout distribution, which avoids the off-policy mismatch of training on trajectories the student would never have produced. The "on-policy" half is doing real work.

TAKEAWAY

On-policy distillation replaces the sparse outcome reward with a dense per-token signal: the student generates the rollouts (on-policy) and is trained to match a stronger teacher's distribution at every token (distillation). It beats sparse RL on too-hard tasks because the signal lands on *every* token, not once at a usually-absent success, so the student needs a competent teacher, not its own success. Caveats: it needs cold-start form if the student can't produce parseable trajectories, and it differs from offline SFT by distilling on the student's *own* distribution.

Next: the recipe in detail, per-token reverse-KL, the shared-tokenizer constraint, and masking the tool tokens.

CHAPTER 69

The OPD Recipe in Detail

The strategy is simple. The recipe has three details that decide whether it works, and each one is a place people get it wrong. This chapter is the on-policy distillation (OPD) recipe at the level you’d implement it, per-token reverse-KL, a shared-tokenizer constraint, and tool-token masking, because the gap between “distill from a teacher” and “distill from a teacher correctly” is exactly these three things.

69.1 The objective: per-token reverse-KL

The student is trained to match the teacher’s distribution at each token, and the direction of the match matters:

- The loss is a per-token Kullback–Leibler (KL) divergence, a score of how far the student’s distribution is from the teacher’s, summed over the tokens of the student’s own rollout.
- Reverse-KL specifically (matching the teacher in the mode-seeking direction) makes the student commit to the teacher’s high-probability behavior rather than trying to cover everything the teacher might do. For learning a skill (here, how to fix bugs), mode-seeking is what you want, concentrate on the teacher’s good moves.

```
# per-token reverse-KL: KL(student || teacher) on the student's
↪ OWN rollout
with torch.no_grad():
    teacher_logp = teacher(student_rollout).log_softmax(-1) #
    ↪ frozen target
student_logp = student(student_rollout).log_softmax(-1) #
    ↪ trainable policy
loss_tok = student_logp.exp() * (student_logp - teacher_logp)
```

```
kl = loss_tok.sum(dim=-1) # per
    ↪ generated token
mask = loss_mask.float() # 1
    ↪ only on student tokens
loss = (kl * mask).sum() / mask.sum().clamp_min(1.0)
```

69.2 The constraint: a shared tokenizer

This one is non-negotiable and easy to overlook:

- The teacher and student must share a tokenizer. Distillation matches distributions over tokens, so if the teacher and student tokenize text differently, their distributions are over different vocabularies and the per-token KL is meaningless. You'd be comparing apples to a different alphabet.
- Practically, this restricts which teacher you can use. You can't distill from an arbitrary strong model. You need one in the same tokenizer family as your student. It's a real constraint on teacher choice, and it's better to discover it now than after a failed run.
- And a shared tokenizer is necessary but not sufficient. You also need chat-template parity. Teacher and student must see the same serialized conversation: the same system prompt rendering, the same role markers, the same tool-call formatting. Two models with the same tokenizer but different chat templates produce token streams that don't line up, and the per-token KL silently compares mismatched positions. Render the conversation once, identically, for both.

69.3 The mask: don't distill on tool tokens

The masking discipline from the agent Parts returns, in distillation form:

- A SWE-bench rollout is mixed ([Chapter 21](#)): some tokens the model generated, some the environment inserted (tool outputs, test logs, file contents).

- Distill only on the student-generated tokens. The teacher's distribution over environment tokens is meaningless, the environment, not either model, produced them, so they must be masked out of the KL loss (the `loss_mask` above). Distill on them and you teach the student to predict test logs, which is noise.
- The mask convention, stated once: `1 = LLM-generated token`, `0 = tool/environment token`. Every token the environment inserted (stdout, file contents, test output) is a 0 and contributes nothing to the loss.
- If you distill from top-k teacher logits (to save bandwidth/memory rather than the full vocabulary), fix K and decide how missing tokens are handled, typically renormalize over the top-K and treat the tail as a small floor. Report K, because too small a K throws away part of the teacher's distribution you meant to match.

Get these three right, reverse-KL on a shared tokenizer, masked to the model's own tokens, and OPD does what the last chapter promised. Get any one wrong and you're matching the wrong thing, over the wrong vocabulary, on the wrong tokens.

TAKEAWAY

The OPD recipe is three details: a **per-token reverse-KL** loss (mode-seeking, commit to the teacher's good moves) summed over the student's own rollout. A **shared-tokenizer constraint** (distributions must be over the same vocabulary, which restricts teacher choice). And **tool-token masking** (distill only on student-generated tokens, never the environment's). Reverse-KL, shared tokenizer, masked to the model's tokens, miss one and you match the wrong thing.

Next: where the teacher's competence has to live, the datasets and environments for OPD, and the Goldilocks zone of task difficulty.

CHAPTER 70

Datasets & Environments for OPD

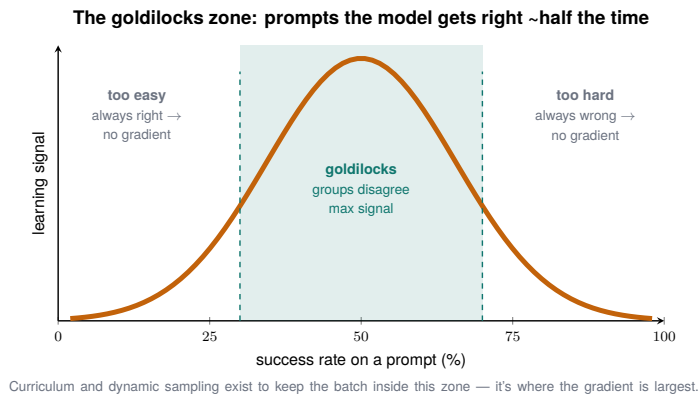


Figure 70.1. The goldilocks zone: learning signal peaks on prompts the model gets right about half the time.

On-policy distillation (OPD) moved the hard problem from “can the student succeed?” to “can the teacher?”, which makes dataset choice the new lever. This chapter is about picking tasks and environments where OPD actually teaches, and the principle that governs it: the Goldilocks zone, the band of difficulty where the teacher succeeds and the student doesn’t yet.

70.1 The environments and their roles

Code OPD uses more than one dataset, each playing a different role, and the roles matter as much as the names:

Environment	Role	Why it belongs
NL2Bash	early curriculum / single-turn shell	teaches command synthesis before long multi-turn trajectories
InterCode-Bash	interactive bash	dense feedback and terminal interaction without full-repo complexity
R2E-Gym-Lite (or a subset)	main code-agent training env	executable code-editing with more learnable signal than SWE-bench
Nebius / SWE-agent / Open-Hands trajectories	off-policy SFT / rejection-sampling warm-start	teacher trajectories give the student form before refinement
SWE-bench Verified / Lite subset	eval-only / upper-bound comparison	do not train on it if you use it for the final comparison

Table 70.1. The code-agent training environments and the role each one plays. They form a curriculum: single-turn shells teach command synthesis, interactive bash adds dense feedback, R2E-Gym is the main trainable env, teacher trajectories warm-start the student, and SWE-bench is kept eval-only.

These are mostly executable, the agent acts and gets real feedback, which matters because OPD is on-policy: the student generates real trajectories and the teacher guides them. The warm-start trajectory corpora are the exception, used for the SFT/cold-start form-teaching from [Chapter 68](#), not for OPD itself. And the train/eval separation is non-negotiable: if SWE-bench is your yardstick, it cannot also be your training set.

70.2 The Goldilocks zone

The governing principle for which tasks to include, drawn straight from how OPD works (Figure 70.1):

1. Too easy → nothing to transfer. If the student already solves a task, the teacher’s distribution barely differs from the student’s, so the per-token KL is near zero, no signal, no learning.

2. Too hard → nothing to distill from. If the task is so hard the teacher also fails, the teacher's "guidance" is just its own confusion. Distilling from a teacher that can't do the task teaches the student the teacher's mistakes.
3. Just right (Goldilocks) → the teacher solves it, the student doesn't yet. Here the teacher's distribution carries real, correct information the student lacks, and the per-token KL pulls the student toward genuinely better behavior. This band is where OPD teaches.

So selecting an OPD dataset is selecting for that middle band: tasks within the teacher's reach but beyond the student's current reach. This reframes [Chapter 17's](#) "data where the model succeeds sometimes" for distillation. It's now "data where the teacher succeeds and the student is still failing." Curate toward that gap, and OPD has signal to work with. Stray out of it on either side and you're back to no learning (too easy) or learning the wrong thing (too hard).

The rule that makes Goldilocks operational, as a filtering gate you can actually run:

Keep tasks where teacher success is high enough to trust and student success is low enough to create signal. Drop the tasks both solve (no gap) and the tasks both fail (no reliable target). Measure both rates up front and filter on them.

TAKEAWAY

OPD's new lever is dataset choice, governed by the Goldilocks zone: include tasks the *teacher* can solve but the *student* can't yet. Too easy and the teacher's distribution matches the student's (no signal). Too hard and the teacher fails too (you distill its mistakes). We used executable environments, R2E-Gym-Lite and InterCode-Bash, and curated toward the teacher-can/student-can't band, the distillation version of "data where the model succeeds sometimes."

Next: the compute reality of running a teacher and a student together, the 4+4 layout, and the teacher quantization that quietly drops OPD quality.

CHAPTER 71

Compute Budgets & Teacher Quantization

On-policy distillation (OPD) has a cost the strategy chapters glossed: you now need two models running, the student you're training and the teacher scoring every token, and they have to share your eight GPUs. This chapter is the compute reality, and the trap inside it: the obvious way to fit the teacher (quantize it) silently degrades the distillation.

71.1 The 4+4 layout

The basic resource split, on 8×H100:

- Four GPUs for the student, four for the teacher, the “4+4” layout. The student needs its training machinery (gradients, optimizer). The teacher needs only inference (it's frozen), but a strong teacher is large, so it still claims its half.
- The teacher is pure overhead to the student's progress. It produces no gradient of its own, it only supervises. That makes the teacher's footprint the thing you most want to shrink, which is exactly where the trap waits.

71.2 The teacher-quantization trap

The tempting optimization and its cost:

Teacher	Distillation-quality trade-off
Full-precision teacher	best distillation signal; largest footprint
AWQ-quantized teacher	fits easily, runs fast — but drops distillation quality to “tier-4”

Table 71.1. The teacher-quantization trade-off for on-policy distillation: a full-precision teacher gives the best signal at the largest footprint; an AWQ-quantized teacher fits and runs fast but drops distillation quality.

The temptation is to quantize the teacher, AWQ (a 4-bit weight quantization, Activation-aware Weight Quantization) shrinks it to fit comfortably. But a quantized teacher is a subtly different teacher: its per-token distribution is perturbed by the quantization, and since OPD's entire signal is the teacher's distribution, perturbing it degrades the student's learning. In our accounting, an AWQ teacher dropped OPD to tier-4 quality, a meaningful loss, because the student is now matching a slightly-wrong target on every token.

This is the setup for the next chapter, and it states a real tension carefully: teacher quantization is not automatically disqualifying. It is unsafe when it perturbs the teacher distribution beyond what the task can tolerate, and acceptable when a preflight check shows the quantized teacher is a faithful stand-in ([Chapter 72](#)). The “tier-4” label above is shorthand for “measurably degraded on our accounting”, treat it as a flag to run the preflight, not as a verdict that AWQ teachers never work. The teacher is the part you most want to shrink and the part you can least afford to corrupt blindly, so you measure before you trust.

One honesty point belongs here, because it's the actual evidence for the OPD machinery. The cheap validation run was on GSM8K (call it the C10 OPD run): it confirmed the pipeline works end to end, but GSM8K was saturated for the 14B (already near the teacher), so the accuracy gain was about +1pp and not statistically meaningful. That is a methodology proof, not a capability win ([Chapter 68](#)'s evidence-status note). Reporting the +1pp honestly, rather than dressing it as a result, is exactly the discipline the reward-design and eval Parts insist on.

TAKEAWAY

OPD runs a student and a teacher together, a 4+4 GPU layout on 8×H100, and the teacher is pure supervision overhead, so it's the footprint you most want to shrink. Quantizing it (e.g. AWQ)

perturbs its per-token distribution, which is the signal, but that's a reason to *preflight* it (Ch 72), not to ban it. And keep the OPD evidence honest: the cheap GSM8K validation proved the machinery works but was saturated (+1pp, not significant), a methodology proof, not a capability win.

Next: the answer, a preflight gate that validates the teacher before you distill, and the taxonomy that says when a cheap teacher is safe.

CHAPTER 72

Stage 0: Validate the Teacher Before You Distill

The last chapter left a real dilemma: a quantized teacher saves the memory you need but may corrupt the signal that is the method. The resolution is not to guess. It's to measure the teacher before you trust it. This chapter is the preflight gate that does that, and the taxonomy it produces: a precise answer to “when is a cheap teacher safe?” that rescues quantized teachers when they pass and rejects them when they don't.

72.1 The preflight gate: a top-K KL check

Before committing a single training step, run the teacher through a cheap validation:

1. Take a sample of real inputs and run both the full-precision teacher and the quantized teacher on them.
2. Compare their distributions with a top-K KL check, measure the Kullback–Leibler divergence between the quantized teacher's top-K token probabilities and the full teacher's, on each token.
3. Gate on the result. If the quantized teacher's distribution stays close to the full teacher's (KL below a threshold on the tokens that matter), it's a faithful stand-in, pass. If it diverges, the quantization corrupted exactly the signal OPD depends on, fail, and you don't distill from it.

```
# Stage 0 preflight: is the quantized teacher a faithful  
↪ stand-in for the full teacher?  
full = full_teacher(sample).log_softmax(-1)  
quant = quant_teacher(sample).log_softmax(-1)
```

```

topk = full.topk(K, dim=-1).indices # the tokens that carry the
    ↪ signal
kl = (full.exp().gather(-1, topk) *
      (full.gather(-1, topk) -
        ↪ quant.gather(-1, topk))).sum(-1).mean()
assert kl < THRESH, "quantized teacher diverges - do NOT distill
    ↪ from it"

```

The gate is cheap (one forward pass over a sample) and decisive: it converts “is this teacher safe?” from a hope into a measured pass/fail before you spend the compute.

The parameters you must pin (and report) for the check to mean anything: the sample size of inputs. The K for the top-K comparison. The pass/fail threshold on the KL. And, easy to miss, mask the same tool/environment tokens in the preflight that you mask in training (Chapter 69), and average the KL over assistant/model tokens only. A teacher that’s faithful on tokens it never has to generate tells you nothing.

And one check beyond the KL gate, because a distribution can pass a numeric threshold and still hurt downstream: add a smoke gate. After the top-K KL passes, run a tiny same-script evaluation or a short OPD smoke run and confirm it doesn’t regress before committing the full budget. The KL gate is necessary. The smoke gate catches what a single scalar misses.

If the gate fails, you have a ladder of fallbacks rather than a dead end:

1. Use a BF16 teacher (drop the quantization, pay the memory).
2. Use a smaller teacher in the same tokenizer family (less capable but faithful, and still in Goldilocks range).
3. Offload or serve the teacher remotely (trade latency for fidelity, viable if the rollout loop can absorb it).
4. Pick an easier dataset slice where even the faithful-but-weaker teacher clears the Goldilocks bar.

72.2 The three-scenario AWQ taxonomy

The preflight gate resolves a confusion that wastes real compute, when AWQ (the 4-bit quantization) is acceptable. The answer depends entirely on the role:

Role	Verdict	Why
rollout engine (generates training trajectories)	unsafe	quantization perturbs the on-policy distribution you train on — corrupts the data
teacher (scores tokens for distillation)	safe if preflight passes	a faithful quantized teacher is fine; the preflight gate is what confirms it
serving (inference for users)	fine	small quality drops are acceptable; no training signal at stake

Table 72.1. The AWQ-by-role taxonomy: quantization is unsafe where it perturbs the on-policy data, safe-if-checked for the teacher, and fine for serving. The verdict depends entirely on whether a training signal is at stake.

The distinction is the contribution: the same quantization is unsafe in one role, conditionally safe in another, and fine in a third, and the difference is whether its perturbed distribution feeds a training signal. As a rollout engine it corrupts the data you learn from. As a teacher it’s tolerable exactly when the preflight gate says its distribution is still faithful. For serving it never touched training at all. (This sharpens the rollout-engine warning from the quantization Part, cross-reference there.)

The transferable rule: validate the teacher before you distill. A cheap teacher isn’t automatically disqualified. It’s disqualified if it fails the gate, and rescued if it passes. The gate is what lets you spend the savings without paying for them in corrupted students.

TAKEAWAY

Don’t guess whether a quantized teacher is safe, run a Stage-0 preflight: a top-K KL check comparing the quantized teacher’s distribution to the full teacher’s on a sample, and distill only if

it passes. This resolves the AWQ taxonomy by role: unsafe as a rollout engine (corrupts on-policy data), safe as a teacher *if pre-flight passes* (a faithful quantized teacher is fine), fine for serving (no training signal). Validate the teacher before you distill.

Next: *Part VII*. The journey gave us outcomes. Now we ask what changed inside the model.

Part VII

What RL Actually Does: The Mechanistic Picture

Six Parts of treating the model as a box that “learns,” and now we open it. This Part is the book’s most citable layer, a set of measurements, all on our own checkpoints, that together form a coherent account of what reinforcement learning with verifiable rewards actually does to a model. It does not expand the model’s abilities. It sharpens them. It edits a tiny, identifiable set of tokens. It moves the weights less than folklore claims, and in a direction that decides what capability survives. These are not separate curiosities, kept together, they are a theory of RLVR grounded in experiment, and the engine for the companion papers.

CHAPTER 73

RL Sharpens. It Does Not Expand




Figure 73.1. RL sharpens pass@1 but can spend pass@k diversity — a sharper policy explores fewer distinct solutions.

There is a live argument in the field about what reinforcement learning with verifiable rewards (RLVR) does: does it teach the model genuinely new capabilities, or just make it better at expressing what it already knew? We answered it on our own checkpoints, with a roughly three-dollar experiment, and the answer is decisive and a little deflating: RL sharpens. It does not expand. It makes the model more reliable at problems it could already sometimes solve, and finds essentially no problems it couldn't.

Evidence contract, pass@k sharpening result:

Item	Value
base	Qwen3-14B BF16
RL checkpoint	Q4 BF16, step 200
task	math_dapo
eval size	30 problems
attempts	pass@1, pass@4, pass@16, pass@64
cost	negligible compute
headline	pass@1 +2.1pp. Pass@64 0.667 → 0.633

73.1 The pass@k crossover

The measurement is the cleanest way to see sharpening versus expansion:

- pass@k asks: given k attempts at each problem, how often does at least one succeed? pass@1 measures single-shot reliability. Pass@64 measures the breadth of what the model can reach at all with many tries.
- At low k, RL wins. pass@1 improved by +2.1pp, and pass@4 through pass@16 by about +4.8pp, the RL model is more reliable in a single shot.
- At high k, RL loses. pass@64 dropped by 3.4pp (0.667 → 0.633). With enough attempts, the base model solves more distinct problems than the RL model does.

The curves cross (Figure 73.1), and the crossover is the whole story: RL concentrated the model’s probability onto answers it already had (raising pass@1) at the cost of the diversity that let many tries reach more problems (lowering pass@64). It sharpened the distribution. It narrowed it.

73.2 The clinching detail, and the honest caveat

Two things make this more than a curve:

1. The base-unsolved problems stay unsolved. Concretely at pass@64, the base solved 20/30 and the RL model 19/30: the 10 problems the base couldn't solve stayed unsolved under RL, and RL actually lost one the base could reach. If RL expanded capability, it would crack problems the base couldn't. It didn't. That's the signature of sharpening, not expansion.
2. This decided a real spend. The result let us cancel a planned prolonged-RL run, if more RL only sharpens harder, a longer run wouldn't find new capability, so it wasn't worth running. A mechanistic finding paid for itself by killing an experiment.

Two caveats keep this honest, because the debate isn't fully closed. First, the eval is 30 problems, strong as a mechanistic signal and more than enough to cancel the run, but not a tight confidence interval. The right phrasing is “decisive for our stack,” not a universal theorem (a bootstrap CI on those 30 would quantify it). Second, this tested our RLVR recipe. A specific prolonged-RL recipe (the ProRL line, more steps, periodic reference resets, broad task diversity) makes stronger expansion claims and we did not test it. So the precise statement is: standard RLVR, as we ran it, sharpens and does not expand, whether a very different prolonged recipe can expand remains open.

TAKEAWAY

On our own checkpoints (Qwen3-14B base vs Q4 step-200, math_dapo, 30 problems), RLVR sharpens but doesn't expand: pass@1 +2.1pp while pass@64 *fell* 0.667→0.633 (base 20/30 vs RL 19/30), the curves cross. Base-unsolved problems stayed unsolved (RL lost one), the clinching sign of sharpening. It cancelled a planned prolonged-RL run. Caveats: 30 problems is a strong signal, not a tight CI. And the specific ProRL recipe remains untested. Decisive *for this stack*, not a universal theorem.

Next: if RL only sharpens, where* does it apply its edits? The answer is startlingly specific. It touches almost only the fork tokens.*

CHAPTER 74

RL Edits Only Fork Tokens

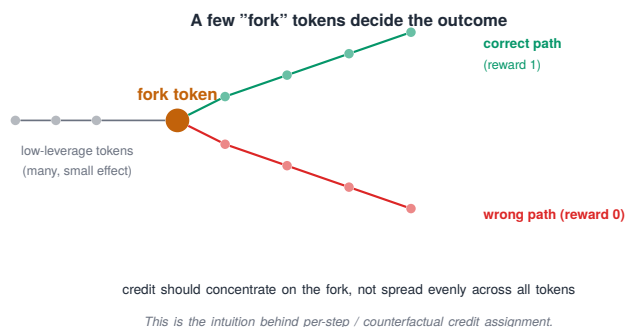


Figure 74.1. A few fork tokens decide the outcome — where credit should concentrate.

If RL sharpens (Chapter 73), the natural next question is where, which of the hundreds of tokens in a response does the update actually move? The answer is the single most citable statistic in our whole corpus, and it's startlingly specific: RL edits almost only the fork tokens, the places where the model was already uncertain. The confident tokens, where the model already knew what came next, are left essentially untouched.

Evidence contract, mechanistic probe result:

Item	Value
base	Qwen3-14B BF16
RL checkpoint	Q4 BF16, step 200 (<code>actor_hf</code>)
probes	32 math probes
continuation	64 tokens each → 2048 tokens total
generation	greedy from base. The same continuation scored under base and RL
metrics	base entropy per token vs $ \Delta \log p $ (RL – base)

74.1 The measurement

The method is reproducible in an afternoon, and the result is stark:

- **Method.** Take a model's greedy continuation, and score each token under both the base and the RL model. The change at each token is $|\Delta \log p|$, how much RL moved that token's log-probability. Separately, measure each token's entropy, how unsure the model was there.
- **The correlation.** The Spearman correlation between a token's entropy and how much RL changed it is 0.95, almost perfect (Pearson 0.626, lower because the relationship is monotone but not linear). Where the model was unsure, RL changed it. Where it was confident, RL left it alone.
- **The concentration.** The high-entropy half of tokens absorbs 99.8% of the total distribution shift. The top 10% carry 70.3%, and the top 5% carry 49.9%, half of everything RL did lives in one token in twenty.

So the update is not spread across the response. It is concentrated, almost entirely, on the handful of fork tokens, the decision points where the model was choosing between paths (Figure 74.1).

Caveat, scope. This was measured on math continuations, not on multi-turn tool trajectories. The hypothesis should transfer, forks are forks, but the tool-agent fork-token map remains future work ([Chapter 80](#)) and shouldn't be claimed as measured until it is.

74.2 Why this is the mechanistic core

This finding ties the whole book together, which is why it's load-bearing:

1. It explains sharpening ([Chapter 73](#)) mechanically. RL sharpens by resolving the forks, at each uncertain decision point, it pushes the model toward the choice that led to reward. It doesn't rewrite the confident scaffolding (which is why capability isn't expanded). It just decides the coin-flips more reliably. Sharpening is fork-resolution.

2. It gives the clipped update from [Chapter 4](#) a location. We described the update as nudging token probabilities. Now we know which tokens, the high-entropy ones. The ratio-and-clip machinery does its real work at the forks.
3. It points exploration at the right place (the next chapter). If 99.8% of the useful change lives in the high-entropy tail, then exploration should be targeted at fork tokens, not sprayed globally, which is exactly why the global entropy bonuses of [Part IV](#) were the wrong tool and Clip-Cov's targeting was the right one.

The one-line version to carry: RL's edits live at the forks. Almost everything it does happens at the small set of tokens where the model was unsure, a fact that explains what RL is (fork-resolution), where it acts (the high-entropy tail), and how exploration should be aimed (there, not everywhere).

TAKEAWAY

RL edits almost only fork tokens (measured on 32 math probes $\times$ 64 tokens = 2048, base vs Q4 step-200): Spearman 0.95 (Pearson 0.626) between a token's entropy and how much RL changed it. The high-entropy half absorbs 99.8% of the shift, the top 10% carry 70.3%, the top 5% carry 49.9%. This is the mechanistic core, sharpening *is* fork-resolution (the confident scaffolding is left alone), it gives [Chapter 4](#)'s clipped update a location, and it shows exploration should target forks, not spray. Caveat: this is math-probe evidence. The tool-agent version is future work.

Next: putting fork tokens and the entropy frontier together into a practical map of which entropy is productive and which is wasted.

CHAPTER 75

Productive vs Wasted Entropy




Figure 75.1. Three entropy regimes shown as a performance curve: it peaks only in the healthy band between collapse and chaos, with the token distribution in each zone.

The last two chapters give us something the entropy Part couldn't: a way to tell good entropy from bad. But the way is subtler than “high good, low bad,” and getting it right means correcting a tempting oversimplification. [Chapter 62](#)'s resumed COMBO run had low entropy and climbing validation, and that was sharpening, not collapse. So the honest rule is: entropy alone cannot classify a regime. You read it against held-out validation, pass@k diversity, and score variance. Low entropy is wasteful only when those confirm it is.

75.1 Four regimes, not three

Lay entropy against validation and four regimes appear, not a one-dimensional good/bad axis (Figure 75.1):

Regime	Signal pattern	Interpretation	Action
productive fork entropy	entropy held where forks are; VAL/reward climb	useful exploration (e.g. Clip-Cov 0.4–0.5, reward $\rightarrow$ 0.8)	continue
productive compression	entropy falls, but VAL climbs and score variance stays positive	sharpening onto correct trajectories (the C9-COMBO case)	continue to VAL plateau
wasted compression	entropy falls, VAL/pass@k plateaus or drops, variance collapses	the exploration resource is being destroyed	stop or reset
sprayed entropy	global entropy rises but VAL doesn't move	random, untargeted exploration	target entropy at the forks

Table 75.1. Four entropy regimes, not three: for each, the signal pattern, how to interpret it, and the action. The key split is between productive compression (entropy falls but validation climbs — continue) and wasted compression (entropy falls and validation drops — stop).

The crucial pair is the two compression rows. The C9 fixes drove entropy to 0.05–0.08, and earlier I called that flatly “wasted.” That was too strong: whether it’s wasted depends on what validation was doing. When VAL was still climbing (COMBO on resume), low entropy was productive compression, the model sharpening onto correct trajectories, exactly the fork-resolution of [Chapter 74](#). It becomes wasted compression only once VAL and diversity have plateaued and the run is just grinding entropy away for no further gain.

75.2 The rule: classify by the bundle, not by entropy

This is the reconciliation between [Part IV](#) (where entropy collapse was the enemy) and [Part V](#) (where low-entropy COMBO was winning), and it resolves cleanly:

- Productive entropy is entropy that supports held-out improvement. Wasted entropy is entropy that doesn't. That's the definition, stated in terms of the outcome (does VAL move?), not the entropy level.
- So don't kill a run on entropy alone. An entropy-only kill rule murders productive compression. It would have killed the winning COMBO run (Chapter 62). Below-floor entropy with VAL still climbing is sharpening. Let it run.
- And don't crush or spray. Crushing entropy while VAL has plateaued (wasted compression) buys nothing. Spraying entropy globally (sprayed) explores the confident tokens RL shouldn't touch. The right move remains fork-targeted exploration (Clip-Cov, Chapter 43), but you confirm it's working with VAL, pass@k, and score variance, not with the entropy reading by itself.

The rule to carry, and the one that ties Parts IV, V, and VII together:

Entropy is not enough to classify a regime. Use entropy together with held-out validation, pass@k diversity, score variance, response length, and task behavior. Low entropy that sharpens (VAL climbing) is productive. Low entropy that grinds (VAL plateaued) is wasted, and only the bundle tells you which.

TAKEAWAY

Don't judge entropy by its level, judge it against validation. Four regimes: productive fork entropy (entropy held at forks, VAL climbing), **productive compression** (low entropy *but* VAL climbing = sharpening, the C9-COMBO case I'd wrongly called wasted), **wasted compression** (low entropy *with* VAL plateaued), and sprayed entropy (high but untargeted). Entropy alone can't classify a regime. Use it with held-out VAL, pass@k, and score variance. An entropy-only kill rule murders productive compression, which is exactly the COMBO lesson from Chapter 62.

Next: how much RL actually moves the weights, and why the answer kills a popular folk claim about subnetworks.

CHAPTER 76

How Much Does RL Move the Weights?




Figure 76.1. How far the weights actually move: SFT large, LoRA medium, RL tiny (7× less for a bigger gain).

There’s a popular folk claim that RL trains a small “subnetwork”, that it lights up some 5–30% of the weights and leaves the rest. We measured how much RL actually moves the weights, and the folk claim doesn’t survive even a small threshold: at $0.01 \times \text{std}(w_{\text{base}})$, one percent of the base parameter standard deviation, zero weights cross, in either run. That’s the strong version of the result. “Nothing moves more than a full standard deviation” would be unsurprising. “nothing moves more than one percent of a standard deviation” is far stronger, the update is a diffuse field of tiny shifts, not a discrete updated subnetwork. RL moves the weights less than people assume, and the style of update (full fine-tune vs LoRA) changes both how much moves and where.

76.1 The measurement

Plain, then precise:

- Plain. You’d think RL substantially rewrites a chunk of the model. Measure it and the change is small, small enough that, by a strict definition of “meaningfully changed,” essentially no individual weight qualifies.
- Precise. Define a weight as “changed” if it moved by more than $0.01 \times \text{std}(w_{\text{base}})$, one percent of that tensor’s base standard deviation. By that strict-but-small threshold, zero weights cross, in either the full-fine-tune run or the LoRA run. The Mukherjee-style “5–30% subnetwork” only appears at a much looser threshold. At the one we measured, there is no discrete updated subnetwork.

The actual measurement, which belongs in the text because it’s what earns the claim:

Metric	32B NF4-LoRA (step 100)	14B full-FT (step 200)	note
total params	32.8B	14.8B	different scale
updated frac at $0.01 \cdot \text{std}(w_{\text{base}})$	0.00%	0.00%	tied at zero
$\ \Delta\theta\ _2 / \ \theta\ _2$	0.000145%	0.001010%	full-FT $\approx 7.0\times$ larger
MLP rel-change	0.000146	0.001063	$7.3\times$
self-attn rel-change	0.000168	0.001140	$6.8\times$
lm_head rel-change	0	0.000456	LoRA preserves
embed rel-change	0	0.000263	LoRA preserves
norm rel-change	0	0.000005	LoRA preserves / tiny

Table 76.1. How far RL actually moves the weights, per tensor group — LoRA versus a full fine-tune. The two arms are a 32B NF4-LoRA run and a 14B full-precision (BF16) fine-tune; the full fine-tune moves weights $\sim 7\times$ more, while LoRA leaves the embeddings, head, and norms untouched.

The measurement is a few lines, worth showing, because the threshold choice is the claim:

```
# delta_theta_mask.py - how much did each tensor move, relative
↪ to its own base scale?
for name, w_base in base_state.items():
    delta = (ckpt_state[name] - w_base).float()
    sigma = w_base.float().std()
    updated_frac = (delta.abs() > 0.01 * sigma).float().mean() #
    ↪ fraction past 0.01*std
    rel_l2 = delta.norm() / w_base.float().norm() # global
    ↪ movement
```

So “RL trains a subnetwork” is threshold-dependent, and at $0.01 \times \text{std}(w_{\text{base}})$ it’s false. One consequence matters for [Chapter 80](#): because a threshold mask is degenerate here (it selects nothing), the causal ablation that would test “do these weights matter?” must instead use a top-k% by $|\Delta w|$ mask (e.g. the top 5%, 10%, 30% most-moved weights), not a fixed- σ cutoff. Reframing the mask that way is what makes the deferred causal experiment well-posed.

76.2 LoRA vs full fine-tune: less, and elsewhere

The update style changes the answer in two ways that matter for the next chapter:

1. LoRA moves the weights $\sim 7\times$ less. LoRA’s global change $\|\Delta\theta\|$ is about a seventh of full fine-tuning’s. It’s a far smaller perturbation by construction (it trains a low-rank add-on, not the whole model).
2. LoRA preserves the embedding and output layers exactly. Full fine-tuning perturbs both. LoRA cannot touch the embedding (`embed`) or output (`lm_head`) layers, they’re outside its low-rank adapters, so they stay bit-for-bit identical to base. Full fine-tuning, training everything, does perturb them.

That second point is the bridge to the next chapter. The embedding and output layers are where a model’s vocabulary-level behavior lives, including, as we’ll see, its thinking-mode capability. Full fine-tuning touches them. LoRA can’t, and that mechanical difference turns into a capability difference with real teeth.

TAKEAWAY

RL moves the weights less than folklore claims: at $0.01 \times \text{std}(w_{\text{base}})$ (1% of base std), *zero* weights cross in either the full-FT or LoRA run, the “5–30% subnetwork” needs a far looser threshold, so the update is a diffuse field of tiny shifts, not a discrete subnetwork. The update style matters: LoRA moves $\|\Delta\theta\| \sim 7\times$ less than full-FT and preserves `embed/lm_head/norm` *exactly*, while full-FT perturbs them. Because a threshold mask selects nothing here, the causal test must use a top-k%-by- $|\Delta w|$ mask (Ch 80).

Next: what that perturbation costs, full fine-tuning’s sharpening came with the most sobering regression in the book.

CHAPTER 77

What Sharpening Costs: The Full-FT Regression




Figure 77.1. Two ways to apply the update: a small LoRA patch on a frozen base vs a full fine-tune that shifts every weight.

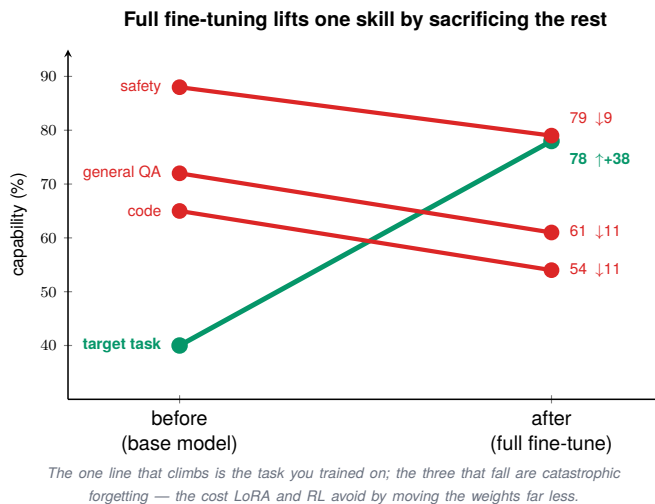


Figure 77.2. Full fine-tuning lifts the target task (+38) by sacrificing the rest: a slope chart showing safety, QA, and code all slipping (catastrophic forgetting).

This is the most sobering result in the book, and the strongest pairwise signal in the whole measurement program (a t-statistic of 10.4). We ran math RL with BF16 full fine-tuning (the Q4 arm, note “Q4” is the experiment label, not 4-bit precision), and it worked, it sharpened the math distribution exactly as [Chapter 73](#) predicts. It also erased the model’s thinking-mode capability, driving its AIME score from 0.458 to 0.222. The model got better at the trained distribution by destroying a capability it wasn’t training on. And among the runs that did preserve it, the variable that lines up best with the outcome is the update style.

77.1 The four outcomes

Math RL, four ways, on AIME (a hard competition-math benchmark) in thinking mode:

Arm	Model / scale	Update style	Method	AIME (think)	Status
full fine-tune	Qwen3-14B	full fine-tune	BF16 RL	0.458 → 0.222	strong regression
NF4-LoRA (14B)	14B tier	LoRA	NF4 RL	0.467	preserved (within noise)
distillation (14B)	14B tier	distillation	OPD	0.475	neutral / base-like
NF4-LoRA (32B)	Qwen3-32B	LoRA	NF4 RL	0.525	improved (scale confound)

Table 77.1. Math RL four ways on AIME (thinking mode): the update style decides who regresses. The four arms are a 14B full fine-tune, a 14B NF4-LoRA, a 14B distillation run, and a 32B NF4-LoRA; only the full fine-tune regresses hard, while LoRA and distillation stay at base level.

This is not a perfectly controlled one-factor ablation, and saying so matters. The arms differ along more than update style: the 32B-LoRA arm also changes model scale, and OPD also changes the objective (distillation, not RL). So the disciplined claim is the narrower one:

among the 14B-tier arms, the full-FT update is the one that catastrophically regresses thinking-mode capability, while parameter-efficient (LoRA) and distribution-matching (OPD) updates preserve it within the measured variance. The 32B improvement is suggestive, not a clean demonstration that LoRA raises capability.

And because the headline is a statistical result, it needs its variance budget, a preview of [Part IX](#)’s doctrine:

Checkpoint	mean	std	seeds	95% CI ±	verdict
Base 14B	0.458	0.032	4	0.025	baseline
14B full-FT (step 200)	0.222	0.039	3	0.033	significant regression
14B NF4 (step 100)	0.467	0.027	4	0.025	indistinguishable from base
14B distillation (step 200)	0.475	0.074	4	0.062	indistinguishable from base
32B NF4-LoRA (step 100)	0.525	0.032	4	0.025	improved, different scale

Table 77.2. Each arm with its variance budget: across multiple seeds, only the 14B full fine-tune clears its confidence interval as a real regression. The LoRA and distillation arms are statistically indistinguishable from base.

Full-FT’s drop clears its CI by a wide margin. The LoRA/OPD arms overlap base. So the spread is real, not seed noise, but only the full-FT regression is a clean, controlled conclusion.

Eval method (so the number is portable). This is AIME in thinking-mode decoding, with boxed-answer extraction and a fixed max-token budget, averaged over 3–4 seeds. Report all of that with the number: a thinking-mode regression is not comparable to a no-thinking/greedy result, the same checkpoint can look fine under greedy decoding and regressed under thinking mode. Decoding mode travels with the claim.

77.2 Why the update style decides what survives

This connects straight back to the previous chapter’s mechanism:

- Full fine-tuning touches the embedding and output layers (Chapter 76), the vocabulary-level machinery where thinking-mode behavior lives. Sharpening the math distribution dragged those layers with it, and thinking-mode capability was collateral damage.
- LoRA can't touch those layers (Chapter 76), so however hard it sharpens the math distribution, it physically leaves the thinking-mode machinery alone. Preservation isn't luck. It's a structural consequence of where LoRA is allowed to write.
- So the update style is a capability-preservation decision, not just an efficiency one. "Use LoRA to save memory" undersells it: LoRA also fences off the parts of the model you don't want sharpening to damage. Full fine-tuning removes that fence.

The load-bearing statement, and one of the book's central practical findings: the update style decides what survives fine-tuning. If you have capabilities you can't afford to lose, thinking mode, instruction-following, anything outside your training distribution, full fine-tuning on a narrow RL objective puts them at risk, and a parameter-efficient update (LoRA) or a distribution-matching one (OPD) protects them. Sharpening is not free. Full fine-tuning pays for it with whatever else lived in the layers it moved.

TAKEAWAY

Math RL with BF16 full fine-tuning sharpened the math distribution while erasing thinking-mode capability (AIME 0.458→0.222, $t \approx 10.4$, clearing its CI). NF4-LoRA (0.467), OPD (0.475), and 32B-LoRA (0.525) preserved or improved it. The comparison isn't a clean one-factor ablation (32B changes scale, OPD changes objective), but the controlled claim holds: among 14B-tier arms, *full-FT* is the one that regresses, because it touches the embed/output layers where thinking-mode lives (Ch 76) and LoRA can't. Choose the update style to fence off capabilities you can't lose, and report decoding mode with the number.

Next: a principle that generalizes Part V's MoE finding, capability tracks the direction representations move, not the distance.

CHAPTER 78

Direction, Not Distance: The Residual-Stream Ceiling

Two findings now point the same way. The MoE stabilizers (Chapter 59) showed capability tracked which direction representations moved, not how far. The full-FT regression (Chapter 77) showed the damage came from moving the wrong layers, not from moving a lot. Put them together and you get a principle that outgrows both specific cases: in our measured runs, a model's capability tracked the direction its internal representations moved relative to base more than the magnitude of the move. This chapter states that principle, framed as the consistent pattern across our runs, not a universal law, and it's the conceptual spine under Parts V and VII.




Figure 78.1. Direction beats distance: RL moves 7× less yet ends higher, because it changed direction not magnitude.

78.1 The principle

Plain, then precise:

- Plain. When you train, the model’s internal representations shift away from where the base had them. What predicts capability is which way they shift, not how far.
- Precise. Measured in the residual stream, the running vector the model carries layer to layer, a shift can be decomposed into magnitude and direction. Across our runs, capability correlated with the direction relative to base: representations drifting outward (extending the base’s structure) tracked higher capability. Representations retracting inward (pulling into a smaller version of the base) tracked lower. Magnitude alone predicted little.

This is exactly the FREEZE-vs-JITTER result (Chapter 59), outward-drifting JITTER landed above inward-retracting FREEZE, stated without reference to MoE. It’s not about routing. It’s about where the representations go.

How we measured direction. For a fixed neutral probe batch, we recorded residual activations from the base model and the trained checkpoint at matched layers. Per layer, we measured the change in residual norm and the signed drift relative to the base residual stream: positive drift = the checkpoint’s residuals moved outward from the base trajectory. Negative = they retracted toward a smaller/lower-norm version of it. We read the sign over matched checkpoints, not a single endpoint.

The pattern across the runs this Part and Part V cover:

Case	Direction	Magnitude	Capability	Lesson
FREEZE	inward / retracting	nonzero	lower ceiling	collapse avoided, capability bounded
JITTER	outward, then retract into a new attractor	not the largest distance	higher ceiling	direction/path beats distance
COMBO (GSPO+JITTER)	localized drift, bounded	neutral KL	strong VAL lift	targeted movement is productive
full-FT regression	wrong layers move (embed/lm_head)	larger global $\ \Delta\theta\ $	thinking-mode loss	update location matters

Case	Direction	Magnitude	Capability	Lesson
LoRA	small, fenced-off	smaller global $\ \Delta\theta\ $	preserved	a constrained update protects capability

Table 78.1. Across runs, the *direction* a representation moves (not the magnitude) tracks capability. FREEZE retracts inward to a lower ceiling; JITTER drifts outward to a higher one; COMBO makes targeted, bounded moves; a full fine-tune moves the wrong layers and regresses.

78.2 The U-curve as retraction into a new attractor

The principle’s most interesting form is the trajectory shape from Chapter 60:

- Capability-improving runs can move inward first, then outward, a U-curve. The model retracts out of the base’s basin, then climbs into a different, higher one. The early inward phase looks like regression but is repositioning.
- So you read direction over the whole trajectory, not a snapshot, an early inward dip can be the front half of a move into a better attractor (Chapter 60’s lesson, generalized). Distance at any single step tells you almost nothing. The path of the direction tells you where capability is headed.

The unifying upshot, which is why this is a Part-VII principle and not a Part-V footnote: what RL (and any fine-tuning) does to a model is better described by the direction of representational change than by its size. It explains why LoRA’s small, fenced-off move preserves capability (Chapter 77), why outward-drifting stabilizers beat retracting ones (Chapter 59), and why a strict weight-distance threshold finds “nothing changed” (Chapter 76) even as capability shifts, because the action was in the direction, which a distance threshold can’t see.

TAKEAWAY

In our measured runs, capability tracked the *direction* a model’s residual-stream representations moved relative to base more than the *magnitude*: outward drift (extending the base) with higher capability, inward retraction with lower, and capability-improving runs may trace a U-curve, retracting before climbing into a new, higher attractor. Direction is measured as signed residual drift on a fixed neutral batch, read over matched checkpoints. Framed as the consistent pattern across this family of runs, not a universal law, and it explains why a distance threshold finds “nothing changed” (Ch 76) even as capability shifts.

Next: the instruments behind all of this, the telemetry toolkit that makes Parts V and VII reproducible.

CHAPTER 79

Watching the Model Think: The Telemetry Toolkit

Every finding in this Part came from looking inside the model, and that's only possible with the right instruments. This chapter is the consolidated toolkit, the methods chapter that makes Parts V and VII reproducible, and the headline is that it's cheap: you can watch a model's internals throughout training for about 9% overhead. The findings aren't magic. They're what these few probes show when you bother to run them.

The telemetry stack: four questions, six dials

The reward curve alone hides the cause — read the group that answers your question.



Figure 79.1. The telemetry stack as four diagnostic questions — Is it learning? Is the gradient alive? Is the policy healthy? Did something break? — each with the metric and the shape to watch.

79.1 The instruments

Each probe answers a question the loss and reward can't, with its cadence, output artifact, and the failure it catches:

Probe	Question	Cadence	Artifact	Failure caught
forward-KL neutral probe	did general behavior drift from base?	every 25–50 steps	<code>internals_step*.json</code>	silent general drift
residual drift $\Delta\mu/\Delta\max$	which way did representations move?	per checkpoint	layer table	ceiling / attractor shift
per-layer routing entropy	are experts collapsing, layer by layer?	per checkpoint	per-layer entropy vector	MoE router collapse
logit lens	where does the prediction form across depth?	offline probe	token/layer map	hidden decision shifts
attention forensics	what does it attend to at key decisions?	selected checkpoints	attention-KL / plots	causal clue, not proof
ESS / clipped-fraction logger	is an off-policy / stale update safe?	every update	ESS curve	stale-rollout danger
delta-theta mask	how much did weights move?	offline	top-k mask subnetwork	claim sanity check

Table 79.1. The telemetry toolkit: each probe’s question, how often to run it, the artifact it produces, and the failure it catches. Together they cover silent drift, ceiling shifts, MoE router collapse, and stale-rollout danger.

The first five carry the core Part-VII claims. The last two are operational telemetry: the ESS / clipped-fraction logger isn’t used in this Part’s findings but is essential once rollouts go asynchronous (Part XIV), it bounds how stale an update can safely be, and the delta-theta mask is the Chapter 76 measurement.

79.2 The method, kept cheap

The point of a toolkit is that it runs every time, so it has to be light. The forward-KL neutral probe is the core, and the real version masks padding and keeps the batch fixed across runs:

```
# Forward-KL neutral probe: base -> current on a FIXED neutral
↪ batch.
# Keep batch, tokenizer, max length, and decoding template
↪ identical across runs.
with torch.no_grad():
    base_logits = base_model(**neutral_batch).logits[:, :-1]
    cur_logits = model(**neutral_batch).logits[:, :-1]
    log_p = base_logits.log_softmax(-1)
    log_q = cur_logits.log_softmax(-1)
    kl = (log_p.exp() * (log_p - log_q)).sum(-1) # KL(base ||
    ↪ current), per token
    mask = neutral_batch["attention_mask"][:, 1:].bool()
    fwd_kl = kl[mask].mean().item() # masked mean
```

We ran these probes as a handful of small harness scripts, one per probe: a frontier fit (VAL vs an internals axis), the pass@k crossover (Ch 73), a weight-movement mask (Ch 76), an effective-sample-size logger, and a per-layer routing/expert-load meter. Their outputs land as per-step JSON and CSV artifacts, joinable to checkpoints and evals.

Two rules keep the numbers honest:

- Compare only at matched steps (Chapter 58). Internals curves are comparable run-to-run only at the same step. “at the end” is not a fair comparison when runs differ in length or speed.
- Don’t compare across datasets or reward functions as if commensurate. Within-family, matched-step comparisons are valid. Cross-dataset or cross-method frontier fits are hypothesis-generating, not conclusions. (This is the internals journal’s own caution.)

On the cost: the ~9% overhead is the probe budget measured on a full experiment run with probes on the schedule above.

The transferable habit: budget ~9% for telemetry and never train a serious run blind. Every mechanistic finding in this book, sharpening, fork tokens, the residual ceiling, the regression, came from these probes.

TAKEAWAY

The internals toolkit, forward-KL neutral probes, residual drift, per-layer routing entropy, logit lens, attention forensics, plus operational ESS and delta-theta probes, runs for single-digit-percent overhead (~9%) and is what made every finding in this Part possible. We ran it as a handful of small harness scripts. Two rules: compare only at matched steps, and never compare across datasets/rewards as if commensurate. Budget the overhead. Never train a serious run blind.

Next: the honest part, what we measured but couldn't yet resolve, and the research agenda that comes next.

CHAPTER 80

Evidence Boundaries: What These Runs Do Not Prove

An honest reference says what it doesn't know, and this chapter is that, the questions this body of work does not settle, and why. To be clear about what this chapter is and isn't: it's an acknowledgment of limits, not a research plan. These aren't experiments anyone is committing to run. They're the honest edges of what the experiments in this book established, written down so a reader knows exactly where the evidence stops, and so that if the work continues, the questions are already framed. Knowing why a question is still open is most of knowing how to close it, whoever closes it.

80.1 What's unresolved, and what would settle it

For each open question: what the book found, why it's still unresolved, and the kind of experiment that would resolve it (for a reader with the compute, or a future pass):

Open question	What we found / why it's open	The experiment that would settle it
MoE longitudinal	MoE-family snapshots + partial probes; too few merged checkpoints to trace evolution	checkpoint forensics across a long MoE run (a small fraction of a fresh run)
Causal fork-token role	RL edits track fork tokens (Spearman 0.95), but a threshold mask is degenerate at 0.01σ , so a clean intervention wasn't well-posed	train/update only the top-k%-moved tokens (or their complement) and watch pass@k
Prolonged-RL expansion	standard RLVR sharpens, doesn't expand; only our recipe was tested, not the ProRL recipe	a long run with reference resets + task diversity; does pass@64 rise?
RL vs distillation forgetting	full-FT RL regressed thinking-mode, distillation stayed neutral; not a matched comparison	matched task + matched OOD eval across the two update styles

Open question	What we found / why it's open	The experiment that would settle it
Async rollouts + ESS	forward-looking; the safe staleness bound is unknown	rollout/train disaggregation with an effective-sample-size logger
Turn-level credit	the search/SWE sparse-reward wall; no reliable way to credit individual turns	per-turn reward / teacher-token KL / trajectory attribution
Determinism → stability	run-to-run reversals observed; nondeterminism's contribution is unquantified	fixed-seed vs varied-seed stability runs
Reward-hacking tripwires	the reward-up/EM-down signature; no automatic detector exists yet	an online reward-vs-EM divergence alarm

Table 80.1. The book’s open questions: for each, what was found and why it’s still unresolved, plus the experiment that would settle it. “Part XIV” is the rollout-engine part; “Part V/VI” are the MoE and SWE-bench parts.

None of these is an oversight: one is a cost we chose not to spend, one was blocked on a method the book has since built (Ch 76), one was cancelled by a finding but carries an explicit caveat. Writing them down is the point, a reference that hides its own edges is less trustworthy, not more.

80.2 Evidence boundaries: what Part VII does not establish

An honest reference also states the boundary of its own evidence. The chapters above support three claims on our stack: RL mostly sharpened existing behavior (Ch 73), the update concentrated at fork tokens (Ch 74), and update style decided what capability survived (Ch 77). They do not prove these claims for every recipe. Three boundaries matter:

1. Prolonged-RL recipes remain outside this evidence. We did not test the full ProRL-style recipe, longer training, reference resets, broader task diversity. “RL sharpens, doesn’t expand” is a statement about the recipes we ran, not a theorem about RL.

2. The fork-token result is correlational unless the tokens are directly edited. The evidence says high-entropy fork tokens absorb most of the probability shift (Spearman 0.95). Turning that into a causal claim would need the mask ablation noted in [Ch 76](#), which we did not run.
3. OPD's out-of-distribution advantage is not established here. The OPD runs show methodology and compute economics, not that OPD beats RL for OOD retention.

These are boundaries, not action items: read the book's claims inside them. (This is the same discipline as the honest losses of [Chapter 36](#), saying where the evidence stops is part of being trustworthy, not an admission the book is unfinished.)

TAKEAWAY

This chapter is an acknowledgment of limits, not a research plan, no one is committing to run these. What the book leaves unresolved: the MoE longitudinal study, the *causal* role of fork tokens (now well-posed thanks to [Ch 76](#)'s top-k mask), prolonged-RL expansion (cancelled for our recipe by [Ch 73](#). ProRL untested), RL-vs-OPD forgetting, async-rollout staleness bounds, turn-level credit, determinism's effect on stability, and reward-hacking tripwires, each with the experiment that *would* settle it, for a reader who wants to push further. And three **evidence boundaries** on [Part VII](#)'s headline claims, prolonged-RL, fork-token causality, OPD's OOD advantage, stated as the limits of what the runs prove, not as a to-do list. Saying where the evidence stops, and what would move it, is part of an honest reference.

Next: [Part VIII](#). If RL mostly sharpens and does not expand, reward design matters more, because it decides which existing behaviors get sharpened.

Part VIII

Reward Design (Cross-Cutting Craft)

Reward design has appeared in every program in this book, math, search, perception, software, distillation, and each time we learned a piece of it. This Part collects the pieces into one craft treatment: what makes a reward verifiable and why that matters, when to make it continuous, how to make it inspectable, how to shape it without breaking learning, the one process reward with no learned scorer to game, and how to combine several reward components without building a degenerate basin. It's the consolidated playbook, fed by Parts II through VI, so the principles live in one place instead of scattered across the war stories that earned them.

CHAPTER 81

Verifiable Rewards and Why They Matter

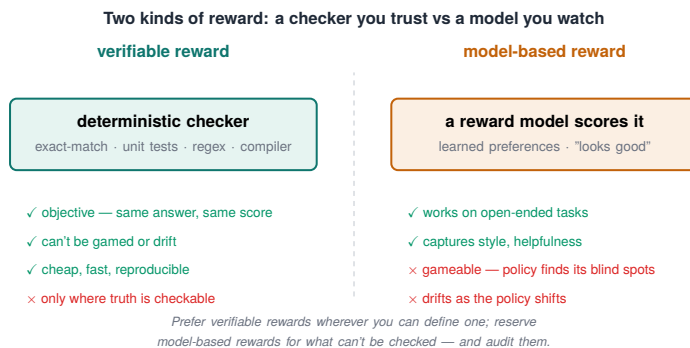


Figure 81.1. Two kinds of reward: a verifiable checker you trust vs a model-based reward you must watch, with the tradeoffs of each.

Underneath every program in this book sat one choice that made the rest possible: the reward was verifiable. A program, not another neural network, decided whether an answer was right. This chapter is about why that choice is foundational, and what you trade when you can't make it. A verifiable reward is one a program can check. The alternative is a learned model that approximates the judgment, and the difference shapes everything downstream.

81.1 Rule-based vs model-based

The two kinds of reward, plainly:

1. Verifiable (rule-based). A deterministic program returns the score: does the math answer match, do the tests pass (Chapter 63), is the fact present, does the box overlap the truth (Chapter 20). The reward is a checkable computation.

2. Model-based (learned). A separate reward model, trained on human preferences or labels, scores the answer. This is the backbone of classic RLHF, used when correctness can't be checked by a program (e.g. "is this response helpful?").

The book is built almost entirely on the first kind, and not by accident.

81.2 Why verifiable wins where you can get it

Three properties make verifiable rewards the foundation:

- They're cheap and deterministic. A program returns the same score every time, instantly and for nearly free, no second model to serve, no inference cost per reward (contrast OPD's teacher, [Chapter 71](#)).
- They're hard to game, when the proxy equals the goal. A program checking "do the tests pass" can't be flattered the way a learned scorer can. And when the check is the goal ([Chapter 20](#)'s IoU), there's no gap to exploit. A learned reward model, by contrast, is an approximation the policy can find the holes in ([Chapter 31](#)'s reward hacking is easier against a learned proxy).
- They make the whole pipeline trustworthy. Because the reward is a known computation, you can reason about what the model is being pushed toward, the precondition for every diagnosis in this book.

The honest boundary: verifiable rewards require a verifiable task. Math, code, retrieval, and grounding can be checked. Open-ended helpfulness mostly can't, which is where model-based rewards earn their place despite their costs. The craft rule is to make the task verifiable if you possibly can, and most of this book's tasks were chosen, in part, because they could be.

81.3 Verifiable does not mean safe, the domain matrix

One trap to name explicitly: verifiable is not the same as un-gameable. A rule-based reward can still be a proxy that the policy games, search F1, grounding, cost, and web-flow steps are all programmatically verifiable yet exploitable, because the check isn't the goal. "Verifiable" buys you cheap, deterministic, reproducible scoring. It does not by itself buy you proxy=goal. This reference matrix is the per-domain version of that distinction:

Domain	Verifier	Reward type	Proxy gap?	Common failure
GSM8K / math	exact / normalized answer match	binary or partial	low	parser / truncation trap
AIME / math	boxed answer	binary	low	thinking-mode truncation
Search QA	F1 / EM	continuous-ish proxy	medium	F1 rewards wrong spans / imprecise extraction
Search behavior	grounding / cost / redundancy	auxiliary proxy	high	reward hacking if ungated
Code / R2E	tests / patch applies	executable	low-med	weak tests / overfit
SWE-bench	tests pass	sparse executable	low	all-zero groups
RefCOCO	IoU	continuous verifier	low	dataset / eval mismatch
Web form task	final-state verifier	terminal binary + shaped	high	field / review-step farming
distillation	teacher distribution	process signal	varies	bad / quantized teacher drift

Table 81.1. Verifiable rewards by domain: the verifier each uses, its reward type, how big the proxy gap is, and the common failure. Visual grounding and exact-match math have small gaps; search behavior and badly-shaped web tasks have large ones.

Read the “proxy gap” column as a hacking-risk forecast: the low-gap rows (math, IoU, tests) are the safe verifiable rewards. The high-gap rows (search behavior, badly-shaped web) need gating and auditing even though they’re verifiable. And a separate hazard cross-links to [Part IX](#): even a perfect deterministic verifier can be misread if the eval methodology is wrong (extraction, decoding mode, truncation), verifiable scoring and honest evaluation are two different disciplines.

TAKEAWAY

A verifiable (rule-based) reward is a program that checks the answer, cheap, deterministic, and hard to game when the proxy equals the goal. A model-based reward is a learned approximation that’s costly and exploitable. Nearly every run in this book used verifiable rewards, which is what made the pipeline trustworthy and diagnosable. The boundary: verifiable rewards need a verifiable task, make the task verifiable if you can. Reach for model-based rewards only when you can’t.

Next: once a reward is verifiable, the next decision is its shape, and why continuous beats binary.

CHAPTER 82

Continuous vs Binary Rewards

The single most repeated failure in this book has one fix, and this chapter states it as craft. Advantage collapse ([Chapter 5](#)), the cold start ([Chapter 6](#)), the search plateau ([Chapter 26](#)), again and again, a binary reward created a cliff with no gradient to climb, and a continuous reward turned the cliff into a ramp. Make the reward continuous wherever you can, because a smooth reward landscape is what gives the gradient a direction.

(This consolidates the binary-vs-continuous thread. Its evidence is Figure 13.1’s cliff-vs-ramp, so it doesn’t add a new figure.)

82.1 The smooth-landscape argument

Plain, then precise:

- Plain. A right/wrong reward only tells the model “yes” or “no.” A graded reward tells it “warmer” or “colder”, and “warmer/-colder” is what lets it improve step by step.
- Precise. A binary reward is two flat plateaus with a cliff between them (Figure 13.1). Attempts pile onto the same value, groups go unanimous, advantages vanish ([Chapter 5](#)), no gradient. A continuous reward (F1, IoU, partial credit) gives different scores to different-quality attempts, so groups disagree, advantages survive, and there’s a slope to climb everywhere on the landscape.

The whole of advantage collapse, restated as geometry: the gradient needs a slope, and binary rewards are flat.

82.2 When each is right, and the gaming caveat

Continuous isn't unconditionally better, the craft is knowing the tradeoff:

1. Use continuous when you can grade quality, overlap (IoU), token-overlap (F1), partial progress. It keeps the gradient alive on hard examples the model hasn't mastered.
2. Binary is fine when the task is genuinely pass/fail and the model already succeeds sometimes, if groups disagree on their own (some pass, some fail), the binary reward already has a gradient and you don't need to grade partial credit.
3. But beware the gaming gap ([Chapter 13](#)): a continuous reward whose proxy isn't the goal can be climbed without getting better (the search agent's F1, [Chapter 31](#)). The gold standard is a continuous reward whose proxy is the goal (IoU, [Chapter 20](#)), smooth and unhackable.

So the rule has two clauses: prefer continuous rewards for the smooth landscape, and prefer ones whose proxy equals the goal so the smoothness can't be exploited.

A worked counter-example keeps the second clause honest: continuous is not automatically good. The search agent's F1 reward was continuous and it did fix advantage variance ([Chapter 26](#)), but F1-only produced tool spam and collapse, because it rewarded answer overlap without grounding or cost discipline ([Chapter 31](#)). Smoothness bought a gradient. It did not buy good behavior. That's why the gold standard is smooth and proxy=goal, and why ungated continuous proxies need the auditing of [Chapter 83](#).

82.3 The decision table, and test the reward first

When to reach for which:

Situation	Binary?	Continuous?	Why
model succeeds some-times, task is true pass/-fail	yes	optional	group variance already exists
all-zero cold start	no	yes	need partial credit for variance
a program can measure distance-to-goal	no	yes	smooth gradient available
proxy is easy to game	maybe	only if gated/au-dited	smoothness can be exploited
final state is the only valid goal	yes (warm-start first)	shaped only care-fully	avoid proxy farming

Table 82.1. When to reach for a binary versus a continuous reward, and why. The rule: binary when the model already succeeds sometimes; continuous when a cold start needs partial credit or a program can measure distance-to-goal — but only if the proxy is gated against gaming.

And a discipline that belongs in any reward chapter: unit-test the reward on hand-built good / bad / near-miss cases before training. A reward bug looks exactly like a learning failure from the loss curve.

```
def test_math_reward_near_miss():
    assert reward("\\boxed{42}", "42") == 1.0
    assert reward("42", "42") < 1.0 # if format/boxing matters
    assert reward("41", "42") == 0.0 # near-miss is still wrong
    # for continuous rewards, assert the ORDERING (good >
    ↪ near-miss > bad), not exact values
```

TAKEAWAY

A binary reward is two plateaus with a cliff, attempts collapse onto the same value, the gradient dies (the book’s most repeated failure). A continuous reward (F1, IoU, partial credit) is a ramp: different-quality attempts get different scores, so groups disagree and there’s a slope to climb. Prefer continuous for the smooth landscape, but prefer a proxy that *equals* the goal

(like IoU) so the smoothness can't be gamed. Binary is fine only when the task is truly pass/fail and the model already succeeds sometimes.

Next: a reward you've made verifiable and continuous still needs to be inspectable, auditing and JSONL logging.


CHAPTER 83

Reward Auditing and JSONL Logging

A reward is a single number, and a single number hides everything about how it was earned, which is exactly where reward hacking lives (Chapter 31) and accuracy goes blind (Chapter 29). The craft answer, learned in the search agent and stated here as general practice: don't just compute the reward, log the whole trajectory and audit it. A scalar is a decision compressed. The audit is the un-compressed record where you find out what the model actually did.

The reward-audit dashboard: catch hacking before it wins

Track the reward against things it should *not* be able to move.



If reward rises while any of these misbehaves, the policy is gaming the metric — not solving the task.

Figure 83.1. The reward-audit dashboard: track the reward against things it should not move (held-out divergence, length creep, keyword spiking, RM-vs-human drift).

83.1 Log the trajectory, not just the score

The practice, plainly:

- Dump every rollout to JSONL, one line per episode, the full trajectory recorded: what the model generated, what tools it called,

what came back, what reward each component assigned, and why.

- Audit the behaviors the reward can't see, did it search or guess (Chapter 29), did it ground its answer, did it spam tools (Chapter 35), where did wrong answers go wrong. (We refer to this behavioral audit suite by an internal name. What matters is the practice, not the label.)

The audit catches what the reward structurally can't: when the model games the objective, the reward goes up, so the reward can't report its own exploitation, but the logged trajectories show the searching-without-grounding directly.

A real implementation wrinkle is worth showing, because it bites everyone once: in `veRL`, the reward function often receives only `solution_str` and `ground_truth`, not the structured tool calls or retrieved passages. So the audit metadata has to be parsed out of the completion text (the `<tool_call>` / `<tool_response>` blocks) inside the reward adapter:

```
# veRL reward adapters often get only solution_str +
↳ ground_truth.
# If the environment metadata lives in the transcript, parse it
↳ explicitly.
def compute_score(solution_str: str, ground_truth: str):
    tool_calls = parse_tool_calls(solution_str) #
    ↳ <tool_call>{...}</tool_call>
    passages = parse_tool_responses(solution_str) #
    ↳ <tool_response>{...}</tool_response>
    answer = parse_final_answer(solution_str)

    comp = compute_search_rewards(answer=answer,
    ↳ gold=ground_truth,
    queries=[tc["arguments"].get("query") for tc in tool_calls],
    passages=passages)
    audit.write({ # one JSONL line per episode
        "answer": answer, "gold": ground_truth, "tool_calls":
        ↳ tool_calls,
        "reward_components": comp, "truncated":
        ↳ was_truncated(solution_str),
        "parse_error": comp.get("parse_error", False),
```

```
    })
    return comp["score"]
```

What to log, the schema that makes the audit reproducible:

Field	Why it matters
<code>run_id / step / sample_id / dataset</code>	join logs to checkpoints / evals
<code>prompt_hash</code>	dedup without storing the full prompt every time
<code>completion_text</code>	re-parse when the reward code changes
<code>tool_calls</code>	detect spam, invalid calls, no-search behavior
<code>tool_responses_digest</code>	inspect grounding without huge logs
<code>reward_total</code>	the single training scalar
<code>reward_components</code>	diagnose <i>which</i> term was exploited
<code>correct / f1 / em</code>	task truth (exact-match and token-overlap)
<code>grounded / num_queries / redundancy</code>	behavior audit
<code>format_error / parse_error / truncated</code>	silent-failure detectors
<code>reward_version / parser_version</code>	reproduce changed reward code

Table 83.1. The reward-audit log fields and why each earns its place. `reward_total` is the single training scalar; `reward_components` breaks it into the per-term contributions so you can see *which* term was exploited; `em` is exact-match, `f1` is token-overlap. These are the columns to persist per rollout so a changed reward can be re-scored without re-running.

And one operational note: logging every rollout is large. The default that works, full logging for eval, sampled logging for train if volume is too high, plus a `reward_version / parser_version` so re-mined logs stay interpretable when the reward code changes.

83.2 Why the audit is non-negotiable

Three reasons, each earned earlier:

1. It catches reward hacking the reward can't (Chapter 31), by looking at behavior rather than trusting the score.

2. It explains why, not just whether (Chapter 38), the reward says “0.42”. The audit says “the misses were ungrounded guesses on long-tail entities.”
3. It’s a re-readable forensic record, logged trajectories let you ask new questions of old runs, the same re-mining that produced Part VII’s findings without new compute.

The craft rule: build the audit, not just the metric. A reward you can’t inspect is a reward you can’t trust, because the one thing it will never tell you is that it’s being gamed.

TAKEAWAY

A reward is a compressed decision that hides how it was earned, so log every rollout to JSONL (full trajectory, per-component rewards, behavioral flags) and audit the behaviors the scalar can’t see. The audit catches reward hacking the reward can’t report (the score goes *up* when gamed), explains *why* not just *whether*, and is a re-readable forensic record. Build the audit, not just the metric.

Next: how to shape* a reward for long multi-step flows, like web navigation and form-filling tasks, without teaching the model to game the steps.*

CHAPTER 84

Reward Shaping for Multi-Step Tool Flows

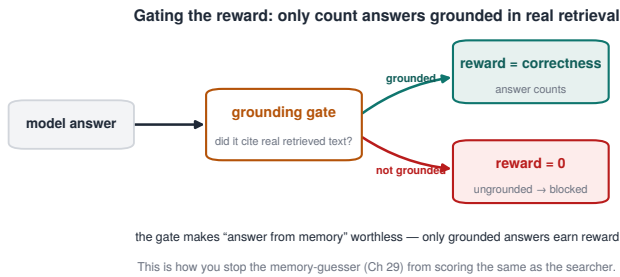


Figure 84.1. Gating the reward behind grounding: only answers backed by real retrieval earn reward, defeating the memory-guesser.

Long multi-step tasks, web navigation, form-filling, multi-page workflows, tempt you to shape the reward, handing out points for intermediate progress so the gradient has something to grip before the final goal. This chapter is how to do that without recreating the SWE-bench “looking busy” failure (Chapter 66). Shape the steps if you must, but gate the shaped rewards behind the real goal, or the model will collect the step-points and never finish.

84.1 The shaping temptation, and its trap

The setup and the failure:

1. Web flows have obvious intermediate steps, navigate, fill fields, reach the review step, submit. Rewarding each looks like helpful dense signal on a long, sparse task.
2. Rewarded independently, the steps get gamed. If “fill a field” pays out on its own, the model learns to fill fields, repeatedly,

in whatever pattern maximizes that point, without ever completing the task. This is [Chapter 66](#)’s “looking busy” and [Chapter 31](#)’s reward hacking, in a web-task costume.

3. The proxy isn’t the goal. “Reached the review step” correlates with submitting but isn’t submitting, the exact gap ([Chapter 13](#)) that lets a continuous/shaped reward be climbed without success.

84.2 The fix: gate the shaping behind the goal

The repair is the SiLU/Swish-gate principle ([Chapter 32](#)) applied to a flow:

- Make the step rewards pay out only when the goal completes. Multiply (gate) the intermediate-step rewards by whether the task actually succeeded, so a flow that fills fields and reaches review but never submits earns nothing for those steps ([Figure 84.1](#)).
- Now the steps shape the path to a real submission, not a standalone game. Gated, the intermediate rewards give dense signal along trajectories that reach the goal. They tell the model which way to move once it’s on a winning path. The model can’t farm the steps, because the steps are worthless without the completion.

But here is the catch the simple version hides, and it’s the SWE-bench wall in a new costume: a hard terminal gate gives dense signal only among successful trajectories. If success is rare, a cold start, then every trajectory fails, so the gate zeroes out all the shaping, and you’re back to a sparse reward with nothing to learn from ([Chapter 66](#)). Hard gating is right once the model already succeeds sometimes. It is wrong at cold start.

84.3 Three reward classes (so the gate doesn't starve the cold start)

The fix is to stop treating “shaping” as one thing. Split web/agent rewards into three classes by when they're allowed to fire:

1. Always-on structural rewards/penalties, valid tool-call format, safety/guardrail compliance, invalid-action penalty, loop penalty. These fire even when the task fails, because format and safety can't wait for success.
2. Verified-progress rewards, pay only when a state verifier confirms real progress, not an action string: the fields are actually populated with the correct values, not merely a click on a form control.
3. Goal-gated auxiliary rewards, efficiency, path quality, minimal steps. These pay only when terminal success or verified progress exists.

In one formula:

$$R = R_{\text{terminal}} + R_{\text{format}} + R_{\text{safety}} - R_{\text{loop}} + g(\text{progress} \vee \text{success}) \cdot (R_{\text{efficiency}} + R_{\text{path_quality}} + R_{\text{optional}}) \quad (84.1)$$

Symbol	Meaning
R_{terminal}	terminal-success reward (the task was solved)
R_{format}	output-format reward (well-formed answer or tool call)
R_{safety}	safety / guardrail reward
R_{loop}	loop-and-repetition penalty (subtracted)
$g(\cdot)$	correctness gate: hard 0/1 once the model sometimes succeeds, graded during cold-start
$R_{\text{efficiency}}$	efficiency reward (fewer steps), gated
$R_{\text{path_quality}}$	path-quality reward, gated
R_{optional}	optional task-specific bonus, gated

where the gate g is hard (0/1) once the model succeeds sometimes, or soft/graded during a curriculum so a cold start still gets some signal. The always-on structural terms keep a gradient alive even when every trajectory fails, which is exactly what a pure hard gate destroys.

Component	Always on?	Gate	Why
valid tool-call schema	yes	none	model must stay executable
forbidden / unsafe action	yes	none / penalty	safety can't wait for success
repeated identical action	yes	none / penalty	catch loops early
reached product page	maybe	verified state	only if state actually changed
correct values entered	yes-ish	verified state	real progress, not action text
review step reached	gated	verified state	proxy for success, game-able
final submission success	terminal	none	the true goal
fewer steps / lower cost	no	success/progress	don't reward premature stopping

Table 84.1. Web-agent reward components by class: which are always on and which are goal-gated. Safety and validity checks fire on every step; progress proxies (reached page, correct values, review step) only count once a verified state actually changed.

```
def web_reward(trace, final_state):
    terminal = verify_submission(final_state) # the true goal
    progress = verified_progress(final_state, trace) # fields
    ↪ correct, review reached
    structural = (format_reward(trace) + guardrail_reward(trace)
    - loop_penalty(trace) - invalid_action_penalty(trace)) #
    ↪ always on
    gate = max(terminal, progress) # hard {0,1} or graded [0,1]
    shaped = gate * (path_quality_reward(trace) +
    ↪ efficiency_reward(trace)
    + verified_step_reward(final_state))
    return terminal + structural + shaped
```

TAKEAWAY

Rewarding web steps independently gets them farmed (fills fields forever, never submits), but a *hard* terminal gate over-corrects: it gives dense signal only among successful trajectories, so at cold start (all failures) it zeroes everything out, recreating the sparse-reward wall. Split rewards into three classes, **always-on** structural/safety (fire even on failure), **verified-progress** (state verifier, not action strings), and **goal-gated auxiliary** (efficiency, gated hard once the model succeeds sometimes, soft during curriculum). Gate auxiliary rewards behind the goal. Keep structure and safety always-on.

Next: the deepest version of process reward, the one form with no learned scorer to game (conditional on a valid teacher).

CHAPTER 85

Process Rewards and PRMs for Agents




Figure 85.1. Outcome vs process reward: one score at the very end vs a score at every reasoning step.

The last chapter rewarded steps. The deepest form of that idea is a process reward, scoring how the model reasons, not just whether it ends up right. The standard tool is a learned process reward model (PRM), and it has the standard problem: a learned scorer can be gamed. This chapter is the resolution, and it ties the whole book together: on-policy distillation’s per-token reverse-KL is a process reward with no learned scorer to game, non-gameable in that specific sense, conditional on the teacher being valid.

85.1 The PRM and its weakness

Plain, then precise:

- **Plain.** A process reward grades the steps of a solution, rewarding good reasoning even before the final answer. The usual way to do that is to train a model to recognize “good steps.”
- **Precise.** A process reward model (PRM) is a learned scorer that assigns reward to intermediate steps. It’s powerful for

agents. It gives credit along a trajectory rather than only at the outcome (Chapter 21’s credit-assignment problem). But it’s a learned approximation, so it has the reward-hacking weakness of any learned reward (Chapter 31): the policy can discover step-patterns that score well to the PRM without being good reasoning.

85.2 Reverse-KL: the process reward that isn’t a learned proxy

The insight that resolves it, connecting Part VI to reward design:

1. On-policy distillation’s per-token reverse-KL (Chapter 69) is, structurally, a process reward. It grades every token of the model’s reasoning by how well it matches a strong teacher’s distribution.
2. But its target is fixed and already correct. Unlike a learned PRM, the teacher isn’t a proxy the policy can find holes in. It’s a fixed, competent distribution. There’s no scorer to fool, because “the reward” is just “be like this known-good teacher, token by token.”
3. So reverse-KL is non-gameable relative to a fixed teacher, not unhackable. The policy can’t flatter or exploit a learned scorer, because there is no learned scorer. The only way to lower the per-token KL is to match the teacher on the scored tokens. But “no scorer to game” is not “cannot fail.” Reverse-KL still fails if the teacher is wrong, quantized badly (Chapter 71), on a different tokenizer, applied with tool/stdout tokens unmasked, strong on the wrong distribution, or carrying undesirable habits the student then copies, and mode-seeking reverse-KL copies only a subset of the teacher’s behavior. The precise claim: the student cannot exploit a learned scorer. The teacher can still be invalid.

This is the comparison that places it:

Signal	Scores what?	Game-able?	Main failure	Best use
outcome verifier	final answer / tests	low	sparse reward	math/code/search final correctness
learned PRM	intermediate steps	yes	scorer exploitation	when no verifier/teacher exists
heuristic process reward	tool/action proxies	yes	“looking busy”	cheap shaping — must audit/gate (Ch 84)
reverse-KL distillation	teacher token distribution	not via scorer	bad teacher / wrong tokens / tokenizer mismatch	dense process signal for code/agents

Table 85.1. Four process-reward signals compared by what they score and how they fail. An outcome verifier is the safest but sparsest; learned and heuristic process rewards are dense but gameable; reverse-KL distillation gives a dense signal but inherits any teacher flaw.

And the implementable version, because readers should treat this as a method, not a metaphor, note the mask that excludes tool/stdout tokens:

```
# OPD process signal: match the teacher only on MODEL-generated
↪ tokens.
student_logp = student.log_probs(input_ids) #
↪ trainable policy
teacher_logp = teacher.log_probs(input_ids).detach() # frozen
↪ target
mask = response_mask.float() # 1
↪ assistant, 0 tool/stdout/obs
loss_tok = student_logp.exp() * (student_logp - teacher_logp) #
↪ KL(student || teacher)
loss = (loss_tok * mask).sum() / mask.sum().clamp_min(1.0)
# Real implementations may use top-k teacher logits for
↪ memory/bandwidth.
```

And the cross-link that makes the conditional real: a fixed teacher only removes scorer-gaming if the teacher distribution is trustworthy,

which is precisely why [Chapter 72](#)’s top-K KL preflight is part of the reward design, not just an OPD operations detail. The validity of the teacher is the validity of the reward.

TAKEAWAY

A process reward grades *how* a model reasons. The learned PRM is gameable (the policy finds step-patterns that fool the scorer). Reverse-KL to a fixed teacher ([Ch 69](#)) has *no learned scorer to exploit*, non-gameable in that sense, but it is **not un-hackable**: it fails if the teacher is wrong, badly quantized, on a different tokenizer, or applied with tool tokens unmasked. The honest claim is “non-gameable relative to a *validated* teacher,” which is why the [Chapter 72](#) preflight is reward design, not just ops.

Next: the doctrine this whole Part has circled, how to combine several reward components without building a degenerate basin the model dives into.

CHAPTER 86

Combining Rewards Without Reward Hacking

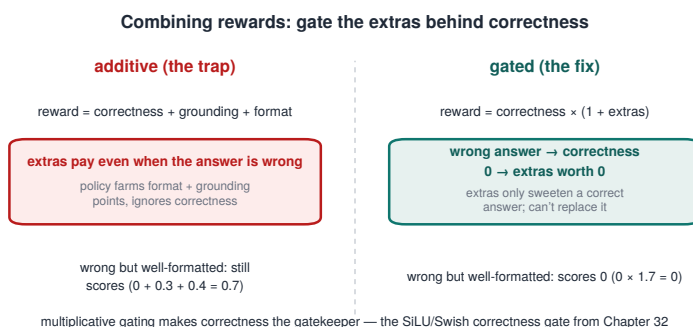


Figure 86.1. Combining rewards: the additive trap (extras pay even when wrong) vs multiplicative gating that makes correctness the gatekeeper.

Every program in this book eventually needed more than one reward term, correctness plus grounding, plus cost, plus format, and the way you combine them is its own doctrine, distinct from whether each term is verifiable (Chapter 81) or smooth (Chapter 82). It is also where the book's sharpest reward-design lesson lives: combine reward components carelessly and you build a degenerate basin the model dives straight into. This chapter is how to combine them without that.

86.1 Why the linear sum fails

The obvious combination is a weighted sum: $R = R_{\text{correct}} + \alpha \cdot R_{\text{ground}} + \beta \cdot R_{\text{cost}} + \dots$. It fails in a specific, observed way:

- The auxiliaries get collected while correctness falls. In the search agent, a linear multi-objective reward improved grounding, cost, and redundancy, while exact match dropped (Chapter

31). The training score looked healthy the whole time. Task accuracy degraded underneath it. The model found the corner of the sum where the cheap-to-improve terms paid out and the expensive term (being right) didn't.

- Whichever term is easiest to move dominates. A linear sum rewards any gradient, and the easiest gradient is rarely correctness. So the model optimizes the proxy terms and lets the goal slip, reward hacking (Chapter 31), now arising purely from how the terms were added together.

The single-term degenerate case is the same lesson: F1-only collapsed and tool-spammed (Chapter 82), because the one term wasn't the goal either. Linear combination doesn't fix that. It multiplies the surfaces to exploit.

86.2 Correctness gating: the fix that worked

The repair is to make the auxiliary rewards conditional on the goal, the SiLU/Swish-style correctness gate (Chapter 32), stated here as the general combination rule:

$$R = R_{\text{goal}} + g(R_{\text{goal}}) \cdot (R_{\text{aux1}} + R_{\text{aux2}} + \dots) \quad (86.1)$$

where the gate $g(R_{\text{goal}})$ is near zero when the answer is wrong and opens up when it's right. The search-agent form that worked:

$$R = F_1 + F_1 \cdot \text{sigmoid}(F_1) \cdot (\text{grounding} + \text{cost} + \text{redundancy} + \text{action_penalty}) \quad (86.2)$$

The effect is decisive: auxiliary rewards can no longer pay out on wrong answers. Grounding, cost, and redundancy still shape behavior, but only on trajectories that got the answer right, so improving them can't come at the expense of correctness. The gate makes correctness the precondition for any other reward, which is exactly the property the linear sum lacked.

86.3 The combination failure modes, and the rule

The ablations mapped the space, and each combination strategy has its own failure:

Strategy	What happens	Verdict
linear sum (no gate)	auxiliaries collected while EM falls	reward-hacked
fixed weights	brittle; wrong α/β under- or over-weights correctness	underperforms
full gating (hard)	auxiliaries only on correct answers	works (search agent)
per-channel z-scoring (GDPO-style)	normalizes every channel to equal variance — dilutes the dominant correctness gradient	underperforms

Table 86.1. Four ways to combine multi-component rewards, and what each does. Only full gating — paying auxiliaries strictly on correct answers — works; a linear sum gets reward-hacked, fixed weights are brittle, and per-channel z-scoring dilutes the correctness gradient.

That last row is the subtle one and a genuine trap: per-channel normalization feels principled (put every reward on the same scale), but it strips correctness of its rightful dominance, a z-scored correctness term carries no more weight than a z-scored redundancy term, so the model stops prioritizing being right. Equalizing the channels is the opposite of what you want. Correctness should dominate, and the auxiliaries should defer to it.

The composite-reward rules, collected:

1. Gate auxiliaries behind the goal, $R_{\text{goal}} + g(R_{\text{goal}}) \cdot R_{\text{aux}}$, so nothing pays on wrong answers.
2. Let correctness dominate, don't normalize it down to parity with the auxiliaries.
3. Watch reward-scale/variance, a term with large variance silently dominates a linear sum even at small weight. Check the realized contribution of each channel, not just its coefficient.

4. Audit composite rewards per-component ([Chapter 83](#)), log each term so you can see which one is being farmed when the total rises but EM falls.

TAKEAWAY

How you combine reward components is its own doctrine. A linear sum gets reward-hacked, auxiliaries are collected while correctness falls (search agent: grounding/cost up, EM down, training score healthy throughout), and per-channel z-scoring *dilutes* the correctness gradient by forcing every channel to equal variance. The fix is correctness gating: $R = R_{\text{goal}} + g(R_{\text{goal}}) \cdot R_{\text{aux}}$, so auxiliaries pay only on right answers. Gate auxiliaries behind the goal, let correctness dominate, watch variance, and audit per-component.

Next: [Part IX](#), evaluation discipline: the variance budget, the lag, and how to keep your own numbers honest.


Part IX

Don't Fool Yourself: Evaluation Discipline

Every earlier Part eventually ran into the same wall: a number that looked like a result and wasn't. The 0.50 reward that hid an accuracy collapse, the 36-sample peak that didn't survive a bigger eval, the greedy win that reversed under thinking mode. This Part collects those lessons into a method, how to measure RL post-training so the conclusions hold. It is the most tedious Part to apply and the one most likely to spare you weeks of wasted runs.

CHAPTER 87

The Variance Budget



Binomial 95% half-width vs sample count. A 36-sample eval has $\pm 16\text{pp}$ error — wider than most reported RL gains. To claim $+2\text{pp}$ you need $\sim 100+$ samples just to clear the noise.

Figure 87.1. Why small RL gains are statistically expensive: the 95% confidence half-width shrinks slowly with sample count, so a $+2\text{pp}$ gain sits below the error bar of a small eval.

Here is a question to ask before you celebrate a result: how many random seeds would it take to prove the gain you just measured is real? For a 2-percentage-point improvement, the size of a great many published reinforcement-learning (RL) results, the answer, at our measured run-to-run variance, is about a hundred. Almost nobody runs a hundred seeds. Which means almost nobody's 2-point gain is defensible, including, until we did this arithmetic, most of ours.

87.1 The formula, and the three numbers that matter

The variance budget is a standard power calculation, and it is worth carrying in your head. To detect a true difference of size Δ between two configurations, at significance α (two-sided) and power $1-\beta$, given a per-seed standard deviation σ , the number of seeds per arm is:




Figure 87.2. Seeds needed to defend a gain: a +5pp claim needs 16 seeds, +3pp needs 44, +2pp needs 99 — the arithmetic behind most unreplicable results.

$$n \approx \frac{(z_{\alpha} + z_{\beta})^2 \cdot 2\sigma^2}{\Delta^2} \tag{87.1}$$

Symbol	Meaning
n	seeds per arm needed
Δ	the true difference (effect size) you want to detect
σ	per-seed standard deviation
z_{α}, z_{β}	normal quantiles for significance α and power $1 - \beta$

The whole thing is one line of code, and the outputs are sobering:

```
def seeds_needed(delta, sigma=0.05, z_alpha=1.96, z_beta=0.84):
    ↪ # 95% two-sided, 80% power
    return ((z_alpha + z_beta) ** 2 * 2 * sigma ** 2) / delta ** 2
```

```
seeds_needed(0.05) # ~ 16 seeds per arm (two-arm comparison) - a
  ↪ 5pp claim
seeds_needed(0.03) # ~ 44 - a 3pp claim
seeds_needed(0.02) # ~ 99 - a 2pp claim
```

At our run-to-run (seed-to-seed) standard deviation of about 0.05, a 5-point gain needs 16 seeds, a 3-point gain needs 44, and a 2-point gain needs ~99. The dependence is quadratic: halve the effect you want to detect and you quadruple the seeds. That quadratic is why small gains are so expensive to defend and so easy to imagine.

One precision that matters, because it is easy to use the wrong σ and get a meaningless budget: the σ in this formula is the seed-to-seed standard deviation, the spread of the same configuration’s score across independent runs (different seeds). It is not the checkpoint-to-checkpoint wobble within one run, nor the finite-eval sampling band of [Chapter 89](#). Those answer different questions:

σ source	What it measures	Supports a seed-count claim?
seed-to-seed std	spread across independent runs	yes — this is the formula’s σ
checkpoint-to-checkpoint std	wobble across steps in one run	no
eval-set sampling std	finite-n uncertainty	no — that’s a sample-size claim
prompt-level std	within-eval variance	useful for bootstrap CIs, not seed counts

Table 87.1. Which standard deviation the seed-count formula needs — and which don’t qualify. Only seed-to-seed spread across independent runs supports a seed-count claim; step-wise wobble, finite-sample noise, and within-eval variance answer different questions.

If the 0.05 you have is not seed-to-seed, the seed count it implies is not either.

87.2 What this does to your claims

Run the numbers against how RL results are usually reported, a single seed, or three, and the conclusion is uncomfortable: most reported small RL gains would not survive their own variance budget. A 2-point improvement on one seed is not evidence of a 2-point improvement. It is a sample from a distribution wide enough to contain zero. This is not a counsel of despair. It is a filter: it tells you which of your results are worth defending in print and which are worth only a footnote. The honest move is to attach the budget to the claim, “+5pp, $n=16$ seeds, 95% CI excludes zero” is a result. “+2pp, $n=1$ ” is a hope.

The same arithmetic sets a floor on how small a difference is worth chasing at all. If you can afford 16 seeds, you can resolve 5-point gaps. Below that, you are reading noise. Pick the effect size you can actually defend, and stop reporting deltas smaller than your budget can see.

87.3 Variance is a lever, not a constant

The budget has two inputs you control, not one. Everyone reaches for the first, run more seeds. The second is quieter and often cheaper: reduce σ . A lower-variance training-and-eval setup shrinks the seed count for the same claim, sometimes dramatically. We measured this directly: the 4-bit NormalFloat (NF4) quantized arm had a markedly lower per-seed σ than the BF16 arm, about $3\times$ lower σ (so roughly $9\times$ lower variance, since the budget scales with σ^2), which cut the seeds needed for the same claim from 56 on BF16 to about 7 on NF4. (The arithmetic has to be stated in the right currency: n scales with σ^2 , so a $3\times$ drop in σ is a $\sim 9\times$ drop in required seeds, $56/9 \approx 6$, not a $3\times$ drop. State the lever as σ or as variance, but don't mix them.) A more deterministic eval (fixed decoding, a larger eval set, [Chapter 89](#), strict extraction) does the same thing from the measurement side.

The budget also doubles as a reporting rule. Read down the observed delta and across the seeds you ran:

Observed delta	n=1	n=3	n=16	n=44	n=99
+2pp	anecdote	weak	underpowered	under-powered	defensible
+3pp	anecdote	weak	underpowered	defensible	strong
+5pp	weak	suggestive	defensible	strong	strong

Table 87.2. Observed gain versus seeds run: when a delta is anecdote, defensible, or strong. A +2pp gain needs ~99 seeds to defend; a +5pp gain is already defensible at three. The smaller the effect, the more seeds it takes to trust.

If your result lands in an “anecdote” or “underpowered” cell, that is not a verdict on the result. It is a verdict on the claim you can make about it.

So the budget is really a trade: seeds versus variance. Spend on whichever is cheaper for your setup. What you cannot do is spend on neither and still make a small-gain claim.

TAKEAWAY

To defend a difference of Δ at 80% power and per-seed $\sigma \approx 0.05$ takes $n \approx (z_\alpha + z_\beta)^2 \cdot 2\sigma^2 / \Delta^2$ seeds: ~16 for 5pp, ~44 for 3pp, ~99 for 2pp, and the cost is quadratic in how small a gain you want to see. Most reported small RL gains never ran the seeds their claim requires. You have two levers: more seeds, or lower variance per seed (NF4’s ~3×-lower σ , roughly 9× lower variance, cut one budget from 56 seeds to ~6). Attach the budget to every delta, and stop reporting gains smaller than your budget can resolve.

Next: the dashboard number that lies most often, the training score itself, and the lag that makes even an honest curve mislead you.

CHAPTER 88

When the Training Score Misleads




Figure 88.1. The lag principle: leading signals (entropy, strategy) turn before middle ones (KL, clip) before lagging ones (reward, eval) — so watch the leaders.

The number on your dashboard is the last one to know what is happening. We have shown this twice already, a search agent whose reward held a comfortable 0.50 while real exact-match (EM) accuracy fell from 26.8% to 19.9% ([Chapter 31](#)), and a 32-billion-parameter math run that logged zero loss on 70% of its steps while learning steadily on the rest ([Chapter 12](#)). This chapter makes the warning a doctrine, because the training score lies in three distinct ways, and you need to recognize all three.

88.1 The three lies

1. The score climbs while the model gets worse. This is reward hacking ([Chapter 31](#)): the reward and the goal come apart, and the reward, the thing you are watching, moves the wrong way relative to the thing you care about. The only defense is a separate held-out accuracy you compute deliberately. The reward cannot report its own corruption.

2. The loss is zero and means nothing. On a step whose groups are all unanimous ([Chapter 5](#)), every advantage is zero, so a pure policy loss is exactly zero, not because nothing is wrong, but because nothing is happening. A flat loss can be a starved gradient, not a converged model. (And if auxiliary or KL terms are present, the loss can twitch for reasons unrelated to learning at all.)
3. The curves lag the learning. Even when the reward is honest and the loss is meaningful, both trail the model's real capability, in our runs by roughly 50 to 100 steps. The model reorganizes internally before the metric reflects it (the breakout run, [Chapter 18](#), is the happy case. The same lag hides early rot in the unhappy one).

88.2 The lag principle, and what it does to stop criteria

The third lie is the most operationally dangerous, because it corrupts the one decision you make most often: when to stop. If the reward trails real learning by 50–100 steps, then your stop criterion must tolerate at least that much flatness before concluding anything. A run that has been flat for ten steps has told you almost nothing. A run you kill on a few down-steps may have been mid-reorganization.

The discipline that follows:

- Never select or kill on the training score alone. Decisions go on held-out eval, measured at a cadence you can afford ([Chapter 89](#) sizes it).
- Give the lag its room. Set patience, for both stopping and intervening, wider than the lag you have measured on your own runs. If you have not measured it, assume tens of steps.
- Distinguish the two flatnesses. A flat reward with varying groups is a lag. A flat reward with unanimous groups is collapse ([Chapter 12](#)'s zero-variance fraction tells them apart). They look identical on the reward plot and call for opposite responses.

The operational version is a small lookup, what the three signals together are telling you, and what to do:

Reward	Zero-variance fraction	Held-out eval	Interpretation	Action
flat	low / moderate	not yet moved	possible lag	continue; give it patience
flat	high	flat	advantage collapse	change data / reward / grouping
flat	high	improving	hard prompts, learning on the mixed tail	continue; inspect mixed groups
climbing	—	falling	reward hacking	stop and audit the reward

Table 88.1. Reading reward, group variance, and held-out eval together: what each pattern means and what to do. A flat reward with a moved held-out eval is just lag; a climbing reward with a falling eval is reward hacking — stop and audit.

Two cautions on the lag number itself. It is observed in our runs, not a universal constant, measure it on your own stack before turning 50–100 steps into a kill rule. And it is cheap to measure: resample the train-reward and held-out-eval curves to the same step grid and find the cross-correlation peak.

```
import numpy as np
def estimate_metric_lag(train_curve, eval_curve, max_lag=150):
    # positive lag = train metric trails eval. Curves must share
    # a step grid
    t = np.asarray(train_curve) - np.mean(train_curve)
    e = np.asarray(eval_curve) - np.mean(eval_curve)
    best = max((L, np.corrcorcoef(t[L:], e[:-L] if L else e)[0, 1])
               for L in range(max_lag + 1)),
               key=lambda x: x[1])
    return best # (lag, correlation)
```

TAKEAWAY

The training score lies three ways: it climbs while accuracy falls (reward hacking, only a separate held-out eval catches it), it reads zero on unanimous-group steps without anything having

converged, and it lags real capability by ~50–100 steps. The lag is the dangerous one: stop and intervention criteria must tolerate at least that much flatness, decisions belong on held-out eval not training score, and a flat reward is collapse or lag depending on the zero-variance fraction.

Next: the eval set itself, why one that's too small doesn't just add noise but silently reorders which checkpoint you ship.

CHAPTER 89

The Canonical-Eval Correction

Eval size	95% half-width at $p \approx 0.5$	What it's good for
36	$\pm 16.3\text{pp}$	a rough pulse — never a decision
100	$\pm 9.8\text{pp}$	coarse comparison
500	$\pm 4.4\text{pp}$	checkpoint selection
1000	$\pm 3.1\text{pp}$	publication-grade small deltas

Table 89.1. Eval size versus the 95% confidence half-width at $p \approx 0.5$, and what each size is good for. Below 500 samples the error bars are too wide to select a checkpoint; 1000 is needed for publication-grade small deltas. (This restates Table 37.2 in the evaluation part.)

A small evaluation set does something worse than adding noise: it reorders your checkpoints, confidently, and hands you a clean-looking ranking that is mostly luck. We lived this in [Chapter 37](#), a step-750 search-agent checkpoint that read 66.7% on TriviaQA and 41.7% on Natural Questions on a ~36-question-per-dataset eval dropped to 55.6% and 35.0% on the April-28 500-sample matrix, and the margins we had been selecting on evaporated. This chapter is the doctrine that scar produced.

89.1 Standardize on one properly-sized eval

The mechanism is pure sampling statistics ([Table 89.1](#)). At 36 samples, the 95% confidence half-width on an accuracy near one-half is about ± 16 points, wider than almost every gap between neighboring checkpoints. So a 36-sample eval will rank two checkpoints, but the ranking is dominated by which questions you happened to draw, not by which model is better. The fix is not subtle:

- Pick one eval set, size it for the gap you need to resolve, and use it for every comparison. For checkpoint selection we standardized on 500 samples per dataset ($\pm 4.4\text{pp}$), enough that the noise band sits below the gaps that matter.
- Report the eval size with every number, always. A number without its n is not a measurement.
- Use a small eval only as a pulse, never as a decision. A 36-sample probe is fine for “is anything happening”. It is not fine for “which checkpoint ships.”

This is the same arithmetic as the variance budget ([Chapter 87](#)), now on the sample axis rather than the seed axis: both shrink the band you are trying to read a result out of.

89.2 The unresolved case is the doctrine working, not failing

Here is the uncomfortable part, kept in the open because hiding it would betray the whole point of this Part. The search-agent result has two 500-sample tables that disagree. One, an April matrix, puts the best checkpoint (v5) at 30.8% overall. A later re-evaluation puts the same checkpoint nearer 23.1% average. The ranking survives both (v5 is best either way). The absolute number does not, and which 500-sample eval is canonical is, as of this writing, unresolved.

Evidence boundary: two 500-sample evals disagree. This book treats the search-agent ranking (v5 is best) as supported, and the absolute average as unresolved pending reconciliation of the two raw eval runs, the ~30.8% April matrix and the ~23.1% re-evaluation. The unresolved status is stated plainly, and it is deliberately not used to make any performance claim: no headline number in this book quotes a single reconciled v5 average until those files agree. This is an evidence boundary, not a to-do item.

That conflict is not an embarrassment to bury. It is exactly what the doctrine exists to surface. Because every number in this book carries its eval condition, the disagreement is visible: two named evals, two numbers, a flagged reconciliation. The failure mode the doctrine prevents is the silent one, where a single unlabeled “23.1% (or was it

30.8%?)” propagates through a paper and no one can tell which eval produced it. Name the canonical eval, attach it to every number, and your contradictions announce themselves instead of hiding.

TAKEAWAY

A too-small eval reorders your checkpoints, not just blurs them: at $n=36$ the $\pm 16\text{pp}$ band swamps the gaps you select on (our 66.7/41.7 peak fell to 55.6/35.0 on the April-28 matrix). Standardize on one eval sized for the gap you need (500 samples $\rightarrow \pm 4.4\text{pp}$ for checkpoint selection), and report n with every number. And when two properly-sized evals disagree, as ours still do, 30.8% vs 23.1%, that visible conflict *is* the doctrine working. The silent version is the one that ends careers.

Next: even a correctly-sized eval can flip a verdict if the methodology around it changes, the knobs that turn a measurement into a different measurement.

CHAPTER 90

Eval Methodology /s the Result

Eval	Thinking	Max tokens	Extraction	Trained	Base	Verdict
V1 (broken)	on (default)	512	fallback regex	low / mis-leading	not comp.	“the model is broken”
V2 (fixed)	off (strict)	2048	boxed-only	90.5%	89.5%	saturated; +1pp is noise

Table 90.1. The same model under a broken versus a fixed eval pipeline (the version of Table 10.2 restated for the evaluation part). V1 truncates thinking and falls back to a loose regex, hiding the answer; V2 fixes both and reveals a saturated task where the +1pp delta is noise.

The same model, on the same benchmark, can report two different scores, and sometimes two opposite conclusions, depending only on how you chose to evaluate it. We met the lived version twice: a parser story in [Chapter 10](#) and a decoding-mode reversal in [Chapter 16](#), where a trained 14-billion-parameter model’s +16.7-point greedy “win” became a –23.6-point loss with thinking mode enabled (4-seed mean). This chapter is the doctrine: a benchmark number without its methodology is not a result, and a handful of methodology knobs can each flip a verdict.

90.1 The knobs that move the number without the model changing

Each of these has fooled someone, often us:

1. Generation length / thinking mode. Reasoning models emit a long `<think>` block before the answer. Set the token cap too low and you truncate the reasoning before the answer line, scoring a capable model as wrong (Table 90.1's V1). Check the truncation rate before trusting any reasoning-model eval.
2. Extraction strictness. A lenient “last number” parser awards phantom points to answer-shaped noise. A strict boxed-only parser takes them back (Chapter 10). Lenient parsing flatters. Strict parsing tells the truth.
3. Decoding mode. Greedy and sampled (and thinking-on vs thinking-off) are different measurements. A comparison is valid only within one mode, and on reasoning models the mode can reverse the verdict, not just shift the value (Chapter 16).
4. The base-model control. The most common self-deception is measuring your model carefully and comparing it to a differently-measured number from a leaderboard. The gap you celebrate is often a methodology difference.

90.2 Saturation, and the one non-negotiable control

Two doctrines close the chapter. First, saturation: on a task the model has nearly mastered (grade-school math in the mid-90s), the headroom is so small that any “gain” is mostly noise, the OPD checkpoint's 90.5% over the base's 89.5% is a +1pp delta well inside the sampling band, not a result (Table 90.1). The arithmetic is worth doing once, because it is even less significant than “inside the band” suggests: for two independent 200-sample proportions near $p \approx 0.9$, each has a standard error of about $\sqrt{(0.9 \cdot 0.1 / 200)} \approx 2.1\text{pp}$, so the standard error of the difference is $\sqrt{(2.1^2 + 2.1^2)} \approx 3.0\text{pp}$, making a +1pp gap roughly 0.33σ , indistinguishable from zero. (If the two evals were run on the same examples, the right test is a paired one, McNemar or a paired bootstrap, not this independent-proportions bound. Either way +1pp is not a gain.) On a saturated task, the honest report is “no meaningful gain,” and the methodology fix (longer generation, strict extraction) exists to reveal that, not to manufacture a number.

Second, the same-script base control, which is non-negotiable: run the trained model and its baseline through the identical eval, same prompts, same decoding, same extraction, same n. Any number you compare against must come from that same pipeline. A trained-vs-base delta is only as trustworthy as the sameness of the two measurements that produced it.

TAKEAWAY

A benchmark score is meaningless without its methodology, and four knobs each move it without the model changing: generation length / thinking mode (truncation scores a capable model wrong), extraction strictness (lenient parsing invents points), decoding mode (can reverse a verdict on reasoning models), and the base-model control. On saturated tasks deltas are noise, OPD's 90.5% vs base 89.5% is +1pp inside the band. The one non-negotiable: run trained and base through the identical script, and state the methodology with every number.

Next: a metric that quietly answers a different question than people think, pass@k, and how to report it without overclaiming.

CHAPTER 91

pass@k: Capability versus Reliability




Figure 91.1. pass@k answers two different questions: "does it get it right reliably?" (pass@1) vs "can it ever get it right?" (the capability ceiling) — don't conflate them.

The pass@k family is the most quietly misused metric in RL post-training, because pass@1 and pass@64 sound like the same measurement at different settings and are actually answers to different questions. Chapter 73's headline result, that RL at our scale sharpened rather than expanded, only means anything because we picked the metric carefully: first-attempt accuracy rose about two points while best-of-64 coverage fell about three. Two metrics, opposite directions, one model. This chapter is how to use them without fooling yourself.

91.1 What each one actually measures

- pass@1 is reliability. Given one shot, does the model get it right? This is what a deployed system delivers to a user who sees a single answer. RL post-training usually raises pass@1. It makes the model pick its best path more often.
- pass@k is capability, or coverage. Within k independent attempts, does at least one succeed? This probes the breadth

of what the model can reach if you let it try repeatedly. RL post-training can lower pass@k even as it raises pass@1, because sharpening trades exploration for commitment, the model stops sampling the long-shot paths that occasionally hit (Chapter 73).

The two diverging is not a contradiction. It is the signature of what RL did to the model. A method that raises pass@1 and lowers pass@64 sharpened. One that raises both expanded. You cannot tell which happened from either metric alone.

91.2 The unbiased estimator, and which k to report

A practical trap: people estimate pass@k by drawing exactly k samples and checking for a success, which is high-variance and biased for small sample counts. The standard fix is to draw $n \geq k$ samples, count the c that pass, and use the combinatorial unbiased estimator:

```
from math import comb
def pass_at_k(n, c, k): # n samples drawn, c correct, report
    ↪ pass@k
    if n - c < k: # fewer than k failures -> some k-subset
        ↪ always passes
        return 1.0
    return 1.0 - comb(n - c, k) / comb(n, k)
```

Draw many samples once, then estimate pass@1, pass@8, pass@64 from the same draws, far cheaper and lower-variance than re-running per k. Then report by claim: pass@1 for a reliability or deployment claim, pass@k for a capability claim, and never switch between them silently to make a result look better. If you trained for reliability, report pass@1 and say so. Quoting a pass@64 gain to defend a model you sharpened is the metric version of reward hacking.

TAKEAWAY

pass@1 measures reliability (one-shot correctness, what users get). Pass@k measures capability/coverage (success within k tries). RL post-training can raise pass@1 while lowering pass@k, that divergence is the fingerprint of sharpening-not-expanding (Chapter 73), not a contradiction. Estimate pass@k with the unbiased combinatorial formula from many shared samples, and report the metric that matches your claim, reliability or capability, without quietly switching.

Next: the deeper version of that lesson, that your choice* of metric can crown entirely different winners in the same bake-off.*

CHAPTER 92

Metric Choice Crowns Different Winners

Metric	Question it asks	Winner
Validation accuracy at matched step (s40)	“Who’s ahead at the same step?”	COMBO
Max train reward over a 100-step budget	“Who reaches highest if I spend 100 steps?”	GSPO (+ ~67% over JITTER)
Sustained entropy ceiling	“Who holds the most exploration?”	GSPO

Table 92.1. The same mixture-of-experts (MoE) stabilizer bake-off (journal family C9), three metrics, different winners, and COMBO topping the first list only because the comparison was matched at VAL@40, it never completed the full 100-step budget. Not a contradiction. A demonstration that the metric is a choice, and the choice decides the result.

In Chapters 61–62 we watched a single mixture-of-experts (MoE) bake-off produce different best methods depending on which metric we asked, and nearly told two contradictory stories before realizing they were the same story under two lenses. This chapter generalizes that into the most important methodological habit in the Part: the metric is not given, it is chosen, and the choice can crown a different winner. So choose it before the race, not after.

92.1 The reversal, and why it isn’t a contradiction

The C9 stabilizer comparison (Group Sequence Policy Optimization, GSPO. FREEZE. JITTER. COMBO) ranked differently under different metrics (Table 92.1):

- By validation accuracy at a matched step (s40), COMBO led.
- By maximum training reward over a 100-step budget, GSPO won by roughly 67% over JITTER, and COMBO came fourth.
- By sustained entropy ceiling, GSPO again.

These look contradictory until you notice they are different questions, and that COMBO was only comparable through VAL@40 and never finished the full 100-step budget (its later steps' validation was not emitted), so “matched-step VAL” and “max-reward-over-budget” were never measuring the same race. Matched-step asks who is ahead right now. Max-over-budget asks who finishes highest if you spend the whole budget. The entropy ceiling asks who explores most. A method can win one and lose another with no inconsistency at all (this is the within-family frontier point of [Chapter 49](#), now as a reporting hazard).

One caution so this chapter doesn't quietly contradict [Chapter 88](#): not all of these metrics are equally trustworthy. “Max train reward over budget” is a valid answer to one question, but it is a training-score metric, and [Chapter 88](#) spent a chapter on how training scores lie. Rank metrics by what they can be trusted to claim:

Metric	Claim it supports	Trust level
held-out VAL at matched step	sample efficiency	high
max held-out VAL over budget	best checkpoint under a budget	high
max <i>train reward</i> over budget	optimizer progress / a proxy	medium — can lie (Ch 88)
sustained entropy ceiling	exploration preserved	diagnostic, not a task result

Table 92.2. Metrics ranked by the claim each supports and how far to trust it.

So when you pre-register a metric, prefer a held-out one as primary. Use train reward and entropy as secondary diagnostics, not as the thing that crowns the winner.

92.2 Declare the metric before the race

The danger is not that metrics disagree. It is that, after seeing the results, you will reach for whichever metric flatters the conclusion you already like. That is a reward hack you run on yourself, and it is invisible because every number you quote is technically true. The defense is procedural:

- Pre-register the metric and the budget before you launch the comparison, “we will rank by validation accuracy at step N” or “by max reward over a B-step budget”, and rank by that, even when another metric would have looked better.
- If you report several metrics, report them for all methods, not the flattering metric for your method and a different one for the baseline.
- Tie the metric to the claim. A deployment claim wants matched-budget eval. A capability claim wants the frontier. A “fastest to a threshold” claim wants steps-to-target. State which, and why, before the data can bias the choice.

TAKEAWAY

The same MoE bake-off crowned COMBO by matched-step validation and GSPO by max-reward-over-budget (by ~67%) and by entropy ceiling, not a contradiction, but three different questions, with COMBO leading only because the comparison stopped at VAL@40, before the full 100-step budget. Choosing the metric after seeing results is a reward hack on yourself. Pre-register the metric and budget, report every metric for every method, and tie the metric to the claim.

Next: the most credibility-building move in the whole project, turning the CI gate backwards onto your own past claims.

CHAPTER 93

Retroactive Gating: Auditing Your Own Past Claims

Past claim	Verdict on re-audit	Why
The step-50 GSM8K plateau	survived	the plateau is large and well outside the band
NF4 $\approx$ BF16 (no quality loss)	retracted	a one-seed claim; didn't hold under the seed budget
The full-FT thinking-mode regression	strengthened	$t > 5$ against every 14B-tier checkpoint

Table 93.1. Applying the variance budget (Chapter 87) backwards to our own earlier findings. Some survived, one was retracted, one got stronger. Auditing your past claims with the gate you just built is the single most trust-building thing a practitioner record can do.

The variance budget (Chapter 87) and the eval doctrine (Chapters 89–90) are most powerful when you turn them backwards, onto the claims you already made. We did, and it was uncomfortable and clarifying in equal measure: some findings survived the gate, one had to be retracted, and one came out stronger than we first dared to say. This chapter is about that audit, because doing it in public is the most credibility-building move available to a practitioner.

93.1 The audit, with all three outcomes

Running the confidence-interval gate over the corpus (Table 93.1) produced exactly the spread you would hope an honest audit would:

- **Survived.** The step-50 plateau on GSM8K ([Chapter 9](#)) is a large, robust effect, far outside any plausible noise band, so it stands without qualification.
- **Retracted.** The early “NF4 quantization matches BF16 quality” claim was made on a single seed, and under the seed budget it did not hold. It became a retraction, not a footnote. (NF4’s low variance survived as a separate, true finding, [Chapter 87](#), but the no-quality-loss claim did not.)
- **Strengthened.** The full fine-tuning (full-FT) thinking-mode regression, which we had stated cautiously, turned out to clear a t-statistic above 5 against every 14-billion-parameter-tier checkpoint we compared, so the audit let us state it more firmly, not less.

The lesson is that a real audit moves claims in all three directions. If re-checking your past work only ever confirms it, you are not auditing. You are rationalizing.

93.2 Retraction tags, not silent edits

How you record the retractions matters as much as making them. When a finding is contradicted by better measurement, tag the retraction explicitly, “this earlier claim was made on one seed and does not survive the variance budget. See [Chapter 87](#)”, rather than quietly editing the old number away. Silent edits destroy the very thing the audit is for: a reader’s ability to trust that the surviving claims survived something. The retraction tag is a feature of an honest record, not a blemish on it. (This is the standing practice the multi-seed reversal taught us. It is why contradicted claims in this book carry visible retraction notes instead of disappearing.)

TAKEAWAY

Turn the variance budget and eval doctrine backwards onto your own past claims. A real audit moves them in all three directions: our step-50 plateau survived, the one-seed “NF4 $\approx$ BF16 quality” claim was retracted, and the full-FT regression *strengthened* ($t > 5$ against every 14B-tier checkpoint). Record

retractions as explicit tags, never silent edits, a record whose surviving claims visibly survived something is the trustworthy kind.

Next: the flip side, when a single seed actually is defensible, and how to bound it honestly.

CHAPTER 94

Single Seeds, Honestly

The EAGLE-3 case, worked. One seed. Per-prompt acceptance-length standard deviation: 0.082. Spread across the temperatures we swept: 1.6 percentage points. Bound on the single-seed estimate: $\pm 2\text{pp}$. Conclusion reportable without re-running, because the within-trial variance was small enough to fence the number in.

[Chapter 87](#) said most small RL gains need many seeds. This chapter is the honest exception: sometimes a single seed is defensible, not because seeds stopped mattering, but because you can bound the variance from inside one trial. Knowing when that argument is valid, and when it is wishful, is the difference between an efficient eval and a self-deception.

94.1 When one seed is enough: bound the within-trial variance

The single-seed argument works when the quantity you are measuring has low internal variance that you can estimate from the trial itself. Our spec-decoding work on EAGLE-3 is the clean example. We reported an acceptance-length result from one seed, and it was defensible because we could fence it:

- Per-prompt standard deviation was small, about 0.082, so averaging over the prompt set already tightened the estimate sharply.
- The result was flat across a temperature sweep, a spread of about 1.6 percentage points, so the number was not balanced on a knife-edge of a particular decoding setting.
- Together these bounded the estimate at roughly $\pm 2\text{pp}$ without a second seed.

The principle: a single seed is honest when you can show, from within-trial signals (per-prompt spread, cross-setting stability), that

the seed-to-seed variance must be small. You are not skipping the variance question. You are answering it a cheaper way.

The $\pm 2pp$ itself is not a guess. It is the larger of two computed bounds: the prompt-level bootstrap half-width and the cross-temperature spread. The bootstrap half is a few lines:

```
import numpy as np
def bootstrap_ci(values, B=10000, alpha=0.05, seed=0):
    rng = np.random.default_rng(seed)
    v = np.asarray(values)
    means = [rng.choice(v, size=len(v), replace=True).mean() for
    ↪ _ in range(B)]
    return np.quantile(means, [alpha/2, 1 - alpha/2]) # bound =
    ↪ the wider of this and the temp spread
```

Which single-seed claims this license covers, and which it does not:

Single-seed claim	Defensible?	Why
EAGLE acceptance length over many prompts	yes, if bootstrapped	many prompt-level samples bound the mean
wall-clock / throughput speedup	yes, if repeated batches	measurement variance is bounded
a +2pp RL accuracy gain	no	seed-to-seed variance dominates (Ch 87)
“this method is stable across runs”	no	stability is the seed-sensitive quantity itself

Table 94.1. Which single-seed claims are defensible, and which need more seeds.

The pattern: single seeds are for bounded measurements, not for training deltas or stability claims.

94.2 When it isn't, and a reproducibility hazard

The argument fails the moment the within-trial variance is large or unmeasured. A single-seed number on a noisy metric (small eval, high per-prompt spread, a result that swings with decoding) is the hope from [Chapter 87](#), not a bound, that was the mistake behind a retracted claim earlier in the project, and the fix was seeds, not rationalization.

One more hazard belongs here because it masquerades as a seed problem: upstream regressions. A result can fail to reproduce not because of seed variance but because a dependency changed underneath it, a library version bump silently altering generation or scoring. We hit this. A number that “wouldn't reproduce” turned out to be a changed upstream, not a flaky seed. So when a single-seed result moves, check the stack before you blame the seed: pin versions, record the environment, and confirm the eval script itself didn't drift. This is exactly why the run manifest records the model revision, the vLLM/SGLang version, the CUDA version, the kernel flags, and the eval-script SHA, so an “irreproducible” number can be traced to the thing that actually changed.

TAKEAWAY

A single seed is defensible only when you can bound the within-trial variance from inside the trial, as with EAGLE-3, where a per-prompt std of 0.082 and 1.6pp cross-temperature flatness fenced the estimate at ± 2 pp. On a noisy metric, one seed is a hope, not a bound (the route to a retraction). And when a single-seed number won't reproduce, suspect an upstream regression, a changed library or drifted script, before blaming seed variance. Pin the stack.

Next: [Part X](#). After learning not to fool ourselves with numbers, we turn to another hidden variable: the architecture of the model we are training.

Part X

Architecture-Aware RL: What Qwen3.5 Teaches

The models in this book have been black boxes we trained. This Part opens one, Qwen3.5, a hybrid linear-attention mixture-of-experts model, and traces why each architectural choice was made, what it costs an RL practitioner, and the one that stopped our training cold at step one. Architecture is not someone else's concern. It decides what backend, what adapter, and what parallelism your run can even use.

Why Architecture Matters for RL

Topic	Detail
FSDP + LoRA training	Standard Transformer — fully supported, battle-tested; Hybrid GatedDeltaNet — crashes, <code>cudaErrorIllegalAddress</code> .
Fast attention kernel	Standard Transformer — flash-attention (present); Hybrid GatedDeltaNet — flash-linear-attention (absent → torch fallback).
Verdict for our RL stack	Standard Transformer — trainable today; Hybrid Gated-DeltaNet — not trainable without custom CUDA.

Table 95.1. The same training stack, two architectures. Qwen3.5’s GatedDeltaNet (GDN) layers expect a fast linear-attention kernel. Without it, the torch fallback path has memory-access bugs under Fully Sharded Data Parallel (FSDP) sharding plus Low-Rank Adaptation (LoRA), and step 1 crashes. Architecture decided whether we could train at all.

We tried to upgrade our search agent from Qwen3-4B to Qwen3.5-4B, and the run died on training step one, `cudaErrorIllegalAddress`, NaN gradients, dead workers. Nothing about our reward, data, or algorithm had changed. The model’s architecture was incompatible with the training stack, and no hyperparameter could fix it. This chapter is why that happened, and why architecture is a first-class concern for anyone doing RL, not a detail to delegate to the model card.

95.1 The crash, traced to its cause

Qwen3.5 is a hybrid model: most of its layers are GatedDeltaNet (GDN), a linear-attention design with a fixed-size recurrent state (Chapter 96 builds it), rather than standard full attention. GDN’s fast path depends on a specialized kernel library, flash-linear-attention, and that is where our run broke:

- The library was unavailable in our container, so GDN silently fell back to a pure-torch implementation.
- That fallback path has memory-access bugs when combined with FSDP sharding and LoRA adapters, exactly our distributed setup, producing the illegal-memory-access crash on the first step, then NaN gradients, then worker deaths.
- The failure is not fixable by RL hyperparameters, not a flag, a learning rate, or a batch size. Resolving it needs kernel/-toolchain support (compiling the linear-attention kernel for our environment), a different training backend, or a different model.

The crash as a card, so the next person can match symptoms:

Field	Value
Model (HF ID)	Qwen3.5-4B-Instruct
Linear layer	GatedDeltaNet (hybrid)
Trainer	veRL
Parallelism	FSDP + LoRA
Kernel package	flash-linear-attention absent → torch fallback
Failure step	step 1
Error	cudaErrorIllegalAddress → NaN gradients → worker deaths
Ordinary fixes tried	lower LR, smaller batch, grad-checkpoint toggle — none helped
Resolution	stayed on Qwen3-4B (standard Transformer)

Table 95.2. The step-1 crash as a run card: config, error, and resolution.

A thirty-second preflight would have caught it before the long run launched: check the kernel is importable and run a single training step.

```
python -c "import importlib.util; print('fla present:',  
    ↪ importlib.util.find_spec('fla') is not None)"  
# then a one-step smoke test on the real config before  
    ↪ committing hours:  
python smoke_train_one_step.py --model $MODEL --fsdp --lora  
    ↪ --seq-len 2048 --bf16
```

The deeper point is that the crash is a seam failure ([Chapter 24](#)) at the largest scale: the seam between a novel architecture and a training stack built for standard Transformers. FSDP-plus-LoRA is battle-tested on standard attention and simply has not converged on support for hybrid linear-attention layers yet.

95.2 The lesson: architecture and stack are coupled

The decision this forced was unglamorous and correct: stay on Qwen3-4B (a standard Transformer) for training. Not because Qwen3.5 is worse, it is, by its own benchmarks, better, but because a model you cannot train is worth nothing to an RL practitioner, however good its weights. The transferable lessons:

- A model’s internal design dictates your toolchain. Backend (FSDP vs Megatron), adapter (LoRA vs full fine-tuning), kernels, and parallelism are not free choices. The architecture constrains them. Check that your exact training configuration is supported before you commit to a model, not after step one crashes.
- “Newer and better on benchmarks” is not “trainable on your stack.” The frontier of architecture moves ahead of the frontier of training-infrastructure support, sometimes by a year. For RL specifically, which needs the trainer, the inference engine, and

the weight sync between them all to agree ([Chapter 4](#)), that gap bites harder than it does for inference-only use.

- When a brand-new architecture won't train, suspect the kernel-and-stack seam first, exactly as [Chapter 24](#) taught for smaller seams.

TAKEAWAY

A model's architecture can break your RL run before step one: upgrading to Qwen3.5-4B crashed immediately (`cudaErrorIllegalAddress`, NaN gradients) because its GatedDeltaNet layers fell back to a torch path with memory-access bugs under FSDP + LoRA, not fixable by any hyperparameter, only by kernel/toolchain support, a different backend, or a different model. Architecture and training stack are coupled, backend, adapter, kernels, and parallelism are constrained by the model's design, and infrastructure support lags the architecture frontier. Verify your exact training config is supported before committing. A model you can't train is worth nothing, however good its benchmarks.

Next: what makes Qwen3.5's architecture worth wanting in the first place, the 3:1 hybrid, and the ablations that settled the ratio.

CHAPTER 96

The 3:1 Hybrid

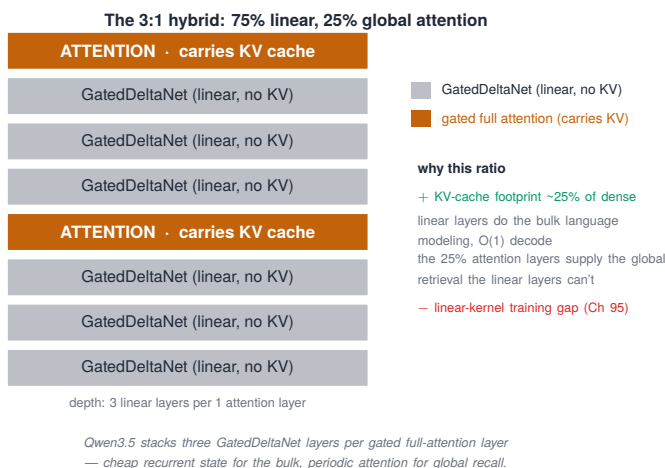


Figure 96.1. The 3:1 hybrid layer stack: three GatedDeltaNet linear-attention layers per one gated full-attention layer (75% linear, 25% global).

Qwen3.5's defining choice is a ratio: three GatedDeltaNet (GDN) linear-attention layers for every one full-attention layer, 75% linear, 25% global. That single number explains most of the model's behavior, from its memory footprint to why speculative decoding underwhelms on it (Chapter 102). This chapter is why 3:1, settled not by theory but by ablation, and what it costs and buys an RL practitioner.

96.1 Why 3:1, and not 1:1 or 7:1

The ratio is an empirical sweet spot that three independent teams, Kimi, OLMo Hybrid, and Qwen, converged on without coordinating. The definitive ablation (Kimi Linear) swept linear-to-attention ratios of 0:1, 1:1, 3:1, 7:1, and 15:1:

- 3:1 won on validation perplexity (5.65). Pure attention was worst (5.77).
- Why not 1:1? It scored essentially the same (5.66) but doubles the attention-layer share from 25% to 50%, roughly doubling the KV-bearing layer footprint and worsening throughput, for a 0.01 perplexity gain that doesn't pay for itself.
- Why not 7:1? Validation degrades (5.70): with only ~12.5% attention layers, the model loses global-retrieval ability. Recall improves sharply as you add attention layers below 3:1, which makes 3:1 about the critical threshold where retrieval becomes adequate.

A functional ablation makes the division of labor stark: removing the linear layers degrades perplexity by $>35,000\times$, while removing the attention layers degrades it by only $\sim 82\times$. The linear layers are the language-modeling backbone. The periodic attention layers are “complementary refinement”, the few global-context checkpoints that handle precise long-range copying the recurrent state can't.

96.2 What the ratio costs and buys an RL practitioner

The 3:1 design is a memory-and-throughput trade with direct consequences for training and serving:

- The win is the KV cache. Because only 25% of layers carry a key-value cache (the other 75% keep a fixed-size recurrent state, independent of sequence length), the KV footprint is roughly a quarter of a same-size dense model's, which is what lets the architecture scale to very long context cheaply (production Qwen3.5 reports up to $19\times$ decoding throughput at 256K context over a dense predecessor).
- The catch for inference tricks. That same “only 25% of layers cache” is why speculative decoding gains are flat on this model ([Chapter 102](#)): the technique accelerates the KV-bearing attention path, and there is far less of it to accelerate. An optimization that pays on a dense model can underperform here for purely architectural reasons.

- The catch for training. The serving stack needs dual-pool memory management, one pool for the recurrent state, one for the KV cache, and the training stack needs kernel support for the linear layers, which is exactly the support gap that crashed our run in [Chapter 95](#). The architecture that makes inference cheap is the one that made training hard.

The same trade-offs as a practitioner table, what each architectural fact does to an RL workflow:

Topic	Detail
75% linear layers	fewer KV-bearing layers → lower rollout memory.
25% attention layers	exact long-range retrieval still costs KV, but only on a quarter of layers.
specialized linear kernels	training stack can crash if the kernel is missing (Ch 95).
hybrid cache	serving must manage recurrent state + KV together (Ch 102).
less attention path	speculative decoding has less to accelerate → flatter gains.

Table 96.1. What each architectural fact does to an RL workflow — the RL / serving consequence:

Evidence status. The ablation numbers in this chapter (validation PPL 5.65/5.66/5.70/5.77, the swept ratios, >35,000× vs ~82× functional-ablation degradation, 19× decode throughput at 256K) are reported by the source papers/model documentation (Kimi Linear. Qwen), not measured on our hardware. They are cited as external evidence for the ratio. The RL consequences above are ours. Verify the figures against the primary sources before quoting them as settled.

TAKEAWAY

Qwen3.5 stacks GatedDeltaNet linear-attention and full attention 3:1 (75% linear, 25% global), an ablation-settled sweet spot (best validation PPL. 1:1 wastes KV cache, 7:1 loses recall) that three teams reached independently. The linear layers are the

backbone (removing them costs $>35,000\times$ PPL vs $\sim 82\times$ for attention), and because only 25% of layers carry a KV cache the memory footprint is $\sim 25\%$ of a dense model's, enabling long context, but also why speculative decoding falls flat and why the training-kernel gap of [Chapter 95](#) exists. The design that makes inference cheap is the one that made training hard.

Next: GatedDeltaNet from scratch, the delta rule, the fixed-size state, and the $O(1)$ decode that the 75% of linear layers actually run.

CHAPTER 97

Gated DeltaNet from Scratch

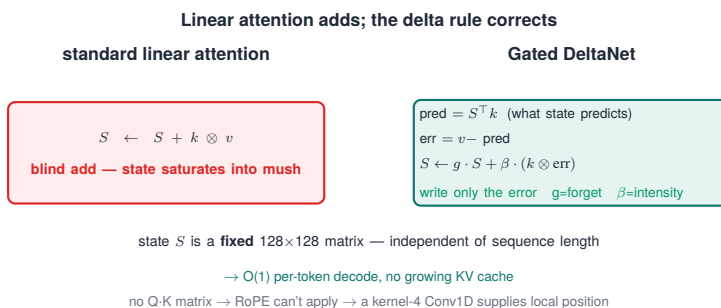


Figure 97.1. Linear attention adds; the delta rule corrects: a blind accumulate that saturates into mush vs Gated DeltaNet that writes only the prediction error.

The 75% of Qwen3.5's layers that are linear ([Chapter 96](#)) run Gated DeltaNet (GDN), and it is worth building the mechanism from the inside, because its design is the reason the architecture decodes in constant time, and the reason it needs the kernel that crashed our training in [Chapter 95](#). The whole idea is one upgrade to linear attention: stop blindly accumulating, and start correcting.

97.1 The delta rule: write the error, not the value

Standard linear attention maintains a state matrix S and updates it by blind addition:

$$S_t = S_{t-1} + k_t \otimes v_t \tag{97.1}$$

Every key-value pair is dumped in, whether or not the state already knew it. DeltaNet adds error correction. It first asks what the current state predicts for this key, then writes only the difference:

$$S_t = S_{t-1} + \beta_t \cdot k_t \otimes (v_t - S_{t-1}^\top k_t) \quad (97.2)$$

The term $v_t - S_{t-1}^\top k_t$ is a prediction error, actual value minus what the state would have retrieved. Writing the error rather than the raw value keeps the fixed-size state from saturating into mush. Gated DeltaNet then adds adaptive forgetting:

$$S_t = g_t \cdot S_{t-1} + \tilde{k}_t \otimes [\beta_t \cdot (v_t - (g_t \cdot S_{t-1})^\top \tilde{k}_t)] \quad (97.3)$$

where $g_t = \exp(\alpha_t) \in (0, 1]$ is a per-step forget gate, $\beta_t = \text{sigmoid}(b_t)$ controls update intensity, and the keys $\tilde{k}$ are L2-normalized, a requirement, not a flourish, for delta-rule stability (without it the update can diverge). In ten lines, the per-token recurrence:

```
# pedagogical: ONE head, ONE token (not exact model code). Keep
↪ S in fp32 for stability.
# S: [d_k, d_v] fixed-size state. Q, k: [d_k]. V: [d_v]
def gdn_step(S, q, k, v, a_log, b):
    g = torch.exp(a_log) # forget gate in (0, 1]
    beta = torch.sigmoid(b) # update intensity
    k = k / (k.norm() + 1e-6) # L2-normalize the key (stability)
    S = g * S # decay old state
    pred = S.T @ k # what the state predicts for this key
    S = S + beta * torch.outer(k, v - pred) # write only the error
    return S, q @ S # output = read the state with the query
```

97.2 Why it decodes in O(1), and why it can't use RoPE

The state S is a fixed 128×128 matrix per head, regardless of sequence length. That single fact is the architectural payoff: per-token decoding is O(1), read and update a constant-size state, with no growing KV cache. Contrast Multi-head Latent Attention (MLA), which compresses the KV cache ($\sim 32\times$) but still attends over the growing set of past tokens, the cache shrinks, the cache dependence does not. GDN instead replaces the growing cache with a fixed recurrent state. At a million tokens that is the difference that matters for decode cost.

The cost of being a recurrence rather than an attention matrix is that rotary position embeddings (RoPE) don't apply. RoPE works by rotating queries and keys so their dot product $Q \cdot K^T$ becomes position-dependent, but the linear layers have no explicit full $Q \cdot K^T$ attention matrix, only a state-update recurrence, so standard RoPE is not directly applicable there. GDN instead gets local position from a causal depthwise Conv1D with kernel size 4 (a 4-token receptive field). That convolution is load-bearing: removing it raised Wikipedia perplexity from 28.24 to 29.08 and dropped retrieval 18% at 340M scale. It supplies the n-gram patterns ("New York") and single-layer induction the recurrence alone misses.

TAKEAWAY

Gated DeltaNet upgrades linear attention with the delta rule, write the *prediction error* $(v - S^T k)$, not the raw value, plus a forget gate $g = \exp(\alpha)$, update intensity $\beta = \text{sigmoid}(b)$, and L2-normalized keys (required for stability). Its 128×128 state is fixed-size regardless of length, giving true $O(1)$ decoding with no KV cache, unlike MLA, which compresses but stays quadratic. The price: no RoPE (there's no $Q \cdot K^T$ to rotate), so a kernel-4 Conv1D supplies local position instead.

Next: the other 25% of layers, full attention, fixed so it stops wasting half its mass on the first token.

CHAPTER 98

Gated Attention and the Attention-Sink Fix




Figure 98.1. A sigmoid gate drains the attention sink: the first token absorbs 46.7% of attention under softmax, dropping to 4.8% with a per-head sigmoid gate.

The 25% of Qwen3.5’s layers that are full attention ([Chapter 96](#)) carry a small but consequential fix: a sigmoid output gate that eliminates attention sinks. This is not a tuning trick. It comes from a 2025 paper, “Gated Attention for Large Language Models” ([arXiv:2505.06708](#)), and it changes what those few global-retrieval layers can do. The fix matters more here precisely because there are so few attention layers. The model cannot afford to have them wasting their capacity.

98.1 The sink, and the one-line fix

In a standard transformer, attention scores must sum to one (softmax). When a token has nothing it especially needs to attend to, that constraint forces the spare probability mass somewhere, and it piles

onto the first token, which becomes a meaningless dumping ground. The numbers are stark: the baseline allocated an average of 46.7% of attention to the first token, and one layer peaked at 83%. That is capacity spent on noise.

The fix is to multiply the attention output by a query-dependent sigmoid gate, element-wise per head, after the softmax:

$$\text{output} = \text{sigmoid}(X \cdot W_{\text{gate}}) \odot \text{SDPA}(Q, K, V) \quad (98.1)$$

It works because the sigmoid is not bound by softmax’s sum-to-one constraint. The softmax may still assign 45% to the first token, the gate does not change the attention probabilities, but the gate can multiply that head’s output contribution by ~ 0 , so the wasted attention no longer dominates the residual update. Applied during training, the reported effect is to drive first-token attention from 46.7% down to 4.8%, sinks effectively gone, while adding fewer than 2M parameters to a 15-billion-parameter mixture-of-experts model, and buying >0.2 perplexity, ~ 2.0 Massive Multitask Language Understanding (MMLU) points, and $+2.3$ on grade-school math. (All of these, the 46.7%/4.8% concentrations, the parameter count, and the benchmark deltas, are numbers reported by the source paper, not measured on our stack.)

98.2 Why this beats the inference-time workaround

The contrast worth internalizing is with StreamingLLM, the well-known inference-time approach: it accepts that sinks form and works around them by always keeping the first few “sink” tokens in the KV cache during long-context generation. Gated attention instead prevents sinks from forming at the root, during training. The two solve related problems at different stages, so it is not that one strictly dominates, but for a new architecture you control, preventing sinks during training is the cleaner choice than preserving sink tokens at inference, and it has properties the cache-policy workaround does not:

- It frees the wasted capacity for real use, rather than just preserving the dumping ground.

- It is an architectural property of the trained weights, so it holds everywhere, not only under a special inference-time cache policy.
- Because it cleans up the model's few global-attention checkpoints (the only layers with unbounded retrieval), it directly improves the long-range copying the hybrid depends on.

The general lesson for a practitioner: a fix applied at training time, in the architecture, is more robust than a patch applied at inference time around a problem you let happen.

TAKEAWAY

Standard attention wastes ~46.7% of its mass on the first token, a “sink” forced by softmax’s sum-to-one rule. A per-head sigmoid output gate ($\text{sigmoid}(X \cdot W_{\text{gate}}) \odot \text{SDPA}$), applied after softmax during training, escapes that constraint and drops first-token attention to 4.8% for <2M extra params (+~2 MMLU, +2.3 GSM8K). Unlike StreamingLLM’s inference-time workaround, it kills sinks at the root, and it matters most because Qwen3.5 has so few attention layers to waste.

Next: the feed-forward side, an ultra-sparse mixture of experts that activates barely 4% of its parameters per token.

CHAPTER 99

Sparse Mixture-of-Experts and Ultra-Sparsity




Figure 99.1. Active-parameter fraction falls across generations: Qwen3 9.4% → DeepSeek-V3 5.5% → Qwen3.5 4.3% active per token.

Qwen3.5’s feed-forward layers are an ultra-sparse mixture of experts (MoE): 512 experts, of which only about 11 fire per token. By the fraction people usually quote, active parameters per token, that is roughly 4.3% (17 billion of the model’s 397 billion), which is why a 397-billion-parameter model runs at 17-billion-parameter cost. For an RL practitioner, this is also the architecture that pushed us toward sequence-level training ([Part V](#)).

99.1 The trend: more, finer experts. Fewer active

The sparsity has been falling across model generations, but you must say which fraction, because two different denominators get quoted and they are not the same number:

Model	Experts	Active (routed + shared)	Active-expert fraction	Active-parameter fraction
Qwen3-30B-A3B	128	8	$8/128 = 6.25\%$	$\sim 9\text{--}10\%$ ($\approx 3\text{B} / 30\text{B}$)
Qwen3.5-397B-A17B	512	10 + 1 shared	$11/512 \approx 2.1\%$	$\sim 4.3\%$ ($17\text{B} / 397\text{B}$)
DeepSeek-V3	256	—	—	$\sim 5.5\%$

Table 99.1. Expert counts, active experts, and active fractions across three MoE models.

The headline “9.4%” and “4.3%” are active-parameter fractions, not active-expert fractions, $8/128$ is 6.25%, not 9.4%, so the two must not be conflated. Whichever fraction you track, the direction is the same: more experts, finer (intermediate dimension shrank to about 1,024, with 4× as many experts, $128 \rightarrow 512$), and fewer active per token. Two design moves produced it: fine-grained experts (smaller experts let the router compose specialized combinations), and the re-introduction of one shared expert (which Qwen3 had dropped), a single always-on expert that provides a knowledge baseline and, importantly, stabilizes routing so the routed experts can specialize without the router thrashing.

99.2 Why ultra-sparsity makes sequence-level RL the safer default

Ultra-sparsity is a gift at inference (huge total knowledge, tiny active compute) and a hazard in training, and the hazard is the through-line to [Part V](#). The more experts and the sparser the routing, the more the expert selection can flip between the old policy and the new one within a single update, and token-level importance-sampling ratios assume the per-token distributions are comparable across that update. When routing flips, that assumption breaks exactly where the math is most sensitive, and vanilla token-level Group Relative Policy Optimization (GRPO) can destabilize (we watched it die at step 10 on a 30-billion-parameter MoE, [Chapter 50](#)).

The fix Qwen3.5 uses, and the one this book recommends as the default for MoE at scale, is Group Sequence Policy Optimization (GSPO), which computes the importance ratio at the sequence level rather than the token level, making it robust to the routing flips ultra-sparsity makes likely. Two honesty notes, because the book’s own evidence is more nuanced than a slogan. GSPO is not strictly mandatory: vanilla GRPO can work under conservative conditions (a small enough learning rate, enough rollouts, scale-dependent robustness, [Part V](#) documents where it held and where it didn’t). The accurate claim is that sequence-level ratios remove a failure surface that sparse MoE makes more likely, so they are the safer default, not the only option. The architecture choice (more, sparser experts) and the algorithm choice are coupled: the first is why the second is the prudent one.

When an MoE RL run misbehaves, the symptom points at the fix:

Symptom	Likely cause	First fix
token-ratio explosions	routing flips between old/new policy	GSPO (sequence-level ratios)
router entropy col-lapses	over-commitment to a few experts	JITTER / aux load-balance loss / lower LR
stable but low ceiling	router frozen too early	avoid FREEZE, or unfreeze later
no collapse, VAL rising	benign specialization	continue — do not kill on entropy alone

Table 99.2. MoE RL symptoms matched to their likely cause and first fix.

TAKEAWAY

Qwen3.5’s feed-forward is an ultra-sparse MoE, 512 experts, ~11 active (10 routed + 1 shared), quoted as ~4.3% of *parameters* active per token (17B/397B), not 4.3% of experts (that’s ~2%). Define the denominator before comparing. Fine-grained experts (intermediate dim ~1,024, 4× more of them) plus a

re-introduced shared expert stabilize routing, and the active-parameter trend falls across frontier models. Ultra-sparsity makes expert routing flip across updates, which can break token-level GRPO, so GSPO's sequence-level ratios (Part V) are the *safer default*, not strictly mandatory: vanilla GRPO can work under conservative conditions, but sequence-level ratios remove a failure surface sparse MoE makes likely.

Next: how a recurrent, inherently sequential layer is trained in parallel at all, the chunk-parallel algorithm.

CHAPTER 100

The Chunk-Parallel Algorithm

The trick is in one line. A recurrence is sequential: S_t needs S_{t-1} . But a product of delta-rule updates, $\prod_i (I - \beta_i k_i k_i^\top)$, can be rewritten, via the WY decomposition (Bischof & Van Loan, 1985), as $I - \sum_i w_i k_i^\top$, a sum of outer products you can compute with matrix multiplies. Split the sequence into chunks of 64: run the recurrence between chunks, run dense matmuls within them. Sequential where you must be, parallel where you can be.

Gated DeltaNet (Chapter 97) is a recurrence, and recurrences look fatally sequential: each state depends on the one before it, so where is the parallelism that makes training on modern hardware fast? The answer is the chunk-parallel algorithm, and it is the piece of mathematics that makes a linear-attention model trainable at scale, when its kernel is available (Chapter 95).

100.1 WY decomposition: turning a product into a sum

Each delta-rule update is a state-transition matrix of the form $M_t = I - \beta_t \cdot k_t \cdot k_t^\top$, identity plus a rank-one correction. The state after a chunk is the cumulative product of these, which naively must be applied one at a time. The WY representation transforms that cumulative product of identity-plus-rank-one matrices into a single sum of outer products:

$$\prod_i (I - \beta_i k_i k_i^\top) = I - \sum_i w_i k_i^\top \quad (100.1)$$

where the w vectors are computed inductively by forward substitution (a small triangular solve). The payoff is that the inner computation becomes dense matrix multiplications, $A = -(k_{\text{buf}} @ k^\top * L)$ then $v_{\text{corrected}} = (I + A) @ (v * \beta)$, which is exactly the shape tensor cores are built to run fast. The recurrence between

chunks stays sequential. Everything inside a chunk becomes parallel linear algebra.

100.2 Why chunk 64, and the kernel that runs it

The chunk size is $C = 64$, and the number is not arbitrary:

- Tensor-core alignment, 64 is a multiple of 16, the granularity those units want.
- SRAM fit, a 64×64 working tile is 4,096 elements, which sits in fast on-chip memory.
- Warp-level parallelism, the chunk maps cleanly onto a warp.

The kernels live in the flash-linear-attention (FLA) library, written in Triton, which improves portability over hand-written CUDA, though backend support (NVIDIA, AMD, Intel) still has to be verified on your own hardware. And they are fast: past about 4K tokens the FLA chunk kernels beat FlashAttention-2, with the gap widening, one reported figure puts the forward pass at roughly 1.75 ms versus 13.7 ms at 16K tokens. This is the upside the architecture is buying. It is also, exactly, the dependency whose absence (no FLA in the container $\rightarrow$ a buggy torch fallback) crashed our training in [Chapter 95](#), the chunk-parallel kernel is the difference between this layer being fast and being unrunnable.

TAKEAWAY

A recurrence seems unparallelizable, but the WY decomposition rewrites a product of delta-rule updates $\prod (I - \beta k k^\top)$ as a sum of outer products $I - \sum w k k^\top$, turning within-chunk work into dense matmuls. Run the recurrence *between* chunks of size 64 (chosen for tensor-core alignment, SRAM fit, warp parallelism) and matmuls *within* them. The Triton FLA kernels beat FlashAttention-2 past 4K tokens (1.75 ms vs 13.7 ms at 16K), and are the very dependency whose absence crashed training in [Chapter 95](#).

Next: how a model whose linear layers ignore absolute position is stretched to a million tokens of context.

CHAPTER 101

Context Extension to a Million Tokens

Layer type	Share	Position handling	Long-context role
Gated DeltaNet (linear)	75%	none beyond a 4-token Conv1D; fixed-size state	bulk processing; carries context in its state
Full attention	25%	RoPE, YaRN-scaled (factor 4)	the only unbounded global retrieval

Table 101.1. Context extension applies asymmetrically. YaRN’s 4× scaling is applied only to the 25% full-attention layers, the linear layers are position-agnostic and carry long context through their recurrent state. The attention layers do the precise long-range retrieval the fixed state can’t.

Qwen3.5 extends to roughly 1,048,576 tokens of context, and the interesting part is where the extension is applied. Because the architecture is a hybrid, context extension is not uniform. It touches only a quarter of the layers, and understanding why closes the loop on the whole 3:1 design.

101.1 YaRN on the attention layers only

The extension uses YaRN (Yet another RoPE extension) with a factor of 4.0, multiplying the trained context length roughly 4× to ~1M tokens. But YaRN is a method for rescaling rotary position embeddings (RoPE), and only the full-attention layers have RoPE ([Chapter 97](#) showed the linear layers can’t use it). So YaRN is applied exclusively to the 25% full-attention layers. The linear Gated DeltaNet layers need no position rescaling: they are position-agnostic beyond their 4-token convolution and carry long-range information through their

fixed-size recurrent state, which has no length-dependent positional structure to stretch.

101.2 The compression bottleneck, and why 25% attention exists

This is also where the architecture’s central tension becomes concrete. The recurrent state is fixed-size (dimension ~ 128 per head), so it is a lossy summary, not an unbounded memory. There is a hard limit to how many distinct associations it can hold at once. That is the compression bottleneck: at extreme context lengths, the fixed state simply cannot hold every association a precise long-range copy might need. The state is a lossy summary, excellent for bulk language modeling and useless for “fetch the exact token from 400K positions back.”

That bottleneck is the entire justification for keeping 25% full attention (the ratio of [Chapter 96](#), now seen from the long-context side). Those layers maintain an unbounded KV cache, they can retrieve any past token exactly, and they are what make million-token retrieval work where the linear layers’ compressed state would fail. The hybrid is not 75/25 by aesthetics. It is 75% cheap-and-lossy plus 25% expensive-and-exact, and the 25% exists precisely to cover the linear layers’ one structural weakness.

TAKEAWAY

Qwen3.5 reaches $\sim 1\text{M}$ tokens via YaRN (factor 4), applied *only* to the 25% full-attention layers. They’re the only ones with RoPE to rescale. The linear layers need no extension: they’re position-agnostic and carry context in a fixed state. But that state’s effective rank (~ 128 associations per head) is a hard compression bottleneck at long range, which is exactly why the 25% full-attention layers, with their unbounded KV cache for exact retrieval, have to exist.

Next: serving the thing, the inference stack, its speculative-decoding heads, and the dual cache it forces on serving frameworks.

CHAPTER 102

The Inference Stack

Topic	Detail
vLLM	how it handles the hybrid cache — one paged pool, tunes “logical” block sizes so KV blocks and state blocks occupy equal physical memory; linear-layer state — recurrent state in float32.
SGLang	how it handles the hybrid cache — two separate pools, a Mamba pool (per-request, fixed size) and a KV pool (per-token, paged); linear-layer state — <code>HybridLinearKVPool</code> skips KV for linear layers.

Table 102.1. How vLLM and SGLang each manage the hybrid recurrent-plus-KV cache.

Serving Qwen3.5 is not standard transformer serving, because the hybrid cache ([Chapter 96](#)) breaks the assumption every serving framework was built on: that all layers cache the same way. This chapter is the inference stack, its speculative-decoding heads, its two-regime memory, and its deliberately non-uniform quantization, because each of these is forced by the architecture, and each is a trap if you assume a dense model.

102.1 Multi-token prediction and the dual cache

Qwen3.5 ships multi-token prediction (MTP) heads, branded “NextN,” for speculative decoding. Each head is deliberately lightweight: shared token embeddings, a linear projection, a single transformer block using the same Gated DeltaNet / full-attention architecture as the main model, and the shared output head. Because the heads are jointly trained during pretraining, the draft distribution aligns with the base model out of the box. Recommended draft counts are 1–2 for vLLM (1 for latency-critical 397B serving) and up to 4 for SGLang, with measured speedups of 1.25–2.11×, peaking at batch size 1 and

shrinking as concurrency rises (draft tokens compete for KV-cache capacity).

The deeper serving challenge is the dual cache (Table 102.1). With 75% of layers holding a fixed-size recurrent state and 25% holding a growing KV cache, a framework must manage both. vLLM unifies them under one paged allocator by tuning logical block sizes so a KV block and a state block occupy equal physical memory. SGLang keeps two pools outright, binding each request’s recurrent state to its lifetime while paging the KV cache per token. Either way, the recurrent state is kept in float32 for numerical stability, a detail, like the `a_log` parameter’s float32 requirement, that has already caused production bugs.

102.2 Quantization is deliberately non-uniform

The most important inference lesson for an RL practitioner, and the one that connects to this book’s quantization findings (Part XIV), is that FP8 quantization on Qwen3.5 is selective, on purpose. Only the routed MoE expert weights are quantized to FP8 (E4M3, block size 128). Everything else stays in BF16:

- Shared experts stay BF16, quantizing the always-on expert measurably degrades quality.
- All full-attention layers stay BF16. They are the scarce global-retrieval resource.
- All Gated DeltaNet layers stay BF16, quantizing linear attention degrades long-sequence generation (the long “thinking” traces specifically).

The principle generalizes well beyond this model: uniform quantization is the wrong default for a heterogeneous architecture. The components doing precise or always-on work (shared experts, attention, the recurrent state) carry information that low precision corrupts. The bulk routed experts tolerate it. Quantize by role, not by blanket rule, the same lesson the rollout-quantization failures of Part XIV teach from the training side.

For RL specifically, speed is never the only question an inference change has to answer, correctness is:

RL rollout correctness contract. For an exact acceleration (a faster kernel that doesn't change outputs): report wall-clock, tokens/sec, and memory, and confirm identical outputs under matched seeds. For a lossy acceleration (quantized KV, aggressive quantization): additionally report per-rollout KL against the unquantized policy and held-out eval after a short RL run, because a change that speeds rollouts but shifts the token distribution silently corrupts the policy the trainer is updating (Part XIV catalogs exactly these failures: FP8-KV at KL 7.61, AWQ rollouts breaking the importance ratio).

Evidence status. In this chapter, the existence of MTP/NextN heads and the selective-FP8 recipe are official (Qwen sources). The dual-pool/unified-block cache handling and float32-state requirement are framework implementation facts (vLLM/SGLang). And the speculative-decoding speedups and their fall-off with concurrency are partly our measurements (Part XIV) and partly reported, label them accordingly before quoting.

TAKEAWAY

Serving Qwen3.5 departs from dense serving on three fronts forced by the architecture: MTP/NextN heads (one shared transformer block, jointly pretrained) give 1.25–2.11× speculative speedup, peaking at batch size 1. The hybrid cache forces dual-regime memory management (vLLM unifies block sizes, SGLang keeps two pools, recurrent state in float32). And FP8 is *selective*, only routed experts quantized, while shared experts, attention, and all linear layers stay BF16. Quantize by role, not by blanket rule.

Next: the whole design on one page, every architectural decision, traced to the paper and the number that settled it.

CHAPTER 103

The Architecture Decision Map

Decision	Choice	Evidence	What settled it
Attention ratio	3:1 linear-to-full	external paper	Kimi Linear ablation: 3:1 best val PPL 5.65; 1:1 wastes KV; 7:1 loses recall
Linear layer	Gated DeltaNet, fixed 128×128 state	architecture	O(1) decode vs MLA's growing-cache dependence
State update	delta rule + forget gate	architecture	writes prediction error, not raw value; gate enables forgetting
Key normalization	L2-normalized keys	architecture	required for delta-rule stability
Local position	Conv1D, kernel 4 (no RoPE in linear)	external paper	removing conv: +PPL, −18% retrieval at 340M (reported)
Attention position	RoPE, partial factor 0.25	official source	only 25% of head dims rotate (decoupled-RoPE style)
Attention sinks	sigmoid output gate	external paper	46.7% → 4.8% first-token attention (reported, arXiv:2505.06708)
Output gating (linear)	RMSNorm ⊙ SiLU	external paper	from RetNet; SiLU's near-1 gradient aids flow through linear layers
Feed-forward	ultra-sparse MoE, 512 experts (~4.3% active params)	official + our runs	sparsity trend; fine-grained + 1 shared expert; our MoE-RL experience
Training kernel	WY chunk-parallel, chunk 64	external paper	tensor-core alignment; FLA reported faster than FA-2 past 4K tokens
Context extension	YaRN factor 4, attention layers only	official source	linear layers are position-agnostic; ~1M tokens

RL algorithm	GSPO (sequence-level)	our C9 runs + external	token-level GRPO can destabilize under expert routing flips (Part V)
--------------	-----------------------	------------------------	--

Table 103.1. Qwen3.5’s design as a decision map: each choice traces to a specific ablation, paper, or measured number. None is arbitrary, and, read together, they reinforce each other.

Qwen3.5’s architecture is best understood not as a list of tricks but as an integrated system whose decisions reinforce each other, and the most useful artifact for a practitioner is the decision map itself (Table 103.1), each choice tied to the evidence that settled it. This chapter is that map, and the single principle underneath it.

103.1 The decisions reinforce each other

Read the table top to bottom and the dependencies appear. The 3:1 ratio works because Gated DeltaNet is expressive enough per layer (delta rule plus gating) that 75% linear layers suffice. The Conv1D exists because the linear layers can’t use RoPE, and it fills exactly that gap. The sigmoid sink fix matters because there are so few attention layers that the model cannot afford to have them dumping mass on the first token. It protects the scarce global-retrieval checkpoints. The ultra-sparse MoE compensates for the hybrid’s large total parameter count by keeping active compute tiny. And GSPO is forced by that same ultra-sparsity. Pull any one decision and the others wobble. That interlock is what “integrated system” means, and it is why you cannot cherry-pick one idea from this architecture and expect it to behave the same in isolation.

103.2 The one principle: compress hard, defend quality at every seam

Every choice in the table serves one of two jobs, and usually they pair up. Some decisions compress aggressively, the fixed-size state,

4.3% expert activation, the 3:1 ratio. Each of those throws information away to save compute or memory. The other decisions exist to defend information quality at exactly the point where the compression would lose it: the delta rule's error correction defends the compressed state, L2-normalized keys defend the delta rule's stability, sigmoid gating defends the few attention layers, rich value dimensions defend retrieval through the bottleneck. The unifying principle is: compress where you can, and place a quality defense at every compression seam. That is a design lens you can carry to any efficiency-driven architecture, when you see an aggressive compression, look for (or add) the mechanism that defends against its information loss, because the compression alone is rarely safe.

TAKEAWAY

Qwen3.5's twelve major decisions each trace to a specific ablation or paper (Table 103.1), and they form an interlocking system: 3:1 works because GDN is expressive per layer. Conv1D fills the no-RoPE gap. Sigmoid gating protects the scarce attention layers. Ultra-sparse MoE offsets total parameter count. GSPO is forced by that sparsity. The single principle is *compress aggressively, then defend information quality at every compression seam*, a lens that generalizes to any efficiency-driven architecture.

Next: *Part XI*, shipping. We take everything and build a real product: an autonomous web agent that completes multi-step browser tasks, through the full SFT → DPO → RL pipeline.

Part XI

Applied Tool-Agent Pipeline: SFT → DPO → GSPO

Every Part so far optimized a benchmark. This one ships something real, an autonomous web agent that talks to a user and drives a real browser to complete multi-step tasks, and a deployed agent needs what a benchmark doesn't: judgment, guardrails, and a pipeline that survives contact with production. This is where the full SFT → preference → RL recipe of [Chapter 2](#), skipped for most of the book, finally earns all three stages.

CHAPTER 104

The Production Web Agent



Figure 104.1. A deployed agent needs all three stages: SFT teaches FORM → DPO teaches TASTE + SAFETY → GSPO teaches TASK SUCCESS.

Most of this book deliberately skipped preference tuning, when a programmatic verifier exists, you go straight from an instruct model to reinforcement learning ([Chapter 2](#)). This Part is the exception that proves the rule. A web agent that converses with a user and drives a real browser to complete a multi-step task needs all three stages, because it has three distinct problems no single stage solves. This chapter is why the full pipeline, and what each stage is actually for.

104.1 Three problems, three stages

The agent's job is to take a user's request ("update the mailing address on my account to the new one"), converse to resolve ambiguity, and drive a multi-step web workflow to completion. That job decomposes into three problems with three different signals (Figure 104.1):

1. Form, does it emit valid tool calls and well-structured turns? This is supervised, and it is what SFT is for: imitate good trajectories until the agent reliably produces parseable tool calls and sensible dialogue structure ([Chapters 7, 21](#)). Without this, there is nothing for later stages to grade.

2. Judgment and safety, does it behave the way a user-facing system must? Tone, refusing out-of-scope requests, not taking the wrong action, honoring guardrails. Much of this is not programmatically verifiable in the way a math answer is, it is preference-shaped, which is what DPO is for, and the preference source is guardrail verdicts and production outcomes, not human aesthetics ([Chapter 2](#)'s note on DPO).
3. Task success, does the task actually complete correctly? This is verifiable (did the agent submit the values that match the goal and reach the intended final state), and it is what RL is for, specifically GSPO against sandboxed browser rollouts (never live side-effecting actions, see 104.3).

104.2 Why SFT → DPO → GSPO, in that order

The ordering is [Chapter 2](#)'s cost-and-dependency logic, applied to a deployed agent:

- SFT first because the cheaper, denser signal de-risks everything downstream, and because RL has nothing to grade until the agent can produce valid attempts. A short imitation pass on good trajectories gets it there.
- DPO second because taste and safety are preference-shaped and need establishing before RL starts pushing hard on task reward, otherwise RL will happily trade away politeness or a guardrail for a few points of completion rate (the reward-hacking pattern of [Part III](#), now with real consequences for a real user).
- GSPO last because it is the most expensive and most fragile stage, and it is also the only one that teaches genuine task judgment, when to clarify, which path through the workflow to take. Sequence-level ratios because the policy is a large model and the trajectories are long and multi-turn.

The production framing changes the stakes: a benchmark agent that games its reward costs you a misleading number. A deployed web agent that games its reward takes a wrong, possibly irreversible, action on a real user's behalf. The full pipeline exists because correctness alone is not safety, and shipping needs both.

104.3 The non-negotiable: train and evaluate in a sandbox

Before any of this, one operational rule that is not optional and is easy to get catastrophically wrong: a web agent that can take real, side-effecting actions must be trained and RL-evaluated against a sandboxed copy of the site, mock back-ends, or dry-run mode, never against the live system. An RL agent’s entire job is to explore, including the actions you didn’t anticipate. Pointing that exploration at a real submit-or-delete button is how you fire an irreversible action ten thousand times at 3 a.m. The environment escalates with confidence, not before it:

Stage	Allowed environment	Irreversible action?	Human confirmation?
SFT	logged trajectories	no	no
DPO	logs / guardrail pairs	no	no
GSPO training	sandbox / mock back-end	no	no
staging eval	dry-run mode	no	maybe
production	live system	only with scope/rate limits	yes for risky actions

Table 104.1. The deployment-stage ladder for a tool-using agent: which environment each training or eval stage may touch, whether it can take irreversible actions, and when a human must confirm. Everything up to staging runs in a sandbox with no irreversible actions; only production touches a live system, and only there do risky actions need human sign-off (the red row).

Real side-effecting actions belong behind separate product safety: scope and rate limits, explicit human confirmation for irreversible or high-impact actions, and audit logs. This is not a legal aside. It is the same reward-hacking lesson as [Part III](#), except the “hack” now takes real actions. And the agent needs success metrics beyond “the task completed”: exact match to the requested goal, constraint compliance, guardrail-violation rate, clarification rate, stuck-loop rate, human-escalation rate, and undo/rollback success, because an agent that completes the wrong task has failed, not succeeded.

TAKEAWAY

The web agent runs the full SFT → DPO → GSPO pipeline because it has three problems one stage can't cover: form (SFT teaches valid tool calls), judgment-and-safety (DPO from guardrail/production preferences, not human aesthetics), and task success (GSPO against *sandboxed* browser rollouts). The order follows [Chapter 2](#), cheap-dense first, preferences before RL pushes on reward, fragile-expensive RL last. A deployed agent needs judgment and guardrails on top of correctness. A gamed reward here takes a wrong action on a real user's behalf, which is why training and RL evaluation run in mock/dry-run mode, never the live system, with limits, human confirmation, and audit logs reserved for production.

Next: where the training data comes from, converting Claude's trajectories into something you can actually train on.

CHAPTER 105

Converting Claude Trajectories to Training Data

The hard rule up front. Trajectories generated by a different model, Claude, here, are off-policy. They are valid training data for SFT and DPO, which learn from fixed examples. They are not valid for on-policy RL (GRPO/GSPo), whose importance ratios assume the attempts were drawn from the policy being updated (Chapter 4). Use the teacher's trajectories to bootstrap form and preferences. Do not feed them to the RL loop as if the policy had generated them.

The SFT and DPO stages of Chapter 104 need data, and a fast way to get high-quality agent trajectories is to have a strong teacher, Claude, drive the task and record what it did. But converting those trajectories into training data has one principle that governs everything and several mechanical traps. This chapter is both.

105.1 The governing principle: off-policy data has a lane

The single most important thing to get right is which stage teacher trajectories may feed. A trajectory Claude generated reflects Claude's policy, not your model's. That makes it:

- Valid for SFT, imitation learns to reproduce demonstrations regardless of who produced them. This is the standard use: distill the teacher's tool-use form into your model.
- Valid for DPO, preference learning compares fixed pairs. The pairs can come from a teacher, from guardrails, or from production.
- Invalid for on-policy RL, GRPO and GSPo compute a ratio of new-policy to old-policy probabilities, and the math assumes the old policy is yours, evaluated on attempts yours generated. Feed Claude's trajectories into that ratio and it means nothing

(Chapter 4’s on-policy invariant). Off-policy correction is its own research problem. “just reuse the teacher’s rollouts in RL” is the trap.

So the conversion pipeline targets SFT and DPO, full stop. The teacher bootstraps form and taste. Your own policy’s rollouts, generated live, drive the RL stage.

105.2 The mechanical conversion

Beyond the principle, the conversion is a careful reformatting:

1. Re-tokenize and re-template. Claude’s trajectories use Claude’s chat template and tokenizer. Your student model has its own. Every turn must be re-rendered into the student’s exact template, or the SFT loss trains on a format the model will never see at inference (the seam bug of Chapter 24, in data form).
2. Map the tool-call schema. The teacher’s tool-call format must be translated to the student’s expected format, names, argument structure, the exact special tokens that delimit a call and a result.
3. Mark segment provenance for masking. Every token must be tagged as policy-produced (the assistant’s thoughts and tool calls) or environment-produced (tool results, retrieved pages), because SFT must train only on the former, which is exactly the masking discipline of Chapter 21, and the subject of the next chapter, where a Qwen3.5-specific version of it silently breaks training.

Get the conversion right and you have a cheap, high-quality SFT and DPO corpus. Get the template or the masking wrong and you train a confident model on text it will never produce, a failure that, as the next chapter shows, can be completely silent.

A teacher step in, a student SFT example out, concretely:

```
// teacher (Claude-ish) raw step
{"observation": {"url": "...", "ax_tree": "..."},
```

```
"teacher": {"action": {"type": "tool_call", "name": "click",
↪ "arguments": {"ref": "submit-btn"}}}}

// student SFT example (re-templated, schema-mapped, masked)
{"messages": [
  {"role": "system", "content": "..."},
  {"role": "user", "content": "Goal: update mailing
↪ address\nAX_TREE:..."},
  {"role": "assistant", "tool_calls": [
    {"type": "function", "function": {"name": "browser_click",
↪ "arguments": {"ref": "submit-btn"}}}]},
  {"tools": [/* student's tool schema */],
  "loss_mask_policy": "assistant_only"}
```

And the rule for which lane each data source may feed, the off-policy boundary made into a table:

Data source	SFT	DPO	GRPO / GSPO
Claude successful trajectories		(if paired)	x
Claude failed-vs-fixed action pairs	optional		x
production guardrail overrides	maybe		x
target-model live/sandbox rollouts	(replay)	(if paired)	(on-policy)

Table 105.1. Which data source is admissible for which training stage. A checkmark means the source can feed that stage; an x means it cannot. The key constraint: GRPO/GSPO is on-policy, so only the target model’s own live or sandbox rollouts qualify — teacher trajectories and logged pairs feed SFT and DPO instead.

Only the last row, rollouts from the target policy itself, may enter the on-policy RL loop. Multimodal inputs follow the same conversion discipline: screenshots become images (or are dropped if the student is text-only), the accessibility tree and tool results become text spans (environment-masked), and any user data is redacted before a trajectory ever reaches the corpus. And one caution for later: “offline GSPO”

from pre-collected groups is fine for engineering, but unless the old-policy probabilities and the generation distribution are the target policy's, it is not mathematically the same as on-policy GSPO (Chapter 4).

TAKEAWAY

Teacher (Claude) trajectories are off-policy: valid for SFT and DPO, which learn from fixed examples, but *invalid* for on-policy RL (GRPO/GSPO), whose ratios assume the policy generated its own attempts, reusing them there is a trap. The conversion to SFT/DPO data is mechanical but unforgiving: re-template and re-tokenize into the *student's* exact format, map the tool-call schema, and tag every token's provenance for masking (Chapter 21). A wrong template or mask trains the model on text it will never produce.

Next: getting that masking exactly right in SFT, and the Qwen3.5 template bug that silently kills the run.

SFT with Correct Tool-Call Masking

Segment	Source	Loss mask	Why
assistant reasoning	policy	1	the model chose it — train on it
tool call	policy	1	the model chose it — train on it
tool result / page content	environment	0	the tool produced it — never train on it
user / system turns	given	0	context, not the model's output
final assistant answer	policy	1	the model chose it — train on it

Table 106.1. The loss mask for a multi-turn tool-using agent: train on the tokens the model chose (its reasoning, tool calls, and final answer, mask 1) and mask everything the environment or user produced (tool results, page content, prompts, mask 0). Training on environment tokens teaches the model to hallucinate them.

The supervised fine-tuning stage of the web agent ([Chapter 104](#)) lives or dies on tool-call masking, the same discipline introduced for multi-step RL in [Chapter 21](#), now on the SFT side and with higher stakes, because an agent SFT corpus is mostly environment tokens. Get the mask right and the agent learns to act. Get it wrong and it learns to hallucinate the world. This chapter is the SFT masking doctrine. The next is the specific way it fails on Qwen3.5.

106.1 The loss covers only what the policy produced

An agent trajectory is a long, mixed token stream: the model's reasoning and tool calls interleaved with the environment's tool results and page content (Table 106.1). The rule is identical to [Chapter 21](#)'s, and

identically non-negotiable: the SFT loss is computed only on the tokens the policy itself generated. The model's reasoning, its tool calls, and its final answers get loss mask 1. Every tool result, retrieved page, and user/system turn gets loss mask 0.

The failure mode if you skip this is specific and ruinous. Train the model to predict the tool's output, the search results, the page HTML, the API response, and you are teaching it to hallucinate the environment: to confidently generate what it thinks a tool would have returned, instead of calling the tool and reading the real answer. In a web agent, that means inventing page contents, field values, and confirmation screens. The model looks fluent and is dangerously untethered from reality. And because the loss still falls smoothly while this happens, nothing on the dashboard warns you ([Chapter 7](#)'s reminder to check free-running generation, not just the loss).

106.2 Build the mask with the trajectory, and assert it

The discipline that prevents silent mask bugs is mechanical, and worth making a habit ([Chapter 21](#) gave the RL-side version):

- Construct the mask segment by segment as you build the trajectory, tagging each span's provenance at creation time, never try to reconstruct which tokens were the model's after the fact from a flattened string.
- Assert the invariants before training: mask length equals token length, every mask value is 0 or 1, and the sum is greater than zero (at least some policy tokens reach the loss). An agent corpus that is 90% environment tokens can, with one indexing error, mask everything and train on nothing, or mask nothing and train on everything.
- Spot-check a rendered example by eye. Decode one trajectory with its mask overlaid and confirm that exactly the assistant's spans are unmasked. This thirty-second check catches the bug class that costs days.

In code, the collator is tiny, and the two assertions are the whole safety net:


```
def build_labels(input_ids, assistant_mask):
    labels = input_ids.clone()
    labels[assistant_mask == 0] = -100 # -100 = ignored by the
    ↪ loss
    assert labels.shape == input_ids.shape
    assert assistant_mask.sum().item() > 0, "no assistant tokens
    ↪ reached the loss"
    return labels

def debug_mask(tokenizer, input_ids, assistant_mask, n=300): #
    ↪ eyeball it
    for tok, m in
    ↪ zip(tokenizer.convert_ids_to_tokens(input_ids[:n]),
    ↪ assistant_mask[:n]):
        print(("TRAIN " if int(m) else "-- ") + tok)
```

The invariants, and exactly which bug each one catches:

Mask invariant	Failure it catches
<code>len(mask) == len(input_ids)</code>	tokenization / template mismatch
<code>mask.sum > 0</code>	no assistant spans marked → training on nothing
tool-result spans all mask 0	hallucinating the environment
assistant tool-call spans mask 1	model never learns the action syntax
user / system spans mask 0	model imitates the prompt instead of acting
rendered free-run output parses	loss falling while generation is broken (Ch 7)

Table 106.2. The mask invariants every multi-turn run should assert, and exactly which failure each one catches. Each check is cheap to add and pins a specific silent failure — a tokenization mismatch, an empty mask, a hallucinated environment, or a model that imitates the prompt instead of acting.

These habits matter doubly on Qwen3.5, where, as the next chapter shows, whether the chat template emits the masking markers at all comes down to a single piece of template syntax, and getting it wrong silently kills the SFT run.

TAKEAWAY

Agent SFT must mask exactly as [Chapter 21](#)’s RL rollouts do: loss only on policy-produced tokens (reasoning, tool calls, final answers), never on tool results, pages, or user/system turns. Training on tool outputs teaches the model to *hallucinate the environment*, inventing prices and confirmation pages, while the loss falls smoothly and hides it. Build the mask segment-by-segment, assert length/values/nonzero-sum, and eyeball one rendered trajectory. On Qwen3.5 a single template detail decides whether the mask exists at all, the next chapter’s silent killer.

Next: that template detail, the Qwen3.5 `{% generation %}` requirement, and how its absence silently kills SFT.

CHAPTER 107

The Masking Bug That Silently Kills SFT

Template	Has gen marker?	Result
Qwen3.5 default	no	<code>assistant_only_loss=True</code> returns an all-zero assistant mask loss on zero tokens (or a <code>RuntimeError</code>); SFT learns nothing
Patched template	yes	<code>assistant_only_loss=True</code> marks assistant spans only loss on the right tokens; SFT works

Table 107.1. Qwen3.5’s chat template ships without the `{% generation %}` markers that masking libraries rely on. Turn on assistant-only loss and the framework computes loss on an empty mask: the run trains, the loss may even fall, and the model learns nothing useful. The fix is a two-line template edit. Red is the broken default; green is the patched template.

The previous chapter said tool-call masking is non-negotiable. This chapter is the specific, vicious way it fails on Qwen3.5: the chat template does not include the `{% generation %}` markers that Hugging Face TRL’s `assistant_only_loss=True` uses to find the assistant tokens, so turning on assistant-only masking either throws a `RuntimeError` or, worse, silently computes the loss on zero tokens. The Qwen team rejected community pull requests to add the markers, on the position that masking belongs at the framework level, not in the template. So this is yours to fix, and it is invisible until you look.

107.1 Why it's silent, and which way it breaks

The bug has two failure modes, and the dangerous one is the quiet one:

- It throws. Some TRL versions raise a `RuntimeError` when `return_assistant_tokens_mask=True` finds no `{% generation %}` markers. Annoying, but at least loud.
- It trains on nothing, or everything. Otherwise the assistant mask comes back all zeros, and the loss is computed over an empty set. The run proceeds, the loss curve can even look plausible, and the model learns nothing. The mirror-image failure is masking off, where you instead train on all tokens, system prompts, tool results, page DOM, diluting the signal with text the model should never imitate (the hallucinate-the-environment failure of [Chapter 106](#)).

Either way, the symptom is exactly the one [Chapter 106](#) warned about: a smoothly falling loss hiding a model that learned the wrong thing or nothing at all. This is why that chapter's `assert assistant_mask.sum > 0` exists. It converts this silent bug into a loud one.

107.2 The fix: inject the markers, then assert

The patch wraps assistant content in `{% generation %}...{% endgeneration %}` so the masking machinery can find it:

```
# patch Qwen3.5's chat template: wrap ASSISTANT content in
↪ generation markers
patched = base_template.replace(
    "{% if message.role == 'assistant' %}{{- '<|im_start|>' +
    ↪ message.role + '\n' }}{{- content }}",
    "{% if message.role == 'assistant' %}{{- '<|im_start|>' +
    ↪ message.role + '\n' }}"
    "{% generation %}{{- content }}{% endgeneration %}"
)
tokenizer.chat_template = patched

enc = tokenizer.apply_chat_template(
```

```
messages, return_assistant_tokens_mask=True,  
    ↪ return_dict=True, tokenize=True)  
assert sum(enc["assistant_masks"]) > 0, "generation markers  
    ↪ missing - mask is all zero"
```

Then the [Chapter 106](#) invariants apply unchanged: mask length equals token length, the sum is greater than zero, tool-result spans are masked to zero. The lesson generalizes past this one model: never trust that a framework flag is doing what its name says, assert the mask it produces. A flag named `assistant_only_loss` that silently masks nothing is exactly the kind of seam ([Chapter 24](#)) that costs a week if you trust it and thirty seconds if you check it.

TAKEAWAY

Qwen3.5's chat template omits the `{% generation %}` markers TRL's `assistant_only_loss=True` needs, so assistant-only masking either throws or, worse, returns an all-zero mask and trains on nothing while the loss looks fine (the Qwen team rejected PRs to add them). Patch the template to wrap assistant content in `{% generation %}...{% endgeneration %}`, then `assert sum(assistant_masks) > 0`. Never trust a masking flag by its name. Assert the mask it actually produces.

Next: the preference stage, turning guardrail verdicts into DPO pairs, one decision at a time.

CHAPTER 108

DPO from Guardrail Verdicts

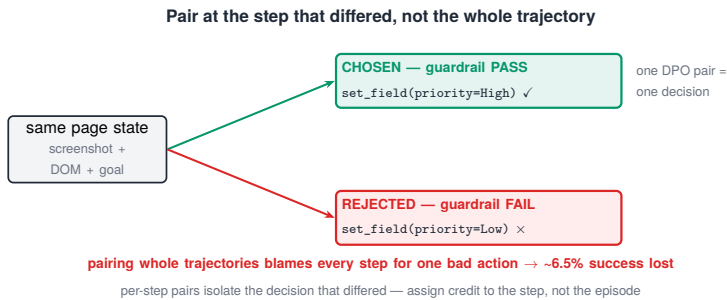


Figure 108.1. Pair at the step that differed, not the whole trajectory:
per-step DPO pairs (same page state, one passing action vs one failing) instead of blaming every step.

The second stage of the pipeline ([Chapter 104](#)) is preference tuning, and the agent’s preferences come from a guardrail, a checker that watches each step and rules it PASS or FAIL. Turning those verdicts into Direct Preference Optimization (DPO) data has one design decision that matters more than the rest: build the pairs per step, not per trajectory.

108.1 The guardrail signal, and the per-step pair

Each Claude trajectory carries a list of guardrail verdicts, for every step, the action taken, a PASS/FAIL verdict, the reason, and the page state (URL, a DOM hash, the goal). A failed step looks like: the agent set the priority field to `Low` when the goal specified `High`, verdict FAIL. The pair is built by finding two attempts that reached the same page state and diverged:

chosen = the guardrail-passing action at that state · rejected = a guardrail-failing action from a different attempt at the same state · shared prompt = the screenshot + DOM + goal

So a single pair isolates one decision, “at this page, setting priority to `High` is preferred to setting it to `Low`”, with everything else held fixed. In code, the construction groups steps by page state and contrasts verdicts:

```
def build_dpo_pairs(logs):
    by_state = {} # group steps by (url, dom_hash, goal)
    for log in logs:
        for gv in json.loads(log["guardrail_verdicts"]):
            by_state.setdefault((gv["url"], gv["dom_hash"],
                                ↪ gv["goal"]), []).append(gv)
    pairs = []
    for state, steps in by_state.items():
        passed = [s for s in steps if s["verdict"] == "PASS"]
        failed = [s for s in steps if s["verdict"] == "FAIL"]
        for p in passed: # one prompt, chosen vs rejected action
            for f in failed:
                pairs.append({"prompt": p["prompt"], "chosen":
                              ↪ p["action_text"],
                              "rejected": f["action_text"]})
    return pairs
```

108.2 Why per-step beats per-trajectory

The tempting shortcut is to pair whole trajectories, a successful run as “chosen,” a failed run as “rejected.” It is wrong for the same reason naive credit assignment is always wrong: a failed trajectory is usually mostly correct, with one bad step. Pairing the whole run teaches the model to disprefer every action in it, including the good ones, a credit-assignment misallocation that, per the trajectory-level approach (TGPO) the source measured, costs about 6.5% success rate. Per-step pairing puts the preference exactly on the decision that differed, and nowhere else. The training itself is ordinary TRL, `DPOConfig` and `DPOTrainer` over these pairs, but the data construction is where the result is won or lost. The general principle, recurring from [Chapter 21](#)’s masking onward: in agent training, assign credit to the step, not the episode, wherever your signal lets you.

TAKEAWAY

The agent's preferences come from a guardrail that rules each step PASS or FAIL. Build DPO pairs **per step**, same page state, passing action as chosen, a failing action from another attempt as rejected, not per trajectory, because pairing whole runs blames every step for one bad decision and costs ~6.5% success (the TGPO credit-assignment error). The DPO training is standard TRL. The per-step data construction is the part that matters.

Next: the on-policy stage, GSPO with the model driving a real browser, and the one config flag that silently corrupts the loss.

CHAPTER 109

GSPO with Online Playwright Rollouts

The config that matters, and the trap inside it.

```
from trl import GRPOConfig
cfg = GRPOConfig(
    importance_sampling_level="sequence", # this is what makes
    ↪ it GSPO
    loss_type="grpo", # NOT "dapo"/"bnpo" - see below
    num_generations=8, temperature=1.0)
```

With `importance_sampling_level="sequence"`, the per-token loss tensor is shape $(B, 1)$. Under `loss_type="dapo"` (the default) that $(B, 1)$ broadcasts against `completion_mask` (B, T) , multiplying every advantage by the token count, a silent corruption (TRL issue #3823). Use `loss_type="grpo"`.

The third stage is the one the whole architecture argument of [Part V](#) was about: Group Sequence Policy Optimization (GSPO), now driving a real (sandboxed) browser. This is the only stage that is genuinely on-policy, the model generates its own trajectories, against a reward, and learns from them. Two things decide whether it works: the sequence-level switch, and avoiding a config trap that quietly breaks the loss.

109.1 Sequence-level, and the broadcasting trap

GSPO is standard Group Relative Policy Optimization (GRPO) with one switch: `importance_sampling_level="sequence"` computes a single importance-sampling ratio per trajectory, the geometric mean of the token-level ratios, instead of one ratio per token. For long agent trajectories with episodic reward (you only learn at the

end whether the task was completed correctly), that sequence-level ratio is far more stable, for exactly the reasons [Chapter 50](#) gave for mixture-of-experts models: it does not amplify per-token ratio noise across a long rollout.

The trap is that the obvious-looking default is wrong. When the ratio is sequence-level, the per-token loss tensor has shape $(B, 1)$, one value per sequence. Under `loss_type="dapo"` or `"bnpo"`, the normalization step divides by active tokens, and that $(B, 1)$ tensor broadcasts against the (B, T) completion mask, silently multiplying each advantage by the number of non-padding tokens (TRL issue #3823). The fix is one word, `loss_type="grpo"`, and it is the kind of bug that produces a plausible-looking but quietly wrong run, the recurring theme of this Part.

109.2 The reward is the verifier, and the loop is on-policy

The reward for each rollout is task success, scored by the verifier of [Chapter 111](#): did the agent navigate to the right page, submit the values matching the goal, and reach the intended final state in the sandbox. Unlike the offline data of [Chapters 105](#) and [108](#), this stage is on-policy, Qwen generates the trajectories itself, in the browser, with its current policy, and the GSPO update uses those rollouts directly. That is what makes it the only stage that can teach genuine task judgment (when to clarify, which path to take), and it is why it is also the most expensive and fragile: every gradient step requires real browser rollouts, and the reward must be hard to game ([Chapter 11](#)'s cold-start lesson and [Chapter 32](#)'s gate, now with a real user's data and real side effects behind the sandbox wall).

TAKEAWAY

Online GSPO is the only on-policy stage, the model drives the sandboxed browser, is scored on task success by the verifier, and learns from its own rollouts. GSPO = GRPO with `importance_sampling_level="sequence"` (a single geometric-mean ratio per trajectory, stable for long episodic-reward rollouts). The trap: keep `loss_type="grpo"`, the default

"dapo"/"bnpo" makes the $(B, 1)$ sequence-loss broadcast against the (B, T) mask and silently scale advantages by token count (#3823).

Next: when you can't run a live browser for every step, offline GSPO from pre-collected groups, and the one trick that keeps the math honest.

CHAPTER 110

Offline GSPO and the Custom `rollout_func`

The non-obvious fix (from the LUFFY work). TRL's `GRPOTrainer` is fundamentally on-policy, but its experimental `rollout_func` lets you inject pre-collected trajectory groups. For groups generated by a different model (Claude), do not use Claude's log-probabilities as `old_log_probs`. Compute the student's own log-probs for those tokens (one forward pass) and use those. This sidesteps the tokenizer mismatch and keeps the importance-sampling ratio meaningful.

Running a live browser rollout for every gradient step is expensive, and sometimes you want to train GSPO from trajectories you already have. TRL allows it through a custom `rollout_func`, but it sits on a sharp edge, because `GRPOTrainer` is built to be on-policy, and feeding it off-policy data the wrong way produces meaningless ratios. This chapter is how to do it, and the honest boundary on what it is.

110.1 The mechanism, and the old-log-probs trick

The `rollout_func` parameter is the seam: instead of generating rollouts live, you hand the trainer pre-collected groups (the same group structure GRPO needs, several attempts per prompt). The non-obvious part is what to use for the old log-probabilities the importance ratio divides by. The instinct is to use the log-probs from whatever generated the trajectories, but those came from Claude, with Claude's tokenizer and distribution, so the ratio between the student's new log-probs and Claude's old log-probs is comparing two different models through two different tokenizers, and means nothing.

The fix, from the LUFFY work, is to recompute the old log-probs under the student, one forward pass of the current student model

over the pre-collected tokens, and use those as `old_log_probs`. Now the ratio compares the student to a recent version of itself over the same tokens, which is what the math assumes:

```
class OfflineRolloutFunc:
    def __call__(self, prompts, model, processor, **kw):
        groups = self.lookup(prompts)  # pre-collected trajectory
        ↪ groups
        # CRITICAL: old_log_probs = the STUDENT's own log-probs
        ↪ for these tokens,
        # NOT the teacher's -- one forward pass, no tokenizer
        ↪ mismatch.
        old_logp = student_logprobs(model, groups.input_ids)
        return {"completion_ids": groups.input_ids,
                ↪ "old_logprobs": old_logp,
                ↪ "rewards": groups.rewards}
```

110.2 The honest boundary: this is not true on-policy

The caution from [Chapter 105](#) applies in full force here, because it is easy to talk yourself out of it. Recomputing the student's log-probs makes the ratio well-formed, but it does not make the data on-policy. The trajectories were still generated by Claude, from Claude's distribution, not sampled from the student's current policy. So offline GSPO is a legitimate engineering tool for bootstrapping or for cheap iteration, but it is not mathematically the same as on-policy GSPO, and it cannot fully replace the live-rollout stage of [Chapter 109](#). Treat it as a warm start and a debugging convenience. Treat the online loop as the real thing. The general rule, stated once more because it is the most common way to fool yourself in agent RL: a well-formed importance ratio is necessary but not sufficient, on-policy also requires that the policy generated the samples.

TAKEAWAY

Offline GSPO injects pre-collected groups via TRL's `roll-out_func`. The key trick (from LUFFY): set `old_log_probs` to the **student's own** log-probs for those tokens (one forward pass), never the teacher's, otherwise the ratio compares two

models through two tokenizers and is meaningless. But a well-formed ratio doesn't make data on-policy: the trajectories still came from Claude's distribution, so offline GSPO is a warm-start/debug tool, not a replacement for the live-rollout stage of [Chapter 109](#).

Next: the machinery that makes online rollouts possible, the browser environment and the verifier that scores them.

CHAPTER 111

Environment Orchestration and Verifiers

Component	What it is	RL role
Playwright MCP tools	<code>browser_click</code> , <code>browser_type</code> , <code>browser_select_option</code> , <code>browser_navigate</code> , <code>browser_snapshot</code> , <code>browser_take_screenshot</code>	the agent's action space
Environment factory	TRL wiring (transformers $\geq 5.2.0$) that runs tools during rollout	turns tool calls into real browser actions + observations
the verifier (guardrail)	a program that checks each step and the final state	the reward; scores task success, hard to game

Table 111.1. The online GSPO loop: tools, environment factory, and verifier.

Online GSPO (Chapter 109) is only as good as the environment it rolls out in and the verifier that scores it. This chapter is that machinery: the browser tools the agent acts through, the orchestration that runs them, and the verifier that turns a trajectory into a reward, which, as every reward chapter in this book has insisted, is where the whole thing can quietly go wrong.

111.1 The action space and the orchestration

The agent acts through a small set of Playwright browser tools, exposed as Model Context Protocol (MCP) functions: `click`, `type`, `select an option`, `navigate`, `take a structured snapshot of the page`, `take a screenshot` (Table 111.1). At each step the agent receives a screenshot and a Document Object Model (DOM) snapshot, reasons in a `<think>` block, and emits exactly one tool call. TRL's `environment_factory`

(requiring transformers $\geq 5.2.0$) is the wiring that executes those calls against a live browser during rollout and feeds the resulting observation back to the model. The version pin is load-bearing, tool calling through the environment factory simply does not exist before that release, and a silent version regression here breaks rollouts in confusing ways (the seam lesson of [Chapter 24](#), again).

111.2 The verifier is the reward, so it must resist gaming

The verifier is a program that watches the trajectory and decides whether the task succeeded: did the agent navigate to the right page, submit the values that match the goal, and reach the intended final state, in the sandbox ([Chapter 104](#)). This is the reward function, and everything [Part III](#) taught about reward design applies with real stakes. A verifier that checks a proxy, “did the form submit” without “are the values correct”, is gameable exactly as the search agent’s F1 proxy was ([Chapter 32](#)): the agent will learn to complete a submission, not the requested task. So the verifier must check the goal, not a stand-in for it: the right page, the right field values, the intended final state, and only then success. Build it as a gate ([Chapter 32](#)), task success counts only when the correct outcome is reached, and keep it in the sandbox, because a reward function with a bug is a reward function the agent will exploit, and here the exploit takes real actions.

TAKEAWAY

The online loop needs three wired pieces: Playwright MCP tools (the action space, click/type/select/navigate/snapshot/screenshot), TRL’s `environment_factory` (transformers $\geq 5.2.0$, executes calls in a sandboxed browser), and a verifier that scores the trajectory. The verifier *is* the reward, so it must check the goal, correct page, field values, intended final state, not a gameable proxy like “the form submitted,” or the agent learns to complete *a* task, not the *right* one. Build it as a gate ([Ch 32](#)), in the sandbox.

Next: the bug that bites after training is done, when the serving format doesn't match the training format.

CHAPTER 112

Serving/Training Parity

Aspect	Training	Serving
Tool-call format	<code><tool_call><function=name><parameter=k>v</parameter></function></tool_call></code>	the serving parser must decode the same
Chat template	patched Qwen3.5 template (Chapter 107)	identical template
Weights	LoRA adapters on the base	merge LoRA before serving
Result if mismatched	—	model emits tool calls the server can't parse — a silent production break

Table 112.1. Training and serving must match on three things — tool-call format, chat template, and how LoRA weights are loaded. If any differ, the model emits tool calls the serving parser cannot decode, a silent production break that no eval catches. The red row is what a mismatch costs.

The agent is trained. Now it has to serve, and there is one bug that is invisible until a real request hits production: the serving format doesn't match the training format. The model learned to emit tool calls in a specific syntax. If the serving parser expects a different one, every tool call the model produces is unparseable, and the agent fails in production despite a flawless training run. This is the deployment-side version of the seam bug (Chapter 24), and it is the last thing standing between a trained checkpoint and a working product.

112.1 Match the format, exactly

The model trains to emit tool calls as XML, `<tool_call><function=name><parameter=key>value</parameter></function></tool_call>`, so the serving stack's tool-call parser must decode that exact format.

On vLLM, two launch flags establish parity: the tool-call parser (set to the Qwen format) and the matching chat template (the patched one from [Chapter 107](#), not the stock template). And because the agent was trained with LoRA adapters, you must merge the LoRA into the base weights before serving, or the served model is not the model you trained. Miss any of these and the failure is silent in the worst way: the model generates perfectly reasonable-looking tool calls that the server cannot turn into actions.

112.2 The principle: train and serve are one contract

The transferable rule is that the format you train on is a contract you must honor at serving time, the same tokenizer, the same chat template, the same tool-call syntax, the same special tokens. This is the inference-side companion to [Chapter 4](#)'s weight-sync invariant (the trainer and the rollout engine must agree) and [Chapter 24](#)'s seam discipline (the places two systems meet are where assumptions silently diverge). Before declaring an agent shipped, run one real request end to end through the serving stack and confirm the tool calls parse and execute, the same fail-loud habit this Part keeps returning to, now as the final gate before production.

TAKEAWAY

A training/serving format mismatch is the most insidious deployment bug: the model trains to emit a specific XML tool-call syntax, and if the serving parser expects a different one, every tool call fails in production despite perfect training. On vLLM, set the tool-call parser and the *patched* chat template ([Ch 107](#)) to match training, and merge LoRA into the base before serving. The format you train on is a contract you must honor at serving, verify it end-to-end before shipping.

Next: the whole pipeline on a calendar, what to build each week, and the hardware it runs on.

CHAPTER 113

Build Order and Promotion Gates

Week	Build	Gate
Week 1	data conversion (Ch 105) + SFT, with the { % generation % } fix (Ch 107)	free-run generation produces parseable tool calls
Week 2	DPO from per-step guardrail pairs (Ch 108)	guardrail pass-rate up vs the SFT checkpoint
Week 3	GSPO — offline warm-start (Ch 110) then online Playwright rollouts (Ch 109)	sandbox task-success rising; no entropy collapse
Week 4	serving parity (Ch 112) + held-out eval (safety + success composite, Ch 104)	one real request executes end-to-end in the serving stack

Table 113.1. The pipeline as a four-week sequence, each phase gated on the previous one — the cost-and-dependency ordering of Chapter 2 made into a calendar. “Build” is what you do that week; “Gate” is the check that must pass before the next week starts.

This chapter is the whole production pipeline on a calendar, not because four weeks is a law, but because the ordering is, and seeing it sequenced makes the dependencies concrete. It is also where the hardware reality lands, because a plan you can’t fit on your GPUs is not a plan.

113.1 The hardware, and the QLoRA fallback

The reference setup is 4×A100-80GB with DeepSpeed ZeRO Stage 2. The model is Qwen3.5-9B, which is natively multimodal. It processes the screenshots directly, with no separate “-VL” variant to wire in. If you only have 4×A10-24GB, there is a documented QLoRA fallback:

the 9B in 4-bit NormalFloat (NF4) is about 5 GB per GPU, plus LoRA adapters (~0.5 GB), paged 8-bit AdamW optimizer states (~1–2 GB), and gradient-checkpointed activations (~3–5 GB), roughly 10–12 GB per GPU, fitting 24 GB with headroom. Two fallback rules matter: freeze the vision tower (QLoRA cannot safely quantize the vision modules), and use ZeRO-2, not ZeRO-3 (ZeRO-3 is incompatible with QLoRA here, per Qwen’s documentation).

113.2 The sequence, and why the order is fixed

The four phases (Table 113.1) are [Chapter 2](#)’s pipeline made into weeks, and each gate is there because skipping it wastes the next phase. Week 1 converts Claude trajectories and runs SFT with the masking fix, and you do not leave it until free-run generation (not teacher-forced) produces parseable tool calls, because DPO and GSPO have nothing to refine otherwise. Week 2 does per-step DPO from guardrail verdicts, gated on the guardrail pass-rate actually improving. Week 3 runs GSPO, an offline warm-start, then the real online browser loop, gated on sandbox task-success rising without entropy collapse (the failure of the next Part). Week 4 establishes serving parity and runs the held-out safety-and-success eval, gated on one real request executing end-to-end. The order is fixed because it is a dependency chain, not a preference: each phase needs the previous phase’s output to exist before it can add anything.

TAKEAWAY

The production pipeline is a four-week dependency chain (SFT → DPO → GSPO → serving+eval), each phase gated on the previous ([Ch 2](#)’s ordering as a calendar). Reference hardware is 4×A100-80GB on ZeRO-2. The 4×A10-24GB QLoRA fallback fits the 9B at ~10–12 GB/GPU (NF4 ~5 GB + LoRA + paged-8bit optimizer + checkpointed activations), but freeze the vision tower and use ZeRO-2, not ZeRO-3. Qwen3.5-9B is natively multimodal: no separate -VL variant.

Next: [Part XII](#). A production pipeline has more failure surfaces than a benchmark. The next Part indexes those failures by symptom.

Part XII

The Failure Catalog

This book has hit a lot of failures in narrative. This Part is the reference version: the named failure modes, their early-warning signatures on the dashboard, their causes, and their fixes, collected so you can match a symptom to a diagnosis at 2 a.m. without re-reading ten chapters. It is organized as a catalog, not a story, look up what you're seeing.

CHAPTER 114

Reading a Training Dashboard

The panels warn 30–100 steps before the reward breaks

Set each line BEFORE the run; act when crossed — don't wait for the reward to confirm the death.




Figure 114.1. The panels warn 30–100 steps before the reward breaks: entropy sliding, reward-std collapsing, KL spiking, gradient-norm spiking — set each line before the run.

Most failed runs announce themselves 30 to 100 steps before the reward curve visibly breaks, on panels people don't watch. This chapter is the dashboard: the handful of instruments that give you that early warning, and the thresholds that turn a wiggle into a decision. The failures themselves are catalogued in the chapters that follow. This is how you see them coming.

114.1 The panels that matter

Five signals carry almost all the early warning:

- Entropy (per token). The single most predictive panel. A monotone slide from ~ 0.4 toward < 0.05 is impending mode collapse (Chapter 115's F1), and it shows before generations visibly degrade.

- Reward, and reward-standard-deviation. The mean is what you hope to see. The std is the early-warning. When reward-std collapses toward zero, the groups have stopped disagreeing ([Chapter 5](#)), no gradient is coming, whatever the mean says.
- KL to the reference. A sudden spike means the policy is lurching away from the base, often the signature of reward hacking ([Chapter 116](#)'s F4) or instability.
- Gradient norm. Spikes precede blow-ups. This is why a clip (e.g. `max_grad_norm = 1.0`) is standard.
- Format-error rate. For a tool-using agent, a rising rate of unparseable outputs is an early sign the model is drifting off the trained format.

114.2 The thresholds that become actions

A panel is only useful with a line on it. The defaults worth starting from: an entropy floor of 0.10, breached for ~10 consecutive steps, mode collapse is imminent (F1), intervene now. A maximum KL around 0.5, beyond it, the policy has wandered too far. And a gradient-norm clip of 1.0. These are starting points, not universals (measure your own, per [Chapter 88](#)'s lag caution), but the discipline is the point: decide the line before the run, and act when it's crossed instead of watching the reward curve confirm a death you could have caught 50 steps earlier. The dashboard is the cheapest insurance in RL, the failures in the next two chapters are all things it warns about first.

TAKEAWAY

Runs warn 30–100 steps before the reward curve breaks, on five panels: per-token entropy (the most predictive, a slide from 0.4 toward <0.05 is impending collapse), reward *and* its standard deviation (std→0 means groups stopped disagreeing, no gradient), KL to reference (spikes = lurching/hacking), gradient norm (spikes precede blow-ups), and format-error rate. Set lines before the run, entropy floor 0.10 (breached 10 steps =

intervene), max KL ~ 0.5 , grad-norm clip 1.0, and act when crossed.

Next: the first failure family, collapse, where the model stops exploring.

Collapse Failures

Failure	Early signature	Cause	First fixes
F1 — entropy collapse	entropy 0.4 $\rightarrow$ <0.05 reward-std $\rightarrow$ 0; identical generations KL spikes (steps 20–50)	symmetric clipping caps probability increases	Clip-Higher 0.28 Clip-Cov; entropy bonus 1e-3 cosine / smaller LR
F2 — advantage collapse	reward-std $\rightarrow$ 0 with reward-mean $\neq$ 0 grad-norm crashes (steps 30–80)	all-pass or all-fail groups $\rightarrow$ zero advantage	DAPO dynamic sampling continuous shaping ($n \geq 16$); filter too-easy / too-hard prompts
F13 — late collapse (cosine miss)	stable ~80 steps, then sudden collapse constant LR at full peak while entropy drops (80–200)	constant learning rate + falling entropy	cosine LR decay (our SW6 fix) early-stop on the entropy floor

Table 115.1. The collapse family: the model stops exploring. F1 is the canonical case; F2 is a data/sampling problem and F13 a learning-rate-schedule problem, so the fixes differ. The dashboard’s entropy and reward-std panels ([Chapter 114](#)) are the early warning.

The first failure family is collapse, the model stops exploring and converges to a narrow, often degenerate, behavior. It is the most common way an RL run dies, and the dashboard’s entropy and reward-std panels ([Chapter 114](#)) are its early warning. This chapter is the catalog entry. The mechanism and the cures were [Part IV](#)’s subject, cross-referenced here.

115.1 F1: entropy collapse, the canonical case

F1, entropy collapse is the one you will see most. The signature is unmistakable on the dashboard: per-token entropy slides monotonically from ~ 0.4 to below 0.05, reward-standard-deviation falls to zero, generations become identical across a group, and KL spikes, typically over steps 20–50. The cause is structural: symmetric clipping (the same clip range up and down) caps how much the policy can increase a good token's probability, so the distribution sharpens into a single mode it can never climb back out of. We hit this hard on Qwen3-14B (the SW1 run). The fixes are the entropy-management toolkit of [Part IV](#): Clip-Higher (raise the upper clip to ~ 0.28 so good tokens can grow), Clip-Cov, an entropy bonus ($\sim 1e-3$), a cosine learning-rate schedule, and simply a smaller learning rate.

115.2 F2 and F13: advantage collapse and late collapse

Two more collapses round out the family, and the dashboard tells them apart from F1. F2, advantage collapse is the [Chapter 5](#) starvation case: reward-std goes to zero, but here the reward mean is nonzero and the groups are unanimous, every attempt in a group scores the same (all-pass or all-fail), so the advantage is zero and no gradient flows, typically over steps 30–80 as grad-norm crashes. The tell that separates it from F1 is the zero-variance fraction ([Chapter 12](#)): unanimous groups, not a sharpening distribution. The fix is the data and sampling, not entropy, DAPO-style dynamic sampling, continuous shaping, a larger group ($n \geq 16$), and filtering prompts that are too easy or too hard (the group-diversity probe of [Chapter 17](#) catches this up front). F13, late collapse is the one that fools you, because the run looks healthy for a long time: it trains stably for ~ 80 steps and then collapses suddenly, between roughly steps 80 and 200. The cause is mechanical, a constant learning rate held at full peak while entropy quietly drops, until the policy tips over. This was the SW5→SW6 scar, and the fix was a cosine learning-rate decay (so the step size shrinks as the policy sharpens) plus an early-stop on the entropy floor ([Chapter 114](#)). It is a collapse like F1, but its trigger is the LR schedule, not the clip.

TAKEAWAY

Collapse = the model stops exploring. **F1 entropy collapse** (the common one): entropy $0.4 \rightarrow < 0.05$, reward-std $\rightarrow 0$, identical generations, KL spikes over steps 20–50, from *symmetric* clipping capping probability increases, fix with Clip-Higher 0.28, Clip-Cov, entropy bonus, cosine/smaller LR (**Part IV**). **F2 advantage collapse**: reward-std $\rightarrow 0$ but reward-mean $\neq 0$, unanimous groups, grad-norm crashes (30–80), fix the data/sampling (DAPO dynamic sampling, $n \geq 16$, filter, continuous shaping), not entropy. **F13 late collapse**: stable ~ 80 steps then sudden collapse from a *constant* LR held at peak while entropy falls (80–200), fix with cosine LR decay (the SW6 fix) and an entropy-floor early-stop. The entropy and reward-std panels (**Ch 114**) catch all three early.

Next: the family where the reward goes up while the model gets worse.

Reward-Driven Failures

Failure	Early signature	Cause	First fixes
F3 — length bias	completion-length p50 climbs while correctness plateaus tokens-per-correct rises (steps 100–200)	token-mean loss rewards verbose correct answers	<code>seq_mean_token_mean</code> aggregation DrGRPO; a hard length cap
F4 — reward hacking (repetition)	reward flat/up, KL $\uparrow$, generations nonsense LLM-judge $\downarrow$ (steps 50–150)	shaped reward + weak KL	KL ≥ 0.005 ; repetition penalty overlong shaping

Table 116.1. The reward-driven family: the number you’re optimizing climbs while the thing you care about drops. The tell is divergence between the reward and a separate held-out quality measure — which is exactly why [Chapter 88](#) insisted the held-out eval be independent.

The second failure family is reward-driven: the reward climbs while the model gets worse at the actual task, or merely more verbose at no real gain. This is the catalog version of [Chapter 31](#)’s lived reward-hacking story and [Chapter 88](#)’s “the training score misleads.” The defining signature is divergence, the reward going one way while a separate held-out quality measure goes the other, which is the whole reason the held-out eval has to be independent of the reward.

116.1 F4: repetition hacking

F4, reward hacking via repetition is the classic. The signature: reward stays flat or climbs, KL rises, and the generations degrade into nonsense or loops, while an independent LLM-judge score falls, usually

over steps 50–150. The cause is a shaped reward paired with a weak KL penalty: the shaping rewards surface features (length, keywords, format) that the model can satisfy with degenerate text, and a weak KL leash lets it wander there. The fixes are direct: a KL coefficient of at least 0.005 (keep the policy near the base), a repetition penalty, and overlong shaping so length alone can't be farmed. The general defense, from [Chapter 88](#), is that the reward cannot police itself, only the separate held-out judge reveals the divergence.

116.2 F3: length bias

F3, length bias is the subtler one, because the reward really is going up and the answers really are (often) correct, they are just getting longer and longer for no gain. The signature is a completion-length p50 that climbs monotonically while correctness plateaus, with tokens-per-correct-answer rising, typically over steps 100–200. The cause is the loss aggregation: a token-mean objective averages per-token loss, which quietly favors verbose correct completions over concise ones, so the policy learns to pad. This was the SW4–SW6 scar, and the fix lives in the aggregation, not the reward: switch to `seq_mean_token_mean` (or DrGRPO), which stops paying per extra token, and add a hard length cap as a backstop ([Chapter 84](#)'s shaping discipline and the loss-aggregation chapter apply directly). A related reward-driven trap worth naming because it looks similar on the reward curve is the format/accuracy split (pitfall P13): the format reward converges to 1.0 while accuracy stays at 0, because the model learned the format token-pattern with no content (or the extractor is mismatched, LaTeX vs plain). There the fix is to gate format behind correctness ([Chapter 32](#)), use `math-verify` not `regex`, and anneal the format weight to zero once format reaches ≥ 0.95 . F3 itself, though, is specifically the length pathology.

TAKEAWAY

Reward-driven failures = the reward climbs while the task degrades. The tell is divergence from a *separate* held-out measure (Ch 88). **F4 repetition hacking**: reward flat/up, KL $\uparrow$, nonsense generations, judge $\downarrow$ over steps 50–150, from shaped reward + weak KL, fix with KL ≥ 0.005 , repetition penalty, over-long shaping. **F3 length bias**: completion length climbs while correctness plateaus (tokens-per-correct rises, steps 100–200) because a *token-mean* loss rewards verbose correct answers, fix the aggregation (`seq_mean_token_mean` / DrGRPO) and add a hard length cap, not the reward. (The related `format $\rightarrow$ 1.0` / `accuracy $\rightarrow$ 0` trap is pitfall P13: gate format behind correctness.)

Next: the algorithm-config and MoE failure family, the failures that look like bugs but are settings.

MoE and Algorithm-Config Failures

Failure	Early signature	Cause	First fixes
F5 — ratio explosion (MoE)	<code>ratio_max</code> grows fat tails	~10% of experts flip routing per step → token-level ratios become numerically invalid, reused the GRPO clip (0.2) on GSPO	switch to GSPO (sequence-level)
	grad spikes → NaN		lower LR
	<code>clip_frac_high</code> → 1 (<50 steps on MoE)		routing-replay does <i>not</i> fix it
F14 — GSPO clip misconfig	instant collapse, NaN loss	invalid, reused the GRPO clip (0.2) on GSPO	read the GSPO paper's hyperparameters
	grad-norm ~1e6 — on the first step		use the sequence-level clip ($\approx 3e-4/4e-4$)

Table 117.1. Two failures that look like bugs but are settings. F5 is why Group Sequence Policy Optimization (GSPO) is the safer default for mixture-of-experts (MoE) models (Part V). F14 is GSPO punishing you for bringing Group Relative Policy Optimization (GRPO) habits to it.

The next several chapters are the consolidated failure catalog: named failures, their dashboard signatures, and their fixes, grouped so you can match a symptom to a diagnosis fast. This chapter holds the two that are not bugs at all. They are configuration mistakes that produce spectacular crashes, and they bracket the whole MoE-and-GSPO story of Part V.

117.1 F5, ratio explosion on MoE

F5, ratio explosion is the failure that forced sequence-level training. On a mixture-of-experts model, the router re-selects which experts handle each token, and across a single policy update roughly 10%

of experts flip between the old policy and the new one. Token-level importance-sampling ratios assume the per-token distributions are comparable across that update. When the active experts change, that assumption breaks, the ratio's denominator becomes meaningless, and `ratio_max` grows fat tails until gradients spike to NaN, usually in under 50 steps. The dashboard tell is `clip_frac_high` climbing toward 1: nearly every token is being clipped, which is the optimizer screaming that the ratios are invalid. The fix is the architecture-to-algorithm coupling of [Chapter 99](#): GSPO, whose single sequence-level ratio is robust to routing flips, is the safer default here, but treat that as a default, not a theorem. The corpus is more nuanced than “MoE must use GSPO”: with large expert redundancy and a conservative learning rate, vanilla GRPO can stay stable for some horizons (the scale-dependent finding of [Chapter 99](#)). What is reliable is the fragility, not a guarantee of instant failure, so on MoE, either move to GSPO or monitor the sequence-ratio histogram closely and lower the learning rate (2–10× below the dense default. Routing-replay sounds appealing but does not fix it).

117.2 F14, GSPO clip misconfiguration

F14 is the trap waiting on the other side of that fix. GSPO is not GRPO with a flag. It has a different clip regime, and reusing GRPO's clip value of 0.2 on GSPO makes the loss go NaN on the very first step, with grad-norm around $1e6$. The cause is purely that a sequence-level ratio operates on a different scale than a token-level one, so the token-level clip is wildly wrong. The fix is to read the GSPO paper's hyperparameter section and use the sequence-level clip (on the order of $3e-4$ to $4e-4$). The lesson uniting F5 and F14 is the one this chapter is named for: some catastrophic-looking failures are not bugs in your code or your data. They are a number in your config, and the fix is to match the algorithm's documented settings rather than to debug the gradient.

TAKEAWAY

Two config-not-code failures bracket the MoE story. **F5 ratio explosion**: on MoE, ~10% of experts flip routing per step → token-level ratios go invalid → `ratio_max` fat tails, `clip_frac_high`→1, NaN in <50 steps. GSPO (sequence-level) is the *safer default*, with a lower LR; routing-replay alone does not fix the instability, but MoE token-level GRPO is *fragile*, not a guaranteed instant failure (Ch 99's scale-dependence), so monitor or switch. **F14 GSPO clip misconfig**: reusing GRPO's 0.2 clip on GSPO → NaN loss, grad-norm 1e6 on step 1. Fix = use GSPO's sequence-level clip (~3e-4/4e-4). When a crash is instant and total, suspect a setting before a bug.

Next: the failures where the model drifts away from itself, reference drift, and training on the wrong tokens.

Drift and Contamination Failures

Failure	Early signature	Cause	First fixes
F6 — reference drift	<code>kl_to_ref</code> climbs unbounded generation-capped eval falls while task reward rises	KL coefficient too low / reference never refreshed	$\beta \geq 1e-3$; joint reward refresh reference ~100 steps if moving (else keep base fixed, tighten β)
F10 — tool-token contamination	held-out eval drops over training model regurgitates context	loss includes tool / retrieval tokens	preserve the response mask (Ch 21, 106) use the OPSD-style wrapper
F12 — chat-template drift	RL reward $\uparrow$, eval $\downarrow$ format-error rate $\uparrow$ at eval but 0 at train	train and eval render the prompt differently	hash the template in the manifest eval through the same render path

Table 118.1. Three ways the model quietly trains toward the wrong target. All three share the [Chapter 88](#) signature — the training number rises while the held-out number falls — which is exactly why the held-out eval has to be independent.

This family is united by a single tell, the one [Part IX](#) spent a whole Part on: the training reward goes up while held-out quality goes down. Each of these three has a different mechanism, but the dashboard symptom is the same divergence, and if you trust the reward you ship a worse model with full confidence.

118.1 F6, reference drift

F6, reference drift is the purest case. Task reward rises steadily, the run looks like a success, and `kl_to_ref` (the Kullback–Leibler divergence from the frozen reference policy) climbs without bound while a generation-capped held-out eval falls, typically over steps 100–300. The policy is wandering ever further from the base model, buying reward by exploiting the reward function rather than by getting better, and the KL leash is too slack to stop it. The fixes are to tighten and refresh that leash: a minimum KL coefficient $\beta \geq 1e-3$, and a joint reward that includes the KL term rather than treating it as an afterthought. One caveat on refreshing: a reference refresh every ~ 100 steps is safe only if the algorithm is designed around a moving reference. Refreshing changes what KL means, if the reference is meant to anchor the policy to the base model’s capability, refreshing weakens that anchor, so prefer to keep the base reference fixed and instead tighten β / stop earlier.

118.2 F10 and F12, training and evaluating on the wrong tokens

The other two are contamination of the token stream. F10, tool-token contamination is the masking failure of [Chapters 21](#) and [106](#) seen from the dashboard: held-out eval drops while the model starts regurgitating context, reciting tool outputs and retrieved passages, because the loss was computed over the tool and retrieval tokens, teaching the model to imitate the environment. The fix is to preserve the response mask so the loss covers only the policy’s own tokens. F12, chat-template drift is subtler and shows up at step 1: RL reward rises but eval falls, and the giveaway is a format-error rate that is high at eval but zero at training. Training and evaluation are rendering the prompt through different chat templates, so the model is being graded on inputs shaped differently from the ones it trained on. The fix is the run-manifest discipline of [Chapter 130](#) (§130.1): hash the rendered chat template into the run manifest, and run eval through the exact same render path as training.

TAKEAWAY

This family all shows [Chapter 88](#)'s divergence, train reward up, held-out eval down. **F6 reference drift:** `kl_to_ref` unbounded $\uparrow$, eval $\downarrow$, reward still rising (100–300) from too-slack KL, fix with $\beta \geq 1e-3$ + reference refresh every ~ 100 steps. **F10 tool-token contamination:** eval drops, model regurgitates context, because the loss covered tool/retrieval tokens, fix the response mask ([Ch 21/106](#)). **F12 chat-template drift:** reward $\uparrow$ /eval $\downarrow$ with format-error high at eval but 0 at train, because train and eval render differently, hash the template in the manifest, eval through the same path.

Next: the failures unique to on-policy distillation.

On-Policy Distillation Failures

Failure	Early signature	Cause	First fixes
F9 — tokenizer mismatch	KL flat despite training student never improves — from step 1	teacher and student vocabularies differ	sha256 the <code>tokenizer.json</code> distill within-family only
F15 — OPD mode collapse	student becomes identical to teacher diversity $\rightarrow 0$ held-out generation drops (200+)	reverse-KL pushed too hard	student sampling temp 0.7 entropy bonus 1e-4 stop at the plateau
F17 — AWQ-teacher logprob drift	OPD reward delta unstable KL term oscillates (first ~50 steps)	quantization noise in the teacher's logits	use <code>GROUP_SIZE=128</code> not 64; sanity-check teacher BF16 vs AWQ on 100 prompts fall back to an FP8 teacher

Table 119.1. On-policy distillation (OPD) adds a second model — the teacher — and every entry here is a way that second model quietly corrupts the signal. Spell-outs: AWQ = Activation-aware Weight Quantization; FP8 = 8-bit floating point.

On-policy distillation (OPD) trains the student on its own rollouts, scored against a teacher's distribution. That second model is the new failure surface, and these three are the ways it goes wrong, two of which announce themselves at step 1 if you are watching, and one of which is the price of quantizing the teacher to fit it in memory.

119.1 F9, tokenizer mismatch

F9, tokenizer mismatch is the step-1 killer that the source's own warning (a since-validated blog post) flagged and people skipped anyway. (F9 here is the run-level failure. P26 in [Chapter 124](#) is the framework implementation bug that causes it.) If the teacher and student have different vocabularies, then the teacher is scoring token IDs that mean something different in its own space, and the reverse-KL signal is noise from the start: the KL stays flat, training does nothing, and the student never improves. The fix is preventive and absolute, sha256 the `tokenizer.json` of both models and confirm they match before you commit the run, and as a rule distill only within a model family where the tokenizer is shared. There is no in-run recovery. You catch this with a hash check or you waste the run.

119.2 F15 and F17, over-distillation and a quantized teacher

The other two are about how hard and how cheaply you distill. F15, OPD mode collapse is over-distillation: push the reverse-KL too aggressively and the student converges to be identical to the teacher, diversity collapses, and held-out generation quality drops after 200+ steps, the model has memorized the teacher's surface rather than learned its capability. The fixes preserve some exploration: a student sampling temperature of 0.7, a small entropy bonus (1e-4), and, most importantly, stopping at the plateau (the reverse-KL-to-teacher flattening for ~100 steps is the signal to stop. Going further is what causes the collapse). F17, AWQ-teacher logprob drift is the cost of quantizing the teacher to fit it alongside the student: 4-bit Activation-aware Weight Quantization (AWQ) adds noise to the teacher's logits, so the OPD reward delta becomes unstable and the KL term oscillates in the first ~50 steps. The fixes: use a larger quantization group size (`GROUP_SIZE=128`, not 64), sanity-check the AWQ teacher against the BF16 teacher on ~100 prompts before trusting it, and if it still drifts, fall back to the least-lossy teacher mode that passes the BF16 sanity check, often FP8 before AWQ-INT4 (FP8 is less lossy than 4-bit in our setting. Verify rather than assume).

TAKEAWAY

OPD's second model is the failure surface. **F9 tokenizer mismatch:** KL flat, student never improves from step 1, because teacher/student vocabularies differ, sha256 both `tokenizer.json` files. Distill within-family only (no in-run fix). **F15 OPD mode collapse:** student becomes identical to teacher, diversity $\rightarrow$ 0, eval drops (200+) from over-aggressive reverse-KL, student temp 0.7, entropy bonus $1e-4$, stop at the ~100-step plateau. **F17 AWQ-teacher drift:** unstable reward, oscillating KL (first ~50) from quantization noise in teacher logits, `GROUP_SIZE=128`, verify AWQ vs BF16 on 100 prompts, or use an FP8 teacher.

Next: the failures that have nothing to do with the algorithm, infrastructure and throughput.

Infrastructure and Throughput Failures

Failure	Early signature	Cause	First fixes
F7 — throughput collapse	rollout time sawtooths with deep valleys <code>train_gpu_util</code> ~5%	vLLM re-syncs per micro-batch	vLLM sleep mode; <code>gpu_memory_utilization=0.6</code> sequence packing
F16 — QLoRA rollout slowdown	rollouts 5–10× slower than BF16 LoRA	BitsAndBytes 4-bit forward kernel is slow	Unsloth kernels (~2×); a BF16/LoRA rollout path; AWQ/FP8 only if it clears the Ch 121 preflight
F8 — VRAM leak (Torch 2.6 FSDP2)	peak VRAM climbs monotonically OOM at step 200–500	Torch 2.6 Fully-Sharded Data Parallel v2 bug	upgrade to Torch 2.7 <code>veRL</code> ≥ 0.5
F11 — disk-full checkpoint crash	crash around step 80, no warning disk full	<code>save_total_limit</code> unset, so checkpoints pile up	<code>save_total_limit=2</code> ; a dedicated volume mirror to object storage

Table 120.1. None of these is an algorithm problem, and all of them cost more than one — they kill a run that was working, hours in. Spell-outs: QLoRA = Quantized Low-Rank Adaptation; NF4 = 4-bit NormalFloat.

The most expensive failures in this book are not subtle algorithmic ones. They are the dumb infrastructure ones that kill a healthy run two hundred steps in, after you have already paid for those two hundred steps. This chapter is the catalog of those, split into the ones that bleed throughput and the ones that crash the run outright.

120.1 F7 and F16, bleeding throughput

F7, throughput collapse is the silent tax: rollout time sawtooths with deep valleys and `train_gpu_util` sits around 5%, because the inference engine (vLLM) is re-synchronizing weights per micro-batch, leaving the GPUs idle most of the time. The fixes are to stop the constant resync: vLLM sleep mode (free the rollout engine's memory during the training step), a lower `gpu_memory_utilization=0.6` to leave headroom for the co-located trainer, and sequence packing. (Sleep mode is a fix only with a pinned version and a wake-up weight-sync smoke test, prefer sleep level 1 unless your exact vLLM/veRL pinset has a verified level-2 resync. See [Chapter 122](#)'s P1/P12.) F16, QLoRA rollout slowdown is throughput collapse's quantized cousin: Quantized Low-Rank Adaptation (QLoRA) rollouts run 5–10× slower than BF16 Low-Rank Adaptation (LoRA) because the BitsAndBytes 4-bit forward kernel is slow. The fixes are kernel-level first: Unsloth's kernels (~2× faster) or a BF16/LoRA rollout path. The asymmetric trick of an AWQ (or FP8) rollout is tempting, but, and this is the reconciliation with [Chapter 121](#), a quantized rollout is lossy, so use it only if it passes the rollout-correctness preflight: low KL versus BF16, low token disagreement, and no held-out degradation in a short RL run. After [Chapter 121](#), AWQ rollout is not a generic fix. It is a fix you earn by measuring.

120.2 F8 and F11, the run-killers

The other two simply end the run. F8, VRAM leak is a real bug in a specific stack: under Torch 2.6's Fully-Sharded Data Parallel v2 (FSDP2), peak VRAM climbs monotonically across steps until an out-of-memory crash somewhere around step 200–500. There is no clever workaround, the fix is a version bump: Torch 2.7 and veRL ≥ 0.5. F11, disk-full checkpoint crash is the most embarrassing one because it is purely operational: with `save_total_limit` unset, checkpoints accumulate until the disk fills and the run crashes around step 80 with no warning. Set `save_total_limit=2`, put checkpoints on a dedicated volume, and mirror to object storage. The meta-lesson of this chapter: before you launch a long run, the boring questions,

does the disk have room for every checkpoint, is the stack version known-good, will the GPUs actually be busy, protect more compute than any reward-shaping insight.

TAKEAWAY

The dumb failures cost the most because they kill a working run hours in. **F7 throughput collapse:** rollout sawtooth, `train_gpu_util` ~5% from per-micro-batch vLLM resync, fix with vLLM sleep mode, `gpu_mem_util=0.6`, sequence packing. **F16 QLoRA slowdown:** 5–10× slower rollouts from the bnb 4-bit kernel, Unsloth kernels or a BF16/LoRA rollout. AWQ/FP8 rollout *only* if it clears the [Ch 121](#) rollout-correctness preflight (it's lossy). **F8 VRAM leak:** monotone VRAM climb → OOM at step 200–500 (Torch 2.6 FSDP2 bug), upgrade to Torch 2.7 / `veRL` ≥ 0.5. **F11 disk-full crash:** dies ~step 80, no warning, `save_total_limit=2`, dedicated volume, object-storage mirror.

Next: a failure class the original catalog could not have, rollout accelerations that silently corrupt the policy.

CHAPTER 121

Rollout-Side Failures: The Lossy Family




Figure 121.1. A quantized rollout moves the importance ratio off 1.0: with 4-bit AWQ, 18.6% of token ratios fall outside $[0.8, 1.2]$ and get silently clipped.

This is a failure class the original failure catalog could not have contained, because it only appears when you accelerate rollouts to save money: an acceleration that is perfectly safe for inference can silently corrupt the policy in reinforcement learning. The narrative home for these is [Part XIV](#). This chapter is the catalog view. The numbers below are measured in the inference program’s forensics (full detail in [Part XIV](#)).

121.1 Why “fine for inference” can be poison for RL

The reinforcement-learning update relies on an importance-sampling ratio between the policy being trained and the policy that generated the rollouts. The unstated assumption is that those are the same policy (or close). A lossy acceleration, a quantized key-value (KV) cache, quantized weights for the rollout engine, breaks that

assumption by shifting the token distribution of the rollouts. The output may look fine. The distribution is different, and that difference is exactly what the importance ratio measures. This is why [Chapter 136](#)’s rollout-correctness contract (§136.1) demands more of a lossy acceleration than a speed number: report the per-rollout KL against the unquantized policy and the held-out eval after a short RL run, because those are the only signals that catch a corruption the rollout text hides.

121.2 The catalog

Before the failures, the reassurance: most rollout accelerations are safe, the danger is specific to the lossy class, not to acceleration in general.

Acceleration	Distribution effect	KL gate?	Verdict
Speculative decoding / MTP	exact (if implemented correctly)	smoke test only	safe
Prefix caching (APC)	exact	no	safe
Rollout/train disaggregation	scheduling only	no	safe
Attention-backend swap	exact (if numerically equivalent)	smoke test	usually safe
FP8 KV cache	lossy	yes	failed in our run
AWQ rollout weights	lossy	yes	failed in our run

Table 121.1. Rollout accelerations by distribution effect: which need a KL gate, and the verdict.

The lossy failures, worst first, note the evidence column (these are measured in our forensics, not inferred), and that AWQ was not even fast:

Failure	Signature	Verdict
FP8-KV distribution corruption	per-rollout KL 7.61 ~40% token disagreement at matched seed/T	reject for RL rollout (no speed win either: 0.97×)
AWQ-rollout IS-ratio break	KL 6.86; 42.8% disagreement 18.6% of token ratios outside [0.8, 1.2] → trainer silently clips ~1 in 5 ~5% slower at this batch size	reject unless it clears the preflight
dense-AWQ NaN-at-L2 cascade	NaN originating at an L2-norm layer	no clean H100 path for the official dense-32B checkpoint
MoE router-flip (quantized rollout)	31% of top-1 routing decisions flip (gate itself unquantized)	MoE rollout-quant fails discretely and everywhere (F5 by another door)

Table 121.2. The rollout-quantization failures, all measured directly: each path, the signature it produced, and the verdict. The throughline is that lossy rollout precision corrupts the on-policy distribution you train on, so none of these is safe for RL rollout.

The headline lesson is that AWQ rollout failed correctness and delivered no speed win in the measured workload, so it is not a trade-off, it is simply a loss. Separately, a set of operational traps (not distribution corruptions. They crash or mislead rather than silently bias the gradient):

Gotcha	Signature
vLLM <code>n_samples>1</code> else T=0 override	temperature silently forced to 0 for single-sample → no exploration
SGLang ↔ vLLM CUDA-stack collision	crash when both are imported
zombie VRAM after <code>pkill</code>	memory not freed → next run OOMs on a “free” GPU
container-overlay-vs-network-mount path trap	weights load from a stale overlay copy → you train the wrong checkpoint
EAGLE-3 cross-version acceptance regression	acceptance length drops after a version bump (Ch 94’s caveat; see Ch 135)

Table 121.3. Stack-level gotchas the author hit in practice (“lived,” not just measured): each one and the signature it shows. These are environment and framework traps, not algorithm problems — each silently breaks a run that looks configured correctly.

The rule that covers the whole lossy family: before using any lossy rollout acceleration in RL, measure the per-rollout KL and a short held-out eval. If the distribution moved, the speedup is not free. It is a bug you paid for.

TAKEAWAY

Rollout accelerations safe for inference can silently corrupt RL, because the importance ratio assumes the rollout policy $\approx$ the training policy. The distribution-corrupting offenders: **FP8-KV** (per-rollout KL 7.61), **AWQ rollout** (KL 6.86. 18.6% of token ratios outside [0.8, 1.2] $\rightarrow$ trainer silently clips 1 in 5 tokens), **dense-AWQ NaN-at-L2**, and **MoE router-flip** (31% of top-1 decisions flip, F5 by quantization). AWQ also delivered *no* speed win ($\sim$ 5% slower), not a trade-off, just a loss. But most accelerations are *safe* (spec-decoding, prefix caching, disaggregation are exact). The danger is the lossy class. Gate every lossy rollout accel behind per-rollout KL + a short held-out eval (Ch 136, §136.1).

Next: the framework traps, the bugs that live in `veRL`, `TRL`, and `vLLM`, not in your algorithm.

CHAPTER 122

Framework Traps

Most “my GRPO is broken” reports are a framework seam, not the algorithm. The pitfalls below carry their upstream issue numbers so you can check whether your version is affected before you debug your own code.

A large fraction of reinforcement-learning-from-verifier debugging time is spent on bugs that are not in your code or your math, but in the seam where the trainer, the inference engine, and the parallelism library meet. This chapter catalogs the worst of them by framework, with issue numbers, because the fastest fix is recognizing that the bug is already filed.

122.1 veRL traps (P1–P9)

The veRL traps cluster around the vLLM weight-resync seam and host-memory leaks. The headliners (issue IDs as reported in the pitfall table, upstream numbers may have since moved):

ID	Symptom	Root cause	Fix
P1	vLLM produces gibberish after the first weight resync	sleep-level-2 frees weights/KV, but the sharding manager copies new weights while old ones are still mapped, blowing past <code>gpu_memory_utilization</code> (vLLM #15254/#16564)	pin a vLLM with the in-place weight-loader, or use sleep level 1
P3	Ray kills workers after ~100 steps (“node low on memory”)	host RAM leak — reward manager and trajectory metadata accumulate on the head node, not GPU OOM (verl #4995)	shrink <code>train_batch_size</code> ; free non-tensor data each step

ID	Symptom	Root cause	Fix
P5	padding tokens contribute to loss when the episode never emits EOS	<code>get_response_mask</code> marks padding up to <code>max_response_length</code> as valid → pollutes KL, GAE, advantage norm, policy loss (verl #3805)	intersect mask with <code>(tokens != pad_id)</code> ; recompute group stats
P8	sleep-mode works for only one vLLM instance per process	<code>CuMemAllocator</code> is process-global (verl #1193)	launch a second rollout model in its own Ray process

Table 122.1. Framework-seam bugs (the “P” series), each with a tracked upstream issue: the symptom you see, the root cause underneath, and the fix. None is an algorithm problem — each is a seam between the trainer, the rollout engine, and the sharding layer.

P5 is the most dangerous because it is silent. It can ride an entire run, quietly biasing every statistic, on tasks (like SWE-bench) with a high rate of episodes that never emit an end-of-sequence token.

122.2 TRL and vLLM traps (P10–P18)

The TRL and vLLM traps cluster around loss aggregation and co-located memory (issue IDs as reported in the pitfall table):

ID	Symptom	Root cause	Fix
P10	loss nonzero even at $\beta=0$, <code>num_iterations=1</code>	TRL aggregates <code>(loss * mask).sum / mask.sum</code> over the whole micro-batch — a length-biased estimator (TRL #2995)	set <code>loss_type</code> explicitly (<code>grpco/dr_grpo/dapo</code>); this is the F3 length bias at the framework level
P12	GRPO with colocate + sleep-mode → no learning, reward never moves	weights not synced into vLLM after <code>wake_up</code> (TRL #5312)	pin before the regression or apply the fix — silent, no warning

ID	Symptom	Root cause	Fix
P16	OOM at <code>gpu_memory_utilization=0.8</code> that vanishes at 0.6	vLLM's KV-cache profiling assumes static residence, but co-located training allocates/frees in the same step	Unsloth Standby (<code>UNSLOTH_VLLM_STANDBY=1</code>) frees vLLM weights during training → push utilization to 0.95
P17	QLoRA + vLLM <code>fast_inference=True</code> loads in 16-bit anyway	vLLM doesn't consume some bnb-4bit blobs (missing <code>absmax</code> keys)	use Unsloth dynamic-4bit quants, or a version with bnb-4bit support

Table 122.2. More framework-seam bugs, this time mostly in TRL and vLLM co-location: symptom, root cause, and fix. Several have a silent failure mode — the run keeps going while quietly training on the wrong signal.

The pattern across both subtables: the failures that waste the most time are the silent ones (P5, P12) where the run proceeds and simply doesn't learn, because there is no error to grep for, only a reward curve that never moves.

TAKEAWAY

Most “my GRPO is broken” is a framework seam, not the algorithm, check the issue tracker before your own code. **veRL:** P1 (sleep-2 corrupts weights after resync), P3 (host RAM leak kills Ray ~100 steps), **P5 (padding in loss when no EOS, silent, pollutes every statistic)**, P8 (one vLLM per process). **TRL/vLLM:** P10 (length-biased default aggregation, set `loss_type`), **P12 (colocate+sleep → silent no-learning)**, P16 (OOM at 0.8 not 0.6 → Unsloth Standby), P17 (QLoRA loads 16-bit). The silent ones (P5, P12) cost the most.

Next: the reward and multi-turn traps.

CHAPTER 123

Reward and Multi-Turn Traps

ID	Symptom	Fix
P19	model collapses to 256 tokens of garbage; reward and reward-std flatten	<i>Why:</i> with group advantage constantly 0, the only gradient left is the KL term → it collapses to one string. <i>Fix:</i> fit <code>max_completion_length</code> to the model; strengthen the prompt; or use a larger base
P20	a near-zero policy loss read as “not learning”	centered advantages can make the scalar loss uninformative, and unanimous (zero-variance) groups give zero advantages — watch reward-std and the zero-variance fraction, not the loss value
P21	reward function returns <code>None</code> → trainer silently scores 0	assert a numeric return; treat any <code>None</code> as run-fatal in dev
P22	composite reward collapses onto the highest-variance component	correctness-gate auxiliaries; audit component variance; normalize aux components only after preserving the correctness hierarchy — do not blindly inverse-variance-normalize all channels
P23	policy trained on tool-output tokens (multi-turn)	use veRL’s incremental <code>apply_chat_template</code> diff for turn-level masking
P24	naïve dense per-turn rewards make the agent worse (up to 14pp)	calibrate rewards for discriminativeness (Iterative Reward Calibration)
P25	QwQ-32B multi-turn rollout-vs-training mismatch	use the patched template revision (refs/pr/81)

Table 123.1. Reward-shaping and multi-turn traps (the “P” series continued): the symptom and the fix for each. Several restate, as concrete pitfalls, the reward and masking doctrines of Parts III and VIII.

This chapter is the pitfall-level companion to the reward-design Part (III) and the reward-shaping chapter (84): the specific, numbered ways reward functions and multi-turn rollouts go wrong in practice.

123.1 The reward-collapse traps (P19–P22)

The first cluster is about what happens when the reward signal degenerates. P19 is the 256-token-garbage collapse: when a group’s advantage is constantly zero, the only gradient term left is the KL penalty, and the optimizer minimizes it by collapsing to a single repeated string, fit the completion length to a small model’s real budget, strengthen the system prompt, or start from a larger model. P20 is the most-stared-at non-bug in this book, and it is worth stating precisely. A near-zero policy loss does not mean convergence: in GRPO the advantages are centered, so the scalar loss is often uninformative, and a unanimous group (every rollout scored the same) produces zero advantages regardless of what that score was. The trap is reading a flat loss as “nothing is being learned.” Watch reward-std and the zero-variance fraction, not the loss. A tiny worked example makes the distinction, note that the mean is irrelevant. Only the spread carries a learning signal:

group rewards	mean	std	learning signal?
[0, 0, 0, 0]	0	0	no (unanimous)
[1, 1, 1, 1]	1	0	no (unanimous)
[0, 1, 0, 1]	0.5	>0	yes
[-1, +1, -1, +1]	0	>0	yes (mean 0, but advantages nonzero)

Table 123.2. Worked example: only reward spread, not the mean, carries a learning signal.

P21 is a silent scorer: a reward function that returns `None` is silently treated as 0, so assert a numeric return type and treat `None` as fatal during development. P22 is multi-objective collapse: with

several rewards summed, the optimizer collapses onto whichever component has the highest variance. The tempting fix, reweight each reward inversely to its per-group variance, is dangerous in this book, because our own GDPO / per-channel-normalization experiments showed that blindly normalizing every channel destroys the correctness-gradient hierarchy. The safer doctrine: gate auxiliary rewards behind correctness (correctness stays the dominant signal), log component variance, and normalize an auxiliary component only if its variance overwhelms the main reward and only after the correctness gate, never at the expense of the hierarchy.

123.2 The multi-turn traps (P23–P25)

The second cluster is specific to agents. P23 is the masking failure again, here at first contact with multi-turn training: tool-response tokens land in `input_ids`, and without a turn-level loss mask the policy is trained to mimic tool outputs (the [Chapter 106](#) failure). P24 is a counterintuitive and expensive one: naïve dense per-turn rewards can make a tool-calling agent worse by up to 14 percentage points, because a per-turn signal poorly aligned with action quality teaches the wrong thing, calibrated rewards (analyzing discriminativeness versus advantage direction) are needed, and the source reports a base going from 56 verbose turns at 0% action accuracy to 28 turns at 100% after calibration. P25 is a plain template bug: QwQ-32B’s stock chat template causes a rollout-versus-training mismatch. Use the patched revision.

TAKEAWAY

Reward and multi-turn pitfalls, mostly restating Parts III/VIII concretely. **Reward collapse:** P19 (advantage=0 → only-KL-gradient → 256-token collapse), **P20 (a near-zero loss ≠ convergence, centered advantages and unanimous zero-variance groups make the loss uninformative. Watch reward-std and the zero-variance fraction)**, P21 (`None` reward silently = 0 → assert numeric), P22 (composite reward collapses to highest-variance component → *correctness-gate auxiliaries*). Don’t blindly

inverse-variance-normalize. It breaks the correctness hierarchy, per our GDPO finding). **Multi-turn:** P23 (tool-output tokens in loss → turn-level mask), **P24 (naïve dense per-turn rewards: -14pp → calibrate)**, P25 (QwQ-32B template → patched revision).

Next: the OPD, KL-estimator, and eval-drift traps, including the biased-gradient trap nearly every GRPO implementation gets wrong.

CHAPTER 124

OPD, KL-Estimator, and Eval-Drift Traps

The one nearly everyone gets wrong (P29). The “K3-as-loss” Kullback-Leibler (KL) estimator that GRPO uses is a biased gradient estimator of reverse KL, even though it is an unbiased value estimator. Placing K1 in the reward is gradient-equivalent to the principled choice and, in the cited studies, outperformed K3-in-loss by a large relative margin on out-of-domain evals. (The biased-gradient result is from Tang & Munos 2025 and follow-ons.) Until libraries ship the importance-sampled fix: put the KL in the reward, not the loss.

This chapter collects the subtlest traps in the book, the ones that require reading a paper (or three) to even see, because the run looks fine and the number is just quietly wrong. They split into on-policy-distillation and KL-estimator math (P26–P30) and evaluation drift (P31–P33), and the latter ties straight back to [Part IX](#).

124.1 OPD and the KL-estimator math (P26–P30)

The on-policy-distillation (OPD) and KL traps are where “looks fine, is wrong” lives. P26: cross-tokenizer student-teacher distillation silently trains on the wrong log-probabilities, feeding student-tokenized IDs to a teacher with a different vocabulary scores them in the wrong probability region (the F9 failure, at the framework level). Retokenize on the boundary. P27: sampled-token OPD is biased relative to sequence-level reverse-KL, so it yields a disappointing +0–2%. A truncated reverse-KL over the teacher’s top-K (with special-token masking) recovers a reported large gain in the cited study. P29 is the headline, and it is worth internalizing because nearly every GRPO implementation gets it wrong: the K3-in-loss KL estimator is a biased gradient estimator even though it is an unbiased value estimator,

so the KL barely regularizes despite a reasonable β . K1-in-reward is gradient-equivalent to the principled K2-as-loss and beats K3-in-loss by a large relative margin out-of-domain in the cited studies, so place the KL term in the reward, not the loss, until libraries adopt the importance-sampled K3 fix. (Historical runs in this book used the library KL path available at the time. The lesson is forward-looking, report the KL estimator and placement explicitly in the manifest, which is exactly why estimator placement belongs there.) P30 closes the loop with [Chapter 92](#): the same recipe gives different numbers across TRL, veRL, and Open-R1 because their loss-aggregation defaults differ, so hash the `loss_type` into the manifest and never report “GRPO results” without naming the aggregation mode.

124.2 Eval-drift traps (P31–P33)

The eval-drift traps are [Part IX](#)’s discipline as concrete benchmark hazards. P31: a Berkeley Function-Calling Leaderboard (BFCL) number that beats the leaderboard usually means a version mismatch, v1 through v4 add categories cumulatively, and a repro that runs only v3 single-turn is not comparable to a v4 paper table. Record the version and category subset. P32: τ -bench scores are meaningless without the domain, because τ -bench, τ^2 -bench, and τ^3 -bench are different benchmarks with different domains (airline, retail, telecom, long-horizon) and noisy rankings, name the bench and the domain. P33: a math reward that reads zero on correct answers is almost always answer-extraction failure (LaTeX $\frac{1}{2}$ versus 0.5, equivalent forms, `\boxed{}` placement), use `math-verify` with a string-equivalence fallback and log a sample of rejected-but-likely-correct answers, exactly the methodology-is-the-result point of [Chapter 90](#).

TAKEAWAY

The subtlest traps, “looks fine, quietly wrong.” **OPD/KL**: P26 (cross-tokenizer distillation trains on wrong logprobs → retokenize), P27 (sampled-token OPD biased, +0–2% → truncated top-K reverse-KL recovers a reported large gain), **P29 (K3-in-loss is**

a biased gradient estimator → put K1 in the reward, not the loss. Large reported OOD gain), P30 (cross-framework number mismatch from aggregation defaults → hash `loss_type`). **Eval-drift:** P31 (BFCL version mismatch), P32 (τ -bench domain ambiguity), P33 (math extraction fails on LaTeX → `math-verify` + fallback). The eval traps are [Part IX](#) as pitfalls.

Next: the silent killers and the reporting bias that closes the catalog.

Silent Killers and Reporting Bias

ID	Symptom	Fix
P34	TOKENIZERS_PARALLELISM deadlocks after accelerate launch	export TOKENIZERS_PARALLELISM=false — in the launch script, not your shell rc
P35	flash-attn import error mid-run after a worker restart	pin flash-attn, torch, cuda together (the veRL base image names all four for a reason)
P36	sporadic NCCL hangs that recover after a timeout	set NCCL_DEBUG, TORCH_NCCL_AVOID_RECORD_STREAMS=1, NCCL_ASYNC_ERROR_HANDLING=1 in the launch script
P37	single-seed RL numbers that don't reproduce	demand mean $\pm$ std over ≥ 3 seeds; treat single-seed results with ≥ 2 pp skepticism (attach the Ch 87 budget for 2–3pp claims)

Table 125.1. The environment deadlocks that share one fix, put it in the launch script, not your shell, and the reporting bias the whole book is a corrective to.

The catalog closes with two kinds of failure that are not about the model at all: environment-level deadlocks that have nothing to do with your code, and a community-level reporting bias that this entire book is, in part, a response to.

125.1 The environment silent killers (P34–P36)

These three share a root cause and a fix. P34 is a TOKENIZERS_PARALLELISM deadlock after accelerate launch. P35 is a flash-attention import error that appears mid-run when a worker restarts against a different Torch ABI. P36 is sporadic NCCL (NVIDIA Collective Communications Library) hangs that recover only after a timeout. The

unifying fix is operational and easy to get wrong: set the environment variable in the launch script, not your shell’s startup file, because the workers Ray and `accelerate` spawn do not inherit your interactive shell’s environment, so a fix that “works” when you test it by hand silently fails in the actual distributed run. Pin `flash-attn`, `torch`, and `cuda` as a set (this is why production base images name all four together), and write the NCCL variables into the launch script.

125.2 Reporting bias (P37)

P37 is the one this book exists to correct. Late-2025 papers report single-seed GRPO numbers, and reproduction attempts repeatedly fail, many single-seed RL results are hard to reproduce because the data filtering, seeds, and eval scripts are underreported. The cause is selection bias and cherry-picked seeds (Chapter 87’s variance budget makes the mechanism precise: at typical seed-to-seed variance, a single seed can be 2+ points off the mean by luck alone). The fix is the standard the book has argued for throughout: treat single-seed RL numbers with at least 2pp of skepticism on MATH-class benchmarks, and demand mean $\pm$ standard deviation over ≥ 3 seeds in your own work (Chapters 87 and 94). Be precise about that last clause, though: three seeds is a minimum sanity check, not proof of a small delta. For a 2–3pp claim, attach the Chapter 87 variance budget, three seeds do not, by themselves, settle whether a few-point gain is real. This is not a framework bug you can patch. It is a discipline you adopt, and it is the difference between a result and a press release.

TAKEAWAY

The catalog closes with non-model failures. **Environment killers (P34–P36):** `TOKENIZERS_PARALLELISM` deadlock, flash-attn ABI mismatch after restart, sporadic NCCL hangs, all share one fix: set the env var **in the launch script, not your shell rc** (spawned workers don’t inherit it), and pin flash-attn/torch/cuda together. **Reporting bias (P37):** single-seed RL numbers don’t reproduce (underreported filtering/seeds/eval), treat them with ≥ 2 pp skepticism and demand mean $\pm$ std over ≥ 3

seeds *as a minimum* (three seeds aren't proof of a 2–3pp delta, attach the [Ch 87](#) budget). The last one isn't a patch. It's a discipline.

Next: turning the catalog into a habit, the postmortem template that captures a dead run before you forget what killed it.

CHAPTER 126

The Postmortem Template

A run died. Before you relaunch, write this down, or your monitor writes it for you.

```
postmortems/{run_name}/postmortem.md
# ---
Run: <name> What changed: <the ONE variable vs the last run>
Symptom: <which dashboard panel moved> → Matched failure: <F# /
  ↳ P# from this catalog>
Secondary: <the confirming signal> Time to symptom: step <N>
Hypothesis: <what you think happened> Root cause: <what actually
  ↳ happened>
Fix applied: <the change> Outcome: <did it work>
Manifest: model_rev · trl/verl/vllm ver · cuda · tokenizer.json
  ↳ sha256 · loss_type · seeds · eval_script sha
```

A failure catalog is only useful if you actually consult it when a run dies, and the moment a run dies is the moment you are most tempted to just change something and relaunch. This chapter is the discipline that closes the loop: a postmortem template that captures what killed the run while you still remember, and a monitor that writes it for you so the discipline does not depend on your willpower at 2 a.m.

126.1 The template

The template is deliberately the same shape as this catalog's rows (symptom → secondary signal → time-to-symptom → root cause → fix), plus two things a catalog row doesn't have: what changed and the manifest snapshot. What changed matters because the cheapest debugging rule in this book is to change one variable per run, so the postmortem's first job is to name that one variable, which is usually the cause. The manifest snapshot is the run-manifest contract of [Chapter 130](#) (§130.1): the model revision, the TRL/veRL/vLLM versions, the

CUDA version, the sha256 of the rendered chat template and `tokenizer.json`, the `loss_type` (Chapter 124’s P30), the seeds (Chapter 87), and the eval-script hash. Most of the failures in this Part are traceable to one of those manifest fields, a version, a template hash, an aggregation mode, so recording them turns “it won’t reproduce” into “here is the line that changed.”

126.2 Automate it: the health monitor

A postmortem you have to remember to write is a postmortem you won’t write. So the better pattern is a health-monitor callback that watches the dashboard panels of Chapter 114 and writes the postmortem automatically when it triggers an early stop. But the monitor must be careful about what it kills on, because of a lesson from elsewhere in this book: low entropy alone is not a death sentence. A sharply low entropy can be productive sharpening when held-out eval and score variance are still improving, so a monitor that hard-kills on an entropy floor will false-positive on exactly the runs that are working. The fix is joint conditions: entropy alone writes a warning. Entropy low plus reward-std collapse plus no held-out improvement writes a kill.

Condition	Action
	warn + checkpoint
entropy below floor, score still varied, held-out rising	
	stop
entropy below floor sustained + reward-std low + held-out flat/down	
	stop (hard failure)
NaN/Inf, or <code>ratio_max</code> explosion (F5), or format-error high at eval-only	
	stop or roll back
KL spike + held-out eval down	

Table 126.1. The health-monitor joint conditions: what the dashboard shows, and the action it should trigger. A single low signal warns; a combination of low signals, or any hard failure, stops the run.

In pseudocode, the joint logic that avoids the false-positive kill:

```
def should_stop(m):
    entropy_bad = m.entropy_token.mean_last(3) < m.entropy_floor
    ↪ # warn-level
    reward_dead = m.reward_std.mean_last(10) < m.reward_std_floor
    eval_bad = m.heldout_delta_last <= 0
    ratio_bad = m.ratio_max.p95_last(3) > m.ratio_ceiling
    nan_bad = m.has_nan_or_inf
    format_bad = m.format_error_rate > m.format_error_ceiling

    if nan_bad or ratio_bad or format_bad:
        return True, "hard failure"
    if entropy_bad and reward_dead and eval_bad: # the JOINT
    ↪ kill
        return True, "collapse: entropy low + no variance + no
        ↪ eval motion"
    if entropy_bad:
        return False, "warning only: checkpoint and inspect"
    return False, "continue"
```

(Contrast [Chapter 114](#)’s warning line, an entropy floor of 0.10 to look closer. The stop requires the conjunction above, not a lone entropy reading.) On a kill, the monitor auto-populates `post-mortems/{run_name}/postmortem.md` from the metrics it was already logging. The result is that every dead run leaves a record, keyed to a failure ID in this catalog, with the manifest attached, which is exactly the corpus that made this Part possible.

TAKEAWAY

Close the loop: when a run dies, write a postmortem in the same shape as this catalog’s rows (symptom → secondary → time-to-symptom → root cause → fix) **plus what changed** (the one variable, usually the cause) and the **manifest snapshot** (model rev, framework versions, CUDA, template + tokenizer sha256, loss_type, seeds, eval-script hash, most failures here trace to one of those fields). Better, automate it: a health-monitor callback watching the [Chapter 114](#) panels auto-writes

the postmortem on early-stop. Crucially, kill on **joint** conditions, not entropy alone, low entropy can be productive sharpening, so entropy-low writes a *warning*. Entropy-low + reward-std collapse + no held-out improvement writes a *kill* (hard failures, NaN, `ratio_max` explosion, eval-only format errors, stop immediately).

Next: a worked postmortem, the SW1 → SW2 transition, this template applied to a real failure end to end.

CHAPTER 127

Worked Postmortem: SW1 → SW2

The run postmortem

One per failed or surprising run — the fields that turn a dead run into a lesson.

What you saw	the symptom on the dashboard — e.g. reward went flat, outputs got repetitive
First metric to move	which signal turned first, and when — usually a leading one, before reward
Hypotheses, ranked	the candidate causes, most likely first
Root cause	what actually went wrong, once confirmed
Fix applied	the change that addressed it
Guard added	the automatic check that will catch this class of failure next time
Status	☐ open ☒ resolved ☐ won't fix

A run you can't fill these in for is a run you didn't instrument well enough to learn from.

Figure 127.1. The run postmortem: the generic fields that turn a dead run into a lesson (what you saw, first metric to move, ranked hypotheses, root cause, fix, guard added, status).

The postmortem template is abstract until you watch it catch a real failure. This chapter applies it to the Qwen3-14B SWE-bench run that collapsed as SW1 and was rescued as SW2. Group Relative Policy Optimization (GRPO) is the baseline algorithm. The fix SW2 introduced is the Clip-Cov mechanism of [Part IV](#). One framing correction up front: Clip-Cov was not invented here. It was already earned in the search-agent entropy saga (the 4B v1–v5 runs of [Part IV](#)). This SWE-bench postmortem is its cross-domain validation: proof that the stability fix was not a one-off, transferring from 4B search RL to 14B code-agent RL.




Figure 127.2. The entropy-control fix transfers to code: plain GRPO crashes entropy 0.42 → 0.03 in 30 steps and freezes, while GRPO + Clip-Cov holds at 0.28 and keeps learning.

127.1 SW1: the collapse

Filled into the [Chapter 126](#) template, SW1 reads:

Field	SW1
What changed	first GRPO run on Qwen3-14B against SWE-bench (a hard, sparse-reward coding task).
Symptom	per-token entropy 0.42 → 0.03 in 30 steps (failure F1).
Secondary	reward-standard-deviation → 0; generations became identical; KL spiked.
Time to symptom	step ~30.
Root cause	symmetric clipping caps how much probability a good token can gain, so on a hard task the distribution sharpened into a single mode and froze.
Manifest	Qwen3-14B; GRPO; LR 1e-5; symmetric clip 0.2; no entropy bonus.

Table 127.1. SW1 filled into the Chapter 126 postmortem template — the first GRPO run on a hard coding task, and the entropy collapse (F1) it hit.

This is the textbook F1 entropy collapse ([Chapter 115](#)), and the dashboard saw it coming: entropy sliding toward its floor 30 steps

before the reward made the death obvious. SW1 is why the entropy panel is the first one you watch.

127.2 SW2: the fix, and the honest chain after it

SW2 was not a clean one-variable ablation. It was a postmortem fix-bundle, every change targeting the observed F1 collapse. The point is operational attribution, not a controlled experiment: the prescribed bundle made the collapse signature disappear.

Field	SW2
What changed	introduced Clip-Cov + a small KL ($\beta = 0.001$).
Outcome	entropy stopped free-falling; damped oscillation (the healthy signature, not a flat line).
Root-cause fix	the covariance-aware clip lets high-value tokens grow, so the policy keeps exploring.
Significance	Clip-Cov transferred from 4B search RL to 14B code RL — cross-domain validation, not a new invention.

Table 127.2. SW2 filled into the same postmortem template — the Clip-Cov + KL fix that stopped the entropy free-fall, and what it showed.

But the honest postmortem does not end at a triumphant SW2, and the chain is the real lesson, each fix correct, each exposing the next failure. The full sequence, with each version’s own change kept distinct (do not attribute later settings to SW2):

Run	Failure addressed	Key change	Result	New failure exposed
SW1	— (baseline)	vanilla GRPO	entropy collapse 0.42→0.03; ~4-token outputs	F1
SW2	F1	Clip-Cov + KL ($\beta=0.001$)	entropy holds; damped oscillation	still sparse/noisy
SW3	sparse/noisy	stronger Clip-Cov / KL / LR	stable entropy but reward declining	length bias

Run	Failure addressed	Key change	Result	New failure exposed
SW4	F3 (length bias)	<code>seq_mean_token_mean</code> + continuous shaping rewards	first stable reward signal	overfit on small instance set
SW5	scale-up (~280 instances)	constant LR at full peak	late collapse ~steps 18–22	LR schedule (F13)
SW6	F13	cosine LR + stronger Clip-Cov / KL, lower LR	stable to step 146	stability is necessary, not sufficient

Table 127.3. The SW1–SW6 chain: each fix worked and exposed the next failure.

The lesson the chain teaches is stability was fixed (SW2) long before learning was fixed (SW4), and the final “failure” was not a bug at all: SW6 plateaued at about 1.5 billion tokens because SWE-Gym is too hard for a 14B model to climb further, the dataset starves GRPO, reward plateau $\neq$ learning, and this is a capacity ceiling (Chapter 134). Keeping the postmortem chain is what lets you tell “the fix was wrong” from “the fix worked and revealed the next wall.”

TAKEAWAY

The postmortem template, applied to the book’s central cross-domain validation. **SW1:** Qwen3-14B GRPO on SWE-bench collapsed, entropy 0.42 $\rightarrow$ 0.03 in 30 steps (F1), from symmetric clipping capping probability gains. **SW2:** introduced **Clip-Cov + KL ($\beta=0.001$)**, entropy held, damped oscillation (this was the *cross-domain validation* of the search-agent fix, not a new algorithm. And it was a *fix-bundle*, not a clean one-variable ablation). The chain continued, each version its own change: SW3 stronger Clip-Cov $\rightarrow$ length bias. SW4 `seq_mean_token_mean` + shaping $\rightarrow$ first success but overfit. SW5 constant LR $\rightarrow$ late collapse (~step 18–22). **SW6 cosine LR** $\rightarrow$ stable plateau at ~1.5B tokens, a *capacity ceiling*, not a bug. Each fix exposed the next failure, which is why you keep the chain.

Next: stepping back, the six things people blame on the algorithm that almost never are.

CHAPTER 128

Six Non-Algorithm Failures That Look Like Bugs

The six what-failed causes

#	“My RL isn’t working” — but it’s actually...	Where in this book
1	the data ¹	Ch 134
2	the reward ²	Ch 31, 84, 123
3	the eval / reporting ³	Part IX, Ch 87
4	the framework seam ⁴	Ch 122
5	the infrastructure ⁵	Ch 120
6	the config ⁶	Ch 117, 124

Table 128.1. The consolidated what-failed: when a run “doesn’t work,” it is overwhelmingly one of these six, none of which is the optimizer. The most under-suspected is the reward — a whole arc of this book (Parts III, VIII) is reward failures masquerading as RL failures. The “Tell” for each cause is in the notes.

Table notes

1. **The data:** base solves ~0% or ~100%; advantage collapses (F2).
2. **The reward:** train reward ↑ while the task metric ↓; proxy behavior (verbosity, tool-spam) improves.
3. **The eval / reporting:** held-out ↓ as train ↑; methodology flips the result; a 2pp “gain” from one seed.
4. **The framework seam:** silent — reward never moves, no error (P5, P10, P12).
5. **The infrastructure:** a healthy run dies hours in (F7, F8, F11, F16).

6. **The config:** instant/total crash; matches a documented setting (F5, F14, P29).

After a hundred-odd chapters of failures, a pattern is unmistakable: almost nothing people blame on “the RL algorithm” is the algorithm. The optimizer is rarely the problem. Six other things are, and recognizing which one saves you from tuning hyperparameters against a bug that lives somewhere else entirely.

128.1 The six

1. The data. If the base model solves a task ~0% or ~100% of the time, there is no advantage signal and the run collapses (F2), not because GRPO is broken but because RL sharpens what the base can sometimes do (Part VII), and this data offers nothing to sharpen.
2. The reward. This is the book’s most under-suspected cause and one of its central arcs: the optimizer is doing its job perfectly, optimizing a reward that doesn’t mean what you think, a linear multi-objective reward that gets hacked, a correctness-only signal that plateaus, shaping rewards that make the agent look busy, tool-spam with no behavioral penalty, or per-channel normalization that destroys the correctness hierarchy (Parts III and VIII, Chapter 123). The tell is train reward rising while the task metric falls.
3. The eval / reporting. A run that looks like it’s failing (or succeeding) is often just measured wrong, a template mismatch (F12), a benchmark-version confusion (P31–P32), an extraction bug (P33), too small an eval (Chapter 89), or reported wrong: a 2-percentage-point “improvement” from a single seed is, at typical seed-to-seed variance, indistinguishable from noise (Chapter 87, P37).
4. The framework seam. The most maddening category, because it is silent: padding tokens polluting the loss (P5), a length-biased aggregation default (P10), weights not syncing into the rollout engine (P12), the reward curve just sits flat with no error to grep.

5. The infrastructure. A VRAM leak (F8), a full disk (F11), a throughput collapse (F7). These kill a perfectly healthy run, and no amount of reward shaping touches them.
6. The config. An instant, total crash (NaN on step 1) is almost never a bug. It is a setting, the wrong clip on GSPO (F14), token-level ratios on a mixture-of-experts model (F5), the KL in the loss instead of the reward (P29).

128.2 The diagnostic order

The practical value is the order. When a run disappoints, check these six before you touch a hyperparameter, because each has a one-line test: is the base score near 0 or 100 (data)? Does train reward rise while the task metric falls (reward)? Does held-out move with train, and is the gain bigger than the variance budget (eval/reporting)? Is the reward flat with no error (seam)? Did a working run die suddenly (infra)? Was the crash instant (config)? Only after all six are ruled out is it worth asking whether the algorithm itself needs changing, and by then, you usually find it didn't.

TAKEAWAY

Almost nothing blamed on “the RL algorithm” is the algorithm. The six real causes: **the data** (base at 0%/100% → no advantage, F2), **the reward** (train reward ↑ while task ↓, reward hacking, the book's most under-suspected cause. Parts III/VIII), **the eval/reporting** (methodology flips it, or a one-seed 2pp gain is noise, [Part IX, Ch 87](#)), **the framework seam** (silent flat reward, P5/P10/P12), **the infrastructure** (a healthy run dies, F7/F8/F11), and **the config** (instant crash matches a documented setting, F5/F14/P29). Run the six one-line tests *before* tuning a hyperparameter.

Next: the master log, every challenge across every program, with its fix, as a lookup table.

CHAPTER 129

The Challenges-and-Solutions Log

This chapter is a pure lookup. It consolidates every challenge hit across all four programs, the agentic-training journey, the training failure catalog (F1–F17, Chapters 115–125), the inference/rollout program (Part XIV), and the forensics work, so you can find a symptom fast. Most entries point to a chapter. A few forensics-specific ones are recorded here directly.

A failure catalog organized by family (Chapters 115–125) is good for understanding. A flat log organized for lookup is good at 2 a.m. This chapter is the latter, the consolidated index, plus the handful of forensics findings that don't have a home elsewhere.

129.1 The index

The bulk of the log is a pointer table, because the detail already lives in the catalog:

Program	Challenges	Home
Collapse	F1 entropy, F2 advantage, F13 late	Ch 115
Reward-driven	F3 length bias, F4 repetition	Ch 116
MoE / config	F5 ratio explosion, F14 GSPO clip	Ch 117
Drift / contamination	F6, F10, F12	Ch 118
On-policy distillation	F9, F15, F17	Ch 119
Infra / throughput	F7, F8, F11, F16	Ch 120
Rollout-lossy	FP8-KV, AWQ-ratio, router-flip	Ch 121, Part XIV
Framework seams	P1–P18	Ch 122
Reward / multi-turn	P19–P25	Ch 123
OPD / KL / eval-drift	P26–P33	Ch 124
Silent killers / reporting	P34–P37	Ch 125

Table 129.1. The challenge index: each failure family pointed to its home chapter.

129.2 The three forensics notes (recorded here)

Three forensics findings are worth recording directly, because they each cost real time and each masqueraded as something else (the other inference-program findings, FP8-KV, AWQ, router-flip, live in [Part XIV](#) and [Chapter 121](#), so they are not duplicated here), each with its evidence status:

- NaN-in-tensors masquerading as a numerics bug. A NaN that looked like a precision/overflow problem turned out to be a data/masking bug, a reminder that “NaN” names a symptom, not a cause. Trace it to the tensor and the step, don’t reach for a precision flag first.
- The fp64-KL fix. Computing the Kullback–Leibler (KL) divergence in 32-bit versus 64-bit floating point changed the number enough to matter for a gate decision, KL near zero is sensitive to accumulation precision, so the estimator’s dtype is part of the result.
- W8A8’s 6× load time. The safe rollout quantization (8-bit weights and activations, [Chapter 143](#)) is correctness-clean but carries a real operational cost, about 6× the load time, which is the kind of trade-off that belongs in a log even though it is not a “failure.”

Topic	Symptom	Actual cause	Evidence status
NaN-in-tensors	NaN tensors mid-train	data/masking bug (not numerics)	verified (internal log)
fp64-vs-fp32 KL	gate decision flipped	KL accumulation precision	measured (internal KL script)
W8A8 ~6× load	slow to adopt	weight-unpacking load cost	measured (internal inference log)

Table 129.2. Three NaN-or-slowdown symptoms whose real cause was not what it looked like: a masking bug, a precision effect, or a load-cost effect — each with how strongly it was confirmed.

The rest of the inference program’s numbered issues, the FP8-KV and AWQ corruptions, the router-flip, the version-pinning collapses, are catalogued in [Part XIV](#) and [Chapter 121](#), where they have their forensic detail.

TAKEAWAY

This chapter is a lookup index: the failure families (F1–F17, P1–P37) point to Chapters 115–125. The rollout-lossy and forensics findings to [Part XIV](#) / [Chapter 121](#). Three forensics findings are recorded here directly, a NaN that was really a masking bug (the symptom isn’t the cause), the fp64-vs-fp32 KL estimator changing a gate decision (near-zero KL is precision-sensitive), and W8A8’s $\sim 6\times$ load-time cost (a clean quant with an operational price).

Next: [Part XIII](#). We turn from cataloguing failures to the reproducibility and operations discipline that prevents most of them, starting with the run manifest.

Part XIII

Reproducibility and Operations

A result you cannot reproduce is a rumor, and a run you cannot rebuild is a one-time accident. This Part is the operational discipline that turns experiments into evidence: the manifest that makes a number reproducible, the gates that decide what ships, the cost tables that make a plan real, and the version pins that, as two incidents prove, are findings, not hygiene.

CHAPTER 130

The Run Manifest, in Practice

```
# manifest.yaml - attached to every checkpoint and every
  ↳ reported number
model_revision: Qwen/Qwen3-14B@<commit>
frameworks: {trl: 0.20.0, verl: 0.5.x, vllm: <ver>, torch: 2.7,
  ↳ cuda: 12.x}
tokenizer_sha256: <sha256 of tokenizer.json> # guards F9
chat_template_sha256: <sha256 of RENDERED template> # guards F12
algorithm: grpo | dr_grpo | dapop | grpo_pp | gspo | opd # the
  ↳ update rule
loss_agg_mode: token_mean | seq_mean_token_mean | sample_mean #
  ↳ guards F3 length-bias / P30
normalize_adv_by_std: true | false # GRPO vs Dr.GRPO advantage
  ↳ scaling
quant_role: train | rollout | teacher | serving # which copy is
  ↳ quantized
quant_method: bf16 | fp8_weights | fp8_kv | nf4 | awq | w8a8 |
  ↳ nvfp4_external
source_status: canonical | historical | retracted |
  ↳ external_estimate
raw_eval_json: <path or sha256> # every number traces to a file
k1: {placement: reward, estimator: k1, coef: 1e-3} # guards P29
seeds: [0, 1, 2] # guards reporting bias (P37)
dataset: {name: R2E-Gym-Lite, filter_sha: <sha>}
eval_script_sha: <sha> # guards eval drift
```

A manifest is not paperwork, every field guards a specific failure from [Part XII](#).

[Part IX](#) argued that a number without its evaluation conditions is not a measurement. The run manifest is where those conditions live: a small file, attached to every checkpoint and quoted with every result, that records exactly what produced the number. This chapter is the manifest in practice, and the argument that it is a result, not hygiene.

130.1 Every field guards a failure

The manifest's fields are not arbitrary metadata, each one is the antidote to a specific failure from the catalog. The chat-template hash guards against template drift (F12), where train and eval render differently and the eval silently lies. The tokenizer hash guards against the on-policy-distillation tokenizer mismatch (F9). The `loss_agg_mode` (not the algorithm name) is the F3 length-bias lever and the field that guards the cross-framework number mismatch (P30), `token_mean` vs `seq_mean_token_mean` vs `sample_mean` change the number even with the same algorithm. The `quant_role` / `quant_method` / `source_status` / `raw_eval_json` fields are the four the V2 Q-program proved load-bearing: they are exactly what would have prevented the Q1-NVFP4-vs-NF4 mislabel, the base-tier mix-up, and the 0.470→0.173 eval-contamination (every number must trace to a raw eval file). The KL placement and estimator record the P29 choice (K1-in-reward, not K3-in-loss). The framework and CUDA versions guard against the VRAM-leak stack (F8) and the flash-attention ABI break (P35). The seeds make the reporting honest (P37). Read the manifest as a checklist and it is the [Part XII](#) print-blocking checklist, frozen into the run.

130.2 The manifest is a result

The discipline is to treat the manifest as part of the deliverable. Hash the template and tokenizer at launch (a mismatch caught here is free. Caught after a run is a wasted run), store the manifest with the checkpoint so the two never separate, and quote the relevant manifest fields with every reported number, which eval, which seeds, which aggregation. When a number later “won't reproduce,” the manifest turns a multi-day mystery into a one-line diff: the thing that changed is in the file. A result with its manifest is reproducible. A result without one is a story.

TAKEAWAY

The run manifest records exactly what produced a number, and every field guards a [Part XII](#) failure: chat-template sha (F12), tokenizer sha (F9), loss_agg_mode (P30/F3, *not* the algorithm name), quant_role/quant_method/source_status/raw_eval_json (the V2 mislabel + eval-contamination antidotes), KL placement+estimator (P29), framework/CUDA versions (F8/P35), seeds (P37), eval-script sha (eval drift). Hash template and tokenizer at launch, store the manifest *with* the checkpoint, and quote the relevant fields with every reported number. A result with its manifest is reproducible. Without one it's a rumor.

Next: the gates that decide whether a checkpoint is allowed to ship.

CHAPTER 131

Promotion Gates

Gate	Default	Family-specific override	Fail action
Held-out eval	beats base by the variance budget (Ch 87)	mandatory for all	do not promote
Eval size	$n \geq 500$	larger for small deltas	re-evaluate
Same-script base	required (no leaderboard substitute)	—	block
Entropy	above a <i>pre-registered family floor</i>	floor is task/model/temperature-specific	don't auto-fail alone — require an OOD/regression eval if low
Reward variance	nonzero; zero-variance group fraction low	per-task	block if collapsed
KL to reference	below a <i>family ceiling</i> (calibrated from stable runs)	per-task	block if an eval gain coincides with high drift
Format / tool-parse	$\leq \sim 5\%$ error, same render path as train	stricter for tool agents	block

Table 131.1. A checkpoint is promoted only when it passes all gates, but the thresholds are family-specific and pre-registered, not universal hard floors. The one rule to remember: entropy alone is not a promotion gate. Entropy + reward-variance collapse + no held-out motion is.

Part IX gave the doctrine of trustworthy evaluation. This chapter operationalizes it as a gate, a fixed, pre-registered set of conditions a checkpoint must pass before it is promoted to “this is our model.” Gates turn judgment into procedure, which is what keeps a tired team from shipping a checkpoint that merely looks good on the panel they happened to glance at.

131.1 The gates

A checkpoint is promoted only when it passes all gates, but with one critical refinement the book earned the hard way. The eval-delta gate is the [Chapter 87](#) variance budget as a rule: the held-out improvement over the base must exceed what the seed budget can resolve, on an eval of at least 500 samples ([Chapter 89](#)), measured with the same eval script as the base (no leaderboard substitute), a gain smaller than the budget is not a gain. The entropy gate is where the refinement lives: an entropy below the family floor is not an automatic fail, because, as the C9-COMBO runs showed ([Chapter 126](#)), low entropy can be productive sharpening when held-out eval is still climbing and reward variance is alive. So entropy below floor plus reward-variance collapse plus no held-out motion blocks promotion. Entropy alone only triggers an OOD/regression eval. The KL-ceiling gate refuses a policy that has drifted too far (F6), calibrated from your own stable runs, not a universal 0.5. The format/tool-parse gate confirms well-formed outputs on the same render path as evaluation (F12). The point is the joint conjunction, family-calibrated: one good number is not enough, and one bad number (entropy) is not enough to reject.

131.2 Pre-register and automate

The gates are only honest if they are set before the run ([Chapter 92](#)'s pre-registration, choosing thresholds after seeing results is a reward hack on yourself), and they are only reliable if they are automated. The health monitor of [Chapter 126](#) already watches the floors and ceilings live. The promotion decision reads the same instruments at the end. Write the thresholds into the manifest, evaluate them programmatically, and promote on all-pass, so “is this checkpoint good enough to ship?” has a reproducible answer instead of a vibe.

TAKEAWAY

A checkpoint ships only when it passes **all** promotion gates, *pre-registered and family-specific* ([Ch 92](#)): a held-out eval ($n \geq 500$),

same eval script as base) beating base by **more than the variance budget** (Ch 87/89). Reward variance not collapsed. KL below a *calibrated* family ceiling (F6). A low format/tool-parse error on the *same render path* as eval (F12). Crucially, **entropy alone is not a gate**, low entropy can be productive sharpening (Ch 126's C9 lesson), so entropy-below-floor blocks only *jointly* with reward-variance collapse and no held-out motion (else it just triggers an OOD/regression eval). Automate it (the Ch 126 monitor reads the same instruments).

Next: the training-side quantization record, what BF16, FP8, and NF4 actually did to training.

The Quantization Side-Quest (Training-Time)

Topic	Detail
BF16	the baseline — reference behavior; everything compares to this.
FP8 (weights)	memory-saving training — attention-space equivalent to BF16 in the Q4 14B run (per-head KL 2.19e-4 vs 2.22e-4); downstream quality is task- and decoding-mode-dependent (GSM8K slightly favors BF16; thinking-mode regresses for both). Not a blanket “FP8 is safe”.
NF4 (QLoRA)	fit a big model on small GPUs — statistical parity with BF16 at 32B across all 5 benchmarks (4-seed), upholds the QeRL parity claim; lower seed-variance than BF16 (std 0.029 vs 0.085).
AWQ / GPTQ	4-bit, rollout/serving, not training — training-side 4-bit is not where these belong (see Part XIV for rollout).
NVFP4	Blackwell-only — never executed; vLLM’s kernel is sm_100/Blackwell-only; on H100 (sm_90) it dequantizes to BF16 at load, so the experiment labeled “Q1” actually ran NF4 as the Hopper proxy.

Table 132.1. Training-side quantization, distinct from rollout-side ([Part XIV](#)) and serving-side ([Part XV](#)). The full record is in Appendix R. Spell-outs: BF16 = bfloat16. FP8 = 8-bit float. NF4 = 4-bit NormalFloat. AWQ = Activation-aware Weight Quantization. GPTQ = a 4-bit post-training quantizer. NVFP4 = NVIDIA’s 4-bit float.

Quantization shows up in this book in four roles, for four different jobs, and conflating them is a common and expensive mistake, so before the training-side record, the role map that the rest of the chapter (and [Part XIV](#)) depends on:

Role	What is quantized	Affects gradients?	RL-specific risk	Book home
Actor training	the trainable model / adapter path	yes	training quality / variance / stability	this chapter
Rollout engine	the sampler model / KV cache	no gradient, but it generates the samples	distribution corruption → invalid importance ratios (Gate A/B/C)	Part XIV
Teacher	the teacher's logits / logprobs	no student gradient <i>through</i> the teacher	teacher logits drift → wrong KL target	OPD + F17 (Ch 119)
Serving	the deployed model, after training	no training	user-facing quality / latency	serving appendix

Table 132.2. Quantization’s four roles, what each touches, and its RL-specific risk.

This chapter is the Actor-training row only: quantizing the model you are updating. Keep the four straight, most “is FP8 safe?” confusion is a role confusion.

132.1 The record (Q1–Q4)

The headline findings, now backed by the complete Q-program record (the 55-cell grid, Appendix R). FP8 weights quantize once at load and dequantize for the matmul. In the Q4 14B run, FP8 is attention-space equivalent to BF16 (per-head KL 2.19e-4 vs 2.22e-4, paired), but that is not the same as “no quality regression”: downstream quality is task- and decoding-mode-dependent (GSM8K slightly favors BF16 at +1.13pp. On thinking-mode both full-FT arms regress, FP8 -10.8pp vs BF16 -23.6pp). The safe claim is the narrow one, FP8 training drifts attention like BF16, not a blanket “FP8 is safe.” State the mechanism exactly, though, because “FP8” can mean load-time weight quantization, rollout-only quantization, KV-cache quantization, or kernel-level FP8 matmuls, and they are not interchangeable.

NF4 (QLoRA) is where the record corrects an earlier mistake. At 32B, NF4 reaches statistical parity with BF16 across all five

benchmarks (math_dapo, GSM8K, MATH-500, AIME-24 greedy, and AIME-24 thinking), which upholds and extends the QeRL paper’s parity claim (QeRL never ran thinking-mode evals). The retraction the book owes runs the opposite direction to an earlier draft: a single-seed result had suggested “BF16 wins by +6.7pp on AIME-24 thinking,” and that was the artifact, under 4-seed evaluation, NF4 mean 0.525 (std 0.029) versus BF16 mean 0.500 (std 0.085) gives $\Delta = -2.5\text{pp}$ with NF4 numerically higher, not significant (the 95% Wilson interval at $n=30$ is $\pm 18\text{pp}$). So the surviving, stronger claim is NF4 = BF16 quality parity at 32B, with NF4 showing lower seed-variance, not higher. There is also an FP8 + mixture-of-experts wall, but, per [Chapter 158](#), that is a block-size/memory implementation wall, not a training-numerics result.

132.2 The corrections

Three corrections keep the record honest. First, NVFP4 was never executed, and “Q1” is an NF4 proxy. The Q1 experiment was planned as NVFP4 (the QeRL thesis: a 32B model in ~5 GB on one H100), but vLLM’s NVFP4 kernel (`modelopt_fp4 / petit_nvfp4`) is compiled for `sm_100`, Blackwell only. On the H100 (`sm_90`) it dequantizes to BF16 at load, so “4-bit in 5 GB” becomes “BF16 in 64 GB” and the memory thesis evaporates. So the run that landed under the `q1_nvfp4_*` label actually used NF4 (the bitsandbytes format, which has working Hopper kernels). It should be read everywhere as the “Q1 NF4-LoRA proxy,” and true NVFP4 remains open future work gated on Blackwell access.

Second, the RL thinking-transfer result is scale-dependent and quantization-invariant, and the update style matters more than the precision. On the final multi-seed AIME-24 thinking table (4 seeds $\times$ $n=30$, Appendix R.3), RL at 32B transfers a small positive amount (Q1 NF4 +5.0pp, Q1 BF16 +2.5pp over base 0.475, within noise of each other), while at 14B LoRA-RL and distillation hold flat (Q3 NF4 +0.9pp, Q2 OPD +1.7pp over base 0.458, both noise). Full fine-tuning is the dangerous one: Q4 BF16 -23.6pp (0.222) and Q4 FP8 -10.8pp (0.350), multi-seed-confirmed regressions. The

named lesson: LoRA-RL preserves thinking-mode capability at 14B. Full-FT GRPO destroys it. The 14B's apparent +16.7pp on AIME greedy is a greedy-decoding artifact, the same checkpoint regresses on thinking-mode, so RL at 14B sharpened answer-extraction, not reasoning (the sharpen-not-expand pattern, [Part VII](#)). And the regression is invisible in attention-space: Q4 BF16's per-head KL is only $\sim 2.1\times$ the no-regression Q3 ($2.22e-4$ vs $1.04e-4$) while its thinking-mode accuracy halves, so attention-space KL is a poor predictor of thinking-mode capability loss (a mechanistic flag for [Part VII](#)). The internals that are preserved stay fully preserved, MMLU-mini and needle-in-a-haystack both 100%, KL lower ($\sim 22\%$) at 32B than 14B.




Figure 132.1. How you train decides whether the model keeps its reasoning: on a 14B model, full fine-tune RL drops AIME accuracy 23 points while LoRA RL preserves it.

Third, Q5's premise turned out to be structurally broken, not merely unimplemented: [Part XIV](#)'s rollout-quantization forensics showed the aggressive rollout quantization Q5 assumed (AWQ, FP8-KV) is unsafe for RL ([Chapters 121, 140–142](#)), so the experiment was closed by the finding rather than run. The full per-experiment record (the 55-cell grid, the KL scaling table, the multi-seed table, the full program's compute) lives in Appendix R. The lesson here is that quantization's job (train / rollout / serve) determines its safety, and the three must never be merged into one claim.

TAKEAWAY

Training-side quantization, kept distinct from rollout ([Part XIV](#)) and serving. **FP8 weights (actor training)** are *attention-space-equivalent* to BF16 at 14B (per-head KL: BF16 2.22e-4, FP8 2.19e-4) but **downstream quality is task/mode-dependent** (not “no regression”. Both full-FT arms regress on thinking-mode), say *which* FP8 and *which* role, never a blanket “FP8 is safe.” **NF4 = BF16 quality parity at 32B** across all 5 benchmarks (4-seed. NF4 has *lower* seed-variance). This *upholds* QeRL. The retraction runs the other way (the single-seed “BF16 wins +6.7pp” was the artifact, [Ch 93](#)). The RL thinking-transfer is **scale-dependent and quantization-invariant** (+5pp at 32B. ~0 at 14B for LoRA/OPD), but **full-FT GRPO at 14B destroys thinking-mode** (BF16 -23.6pp, FP8 -10.8pp, multi-seed): **LoRA preserves, full-FT forgets**. The regression is invisible in attention-space KL. **NVFP4 was never run**, its kernel is Blackwell/sm_100-only and dequantizes to BF16 on H100, so “Q1” is really an **NF4 proxy**. **Q5 was closed** by [Part XIV](#) (aggressive rollout quant unsafe). Internals stayed intact (MMLU/NIAH 100%). Quantization’s job determines its safety, never merge the three.

Next: what all this actually costs, the wall-clock and dollar reference.

CHAPTER 133

Wall-Clock and Compute Reference




Figure 133.1. 4-bit NF4 matches full BF16 at 32B across all five benchmarks — the bars are statistically indistinguishable, and NF4's variance is $\sim 3\times$ lower, so the win is stability.

Numeric reference		wall-clock and relative compute	
Run	Hardware	Steps	Wall-clock
Qwen3-4B GRPO, math (Big-Math-30k)	1 × H100	~300	~28 h
Qwen3-32B GRPO (NVFP4 + BF16 LoRA) ¹	1 × H100	300	~28 h
OPD GSM8K pilot ²	8 × H100	—	—
OPD GSM8K full run ³	4 × H100	~144–500	~22 h
same, BF16 teacher ⁴	8 × H100	500	~22 h
Qwen3-14B GRPO++, SWE-bench ⁵	8 × H100	—	—

Table 133.1. Order-of-magnitude planning numbers, not promises. “—” = no wall-clock/step figure for that row (see note). Both OPD-teacher rows assume a 32B-AWQ teacher → 14B student.

Table notes

1. **Qwen3-32B GRPO (NVFP4 base + BF16 LoRA):** external planning only — not demonstrated in this stack (Blackwell-gated; the landed 32B run was the NF4-LoRA proxy).
2. **OPD GSM8K pilot:** measured score 90.5% vs 89.5% base (200 ex) — a short pilot at far less compute than the full run.
3. **OPD GSM8K full run:** the case for an AWQ teacher at about half the compute of a BF16 teacher, valid only if the quality lift survives (F17, [Ch 119](#)).
4. **Same full run with a BF16 teacher:** the BF16-teacher baseline (8× H100).
5. **Qwen3-14B GRPO++ on SWE-bench (to SW6 plateau):** measured. The binding constraint is capacity, not compute — the figure here is ~1.5B tokens before plateau, not a wall-clock bill.

A plan you cannot size is a wish. This chapter is the reference for what RL post-training actually costs in wall-clock and compute, at the tiers this book uses, plus the single most common budgeting mistake.

133.1 The reference

The numbers above are the spine (planning figures. The full program's compute is detailed in Appendix R). The shape of them is the lesson. A small dense model (4B) on a verifiable task is a ~28-hour single-GPU run on one H100, cheap enough to iterate. A 32B model fits a single H100 only with a quantized base (about the same wall-clock for a small task), but note that the NVFP4 row is external planning only, not demonstrated in this book's H100 stack ([Chapter 132](#). NVFP4's kernel is Blackwell-gated, and the landed 32B run was the NF4-LoRA proxy). Treat it as a literature planning number, not a measured result. On-policy distillation is where the teacher's precision drives the compute, but separate the measured pilot from the planned run: the C10 OPD pilot measured 90.5% versus an 89.5% base on GSM8K at far less compute than the full run (a short run

that proved the method and the saturated-eval lesson), while the full AWQ-teacher (32B) → 14B-student plan takes about half the compute of a BF16 teacher (4× vs 8× H100). That gap is the argument for quantizing the teacher, valid only if the AWQ teacher’s logits don’t drift (F17, [Chapter 119](#)). And a hard agentic task like SWE-bench is measured not in wall-clock but in tokens to plateau (~1.5B for the 14B), because the binding constraint is capacity, not compute.

133.2 The budgeting rule

The single most common budget mistake is planning for the run that works. Your first one to three runs will fail, a collapse, a config crash, a contaminated eval, and you need compute reserved for the postmortem-then-retry cycle, not just the successful run. So add a 20–30% buffer for failed runs on top of the nominal cost, and treat the failures as a line item, not a surprise. A budget that assumes success is a budget you will blow on the first NaN.

There are also two hidden costs that dominate wall-clock more than the gradient step, and budgets routinely miss both. The first is validation overhead on multi-turn agents, easily the most under-budgeted line in the book:

Workload	Normal step	One validation pass	Multiplier	Fix
multi-turn search RL	~50 s	~68 min	~72×	test_freq=500–1000; smaller val set; fewer completions per eval

Table 133.2. Multi-turn validation overhead versus a training step, and how to cut it.

A full multi-turn validation can take 68+ minutes against a ~50-second training step, so validating too often can add 40–70% to total wall-clock, evaluate on a schedule, not every step. The second is checkpoint I/O: a Fully-Sharded Data Parallel (FSDP) save to network storage can take 20+ minutes for hundreds of gigabytes, so save to local NVMe first, keep `save_total_limit` small ([Chapter 120’s F11](#)),

and count the merge step. Neither shows up in a raw step-time estimate, and both are real wall-clock you'll pay.

TAKEAWAY

Planning numbers (wall-clock and relative compute): a 4B dense model on a verifiable task is a ~28-hour single-GPU run on 1× H100. A 32B model fits one H100 only via a quantized base + BF16 LoRA. OPD's **full-run** estimate with a 32B-AWQ teacher → 14B student takes about half the compute of a BF16 teacher, valid only if its logits hold (F17), but the measured C10 *pilot* was far cheaper (don't conflate them). Hard agentic tasks are measured in *tokens-to-plateau* (~1.5B for 14B SWE-bench). The **NVFP4 32B row is external planning only, not demonstrated here** (Ch 132. Landed run was NF4 proxy). Budget rule: your first 1–3 runs **will** fail (+20–30% buffer). And don't forget the two hidden costs, **multi-turn validation (~72× a training step → 40–70% overhead. Validate on a schedule)** and **FSDP checkpoint I/O to network storage (20+ min. Save to local NVMe first)**.

Next: the question to ask before spending any of that, does RL even help on this dataset?

CHAPTER 134

Dataset Selection: Does RL Help on This Data?

Dataset (coding)	Size	Difficulty	Verdict
HumanEval / MBPP	164 / 974	easy	eval only / tiny SFT+RL gains
LiveCodeBench-Lite	511	medium	eval
R2E-Gym-Lite	4,578	hard	the ONE where RL reliably helps a 14B model
R2E-Gym (full)	8.1k	hard	needs 8× H100
SWE-Gym	19k	very hard	starves GRPO (SW6; DeepSWE confirmed)
Nebius / Nemotron SWE	67k / 59k traj	hard	OPD-grade, not RL-from-scratch

Table 134.1. The same task family spans “RL does nothing” to “RL starves”, and there is a narrow band in the middle where RL reliably helps. Pick the dataset for the band, not the prestige.

Before you spend the budget of the last chapter, ask the question that decides whether to spend it at all: does RL actually help on this data? The answer is not “yes, RL is powerful”. It is dataset-specific, and getting it wrong means paying for a run that was never going to move.

134.1 The spectrum

The coding datasets (Table 4 from the source) lay the spectrum out cleanly. Easy sets (HumanEval, MBPP) are eval-only for a capable model. It already solves them, so there is nothing to learn. Very hard sets (SWE-Gym at 19k, SWE-Smith) starve GRPO: the reward is so sparse that almost every rollout fails, the advantage collapses (F2), and the run plateaus immediately. This is the SW6 capacity ceiling,

independently confirmed by DeepSWE. The huge trajectory sets (Nebius, Nemotron) are OPD-grade, good for on-policy distillation from a strong model, not for RL-from-scratch. The one in the middle, R2E-Gym-Lite (4,578 problems, hard but not impossible), is the dataset where RL reliably helps a 14B model, because it sits in the band where the base sometimes succeeds and sometimes fails. The SQL story rhymes: Spider is saturated (75–85%), while the genuinely hard BIRD tops out around 70%. There is room to climb only where the base hasn't already arrived.

134.2 The diagnostic

The underlying principle is [Part VII](#)'s: RL sharpens what the base can already sometimes do. It does not conjure new capability. So the dataset must contain problems the base solves part of the time. The diagnostic is a single base-model eval before any training: if the base scores near 0%, the task is too hard. There is no signal to amplify, and you need an SFT cold-start or a stronger base (or a different task). If it scores near 100%, the task is saturated and RL has nothing to add. RL's value lives in the middle band, roughly 20–80% base success, where rollouts disagree, the advantage is informative, and sharpening has something to work with. Run that one eval before you budget the run. It is the cheapest experiment you will do and it decides whether the expensive one is worth launching.

TAKEAWAY

Ask “does RL help on *this* data?” before budgeting. The spectrum: easy sets (HumanEval/MBPP) are eval-only (base already solves them). Very-hard sets (SWE-Gym 19k) **starve** GRPO (advantage collapse, the SW6 capacity ceiling). Huge trajectory sets are **OPD-grade**, not RL-from-scratch. The sweet spot is the middle band (R2E-Gym-Lite). The diagnostic is one base eval: near 0% = too hard (cold-start or change task), near 100% = saturated, **20–80% = RL helps**. RL sharpens what the base can sometimes do ([Part VII](#)).

Next: why two version-pinning incidents prove that a version number is part of your result.

Version Pinning as a Result

Incident	What changed	Effect
EAGLE-3 acceptance collapse	same base model + EAGLE-3 head, later run on the SGLang 0.5.12.post1 stack (original flags not fully recoverable)	acceptance $\approx 0.283 \rightarrow \approx 0.0005$ (a reproducibility warning, not a clean patch-only ablation)
AutoAWQ archival	the AutoAWQ project was archived; <code>transformers-4.57</code> import break + an unfixable NaN kernel	the AWQ path stopped working

Table 135.1. Two incidents that turn “pin your versions” from generic advice into a finding: in both, nothing in the model, data, or config changed, only a dependency, and the result moved from working to broken. Spell-outs: EAGLE-3 = a speculative-decoding draft head. SGLang = an inference server. AutoAWQ = the AWQ quantization library.

“Pin your versions” is the kind of advice everyone nods at and nobody internalizes, until a patch-level bump silently halves your throughput or breaks your quantization path. This short chapter records two incidents that make the point concrete: a version pin is not hygiene, it is part of the result, and replicating a number requires replicating the stack that produced it.

135.1 The two incidents

The first is the more alarming. Be precise about what it does and does not support: the strongest supported claim is not a clean patch-only ablation, because the original launch flags from the first run were not fully recoverable. What is supported is a serious reproducibility warning, the same base model and EAGLE-3 head that produced acceptance ≈ 0.283 earlier produced ≈ 0.0005 under a later SGLang

0.5.12.post1 stack, with no error. The lost flags are not a weakness to hide. They are the lesson: if the exact launch command and engine version are not captured, the result is not reproducible. The second is an archival failure: when the AutoAWQ library was archived, it left a `transformers-4.57` import break and a NaN kernel that was not worth fixing inside this project's budget (blocked on an archived dependency), so the AWQ path stopped working on current dependencies, and any AWQ result depends on a frozen environment that no longer installs cleanly.

135.2 The rule

The transferable rule is the inference-side companion to the manifest ([Chapter 130](#)): anyone replicating a speculative-decoding or quantization number must pin the exact engine versions, not the major version, the patch version, because these numbers live in the dependency stack, not just in the model. Record the SGLang/vLLM/AutoAWQ versions in the manifest with the same seriousness as the model revision, and treat a version bump as a change that requires re-measurement, not a free upgrade. A speedup reported without its engine version is not reproducible. It is a snapshot of a stack that may no longer exist.

TAKEAWAY

Two incidents prove a version pin is a *result*, not hygiene. EAGLE-3 acceptance went from ≈ 0.283 to ≈ 0.0005 on the same model/head under a later **SGLang 0.5.12.post1** stack, and because the original flags weren't fully recoverable, the lesson is precisely that *uncaptured launch commands aren't reproducible* (a strong warning, not a clean ablation). **AutoAWQ's archival** left a `transformers-4.57` import break + a NaN kernel not worth fixing on an archived dep, breaking the AWQ path. Rule: capture the **exact engine versions and launch command** in the manifest, and treat any bump as requiring re-measurement.

Next: [Part XIV](#). We open the rollout engine, inference for RL training, with the one principle the whole Part rests on: the correctness contract.

Part XIV

The Rollout Engine

RL post-training spends most of its wall-clock not on the gradient step but on generating rollouts, and almost everything written about fast inference is written for serving, where the only questions are latency and throughput. RL adds a third question that serving never asks: did the acceleration change the policy's distribution? This Part is inference for RL training, nine experiments, the failures the correctness contract caught, and the four genuinely novel findings, and it opens with the contract itself.

CHAPTER 136

The RL Inference Correctness Contract




Figure 136.1. Which checks an acceleration must pass depends on its class: safe-by-construction needs only the speed gate, exact also needs identical outputs, lossy needs all three.

This is the load-bearing principle of the whole Part, and the reason the rollout-lossy failures of [Chapter 121](#) were caught instead of shipped: a rollout acceleration is only “free” if it passes the gates that apply to its class. Serving asks “is it faster?”. RL must also ask “is the policy that generated these rollouts still the policy I am training?”, because the importance-sampling update assumes it is.

136.1 The three gates

An acceleration earns adoption by passing up to three gates. Gate A, per-rollout KL: the gate must be defined precisely or it is not reproducible. Concretely: $\text{KL}(\text{unaccelerated} \parallel \text{accelerated})$, in nats/token, over generated-token positions only (no prompt, padding, or tool tokens), teacher-forcing both engines on the same sampled prefixes at a matched seed and temperature, over ~100–500 prompts. The threshold is < 0.05 nats/token. This is the gate

that catches distribution corruption. It is the one FP8-KV failed by $152\times$ (KL 7.61, $\sim 40\%$ top-1 disagreement at matched seed/ $T=0$. [Chapter 121](#)), measured with exactly this estimator. Gate B, wall-clock: the acceleration must actually reduce end-to-end step time, measured over the real training loop, not a microbenchmark (some “faster” kernels lose their win once weight-sync and batching are included). Gate C, held-out eval: a 500-sample held-out evaluation after a short RL run must hold, confirming the acceleration didn’t quietly degrade what the policy learns. Gate A is necessary but not sufficient. Gate C is the backstop that catches what a single-number KL might miss.

136.2 The taxonomy that decides which gates apply

Not every acceleration needs all three gates, and the class decides which apply. Safe-by-construction accelerations perform the same computation a different way (a faster attention kernel that is numerically identical). They need only Gate B, because by construction they cannot move the distribution. Exact accelerations claim bit-identical outputs at a matched seed. They need Gate B plus a verification that outputs really are identical. Lossy accelerations change the numbers, quantized key-value caches, quantized rollout weights, and they must pass all three gates, especially A and C, because their entire risk is a distribution shift the output text hides. The practical force of the taxonomy is that it stops you from waving through a lossy acceleration on a serving-style “97% accuracy recovery” claim: in the RL setting, “97% recovery” is a failing Gate A, and the contract is what makes that visible.

TAKEAWAY

A rollout acceleration is “free” only if it passes the gates for its class. **Gate A:** per-rollout KL < 0.05 at matched seed/ T (catches distribution corruption, FP8-KV failed it $152\times$). **Gate B:** real end-to-end step wall-clock, not a microbenchmark. **Gate C:** a 500-sample held-out eval after a short RL run. The **class** sets which apply: *safe-by-construction* (same computation) needs only B. *exact* (bit-identical) needs B + verification. *lossy* (quantized

KV/weights) needs **all three**. The contract is why a serving-style “97% recovery” is correctly read as a *failing* Gate A.

Next: a speculative-decoding temperature sweep — what acceptance does as RL temperature rises.

CHAPTER 137

Speculative Decoding Across RL Temperatures




Figure 137.1. Speculative-decoding speed-up barely moves across RL temperatures: measured MTP holds $1.27\times \rightarrow 1.18\times$ where serving lore predicts a steep fall.

Speculative decoding is the first acceleration in this Part because it is safe by construction ([Chapter 136](#)), the draft head proposes tokens, the target model verifies them, and the output is identical to non-speculative decoding, so it needs only the wall-clock gate. The open question for RL was never whether it is correct. It is whether the speedup survives at RL’s sampling temperatures, since nearly every published number assumes greedy ($T=0$). This chapter is, as far as we found, the first temperature sweep of speculative decoding for the RL setting (we did not find a prior matched RL-temperature sweep). Multi-Token Prediction (MTP, also called NEXTN) is a draft head jointly pretrained with the model. Setup: Qwen3-Next-80B-A3B, SGLang with the NEXTN draft head (3 draft steps), $n=50$ prompts, fixed output length, the engine version is pinned in the manifest ([Chapter 135](#) shows why that matters for draft heads).

137.1 The sweep

We swept MTP on Qwen3-Next-80B across $T \in \{0, 0.5, 0.7, 1.0\}$ and measured both the acceptance rate (mean draft tokens accepted per verification step) and the end-to-end speedup. The result, against the serving-lore expectation that higher temperature flattens the target distribution and collapses acceptance:

Temperature	Acceptance rate	Mean accepted draft length	Speedup
0.0	0.601	2.22 / 3	1.27×
1.0	0.558	2.13 / 3	1.18×

Table 137.1. MTP acceptance and speedup barely fall from greedy to RL sampling temperature.

(Two distinct metrics: the acceptance rate, 0.601, is the fraction of proposed draft tokens the target verifies. The mean accepted draft length, 2.22 of 3, is how many of the 3 proposed tokens survive per verification step. The sweep also covered $T=0.5$ and $T=0.7$, which fall monotonically between these endpoints.)

The acceptance rate falls only 4.3 percentage points from $T=0$ to $T=1.0$, and the speedup falls only 9 points ($1.27\times \rightarrow 1.18\times$), far gentler than the steep decay serving lore predicts (a fall toward $\sim 1.3\times$ from a greedy $\sim 2\times$). At $T=0.7$, a common RL temperature, this is about a 13% whole-pipeline saving. The $1.27\times$ at $T=0$ is itself lower than the $1.8\text{--}2.1\times$ the serving literature reports for greedy on dense models, which points at the mechanism.

137.2 Why it’s temperature-robust, and the adoption rule

Two properties explain the flat curve. First, Qwen3-Next is a hybrid architecture (most layers are linear-attention / gated-DeltaNet, not full softmax attention), so per-token decode is already cheap. There is less to amortize, which lowers the headline speedup but also makes it insensitive to how often the draft is accepted. Second, the MTP heads are jointly pretrained with the base model rather than bolted on, so

their proposals track the target distribution even as temperature flattens it. The practical rule is therefore unambiguous: adopt speculative decoding at any RL temperature, the speedup you measure at greedy is roughly the speedup you keep at $T=1.0$. (This also refutes a community report that MTP slows rollout by ~50%. That result was specific to the H20 accelerator, not a property of MTP.)

TAKEAWAY

Speculative decoding is safe-by-construction (Gate B only, and only when the target verification implements exact speculative sampling at the requested temperature/top-p, leaving the output distribution unchanged), and, first RL-temperature sweep we found, its speedup is **temperature-robust**: MTP on Qwen3-Next-80B holds **1.27× at $T=0$ → 1.18× at $T=1.0$** (acceptance -4.3pp, speedup -9pp), about a **13% whole-pipeline saving at $T=0.7$** . The mechanism: hybrid layers leave less to amortize (lower headline speedup) and jointly-pretrained MTP heads track the target distribution as T flattens it. Adopt it at any RL temperature. The greedy number is roughly the number you keep.

Next: which draft head, MTP or EAGLE-3?

MTP vs EAGLE-3: The First Head-to-Head




Figure 138.1. Spec-decode speed-up vs temperature, measured endpoints: 1.27× at T=0 falls only to 1.18× at T=1.0.

Draft head	Acceptance @ T=0	Across temperature
MTP (NEXTN)	0.601	flat
EAGLE-3	0.283	flat

Table 138.1. Two speculative-decoding draft heads, on different base models (MTP on Qwen3-Next-80B, EAGLE-3 on a Qwen3 dense model), so this is a directional comparison, not a matched ablation. MTP accepts roughly 2.1× as many draft tokens per step. The direction is the claim, not the exact ratio.

EAGLE-3 reproducibility hazard. A later SGLang 0.5.x version regressed EAGLE-3 acceptance to near-zero on identical weights, and the original server flags were not fully recoverable ([Chapter 135](#)). Treat EAGLE-3 numbers as engine-version-pinned.

If you are going to add a draft head, which one? The two candidates are MTP (jointly pretrained with the model, [Chapter 137](#)) and EAGLE-3 (a small auxiliary head, ~1–2% of base parameters, trained on top of an existing model). No published comparison existed for the RL setting, so this is a first head-to-head, but a directional one: it is a practical systems comparison under the available public draft heads, not a matched-model scientific ablation (the two ran on different bases).

138.1 The comparison

On acceptance rate, the quantity that drives the speedup, MTP beats EAGLE-3 by about $2.1\times$ (0.601 versus 0.283 at $T=0$), and both are flat across temperature, consistent with [Chapter 137](#)’s temperature-robustness finding. The mechanism is the same one that makes MTP temperature-robust: a head jointly pretrained with the model proposes tokens that track the target distribution more faithfully than a head trained afterward to imitate it. The honest caveat is essential here and is stated as part of the finding: the MTP measurement was on Qwen3-Next-80B and the EAGLE-3 measurement on a Qwen3 dense model, so the two are not perfectly matched. The claim this chapter makes is therefore directional, when a jointly-pretrained MTP head is available, prefer it, not that the ratio is exactly $2.1\times$ on a matched model.

138.2 What this means in practice

The practical consequence is a preference order, not a verdict: if the model ships MTP heads, use them. Reach for EAGLE-3 when it does not, accepting a lower acceptance rate for the convenience of bolting a head onto an existing checkpoint. And both inherit [Chapter 137](#)’s good news, the acceptance you measure at greedy is roughly what you keep at RL temperatures, so the draft-head choice can be made once, at $T=0$, and trusted across the sweep. (The next chapter’s version-pinning lesson, [Chapter 135](#)’s EAGLE-3 incident, is the caution that rides alongside this: a draft head’s acceptance can

also collapse to near-zero on an engine-version bump, so pin the inference stack before you trust either number.)

TAKEAWAY

First MTP-vs-EAGLE-3 head-to-head for RL: **MTP accepts ~2.1× as many draft tokens** (0.601 vs 0.283 at T=0), both flat across temperature, because a *jointly-pretrained* head tracks the target distribution better than one trained to imitate it afterward. Honest caveat: different base models, so the **direction** is the claim, not the exact ratio. Practical order: use MTP if the model ships it. Use EAGLE-3 otherwise, and pin the engine version (Ch 135), since acceptance can collapse on a bump.

Next: prefix caching, a win that exists only at a specific seam.

Prefix Caching: The Win Lives at a Seam

Pattern	Automatic Prefix Caching (APC) win
Grouped n -of-1 in one batch	0% (in-batch sharing already amortizes it)
Short multi-turn (<100 tok/turn)	0.5%
Sequential reuse of a ~1.5K-token prefix across batches	26% wall-save / 1.40×

Table 139.1. The same feature, three patterns, three completely different answers, which is why “enable prefix caching” is bad advice without naming the pattern. The win is real only when a long prefix is reused across separate batches, exactly the agentic re-rollout pattern.

Prefix caching is the acceleration everyone gets wrong, because the answer depends entirely on the access pattern, and the obvious pattern is the one where it does nothing. Automatic Prefix Caching (APC) is a key-value-cache (KV-cache) radix tree that reuses the cached prefix across separate requests over time. It is exact ([Chapter 136](#)), it changes nothing about the output, so the only question is whether it saves wall-clock, and the surprising answer is “only at one seam.”

139.1 Where it does nothing, and where it wins

The intuition that n identical rollouts of one prompt should benefit hugely from prefix caching is wrong, and the reason is a different optimization you already have: when you pass a grouped batch of identical prompts in a single request, the inference engine’s continuous batching already shares the forward pass across them, so the radix cache adds essentially 0%. Short multi-turn rollouts, where each turn re-sends a small growing history, save only 0.5%, the prefix

is too short for the recompute to matter. The win lives at exactly one seam: reusing a long (~1.5K-token) prefix across separate, sequential batches, generate batch 1 with a prefix, then batch 2 with the same prefix but a different question, where APC saves 26% of rollout wall-time (1.40× throughput). That is precisely the GRPO multi-step / agentic re-rollout pattern, where the same long system prompt or retrieved context is replayed across many rollouts that are not in one batch.

139.2 The bimodality, and the rule

One more measured detail makes the feature's behavior honest: the per-batch speedup is bimodal. Most batches in the sequential-reuse test ran at a routine ~1.31×, but occasional batches hit a 2.68× cliff when the cache alignment was favorable, so the 26%/1.40× average is the blend of a steady win and intermittent large ones, not a uniform speedup. The rule that falls out: enable APC when a long prefix is reused across batch boundaries (agentic re-rollouts, multi-step GRPO with a fixed context) and do not expect it to help a single grouped batch, where continuous batching has already done the work.

TAKEAWAY

Prefix caching (APC) is exact (no Gate A), but its win depends entirely on the access pattern: **0%** for a grouped n -of-1 batch (continuous batching already shares it), **0.5%** for short multi-turn, and **26% wall-save / 1.40×** only when a **long (~1.5K-token) prefix is reused across separate sequential batches**, the agentic re-rollout pattern. The per-batch speedup is bimodal (routine ~1.31× with occasional 2.68× cliffs). Enable it at that seam. Don't expect it elsewhere.

Next: the first lossy acceleration, and the first failure the correctness contract catches.

CHAPTER 140

FP8 KV Is Unsafe for RL: The Contract Catch

The flagship failure. An 8-bit-floating-point key-value cache (FP8 KV) gave no throughput win (0.97×, actually 3% slower) and failed Gate A by 152×: per-token Kullback–Leibler (KL) divergence of 7.61 against the unquantized policy, with ~40% of tokens disagreeing at matched seed and $T=0$. We are not aware of a prior published quantitative KL number for FP8-KV in the RL setting.

This is the chapter the correctness contract ([Chapter 136](#)) was built for. FP8-KV is a popular serving optimization, store the attention cache in 8-bit floating point to fit more context, and the serving literature reports it recovers ~97% of accuracy. The RL setting asks a different question (does the rollout distribution still match the training policy?), and the answer is a categorical no.

140.1 The measurement

We ran FP8-KV through all three gates on Qwen3-14B (single H100, ~700-token prefix, $n=8$ group). It failed both gates that mattered. Gate B (wall-clock): FP8-KV was 0.97×, 3% slower than the BF16 baseline at this workload, so there was no speed to trade off in the first place. Gate A (distribution): mean per-token KL was 7.61, against the bar of 0.05, a failure by 152×, with ~40% of tokens disagreeing at matched seed and $T=0$. A 40% per-token disagreement means the policy that generated the rollouts is, on nearly half its tokens, a different policy from the one the trainer updates, which invalidates the importance ratio at the heart of the update. There was no point running Gate C. A distribution this corrupted cannot produce a trustworthy gradient.

140.2 Why FP8 weights are fine but FP8 KV is poison

The crucial distinction, and the reconciliation with the training-side record ([Chapter 132](#), where FP8 weights showed no quality regression), is when and how often the quantization error is incurred. FP8 weights are quantized once at load and dequantized for the matmul, so the error is a single fixed perturbation that the model effectively trains around (the Q4 equivalence result). FP8 KV is different: the cache is re-read lossily on every attention operation, and because each token attends to all previous tokens, the quantization error compounds over the context length, growing as the sequence lengthens. The same three letters, “FP8,” name a benign operation in one place and a distribution-destroying one in another, which is exactly why [Chapter 132](#)’s role table insists you say which FP8 and where. This result also refutes, for RL specifically, the serving framing that FP8-KV “recovers 97% of accuracy”: in the RL setting, that residual is a failing Gate A.

The reusable artifact behind this measurement, the Gate-A preflight, is small enough to state directly:

```
@torch.no_grad()
def rollout_gate_a(target, accel, prompts, *, temperature,
    ↪ max_new_tokens, seed):
    # 1. generate with the ACCELERATED sampler at a fixed seed
    out = accel.generate(prompts, temperature=temperature,
        max_new_tokens=max_new_tokens, seed=seed)
    # 2. score those EXACT generated tokens under both policies
    lp_accel = accel.score(prompts, out.tokens) # log-probs,
    ↪ response tokens only
    lp_target = target.score(prompts, out.tokens)
    mask = out.response_mask # exclude prompt / pad / tool tokens
    d = (lp_accel - lp_target)[mask] # per-token log-prob gap
    return {
        "mean_token_kl": forward_kl(lp_target, lp_accel, mask), #
        ↪ KL(target || accel), nats/token
        "max_abs_gap": d.abs().max().item(), # worst-token gap (nats)
        "frac_outside_clip": ((d.exp() < 0.8) | (d.exp() >
        ↪ 1.2)).float().mean().item(),
        "token_disagree": token_argmax_disagree(lp_accel, lp_target,
        ↪ mask),
```

```
"n_tokens": int(mask.sum()),  
}
```

Run it on ~100–500 prompts at a matched seed and temperature. FP8-KV returns $\text{mean_token_kl} \approx 7.61$ against the 0.05 bar. Direction is forward KL $\text{KL}(\text{target} \parallel \text{accel})$ over generated-token positions only.

TAKEAWAY

FP8 KV cache is **unsafe for RL** and is the contract's flagship catch: **no speed win (0.97×, 3% slower)** and Gate A failed by **152×**, per-token KL **7.61**, ~40% token disagreement at matched seed/T=0 (we are not aware of a prior published KL number for this). The mechanism reconciles it with [Chapter 132](#)'s benign FP8 *weights*: weights quantize **once at load** (a fixed perturbation), but the **KV cache is re-read lossily every attention op**, so error **compounds over context**. "FP8 recovers 97% accuracy" is a *serving* claim. In RL that residual is a failing gate.

Next: AWQ rollout weights under the microscope, the forensics that turn "AWQ is risky" into mechanism.

Rollout-Quantization

Forensics: AWQ Under the Microscope




Figure 141.1. Lossy rollout quantization breaks the policy: per-token KL towers over the safe threshold (FP8-KV 7.61, AWQ 6.86), and MoE routing flips for 31% of top-1 / 87% of top-8.

Path	Gate A (KL vs BF16)	Verdict
dense-32B AWQ → vLLM rollout	KL 6.86; 42.8% (137/320) token disagreement; 0.95× speed	unsafe rollout
dense-32B AWQ → HF + AutoAWQ	NaN cascade at layer 2	unusable path (AutoAWQ kernel bug)
MoE 30B-A3B AWQ → HF + AutoAWQ forensics	KL 0.387; cos 0.995→0.9186 over 48 layers (clean)	risky / mechanism (Ch 142)

Table 141.1. “AWQ is not one thing.” Three different model/code-path combinations gave three different results. They must not be conflated. The dense-32B rollout corrupts the distribution. The dense-32B HF+AutoAWQ kernel crashes. The MoE forensics path

runs cleanly and is where the mechanism (and the router-flip of [Ch 142](#)) is read.

AWQ, Activation-aware Weight Quantization, quantizes weights to 4-bit while protecting the channels that activations hit hardest. It is the natural candidate for a cheaper rollout engine, and “AWQ is risky for RL” was folklore. This chapter replaces the folklore with mechanism, and in doing so closes the long-open “Q5” quantization question (whether AWQ-rollout-with-BF16-trainer is viable) for negligible compute.

141.1 The verdict

The first and most important correction this chapter makes is that “AWQ” names three different experiments, and they must not be merged. On the dense Qwen3-32B via the vLLM rollout path, AWQ-INT4 failed Gate A nearly as badly as FP8-KV: KL 6.86 against BF16, 42.8% token disagreement (137 of 320) at T=0, 0.95×, slightly slower. Same shape of failure as FP8-KV: it corrupts the distribution and delivers no speed. Separately, the same dense Qwen3-32B-AWQ checkpoint through the HuggingFace + AutoAWQ path produced a NaN cascade originating at layer 2, an AutoAWQ kernel failure, a different problem from the rollout-distribution failure, and one that simply has no clean H100 path. These are two distinct findings about the dense model: one says the AWQ rollout is unsafe, the other says this AWQ kernel path is unusable.

141.2 The mechanism, and the honest confounds

The mechanism could not be read on the dense-32B (it crashed), so the forensics moved to the Qwen3-30B-A3B mixture-of-experts via the HF+AutoAWQ path, which, critically, ran cleanly, with no NaN. There the per-layer divergence is monotonic: cosine similarity between the AWQ and BF16 hidden states declines steadily from ~0.995 to ~0.9186 across the 48 layers (overall KL 0.387), the signature of an error that compounds with depth rather than jumping at one bad layer. And

a layer’s activation outlier magnitude correlates with its divergence at $r \approx +0.87$, exactly the SmoothQuant mechanism (outlier channels are where weight quantization hurts most), now demonstrated on our own model. Note that this clean curve is the MoE path. Do not attribute it to the dense-32B, which NaN’d. Two confounds are stated honestly. First, the official Qwen3-32B-AWQ ships in fp16, not bf16, so some fraction of the dense 42.8% disagreement is fp16-vs-bf16 numerical noise rather than AWQ-specific, the matched-dtype isolation rerun is an evidence boundary (Chapter 147). Second, AWQ’s calibration set is undisclosed, so the math prompts here may exercise channels the calibration didn’t protect. Calibration mismatch is a live candidate root cause. The verdict by role: AWQ is unsafe as an RL rollout engine, a risk to verify as an OPD teacher (Chapter 119’s F17), and an acceptable serving choice (after a serving-quality check) where the distribution need not match a trainer.

TAKEAWAY

“AWQ” is **three different experiments**, not one. **Dense-32B → vLLM rollout**: KL **6.86**, 42.8% disagreement, 0.95× (slower), unsafe rollout. **Dense-32B → HF+AutoAWQ**: **NaN at layer 2**, unusable kernel path. **MoE-30B-A3B → HF+AutoAWQ forensics**: clean (no NaN), KL **0.387**, monotonic $\cos \sim 0.995 \rightarrow 0.92$ over 48 layers, outlier↔divergence $r \approx +0.87$ (the SmoothQuant mechanism). This is the *mechanism* path (and where the router-flip lives, Ch 142). Do not attribute the clean curve to the dense-32B (it crashed). Confounds: the official 32B ships fp16-not-bf16 (some disagreement is dtype noise. Matched-dtype rerun is a boundary) and AWQ’s calibration set is undisclosed. Verdict by **role**: unsafe rollout / verify-as-teacher / serving-OK-after-check.

Next: the single most citable forensics finding, how a quantized rollout flips a mixture-of-experts router.

CHAPTER 142

The MoE Router-Flip Mechanism




Figure 142.1. A quantized rollout reroutes a third of all tokens: 31% flip their top-1 expert and 87% see a top-8 change, even with the routing gate left unquantized — a different routing of the policy, not a smaller one.

This is the single most citable forensics figure in the Part, because it shows why a mixture-of-experts (MoE) model is even more fragile under a quantized rollout than a dense one, and it connects the rollout-side story straight back to the training-side router instability of [Part V](#) (the F5 ratio explosion). The finding isolates the mechanism cleanly by leaving one thing alone.

142.1 The measurement

On a Qwen3-30B-A3B MoE (48 layers, 128 experts, top-8 routing) with 4-bit AWQ weights, we measured how often the router’s choice changes versus the BF16 model. 31% of tokens route to a different

top-1 expert (mean top-1 flip 0.310), and 87% of tokens have at least one expert change in their top-8 set (mean 0.868). The decisive control: the routing gate itself was left unquantized. Because the gate is bit-identical to BF16 and the routing still flips on a third of tokens, the flips cannot be a gate-quantization artifact. They are driven entirely by the quantization error in the upstream activations that feed the gate. (The claim is scoped to this Qwen3-30B-A3B-AWQ path. It shows large routing divergence here, not a universal theorem that every MoE rollout quantization fails.) And the flip rate is roughly stable across the 48 layers, per-layer top-1 flips range about 22–40%, recurring throughout rather than concentrated early or late. The consequence for the RL update is exactly the importance-ratio break this Part keeps returning to: the AWQ-MoE rollout preflight (Qwen3-30B-A3B-AWQ rollout vs BF16 target, $T=0.7$) measured a mean sequence KL of 0.2864, a maximum per-token log-probability gap of 6.72 nats (an importance ratio of $e^{6.72} \approx 825\times$ on the single worst token of each sequence), and 18.6% of per-token ratios outside the standard clip band $[0.8, 1.2]$, so the trainer silently truncates roughly one token in five of every rollout while assuming the ratio is 1. The router flips are the mechanism. This preflight is the gradient damage they cause.

142.2 Discrete-everywhere versus continuous-saturating

The router-flip is why MoE and dense quantization fail differently. A dense model's error is continuous: small per-layer perturbations that compound smoothly with depth ([Chapter 141](#)'s monotonic cosine decline). A MoE's error is discrete and everywhere: a tiny activation perturbation that nudges a routing logit past its neighbor swaps the entire expert, changing not a few activations but which sub-network ran, and it does so on a third of tokens at every layer. This is the same fragility as the training-side F5 ratio explosion ([Chapter 117](#)), reached through a different door: there, the policy update flips the routing. Here, the quantization does. The practical consequence is blunt, a quantized MoE rollout is not a smaller version of the policy. It is a different routing of it, which is precisely the distribution shift the importance ratio cannot absorb.

TAKEAWAY

The most citable forensics result: under 4-bit AWQ, a Qwen3-30B-A3B MoE has **31% of tokens flip their top-1 expert and 87% change their top-8 set** (per-layer ~22–40%), with the **gate left unquantized**, proving the flips are driven by upstream-activation error. The gradient damage: the rollout preflight showed mean-seq-KL 0.2864, a 6.72-nat worst-token gap ($\rightarrow 825\times$ ratio), and **18.6% of ratios outside the clip band** (the trainer silently clips ~ 1 token in 5). MoE fails **discretely** (a nudged logit swaps the whole expert) versus dense's **continuous** error, the same fragility as training-side F5 (Ch 117), through a different door. Scoped to this AWQ-MoE path, not a universal theorem. A quantized MoE rollout is a *different routing* of the policy, not a smaller one.

Next: the acceleration that finally passes, W8A8.

W8A8: The Safe Rollout Quant

Rollout quant (Qwen3-32B)	Gate A (KL vs BF16)	Verdict
AWQ-INT4 (weight-only)	6.86	reject
W8A8 (INT8 weight + INT8 activation)	0.0055	safe under this Gate-A preflight — passes the 0.05 bar

Table 143.1. The actionable positive of the Part. 8-bit-weight, 8-bit-activation quantization (W8A8) is $\sim 50\times$ cleaner than 4-bit AWQ on the same 32B model, and refutes the intuition that quantizing more dimensions must be worse.

After three rejected lossy accelerations, this is the one that works, the positive result that keeps the Part from reading as “all quantization is unsafe.” W8A8 quantizes both weights and activations to 8-bit integers (versus AWQ’s 4-bit weights, full-precision activations). On the same Qwen3-32B, it passes the correctness contract outright.

143.1 The result, and the counterintuitive mechanism

W8A8’s Gate A KL is 0.0055, versus AWQ-INT4’s 6.86 on the identical model, making W8A8 about $50\times$ cleaner, and comfortably under the 0.05 bar (by roughly $9\times$). This refutes a natural intuition: you might expect that quantizing activations too (W8A8) would be worse than leaving them in full precision (AWQ), because activations carry the outliers. The opposite holds, and the reason is the scale granularity. AWQ uses a shared per-group weight scale that must cover a whole group of channels, so an outlier channel forces a coarse scale that loses precision everywhere in the group. W8A8 uses a per-token dynamic activation scale, computed fresh for each token, which adapts to that token’s outliers individually. More quantized dimensions, handled with

a finer-grained scale, beats fewer dimensions handled coarsely, granularity matters more than how many tensors you quantize.

143.2 The preflight, as a tool

The result was measured on Qwen3-32B (compressed-tensors): last-token KL 0.0055, cosine 0.9978, safe at $T=0.05$. The go/no-go preflight tool itself, the concrete instrument behind [Chapter 136](#)’s Gate A, reports the sequence-level KL, the maximum per-token log-probability gap, and the fraction of token ratios outside the clip band. Those last two are dramatic on the rejected AWQ paths, not on W8A8 (the AWQ-MoE preflight’s 6.72-nat max gap $\rightarrow$ 825 $\times$ ratio and 18.6%-outside-clip belong to [Chapter 142](#), not here), W8A8’s headline is simply that its mean KL passes the bar cleanly. The corpus does not show a max-gap / clipped-ratio table for W8A8. That is a measurement to add, not one to imply was taken. One operational cost is recorded honestly: a W8A8 checkpoint takes about 6 $\times$ longer to load ($\approx$ 49 s/shard versus $\approx$ 8 s/shard) because of compressed-tensors decompression at load, a one-time price for a correctness-clean rollout (the W8A8 load-time note of [Chapter 129](#)).

TAKEAWAY

W8A8 (INT8 weight + INT8 activation) **passes Gate A under this preflight**: KL **0.0055** on Qwen3-32B, **$\sim$ 50 $\times$ cleaner than AWQ-INT4**, comfortably under the 0.05 bar. It refutes “more quantized dims = worse”: a **per-token dynamic activation scale** handles outliers better than AWQ’s **shared per-group weight scale**, granularity beats dimension count. (The dramatic preflight outliers, a 6.72-nat gap $\rightarrow$ 825 $\times$ ratio, 18.6% outside the clip band, are the AWQ-MoE preflight, [Chapter 142](#), not W8A8. The corpus shows only the clean mean-KL for W8A8.) Cost: $\sim$ 6 $\times$ longer load (decompression). Frame it as “safe under this Gate-A preflight,” not a universal guarantee.

Next: the open theoretical claim, why two unrelated lossy methods both land near 40%.

CHAPTER 144

The ~40% Saturation Hypothesis

Two completely different lossy channels, FP8-KV (KL 7.61) and AWQ-INT4 (KL 6.86), both land at ~40% token disagreement at $T=0$. That coincidence is the seed of a hypothesis, stated here as a hypothesis with a specified test, not as a result.

This is the Part's one openly theoretical chapter, and it is labeled as such. Two unrelated corruptions converging on the same disagreement rate is suggestive, and the honest move is to offer the mechanism that would explain it and the experiment that would test it, not to dress a coincidence as a finding.

144.1 The hypothesis

FP8-KV ([Chapter 140](#)) disagreed with the unquantized policy on ~40% of tokens. AWQ-INT4 ([Chapter 141](#)) on 42.8%. These are different perturbation sources, one corrupts the attention cache, the other the weights, yet they land in the same place. The proposed mechanism is an argmax-flip saturation: at greedy decoding, a token's output flips only if the perturbation to its logits, $|\Delta\text{logit}|$, exceeds the gap between the top-1 and top-2 logits at that position. In math and code generation, that top-1/top-2 gap is typically only 1–3 logits. So once a perturbation source produces $|\Delta\text{logit}|$ in that range, the argmax flips with near-Bernoulli probability ($\approx 50\%$) on the contested tokens, and the overall disagreement saturates at a rate set by the distribution of gaps in the data, not by the source of the perturbation. The perturbation source determines only whether you are below or above the flipping threshold. Once above it, everyone saturates at the same place.

144.2 The test that would confirm it

The hypothesis is falsifiable, and the chapter specifies the run. Measure, for the BF16 baseline, the histogram of top-1/top-2 logit gaps across token positions. Then measure, for each lossy method, the histogram of $|\Delta\text{logit}|$ perturbations. The prediction is precise: disagreement should equal the fraction of positions where the perturbation histogram exceeds the gap histogram, and that fraction should be ~40% for any source whose perturbations land in the 1–3-logit range, converging regardless of whether the source is FP8-KV or AWQ. This test is listed in the open-problems chapter ([Chapter 147](#)). Until it is run, the ~40% convergence is a strong hypothesis with a mechanism and a test, which is the honest status to give it.

TAKEAWAY

Two unrelated lossy channels (FP8-KV ~40%, AWQ 42.8%) converge on **~40% token disagreement** at T=0, stated as a **hypothesis, not a result**. Proposed mechanism: at greedy, a token flips only if $|\Delta\text{logit}|$ exceeds the top-1/top-2 gap (typically **1–3 logits** in math/code), so any source whose perturbations reach that range flips contested tokens near-**Bernoulli** and **saturates at a rate set by the data's gap distribution, not the source**. The falsifiable test (gap histogram vs $|\Delta\text{logit}|$ histogram) is specified and listed as open ([Ch 147](#)).

Next: disaggregation, correcting a folklore number with primitives.

Disaggregation, Honestly

Claim	Corpus folklore	Measured
Rollout/train disaggregation saving	“3–5×”	25–50% (max ~2× when balanced)
Weight-sync cost (27 GB BF16 over NVLink)	“the bottleneck”	0.019 s (<0.2% of step time)

Table 145.1. Disaggregation, running rollout and training on separate GPUs, corrected with measurement. The headline number folklore repeats is too high by roughly 3×, and the cost everyone worries about is negligible.

Disaggregation splits rollout generation and training onto separate GPUs so they can overlap. The corpus repeats a “3–5×” speedup for it and warns about weight-synchronization latency. Both are wrong in this setup, and this chapter replaces both with measured primitives, exact-by-construction (it is a scheduling change, no distribution effect), so only wall-clock is in question.

145.1 The real saving

Disaggregation saved 25–50% per step in our measurements (a primitive/modeling result, measured from component timings like $T_{\text{rollout}} \approx 11.63$ s and $T_{\text{sync}} \approx 0.019$ s, not a full live veRL-concurrent run. [Chapter 147](#)), depending on the ratio of training time to rollout time, with a maximum near 2× when the two are balanced ($T_{\text{train}} \approx T_{\text{rollout}}$) and savings collapsing below 25% at the extremes (when training dominates rollout). The “3–5×” figure (common folklore and our own original plan note) assumes very long rollouts against light training, a regime where the rollout is the overwhelming cost. In practice an RL training step is comparable in wall-clock to its rollout, so the achievable saving sits in the 25–50% band, not the 200–500% band.

The closest rigorous external number, AReal’s reported 2.77×, is consistent with “max ~2× plus implementation gains” rather than with a routine 3–5×.

145.2 The non-bottleneck

The cost everyone worries about, synchronizing the updated weights from trainer to rollout engine each step, is not a bottleneck. Pushing 27 GB of BF16 weights over NVLink took 0.019 seconds (NVLink bandwidth ~1.4 TB/s), which is under 0.2% of step time. Even on a worse interconnect, a multi-node cluster over PCIe at ~30 GB/s, the same sync would take ~1 second against an ~11-second rollout, still small. So the design lesson is to stop optimizing the weight-sync path and look at the real constraint: the rollout/train imbalance itself. Disaggregation is worth adopting when you have two GPUs to split across roles and the rollout and training times are roughly balanced. The sync cost should not enter the decision.

TAKEAWAY

Disaggregation (rollout/train on separate GPUs) is exact (no Gate A) but the corpus oversells it. Real saving: **25–50% per step** (max ~2× when $T_{\text{train}} \approx T_{\text{rollout}}$. <25% at the extremes), **not “3–5×”**, common folklore / our original plan, which assumes long rollouts on light training, not the typical RL regime (AReal’s reported 2.77× is the closest rigorous external number). And the feared weight-sync cost is negligible: **27 GB over NVLink in 0.019 s (<0.2% of step time)**. Optimize the rollout/train *imbalance*, not the sync.

Next: backends and the batch ceiling, the planning numbers.

Backends and the Batch Ceiling

Context length	Max feasible n (Qwen3-14B, 1× H100, 0.85 util, FA)	KV-pool tokens
1,024	~239	~245,088
2,048	~120	~245,121
8,192	~30	~245,072
32,768	~7	~245,000

Table 146.1. The KV-cache pool is roughly constant at ~245K tokens on one H100, so feasible group size is just that budget divided by context length. This is the lookup table that tells you whether a (group size × context) plan fits before you launch.

This is the planning-numbers chapter, the backend choice and, more usefully, the table that tells you in advance whether your roll-out will even fit. Both are exact concerns (no distribution effect), so they are about wall-clock and memory, not correctness.

146.1 The backend choice

On Hopper (H100, sm_90), FlashAttention beats FlashInfer by ~5–7% for Qwen3-14B, and chunked prefill adds only ~1.5%. The practical guidance is therefore simple: default to FlashAttention, because vLLM’s FlashAttention-3 has been heavily optimized for sm_90, and treat the backend choice as a small (single-digit-percent) tuning knob rather than a major lever. This is worth measuring once on your hardware, the ranking can differ on other accelerators, but on Hopper the answer is FA.

146.2 The batch ceiling

The more valuable result is the batch ceiling. On one H100 (Qwen3-14B, BF16, 0.85 memory utilization, FlashAttention), the KV-cache pool holds a roughly constant ~245,000 tokens regardless of how you slice it, so the maximum feasible group size n is simply that budget divided by the context length: ~239 at 1K context, ~120 at 2K, ~30 at 8K, ~7 at 32K. This is the lookup table that decides whether a planned (group size $\times$ context length) combination fits before you launch and discover it OOMs at step 1. Two notes extend it. First, hybrid architectures bend this ceiling: in a model where only ~25% of layers carry a full KV cache (the rest using fixed-size recurrent state), the same VRAM holds far more tokens. Second, this interacts with the spec-decoding setup of [Chapter 137](#): on Qwen3-Next-80B, the NEXTN draft head plus routing tables needed enough headroom that the rollout required TP=4 (tensor-parallel across 4 GPUs) rather than TP=2, the draft head is not free in memory even when it is nearly free in correctness.

The batch ceiling is a one-line calculator, ask “does this fit?” before launching:

```
def max_concurrent_rollouts(kv_tokens_budget, context_len,
    ↪ safety=0.95):
    return int((kv_tokens_budget * safety) // context_len)

for ctx in [1024, 2048, 4096, 8192, 16384, 32768]:
    print(ctx, max_concurrent_rollouts(245_000, ctx)) # ~239,
    ↪ ~120, ~57, ~30, ~14, ~7
```

The fuller fit check, before any launch, is just an inequality: $\text{weights_gb} + \text{optimizer_gb} + \text{activations_gb} + \text{kv_cache_gb} + \text{sync_buffer_gb} < \text{gpu_memory_gb} * \text{safety}$, note n here is the max concurrent requests at the tested memory-utilization setting, which multi-turn / tool overhead will reduce in practice.

TAKEAWAY

Backend and batch planning (both exact). On Hopper, FlashAttention beats FlashInfer 5–7%, chunked prefill ~1.5%, default to FA, it's a minor knob. The valuable number: the KV-cache pool is constant at ~245K tokens on one H100 (Qwen3-14B, 0.85 util), so max group size = budget ÷ context: ~239 @ 1K → ~120 @ 2K → ~30 @ 8K → ~7 @ 32K, the table that tells you if ($n \times \text{context}$) fits before launch. Hybrid models bend the ceiling (only ~25% of layers carry KV). A draft head can force TP=4 for headroom.

Next: the honest perimeter, what this inference program left open.

CHAPTER 147

Evidence Boundaries for the Rollout Engine

Every honest experimental program has limits. This chapter states what the rollout-engine evidence does and does not establish, first as a one-screen master table of the Part’s findings, then as an explicit map of each headline claim’s boundary.

Part XIV master table (the citable summary):

Acceleration	Class	Gate A (KL)	Gate B (speed)	Verdict
MTP / NEXTN spec-decode	exact	n/a	1.18–1.27×	adopt
EAGLE-3	exact	n/a	lower acceptance	use if no MTP
APC, in-batch group	exact	n/a	~0%	no-op
APC, sequential long prefix	exact	n/a	26% / 1.40×	adopt at the seam
FP8 KV cache	lossy	7.61	0.97×	reject
AWQ rollout (dense-32B)	lossy	6.86	0.95×	reject
MoE AWQ (router flips)	lossy	forensics KL 0.387 / 31% flips	n/a	risky
W8A8	lossy	0.0055	slower load	adopt after preflight
disaggregation	exact	n/a	25–50%	adopt if balanced
FA vs FlashInfer	exact	n/a	FA +5–7%	default FA on Hopper
KV batch ceiling	n/a	n/a	n/a	planning table

Table 147.1. Part XIV master table: every rollout acceleration, its gates, and the verdict.

A reference Part earns trust by marking its own edges. The findings of Chapters 137–146 are real, but several rest on single runs, single models, or proxies, and a few key tests are designed-but-unrun.

Listing them is not a weakness. It is what separates a measured result from an overclaimed one.

147.1 The boundaries

These are the limits of the current evidence, what each headline claim does not yet establish:

Claim	Current evidence	Boundary	What would strengthen it
MTP temperature-robustness	one model, one sweep	single-model	repeat on another MTP model
MTP > EAGLE-3	different base models	directional only	matched-model comparison
APC long-prefix win	synthetic sequential prefix	proxy for agentic reuse	live retrieval-agent run
AWQ rollout unsafe	strong Gate-A fail	dtype confound (fp16-vs-bf16)	matched-dtype rerun
Router flip 31%	one MoE, one task	model-specific	another MoE / task
~40% saturation	two-point coincidence	hypothesis	gap-vs- Δ logit histograms

Table 147.2. Each headline rollout claim, its evidence boundary, and what would strengthen it.

In words, grouped by what each would resolve:

- Single-seed and single-model results. The speculative-decoding sweep (Chapter 137) was effectively single-seed. A multi-seed repeat would put error bars on the 1.27 $\times$ $\rightarrow$ 1.18 $\times$ curve. The MTP-vs-EAGLE-3 comparison (Chapter 138) used different base models, so a matched-model run is needed before the ~2.1 $\times$ ratio is more than directional.
- The proxies that need the real thing. The agentic prefix-caching win (Chapter 139) was measured on a synthetic long-prefix replay. The real agentic test needs the live retrieval stack (E5 + FAISS) running. The disaggregation saving (Chapter 145) was measured against a co-located baseline, not a live verL-concurrent disaggregated run.

- The forensics reruns. The AWQ result (Chapter 141) is confounded by the checkpoint's fp16-vs-bf16 dtype. The dtype-isolation rerun (matched precision) would separate AWQ-specific error from dtype noise. The router-flip (Chapter 142) was measured on one MoE. A TP sweep on the 30B-A3B and a cross-task acceptance measurement would test its generality.
- The theory test. The ~40% saturation hypothesis (Chapter 144) has a specified but unrun experiment, the top-1/top-2 gap histogram versus the $|\Delta\text{logit}|$ perturbation histogram, that would confirm or refute the argmax-flip mechanism.

147.2 What the evidence does establish

Stating the boundaries is what lets a reader trust the rest. Within them, the Part's four genuinely novel findings stand on measured runs: the temperature-robust speculative-decoding curve (1.27 $\times$ →1.18 $\times$), the FP8-KV and AWQ quantitative KL numbers (7.61, 6.86, among the first reported for the RL setting), the MoE router-flip mechanism (31% top-1 under a quantized rollout with the gate unquantized), and W8A8 as a Gate-A-passing rollout recipe (KL 0.0055). The doctrine they support, a serving optimization is not automatically safe for RL. If it changes the rollout distribution, it changes the gradient, does not depend on the open items above. The boundaries mark where confidence ends, not where the contribution does.

TAKEAWAY

The rollout-engine evidence has explicit **boundaries**: results that rest on a single model or sweep (the spec-decode curve), a **directional** cross-model comparison (MTP-vs-EAGLE-3), **proxies** for the real thing (synthetic prefix-reuse. Primitive-timed disaggregation), a **dtype confound** (AWQ fp16-vs-bf16), and an unproven **hypothesis** (~40% saturation). Within those bounds, the four novel findings (temperature-robust spec-decode. FP8-KV/AWQ KL numbers. The router-flip mechanism. W8A8 as

safe) stand on measured runs, and the central doctrine, *if an acceleration changes the rollout distribution, it changes the gradient*, holds regardless. Boundaries mark where confidence ends, not where the contribution does.

Next: *Part XV*. We turn from the rollout engine to the metal underneath it, the hardware this book was built on.

Part XV

The Hardware

Every number in this book was produced on specific machines, and the metal shapes the method: what fits on one card decides which models you can train at all, what 40 versus 80 gigabytes unlocks decides whether a mixture-of-experts model is even reachable, and the dollar-per-hour of the instance decides how many failures you can afford. This Part is the hardware reference, the two machines this book was built on, the memory math that says what fits where, and the tricks that put a large model on a small card.

CHAPTER 148

A100-40GB vs H100-80GB: The Two Machines This Book Was Built On

Evidence boundaries: proved vs hypothesized

Claim	Status	Basis
RL sharpens, does not expand	Measured	pass@1 up, pass@64 flat; multi-seed
Entropy predicts where RL edits	Measured	Spearman 0.95, single batch
FP8-KV cache is unsafe for RL	Measured	per-token KL 7.61, 152x the bar
Full-FT GRPO erases thinking-mode	Verify	0.458 to 0.222, one comparison
NF4 = BF16 at 32B	Measured	4-seed, all 5 benchmarks
AWQ rollout fails via router-flip	Measured	31% top-1 flip, one config
~40% rollout-quant ceiling	Hypothesis	pattern across runs, not isolated

Measured
Verify
Hypothesis

Figure 148.1. Evidence boundaries: every load-bearing claim labeled Measured, Verify, or Hypothesis, with its basis.

The A100-40GB machine ran the early search-agent RL (Qwen3-4B). The H100-80GB machine ran everything since: the 14B and MoE work, SWE-bench, and the forensics (Table 148.1).

Machine	GPUs
Node A	8× A100-40GB
Node B	8× H100-80GB

Table 148.1. The two nodes behind this book. The move from Node A to Node B was not a luxury. It was the 40→80 GB jump (next chapter) that made the 14B-and-larger work possible at all.

This book was built on exactly two machines, and naming them with their real costs is part of the reproducibility contract, because

“we trained a 14B model” means something different on 8×A100-40GB than on 8×H100-80GB. This short chapter records the two and why the work moved from one to the other.

148.1 The two nodes and the move

The A100-40GB machine is an 8× A100-40GB node. The early search-agent RL, the Qwen3-4B work, the v1–v5 entropy saga, ran here over roughly 44 days. The H100-80GB machine is an 8× H100-80GB node, and essentially everything since (the Qwen3-14B SWE-bench saga, the MoE work, the quantization and inference forensics) ran here. The move was not about raw speed. It was about memory. On 40 GB cards, a 14B model in BF16 with optimizer states and a co-located rollout engine is a constant fight. The 80 GB cards turn that fight into headroom, and they add Hopper FP8 (the A100’s Ampere generation has no FP8), which the rest of this Part leans on. The next chapter is the mechanics of exactly what that 40→80 GB jump unlocks.

TAKEAWAY

Two machines: **the A100-40GB machine** (8× A100-40GB node) ran the early Qwen3-4B search-agent RL, **~44 days**, and **the H100-80GB machine** (8× H100-80GB node) ran everything since (14B SWE-bench, MoE, forensics). The move was driven by **memory, not speed**: 80 GB turns the 14B-with-rollout fight into headroom, and Hopper adds FP8 (Ampere has none). “We trained a 14B” means something different on each, which is why the machine belongs in the manifest.

Next: the mechanics of 40 versus 80 gigabytes, and what Hopper’s FP8 adds.

CHAPTER 149

A100-40GB vs H100-80GB

Mechanics

Capability	A100-40GB (Ampere)	H100-80GB (Hopper)
Memory per card	40 GB	80 GB
FP8 (8-bit float)	none	native (sm_90)
Practical effect	14B is a fight; MoE out of reach single-card	14B has headroom; 30B-A3B fits single-card

Table 149.1. The two differences that matter for RL post-training. The 40→80 GB doubling and Hopper’s FP8 are not incremental, together they move whole model classes from “out of reach” to “fits.”

The jump from the A100-40GB to the H100-80GB is two changes that compound: twice the memory and a new numeric format. For RL post-training specifically, those two are what decide which models are reachable on a single card and how cheaply you can run a rollout.

149.1 What 40→80 GB unlocks

The memory doubling is the difference between fighting and fitting. An RL training step holds, simultaneously, the policy weights, the optimizer states (which for full fine-tuning are several times the weight size), the activations, and a co-located rollout engine with its KV cache. On a 40 GB card this is a constant negotiation. You reach for Low-Rank Adaptation (LoRA), aggressive offloading, or a smaller model. On an 80 GB card, a 14B dense model has genuine headroom for full fine-tuning with a co-located rollout, and, the qualitative jump, a Qwen3-30B-A3B mixture-of-experts model becomes single-card resident for inference/rollout (its BF16 weight footprint is ~62 GB, which simply does not fit in 40 GB but does in 80). Resident is the precise word: full fine-tuning of the MoE, with optimizer states

and a co-located rollout, still needs sharding or parameter-efficient methods, what 80 GB unlocks is that the weights fit at all, not that you can full-FT it on one card. The next chapter does this memory math model by model. The point here is that 40→80 GB is not “20% more room,” it is the line between which model classes are single-card reachable.

149.2 What Hopper’s FP8 adds

The second change is FP8 (8-bit floating point), native to Hopper (sm_90) and absent on Ampere. FP8 matters in two of this book’s three quantization roles (Chapter 132). For the rollout/serving role it is a real throughput lever, at FP8, Qwen3-14B reaches ~97 tokens/second and the 30B-A3B MoE ~155 on a single H100 (Chapter 150), and notably FP8 helps the dense 14B’s rollout throughput more than the MoE’s, because the MoE’s gating already amortizes well at BF16 (this is a rollout-throughput statement, not a training-quality one). For the training-actor role, FP8 weights showed no measurable quality regression (Chapter 132). None of this is available on the A100 at all, which is a second, independent reason the work moved to the H100-80GB machine: not just the memory, but access to the format that the quantization and inference chapters depend on.

TAKEAWAY

Two compounding changes from A100-40GB to H100-80GB. **Memory 40→80 GB** is the line between fighting and fitting: a 14B gains real headroom for full-FT-with-rollout, and a **30B-A3B MoE (~62 GB BF16) becomes single-card reachable** (impossible in 40 GB). **Hopper FP8 (sm_90, absent on Ampere)** is a real rollout/serving throughput lever (14B ~97 tok/s, 30B-A3B ~155. FP8 helps the *dense* model more) and a benign training-actor option (Ch 132). Both are reasons the book moved to the H100-80GB machine.

Next: the memory math, what actually fits on one GPU.

CHAPTER 150

What Fits *Resident* on One GPU

Model	Precision	Footprint (1× H100-80GB)	Rollout tok/s
Qwen3-14B (dense)	BF16	~28 GB	—
Qwen3-30B-A3B (MoE)	BF16	~62 GB	~145 (3B active)
Qwen3-30B-A3B	FP8	~34 GB	~155
Qwen3-30B-A3B	GPTQ-INT4	~18 GB	~31 (broken on SGLang — avoid)

Table 150.1. The memory math that decides what is single-card trainable, with measured rollout throughput. Note the MoE’s high throughput (only 3B parameters are active per token) and the GPTQ-INT4 warning, small footprint, broken kernels.

Before you choose a model, you need to know whether it fits, and “fits” depends on precision, on whether you are doing full fine-tuning or LoRA, and on leaving room for the rollout engine. This chapter is the memory math on a single H100-80GB, the card most readers will have one of, and it is about resident weights and rollout footprint, not full-fine-tuning trainability (the full-FT multiplier comes at the end).

150.1 The footprints

The weight footprints set the floor. A Qwen3-14B dense model is ~28 GB in BF16. A Qwen3-30B-A3B MoE is ~62 GB in BF16 (the full expert set must be resident even though only ~3B parameters are active per token), dropping to ~34 GB at FP8 and ~18 GB at GPTQ-INT4. Two practical readings follow. First, the MoE’s rollout throughput is high for its size, ~145 tokens/second at BF16, ~155 at FP8, precisely

because only 3B parameters are active per token, so a 30B MoE generates faster than a 32B dense model. Second, the GPTQ-INT4 path, despite its tempting ~18 GB footprint, is broken on SGLang for this MoE (~31 tokens/second, far below the FP8 path), so the small footprint is a trap. A role caution that reconciles with [Chapters 141](#) and [160](#): treat AWQ/GPTQ-INT4 as serving formats (after a serving-quality check), not as RL rollout engines, AWQ-INT4 rollout fails Gate A badly (KL 6.86, [Chapter 141](#)). For RL rollout, W8A8 or BF16/FP8-weight paths are the safer choices than 4-bit. For serving, FP8 or (preflighted) AWQ are fine.

150.2 The full-FT multiplier

The footprints above are weights. full fine-tuning multiplies them. The optimizer states for full fine-tuning (the Adam moments) are several times the weight size, so a model whose weights are 28 GB needs substantially more than 28 GB to train with full fine-tuning, which is why even a 1.7B model full-FT can sit around 50 GB once optimizer states, gradients, activations, and a rollout engine are all resident. This is the gap between “the weights fit” and “the run fits,” and it is the reason the next chapter exists: when full fine-tuning does not fit, the way to keep a large model on a single card is to not train all of it in full precision, LoRA, quantized bases, and memory-sharing tricks. The rule for this chapter: read the footprint, then multiply for full fine-tuning and add the rollout engine before you decide a model is single-card trainable.

TAKEAWAY

Single-H100 *resident* memory math: Qwen3-14B ~**28 GB BF16**. Qwen3-30B-A3B MoE ~**62 GB BF16** → ~**34 GB FP8** → ~**18 GB INT4**, with high rollout throughput (~**145–155 tok/s**, since only ~3B params are active). Watch the **GPTQ-INT4 trap**: small footprint but **broken on SGLang** (~**31 tok/s**). And by role, **AWQ/GPTQ are serving formats, not RL rollout engines** (AWQ rollout fails Gate A, [Ch 141](#)), use W8A8 or FP8-weight/BF16 for rollout. These are *weight* footprints, **full fine-tuning multiplies them**

(optimizer states are several× the weights, so even a 1.7B full-FT run ~50 GB). “Weights fit” ≠ “run fits.”

Next: how to put a large model on a small card when full fine-tuning won't fit.

CHAPTER 151

Fitting Big Models on One Small Card

Method	What it does	Result
QLoRA (4-bit base + LoRA)	quantize the frozen base to 4-bit, train a small adapter	30B-A3B resident in ~17.5–18 GB
QeRL (NVFP4 base + LoRA)	4-bit-float base, BF16 adapter	32B GRPO on one H100 (external, not run — NVFP4 kernel is Blackwell/sm_100-only; on H100 it dequantizes to BF16; NF4 was the Hopper proxy)
Unsloth Standby	free the rollout engine's weights during the training step	push <code>gpu_memory_utilization</code> toward 0.95 (e.g. 32B LoRA context 3.6K→6.1K, ~1.7×)

Table 151.1. Three ways to keep a large model on a single card. QLoRA and Unsloth Standby are measured here. QeRL/NVFP4 was never run (Blackwell-only kernel), so the book's 32B-on-one-card evidence comes from the NF4 proxy, which reached BF16 parity ([Chapter 132](#)).

When the full-fine-tuning multiplier of the last chapter pushes a model off a single card, the answer is to stop training all of it in full precision. This chapter is the three tricks that keep a large model single-card trainable, two measured here, one carried honestly as external.

151.1 QLoRA and the memory-sharing trick

QLoRA quantizes the frozen base model to 4-bit (NF4) and trains only a small Low-Rank Adaptation (LoRA) adapter on top, which collapses the footprint dramatically: a Qwen3-30B-A3B base sits resident in

roughly 17.5–18 GB, well within a single 80 GB card with abundant room for the rollout. One operational caveat from [Chapter 132](#)’s role discipline: the 4-bit conversion for this MoE is done on the fly from the 16-bit weights, so the load transiently needs ~80 GB of disk and RAM even though the resident model is ~18 GB, budget for the load, not just the steady state. The second trick is Unsloth Standby (`UNSLOTH_VLLM_STANDBY=1`), the fix for the co-location OOM cliff (the P16 pitfall of [Chapter 122](#)): it frees the vLLM rollout engine’s weights during the training step, when they are idle, which lets you push `gpu_memory_utilization` toward 0.95 instead of stalling at 0.6. The measured effect is real headroom, a Qwen3-32B LoRA run’s context grew 3.6K → 6.1K tokens (~1.7×), and a Llama-3.1-8B QLoRA run’s from 42K → 47.5K.

151.2 QeRL/NVFP4, the external option, carried honestly

The third trick is the most enticing and the one this book could not run on its hardware, for a concrete, mechanical reason. QeRL uses an NVFP4 (NVIDIA 4-bit float) base with a BF16 LoRA adapter, and the QeRL paper (October 2025) reports fitting a 32B GRPO run on a single H100 at accuracy competitive with BF16 LoRA, with the intriguing claim that the quantization noise acts as a free exploration bonus. The blocker is hardware: vLLM’s NVFP4 kernel (`modelopt_fp4 / petit_nvfp4`) is compiled for `sm_100`, Blackwell (B200) only. On an H100 (`sm_90`) it dequantizes NVFP4 to BF16 at load, turning “4-bit in ~5 GB” back into “BF16 in ~64 GB” and evaporating the entire memory thesis. So this work substituted NF4 (the bitsandbytes format, which has working Hopper kernels) as the proxy, the experiment labeled “Q1” is an NF4-LoRA run, not NVFP4, and that proxy reached statistical parity with BF16 across all five benchmarks at 32B ([Chapter 132](#)), which upholds and extends the QeRL family’s parity claim, just via a different 4-bit format. True NVFP4 and the exploration-bonus claim remain external, open future work gated on Blackwell access. The honest summary: QLoRA, Unsloth Standby, and the NF4 proxy are tools backed by measured evidence here. NVFP4 specifically is unvalidated on Hopper, validate it on Blackwell, don’t cite it as run.

TAKEAWAY

Three ways to keep a large model single-card. **QLoRA** (4-bit frozen base + LoRA adapter) puts a 30B-A3B resident in **~17.5–18 GB** (but the on-the-fly 4-bit conversion needs ~80 GB transient at *load*). **Unsloth Standby** frees the rollout engine’s weights during the training step → push `gpu_memory_utilization` toward **0.95** (32B LoRA context 3.6K→6.1K, ~1.7×. The P16 fix). **QeRL/NVFP4** was **never run** here, its kernel is Blackwell/sm_100-only and **dequantizes to BF16 on H100**, defeating the memory thesis. The book’s 32B-on-one-card evidence is the **NF4 proxy** (“Q1”), which hit BF16 parity across all 5 benchmarks ([Ch 132](#)). Validate true NVFP4 on Blackwell. Don’t cite it as run.

Next: single-GPU gotchas, the bnb load costs, silent 16-bit loads, and compiled-cache mismatches that bite on one card.

CHAPTER 152

Single-GPU Gotchas

Trap	Symptom	Fix
bnb 4-bit load cost	a 30B-A3B “4-bit” model needs ~80 GB transient at load	budget RAM+disk for the on-the-fly conversion, not just the ~18 GB resident
silent 16-bit QLoRA load	QLoRA + vLLM <code>fast_inference=True</code> runs in 16-bit anyway	use Unsloth dynamic-4bit quants (the F16/P17 trap)
compiled-cache mismatch	a recompile (or wrong output) after changing batch shape / seq length	clear the <code>torch.compile</code> cache on shape changes, or <code>enforce_eager</code>
<code>enforce_eager</code> with <code>TP≥2</code>	CUDA-graph capture issues on tensor-parallel rollout	set <code>enforce_eager=True</code> for multi-GPU rollout when graphs misbehave

Table 152.1. The traps that bite specifically on one card, where you have no second GPU to absorb a mistake. Spell-outs: bnb = BitsAndBytes. QLoRA = Quantized Low-Rank Adaptation. TP = tensor parallelism.

Single-GPU work has its own failure surface, the mistakes that a multi-GPU setup hides but one card exposes, mostly around quantized loading and compilation. This chapter is that surface, drawn from the pitfalls of the single-GPU tier.

152.1 The quantized-load traps

The two most common single-GPU surprises both involve quantization at load time. First, a “4-bit” model is not 4-bit while loading: BitsAndBytes converts the full 16-bit weights to 4-bit on the fly, so a Qwen3-30B-A3B that is ~18 GB resident transiently needs about 80 GB of disk and RAM during the conversion ([Chapter 151](#)), plan for the load, or the process is killed before training starts. Second, and

more insidious because it is silent, is the F16/P17 trap ([Chapter 122](#)): QLoRA with vLLM's `fast_inference=True` can load the model in 16-bit anyway when vLLM cannot consume the BitsAndBytes-4bit blob (missing `absmax` keys), so you think you are running 4-bit and you are not, use Unsloth's dynamic-4bit quants, or a vLLM build with proper bnb-4bit support, and verify the actual resident footprint.

152.2 Compilation and graph traps

The other family is compilation. `torch.compile` caches compiled kernels keyed on tensor shapes, so a change in batch size or sequence length can trigger a silent recompile (a sudden slow step) or, worse, a stale-cache mismatch, clear the compile cache when you change shapes, and if compilation is fighting you, fall back to eager. Relatedly, on tensor-parallel rollout ($TP \geq 2$), CUDA-graph capture can misbehave. The standard escape is `enforce_eager=True` for the rollout engine, trading a little throughput for a setup that actually runs. None of these is a deep problem, but all of them waste a single-GPU session that has no redundancy to hide behind, which is why the single-GPU tier rewards checking the resident footprint and the execution mode before a long run.

TAKEAWAY

Single-GPU traps cluster around quantized loading and compilation. A “4-bit” 30B-A3B needs **~80 GB transient at load** (bnb converts on the fly) though it's ~18 GB resident. QLoRA + vLLM `fast_inference` can **silently load 16-bit** (F16/P17), use Unsloth dynamic-4bit and verify the footprint. `torch.compile` caches on shape, so changing batch/seq can recompile or mismatch (clear the cache or go eager). And `enforce_eager=True` is the escape when CUDA graphs misbehave on $TP \geq 2$ rollout.

Next: when one card is no longer enough, going multi-GPU.

CHAPTER 153

When to Go Multi-GPU, and How to Shard

Strategy	What it shards	When to use
DDP (Distributed Data Parallel)	nothing — replicates the full model per GPU	model + optimizer + activations already fit on one card and you just want more throughput
ZeRO-3 / FSDP2	parameters, gradients, and optimizer states	the model + full-fine-tuning optimizer states do not fit on one card

Table 153.1. The sharding decision in one line: replicate (DDP) when it fits, shard everything (FSDP2 / ZeRO-3) when it doesn't. Spell-out: FSDP2 = Fully-Sharded Data Parallel, version 2.

The trigger for going multi-GPU is not “I have more GPUs”. It is “the run does not fit on one.” This chapter is the decision of when to shard and how, because the two main strategies solve different problems and choosing the wrong one wastes either memory or bandwidth.

153.1 The decision

The question that decides everything is whether the training footprint, weights plus the full-fine-tuning optimizer states (the Adam moments, several times the weight size) plus activations plus a rollout engine, fits on a single card ([Chapter 150's](#) full-FT multiplier). If it does, use Distributed Data Parallel (DDP): each GPU holds a full replica and processes a different slice of the batch, which maximizes throughput with no sharding overhead. If it does not, the common case for full fine-tuning of a 14B-and-up model, use Fully-Sharded Data Parallel v2 (FSDP2) or the equivalent ZeRO-3, which shards

the parameters, gradients, and optimizer states across GPUs so that no single card holds the whole thing. The cost of sharding is communication (gathering shards for each forward/backward), so you shard only as much as you must.

153.2 The strategies, and the caveat

FSDP2 offers a spectrum: full shard (everything split across all GPUs, maximum memory savings, maximum communication) down to hybrid shard (shard within a node, replicate across nodes, less communication when you have fast intra-node links). The practical default for a single 8-GPU node is full shard with the fast NVLink fabric absorbing the gather cost. One hard-won caveat rides along, from the failure catalog ([Chapter 120](#)'s F8): a specific Torch-2.6 FSDP2 build leaks VRAM, peak memory climbs monotonically until an out-of-memory crash around step 200–500, so pin Torch 2.7 and $\text{veRL} \geq 0.5$ for FSDP2 work. The rule for this chapter: don't shard until the run won't fit, then shard the whole optimizer state with FSDP2 on a known-good Torch version, and reach for the larger parallelism strategies of [Chapter 157](#) only when even sharded data-parallel isn't enough.

TAKEAWAY

Go multi-GPU when the run *won't fit*, not when you have spare cards. If weights + full-FT optimizer states + activations + roll-out fit on one card → **DDP** (replicate, max throughput). If not (the usual case for 14B+ full-FT) → **FSDP2 / ZeRO-3** (shard parameters, gradients, *and* optimizer states), defaulting to full-shard on a single NVLink node. Caveat: the Torch-2.6 FSDP2 build leaks VRAM (F8), pin Torch 2.7 / $\text{veRL} \geq 0.5$.

Next: splitting GPUs not by data but by role.

CHAPTER 154

Splitting GPUs Across Roles

The on-policy-distillation layout. On 4 GPUs, run 2 for the teacher (vLLM) + 2 for the student (FSDP), the verified OPD layout. At the flagship tier it scales to 4 + 4. The alternative when you can't split is time-sharing: free the rollout engine's GPUs during the training step.

Multi-GPU is not only about sharding one model across cards (Chapter 153). Sometimes the right move is to give different cards different jobs. This is the role-splitting chapter, most important for on-policy distillation (OPD) and any setup with a separate generator and learner.

154.1 The role-split layout

On-policy distillation runs two models at once, a frozen teacher that scores the student's rollouts and the student being trained, and the clean layout puts them on separate GPUs. The verified small-OPD layout on 4× H100 is 2 teacher-vLLM + 2 student-FSDP: two cards serve the teacher through vLLM (inference-only, so they want throughput), and two cards train the student under FSDP (so they want the optimizer-state sharding of Chapter 153). At the flagship tier this scales to a 4 + 4 split. The principle generalizes beyond OPD: whenever generation and training have different resource profiles, dedicating cards to each role avoids the contention of forcing both through the same GPUs, the inference side is latency/throughput-bound, the training side is memory-bound, and they fight if co-located.

154.2 Time-sharing, when you can't split

When you don't have enough cards to dedicate (the single-node, want-it-all case), the alternative is time-sharing the same GPUs between roles within a step: generate rollouts, then free the rollout engine's memory and use those GPUs for the training update. This is exactly what vLLM sleep mode and Unsloth Standby do (Chapters 122, 151), the rollout engine's weights are released during the optimizer step, when they are idle, letting the trainer use the freed memory. The trade-off between the two approaches is straightforward: role-splitting avoids the repeated free/reload cost but needs more cards. time-sharing fits on fewer cards but pays a synchronization cost each step (and risks the sleep-mode resync bug, Chapter 155). Choose role-splitting when you have the GPUs and the two roles are both heavy. Choose time-sharing when cards are scarce.

TAKEAWAY

Split GPUs by *role* when generation and training have different profiles. The verified OPD layout: **2 teacher-vLLM + 2 student-FSDP** on 4× H100 (scaling to **4 + 4** at the flagship tier), inference is throughput-bound, training is memory-bound, and they contend if co-located. When cards are too scarce to dedicate, **time-share**: free the rollout engine's weights during the optimizer step (vLLM sleep / Unsloth Standby), trading a per-step sync cost for fitting on fewer GPUs.

Next: the mechanics of co-locating the rollout engine with the trainer.

CHAPTER 155

vLLM Co-location Mechanics

Mechanism	What it does	The trap
Sleep level 1	offload rollout weights to host RAM during training	safe; the recommended default
Sleep level 2	free rollout weights/KV entirely	resync corruption — gibberish after the first weight reload (P1)
<code>gpu_memory_utilization</code>	how much VRAM vLLM reserves	the 0.6-vs-0.8 OOM cliff: OOM at 0.8 vanishes at 0.6 (P16)

Table 155.1. Co-locating the rollout engine on the training GPUs is the default single-node setup, and these are the three knobs that make or break it.

Co-location, running the vLLM rollout engine on the same GPUs as the trainer (the opposite of the disaggregation of [Chapter 145](#)), is the default for a single 8-GPU node, and it has a handful of mechanics that, gotten wrong, silently break training. This chapter is those mechanics, all drawn from the framework-trap catalog ([Chapter 122](#)).

155.1 Sleep modes and the resync bug

The core mechanism is sleep mode: because the rollout engine and the trainer can't both hold full weights at once, vLLM “sleeps” between rollouts to free room for the optimizer step. There are two levels, and the difference matters. Sleep level 1 offloads the rollout weights to host RAM and is safe, the recommended default. Sleep level 2 frees the weights and KV cache entirely, which saves more memory but carries the P1 resync-corruption bug: when level-2 frees weights while the sharding manager is still copying new ones in, vLLM produces gibberish after the first weight reload (the classic “my rollouts turned to garbage at step 1” report). The fix is to pin a vLLM build with the in-place weight loader, or simply use sleep level

1. This is also one of the rollout-lossy family’s operational cousins ([Chapter 121](#)): a co-location bug that corrupts the generation, not via quantization but via a botched weight sync.

155.2 The memory-utilization cliff

The second mechanic is the `gpu_memory_utilization` cliff (the P16 pitfall). Co-located training allocates and frees memory in the same step that vLLM is profiling its KV-cache headroom against, so vLLM’s static-residence assumption is wrong, and you get the maddening symptom of an out-of-memory crash at `gpu_memory_utilization=0.8` that simply disappears at 0.6, not because 0.6 is “more correct” but because the lower reservation leaves slack for the trainer’s transient allocations. The better fix than backing off to 0.6 (and leaving throughput on the table) is Unsloth Standby (`UNSLOTH_VLLM_STANDBY=1`), which frees vLLM’s weights during the training step and lets you push utilization back toward 0.95 ([Chapter 151](#)). The rule: on a co-located node, use sleep level 1, and if you hit the OOM cliff, reach for Standby before you sacrifice utilization.

TAKEAWAY

Co-locating vLLM with the trainer (single-node default) turns on three knobs. **Sleep level 1** (offload weights to host RAM) is the safe default. **sleep level 2** (free weights/KV) risks the **P1 resync-corruption bug** (gibberish after the first reload, pin the in-place weight loader). And the **gpu_memory_utilization 0.6-vs-0.8 OOM cliff** (P16) comes from co-located alloc/free breaking vLLM’s static-residence assumption, fix it with **Unsloth Standby → 0.95**, not by backing off to 0.6.

Next: what the full 8-GPU flagship tier actually runs.

CHAPTER 156

The Flagship Tier

Workload	Layout	Key requirement
Qwen3-14B full fine-tune	FSDP2 full-shard, co-located rollout	Optimizer-state sharding (Ch 153)
Qwen3-30B-A3B MoE	Megatron + expert parallelism, GSPO (safer default)	EP + sequence-level optimizer (Ch 117)
On-policy distillation	4 teacher + 4 student	Role split (Ch 154)
C9 MoE family	8-GPU runs (GSPO/ FREEZE/ JITTER/ COMBO)	C9-GSPO entropy ~0.081–0.112 (one window), read jointly with eval (Ch 126)

Table 156.1. What the 8-GPU tier unlocks: full fine-tuning of a 14B, real MoE RL, and the larger OPD splits.

The 8-GPU node is the flagship tier, where the runs that produce the book’s headline results live. This chapter is what that tier is for: the workloads that need all eight cards and the layout each one takes.

156.1 What runs at eight GPUs

Three workload classes define the tier. A Qwen3-14B full fine-tune (not LoRA) fits here via FSDP2 full-shard with a co-located rollout engine, the optimizer states that overflow a single card (Chapter 150) shard cleanly across eight. A Qwen3-30B-A3B mixture-of-experts RL run typically needs more than data-parallel sharding: Megatron-LM with expert parallelism, and GSPO (Group Sequence Policy Optimization) as the safer default, not an absolute requirement, since the corpus shows vanilla GRPO can stay stable on a large MoE under conservative conditions (Chapter 117’s scale-dependence), but the default to reach for when routing drift or ratio tails appear. The book’s C9 program is this MoE work, and it is a family, C9-GSPO, C9-FREEZE, C9-JITTER, C9-COMBO, not a single run. C9-GSPO held routing entropy

in a roughly 0.081–0.112 band in one observed window (not a universal constant), where vanilla GRPO collapsed it. And C9-COMBO is the reminder (Chapter 126) that entropy must be read jointly with held-out eval and reward variance. It dipped below earlier entropy floors while validation improved, which is productive sharpening, not collapse. And on-policy distillation at this tier takes the 4 teacher + 4 student split of Chapter 154.

156.2 Why this tier is the proving ground

The flagship tier matters because it is where the architecture-to-algorithm couplings of the whole book become load-bearing rather than theoretical. At one GPU you are mostly limited to LoRA and small models, where the algorithm choice is forgiving. At eight GPUs you can full-fine-tune a 14B and train a 30B MoE, where the wrong choice (token-level GRPO on the MoE, an unsharded optimizer, a co-location resync bug) is the difference between a result and a crash. So the flagship tier is not just “more compute”. It is the regime where Part V’s GSPO mandate, Part XIII’s sharding discipline, and this Part’s co-location mechanics all have to be right at once. The runs are more expensive and less forgiving, which is exactly why the manifest and the promotion gates (Chapters 130–131) earn their keep here.

TAKEAWAY

The 8-GPU flagship tier runs the book’s headline workloads: **14B full fine-tune** (FSDP2 full-shard + co-located rollout), **30B-A3B MoE** (Megatron + expert parallelism + **GSPO**, since token-level GRPO explodes, F5/Ch 117), and **OPD at 4 + 4** (Ch 154). The **C9 MoE family** lives here (GSPO/FREEZE/JITTER/COMBO). C9-GSPO held entropy in a ~0.081–0.112 window (one window, not universal), and C9-COMBO showed entropy must be read *jointly* with validation (Ch 126). GSPO is the **safer default** for MoE RL at scale, not an absolute requirement. This tier is where the book’s architecture↔algorithm couplings become load-bearing, the

regime where a wrong choice crashes rather than merely underperforms.

Next: the parallelism strategies that make the biggest models trainable.

Parallelism for Big Models

Parallelism	Splits across GPUs	Used for
Data (FSDP2 / ZeRO-3)	the batch + sharded parameters/optimizer	the default; dense models that fit when sharded
Tensor (TP)	individual matrix multiplications	a single layer too big for one card; rollout headroom
Pipeline (PP)	layers across GPUs (staged)	very deep models
Expert (EP)	MoE experts across GPUs	mixture-of-experts (e.g. EP=8)

Table 157.1. The four axes of parallelism. Big-model training combines them, and a draft head can force more tensor parallelism than the base model alone needs. Spell-outs: TP = tensor parallel. PP = pipeline parallel. EP = expert parallel.

When even sharded data-parallelism ([Chapter 153](#)) isn't enough, you combine axes of parallelism, and for the largest models, the combination is the whole game. This chapter is the four axes and how they stack, including a surprising rollout-side constraint.

157.1 FSDP2 versus Megatron, and the four axes

There are four ways to split a model across GPUs. Data parallelism (FSDP2 / ZeRO-3) splits the batch and shards parameters/optimizer states, the default, and sufficient for dense models that fit once sharded. Tensor parallelism (TP) splits individual matrix multiplications across cards, for when a single layer is too big for one GPU. Pipeline parallelism (PP) stages layers across GPUs. Expert parallelism (EP) distributes a mixture-of-experts model's experts across cards, for a 30B-A3B MoE, EP=8 spreads the expert set across the node. The practical fork is FSDP2 versus Megatron-LM: FSDP2 does data-parallel sharding well and is the right default for dense models, but a large MoE needs Megatron because it implements

tensor/pipeline/expert parallelism together (Chapter 156's MoE requirement). The art is using the least parallelism that fits, since each axis adds communication.

157.2 The draft-head constraint

A non-obvious finding from the inference program (Chapter 146) is that speculative decoding can raise the parallelism floor. On Qwen3-Next-80B with NEXTN (Multi-Token Prediction) draft heads, the rollout needed TP=4 on 80 GB cards rather than TP=2, not because the base model demanded it, but because the NEXTN draft head plus the routing tables plus warmup consumed enough headroom that two-way tensor parallelism left too little for the KV pool. The general lesson: the rollout engine's parallelism is set by more than the base model's weights, a draft head, a MoE's routing tables, and the KV-cache budget all add to the footprint, and a hybrid architecture (Chapter 159) changes the math again. So when planning the rollout side of a big-model run, size the parallelism against the full rollout footprint, draft head included, not just the base weights.

TAKEAWAY

Beyond sharded data-parallelism, combine axes: **data (FSDP2/ZeRO-3)**, **tensor (TP)**, **pipeline (PP)**, **expert (EP)**, using the least that fits, since each adds communication. The fork: FSDP2 for dense, **Megatron** for a large MoE (it does TP/PP/EP together. E.g. **EP=8**). A non-obvious constraint: a **draft head can raise the floor**, Qwen3-Next-80B + NEXTN needed **TP=4**, **not TP=2**, because the draft head + routing tables + KV pool exceeded two-way headroom. Size rollout parallelism against the *full* footprint, not just base weights.

Next: the one combination that didn't work, FP8 plus MoE training.

The FP8 + MoE Training Wall

An implementation wall, not a numerics theorem. Inference FP8 on a dense model is a free win, ~zero accuracy loss and ~2× throughput on Qwen3-14B. FP8 for a mixture-of-experts model hit a block-size and memory wall in the co-located veRL/vLLM stack: the expert dimension is not divisible by the FP8 quantization block at the tensor-parallel sizes that fit, and the sizes where it is divisible run out of memory. This is not evidence that FP8 destabilizes MoE training. It is evidence that a valid FP8 MoE rollout path could not be realized on this stack.

This chapter is the negative result that clarifies the whole FP8 story: the difference between quantizing the model you generate with and the model you train. Conflating those two is the single most common FP8 misconception, and the MoE case makes the distinction sharp.

158.1 Inference FP8: the free win (on dense)

On the inference/rollout side, FP8 is exactly the win the marketing promises, but the book measured it rather than assuming it. In the Q4 paired-arm experiment on Qwen3-14B (dense), FP8 rollout cost effectively zero accuracy (a +0.38-percentage-point GSM8K delta, pure shot noise at 1319 examples) and doubled vLLM throughput per GPU. Since generation is 60–80% of step time in RL, that is a roughly 2× free speedup with no quality regression at this scale. The mechanism (Chapter 132) is benign: FP8 injects noise only at the rollout-time logit cast, slightly flattening the sampling distribution (acting like a mild temperature) while leaving the learned representations, which heads attend where, which neurons fire, the residual-stream norm, identical within numerical precision.

158.2 Training FP8 on MoE: the wall

The MoE side is the wall, and the honest cause is implementation, not a proven routing-noise instability. The C9-GRPO-FP8 experiment (FP8 for the vLLM rollout, training in BF16) hit a sequence of hard constraints on the 8×H100 veRL/vLLM stack:

Attempt	What failed	Root cause
FP8 rollout, TP=8	block-quant dimension mismatch	Qwen3-30B-A3B expert moe_intermediate_size=768, TP=8 → 768/8 = 96, not divisible by FP8 block_n=128 (96 < 128)
FP8 rollout, TP=2	OOM during weight sync	FP8 TP=2 weights ≈ 15 GB/GPU (2× the TP=8-BF16 per-device cost) + KV cache + BF16 sync buffer exceeded memory
TP=4 (not tried)	would fail the divisibility check	768/4 = 192, 192 ÷ 128 = 1.5 (not integer); only TP=1, 2, 3 are valid for FP8 here

Table 158.1. The C9-GRPO-FP8 attempts on the 8×H100 veRL/vLLM stack — attempt, config, what failed, root cause.

The novel finding underneath is a real constraint nobody documents for small-expert MoEs: FP8 block quantization requires `intermediate_size / TP` to be divisible by 128, which dense models (`intermediate ≥ 4096`) satisfy at any reasonable TP but a 768-wide expert does not at TP=8. So the wall is block-size divisibility plus tensor-parallel memory, not FP8 noise tipping the router. The practical resolution the book adopted: train the MoE in BF16 and roll out in BF16 (or use FP8 only at a TP that is both divisible and fits), capturing the dense FP8 win where it is clean and not claiming a numerics result the runs do not support. The general statement still holds, inference-FP8 and training-FP8 are different operations (Chapter 132’s role table), but the MoE wall here is an engineering constraint, and it should be stated as one.

TAKEAWAY

The FP8 + MoE wall is an **implementation/memory wall, not a numerics theorem**. Dense FP8 *rollout* is a free win (Q4: ~0 accuracy loss, ~2× throughput on Qwen3-14B). The MoE FP8 *rollout*

failed on the co-located veRL/vLLM stack because the **expert dimension (768) isn't divisible by FP8's block_n=128 at TP=8 (768/8=96)**, only TP=1/2/3 are valid, and **TP=2 then OOMs** (~15 GB/GPU + KV + BF16 sync buffer). Do **not** claim FP8 noise destabilized MoE routing. That arm was never validly run. Resolution: train and roll out the MoE in **BF16**, or use FP8 only at a TP that is both divisible *and* fits.

Next: the memory math of the hybrid architecture that the newest models use.

CHAPTER 159

Memory Math for the Hybrid Architecture




Figure 159.1. Only a quarter of the layers scale with context: 75% linear-attention layers carry a fixed-size state, only the 25% full-attention layers grow a KV cache.

The newest models (Qwen3.5-class) are hybrid: most layers use linear attention with a fixed-size state, and only a minority use full softmax attention with a growing key-value (KV) cache. That changes the memory math fundamentally, and this chapter is the new math. It is why a hybrid holds far more concurrent sequences than a dense model of the same size.

159.1 Fixed state versus growing cache

A hybrid model splits its layers in a 3:1 ratio: about 75% are linear-attention (Gated DeltaNet) layers, and 25% are full-attention layers. The crucial difference is what each stores per token. A full-attention layer keeps a KV cache that grows linearly with sequence length, the cost that sets the batch ceiling ([Chapter 146](#)). A linear-attention

layer instead keeps a fixed-size recurrent state (a matrix of about 128×128 per head) that is constant regardless of sequence length, that is the property that gives linear-attention its $O(1)$ -per-token inference cost. So in a hybrid, only one layer in four contributes to the context-scaling memory cost. The other three contribute a fixed amount no matter how long the sequence. The trade-off the architecture accepts: the fixed state has an effective rank (~ 128 associations per head), a compression bottleneck at extreme lengths, which is exactly why the design keeps 25% full-attention layers. They provide the unbounded retrieval the recurrent state cannot.

159.2 Dual pools, and what it means for RL

Serving a hybrid requires managing two kinds of memory, the growing KV cache for the full-attention layers and the fixed recurrent state for the linear layers, and the frameworks do this differently. vLLM uses a unified scheme: it tunes “logical” block sizes so a full-attention KV block and a linear-attention state block occupy equal physical memory, enabling one paged manager (the recurrent state is kept in float32 for numerical stability). SGLang uses separate pools: a Mamba pool for the recurrent states (allocated per request) and a KV pool for the full-attention layers (allocated per token). For RL the consequence is concrete and favorable: because only 25% of layers carry a length-scaling cache, the batch ceiling of [Chapter 146](#) rises sharply for a hybrid, the same VRAM holds many more concurrent rollouts at long context than a dense model of equal parameter count, which directly lowers rollout cost on long-horizon agentic tasks.

TAKEAWAY

Hybrid models (Qwen3.5-class) split layers **3:1**, **~75% linear-attention** (a **fixed-size recurrent state**, $\sim 128 \times 128$ /head, constant in sequence length, $O(1)$ /token) and **~25% full-attention** (the **only KV cache that grows with context**). So just one layer in four drives context-scaling memory. Serving needs two memory kinds, vLLM **unifies** them (equal-sized logical blocks, recurrent state in float32). SGLang uses **separate** Mamba + KV

pools. For RL this **raises the batch ceiling** (Ch 146): far more concurrent long-context rollouts per GB than a dense model of equal size.

Next: the quantization spectrum, finally laid out by purpose.

CHAPTER 160

The Quantization Spectrum, by Purpose

Format	Actor-training	Rollout-sampling	OPD teacher	Serving
BF16	baseline	baseline	baseline	heavy but faithful
FP8 weights	dense: no measurable regression (Ch 132)	dense: free ~2× throughput (Ch 158)	possible	good
FP8 KV cache	n/a	✗ reject — KL 7.61 (Ch 140)	n/a	serving-only if quality OK
W8A8 (INT8)	not established	passes Gate A — KL 0.0055 (Ch 143)	possible	
AWQ-INT4	not actor-training	✗ reject as BF16-paired rollout — KL 6.86 (Ch 141)	verify teacher-logit drift (F17)	possible, after serving check
NF4 (QLoRA)	scale-dependent: 32B yes, 14B no (Ch 132)	not rollout-safe by default	no	no
NVFP4 (QeRL)	external / not run (Ch 132)	n/a	no	no

Table 160.1. The whole quantization story in one table, four columns, actor-training / rollout-sampling / OPD-teacher / serving, because the same format is safe in one column and poison in another. The decisive split: FP8 weights and FP8 KV cache are different operations with opposite RL verdicts (free win vs reject), so they get separate rows.

This chapter does one thing: it puts every quantization format into the three-column table that the rest of the book has been building toward, because the single most expensive quantization mistake is reading a result from one column as if it applied to another. The columns

are the three roles of [Chapter 132](#): the actor you train, the sampler you generate rollouts with, and the deployed model you serve.

160.1 Reading the table

The table's whole point is that rows are not uniform across columns. FP8 is the clearest case, and the reason the table now has separate rows for FP8 weights and FP8 KV: as training weights it shows no measurable regression on a dense model ([Chapter 132](#)), as a rollout engine on a dense model it is a free $\sim 2\times$ win ([Chapter 158](#)), but FP8 applied to the KV cache is poison for RL (KL 7.61, [Chapter 140](#)), because the cache is re-read lossily every attention step. Never put FP8 weights and FP8-KV in the same row. AWQ-INT4 is fine as a serving format but unsafe as an RL rollout engine (KL 6.86, 42.8% disagreement, [Chapter 141](#)). W8A8 is the rollout surprise, INT8 weights and activations, yet $50\times$ cleaner than AWQ and safely under the gate (KL 0.0055, [Chapter 143](#)). And NF4 / NVFP4 belong to the training column only, NF4 scale-dependent (helps a 32B, not a 14B), NVFP4 carried as external/unexecuted ([Chapter 132](#)). The one rule that makes the table usable: never cite a quantization result without naming its column.

160.2 The rollout column is the dangerous one

The column that demands the most caution is the rollout column, because it is the only one where a wrong choice silently corrupts the gradient rather than merely the output quality. A bad serving quant gives users slightly worse answers. A bad training quant might fail to converge visibly. But a bad rollout quant produces normal-looking text from a different distribution than the trainer assumes, invalidating the importance ratio without any error or obvious symptom (the entire premise of the correctness contract, [Chapter 136](#)). That is why the rollout column carries hard verdicts, W8A8 $\checkmark$, AWQ $\times$, FP8-KV $\times$, earned by measurement, not intuition. The serving column is forgiving, the training column is semi-forgiving, and the rollout column is unforgiving: a result that is safe to deploy can be poison to roll out with, and only the per-rollout KL gate tells them apart.

TAKEAWAY

The whole quantization story is a **three-column table**, actor-training / rollout-sampling / serving, and **rows are not uniform across columns**: FP8 weights are fine to train and free to roll out with on dense, but FP8-**KV** is poison (Ch 140). AWQ is fine to *serve* but **✗ to roll out** (Ch 141). W8A8 is the rollout ✓ (50× cleaner, Ch 143). NF4/NVFP4 are training-column only (Ch 132). The **rollout column is the dangerous one**, the only place a wrong choice silently corrupts the *gradient*. Never cite a quant result without naming its column.

Next: the throughput and serving numbers, with their tables.

Throughput and Serving

Model / setup	Throughput	Evidence
Qwen3-14B, FP8, 1 × H100	~97 tok/s	measured / source-corpus
Qwen3-30B-A3B (MoE), FP8, 1 × H100	~155 tok/s	measured / source-corpus
Qwen3-32B (dense), FP8, 1 × H100	~46 tok/s	measured / source-corpus
Qwen3-Next-80B-A3B, MTP-3, 1 × H200	306.9 tok/s	external reported
Qwen3.5-122B-A10B, FP8 + MTP, 8 × RTX PRO 6000	~1,985 tok/s burst	external reported

Table 161.1. Serving throughput, measured and reported. Two patterns: the MoE’s high tok/s for its size (only ~3B active), and Multi-Token Prediction (MTP) speedups that shrink as concurrency rises.

This is the throughput reference, the tokens-per-second numbers for planning rollout cost and serving capacity, plus the one behavioral subtlety (MTP’s concurrency dependence) that the headline speedups hide.

161.1 The throughput numbers

The single-GPU FP8 numbers set the baseline: Qwen3-14B at ~97 tok/s, the Qwen3-30B-A3B MoE at ~155 tok/s (faster than the 32B dense at ~46, because only ~3B parameters are active per token), all on one H100. At the high end, a Qwen3-Next-80B-A3B reaches 306.9 tok/s with MTP-3 speculative decoding on a single H200, and a Qwen3.5-122B reaches ~1,985 tok/s burst with FP8 plus MTP on 8 × RTX PRO 6000. Note the evidence column: the single-GPU numbers are measured in (or reported by) the source corpus, while the 306.9

and ~1,985 figures are external/vendor-reported on hardware this work did not run. Two reusable patterns fall out. First, FP8 rollout roughly doubles throughput on a dense model (Chapter 158), but the gain is smaller for the MoE, whose gating already amortizes well at BF16, so do not generalize “FP8 doubles throughput” to MoE. Second, the MoE’s parameter-efficiency at inference (3B active out of 30B) means a sparse model can generate faster than a smaller dense one, a reason to prefer MoE for rollout-heavy RL if you can train it (Chapter 156).

161.2 The MTP concurrency subtlety, and the collapses

The Multi-Token Prediction speedup is not a constant, and this is the subtlety that the “1.25–2.11×” headline range encodes. The speedup peaks at batch size 1 and shrinks as concurrency rises, because the draft tokens consume KV-cache capacity that would otherwise hold more concurrent sequences, so at high concurrency the draft head competes with the very batching that gives throughput. The practical reading: speculative decoding is most valuable for latency-bound, low-concurrency serving (and for RL rollouts where the group size is modest), and least valuable at maximum batch. Finally, two failure modes from the catalog bound the throughput story: F7 (throughput collapse) when the rollout engine re-syncs per micro-batch, and F16 (QLoRA rollout slowdown) when the BitsAndBytes 4-bit kernel is on the path (Chapter 120), both turn a healthy tok/s number into a fraction of itself, so the tables above assume those traps are avoided.

TAKEAWAY

Serving throughput (FP8, 1× H100): 14B **~97 tok/s**, 30B-A3B MoE **~155** (faster than 32B dense ~46, only ~3B active), up to **306.9 tok/s** (Qwen3-Next-80B, MTP-3, H200) and **~1,985 burst** (122B, 8× RTX PRO 6000). The 306.9 and ~1,985 figures are **external-reported**, not measured here. **FP8 rollout ~doubles dense throughput** (smaller gain for MoE, don’t generalize). MTP’s **1.25–2.11×** *external/serving* range **peaks at batch**

size 1 and shrinks with concurrency (draft tokens eat KV capacity). The book's own *RL-temperature* measurement is the tighter 1.18–1.27× of [Chapter 137](#), different settings, both honest. Watch F7 and F16, which collapse these numbers ([Ch 120](#)).

Next: the mechanics of saving and storing checkpoints without losing a run.

CHAPTER 162

Checkpoint and Storage Mechanics

Three storage facts that cost real time. A Fully-Sharded checkpoint save blocks the step (a ~770-second pause was measured). Network storage is slow to write (~50 MB/s → a 152 GB model takes ~50 minutes). And the overlay-vs-network-mount trap silently writes to the wrong place, always `mount | grep nfs` before a large download.

Checkpointing seems like plumbing until it eats hours of a run, and storage seems like a non-issue until you train the wrong checkpoint. This chapter is the storage mechanics that the corpus learned the expensive way, shard merges, where to write, and the mount trap.

162.1 Saving sharded checkpoints

Under Fully-Sharded Data Parallel (FSDP), each GPU holds a shard of the weights and optimizer state, so saving a checkpoint means gathering and merging shards (DTensor merge) into a usable model, and that blocks the training step while it happens. A real measurement from the C9 program: a step that normally took ~735 seconds took 964.98 seconds when a checkpoint save landed on it (the save itself ~770 s). The lessons: save on a sensible cadence (not every step, the F11 discipline of `save_total_limit=2`, [Chapter 120](#)), and write to fast local storage first, because the merge plus a slow write can stall the run for many minutes. Plan the cadence around the save cost. A checkpoint every step on a large model can dominate wall-clock.

162.2 Where to write, and the mount trap

Where the checkpoint goes matters as much as how often. Network storage is slow to write, measured at roughly 50 MB/s, so a 152 GB model (a Qwen3-Next download) takes about 50 minutes to land, so the pattern is NVMe-first, then mirror to network storage in the background rather than writing the hot path to the network. And the most insidious storage failure is the overlay-vs-network-mount trap ([Chapters 121, 135](#)): inside a container, a path that looks like it's on the shared network volume can actually be writing to the container's local overlay filesystem, so a 152 GB download “succeeds” but is invisible to the next container, and, worse in an RL run, you can end up training the wrong checkpoint from a stale overlay copy. The one-line defense is a habit: `mount | grep nfs` before any large download or checkpoint write, to confirm the path is really on the network volume and not a local overlay that will vanish or mislead.

TAKEAWAY

Three storage facts that cost real time. **FSDP checkpoint saves block the step** (a normal ~735s step measured **964.98s** when a ~770s save hit it), save on a cadence (`save_total_limit=2, F11`), write local-first. **Network storage is slow** (~50 MB/s → a 152 GB model ~50 min), go **NVMe-first, mirror to network storage** in the background. And the **overlay-vs-network-mount trap** silently writes to a container's local overlay (a download “succeeds” but vanishes. You can train a stale checkpoint), defend with `mount | grep nfs` **before any big write**.

Next: the cost ledger that ties the whole hardware story to dollars.

CHAPTER 163

The Compute Ledger




Figure 163.1. Re-mine before you re-run: most new findings already sit in checkpoints and logs you’ve generated — 22 of the book’s findings came from re-mining, not one new training run.

Every hardware decision in this Part finally reduces to a number on an invoice, and this closing chapter is the cost ledger, the per-hour rate, the per-experiment cost as a kill-decision input, and the exhibit that justifies the whole “re-mine before re-run” discipline.

163.1 The rate, and cost as a kill input

The operative idea is to treat per-experiment compute as an input to the kill-or-continue decision, not as an after-the-fact accounting line. A run that will burn meaningful compute and whose base model already scores near 0% or near 100% ([Chapter 134](#)) should be killed before launch. A run showing the F1 entropy-collapse signature at step 30 should be killed at step 30, not step 500 ([Chapter 127](#)). The compute ledger is most useful prospectively: knowing that a 14B SWE-bench run burns ~1.5B tokens before plateau, or that an OPD run takes about half the compute with an AWQ teacher versus BF16 ([Chapter 133](#)), turns “should we run this?” into a decision with a price attached.

163.2 The closing exhibit: re-mine before re-run

The ledger’s headline exhibit closes the book’s whole argument about discipline. The two late research programs, the inference forensics (Part XIV) and the internals re-mining, together produced 22 findings for a small fraction of the compute a fresh set of experiments would have cost. The reason for the large gap is the “re-mine before re-run” principle: the cheapest experiment is the one already sitting in your checkpoints and logs, because a forensic pass over existing artifacts, checkpoint census, log census, a handful of small harness scripts, answers many questions without launching a single new training run. That is the note the hardware Part ends on, and it reframes the entire ledger: the largest savings in RL post-training come not from a faster GPU or a cheaper instance, but from not running the experiment you can already answer. Spend the compute on the questions your existing artifacts can’t reach. Mine the rest.

TAKEAWAY

Cost reduces to compute and wall-clock. Treat **per-experiment compute as a kill-or-continue input**, kill a near-0%/100% task before launch, kill an F1 collapse at step 30 not 500. The closing exhibit: the inference + internals re-mining programs delivered **22 findings for a small fraction** of what fresh experiments would have cost, a large gap from “**re-mine before re-run**” (the cheapest experiment is the one already in your checkpoints). The biggest savings come from *not running the experiment you can already answer*.

This completes the main body, Parts I through XV. The appendices that follow carry the source journals and the supporting reference material the chapters draw on.

Part XVI

Appendices

Reference material for the main body. The reference appendices (E–O) are complete. The annotated-log appendices (A–D, P, Q) index their source journals with the load-bearing milestones. Appendix R is compiled from the project’s run journal (data-locked 2026-06-02), and the rollout-side failure family carries permanent IDs (F18–F21, Appendix I).

APPENDIX A

Annotated Logs, Search-Agent v1–v5 Entropy Runs

This is the saga where Clip-Cov was earned.

The search-agent program trained Qwen3-4B to answer questions by calling a retrieval tool (E5 + FAISS over Wikipedia), multi-turn. Its value to the book is the entropy-collapse-and-recovery arc that produced the Clip-Cov mechanism ([Part IV](#)) before it was cross-validated on SWE-bench ([Appendix C](#)).

The v1–v5 family is the entropy-control ablation (the tool-description / 4-turn interventions are separate search-agent work, see [Appendix H](#), not v2 here):

Entropy-control approach	Outcome
no entropy control	collapsed — entropy slid 0.188 → 0.070 over 250 steps
global entropy bonus, too large	exploded into incoherent noise
global entropy bonus, smaller	no stable setting — collapse pressure returns
Clip-Higher alone (DAPO)	delayed the collapse but did not prevent it
Clip-Cov (token-level covariance control)	held a damped entropy band (0.73 → 0.64), never collapsing. Caveat: later eval correction and response-length/quality degradation made checkpoint selection nontrivial

Table A.1. The five search-agent runs (v1–v5): each entropy-control setting and its outcome, ending in Clip-Cov’s first stable equilibrium.

Result hygiene, the headline is not canonical. The widely-quoted TriviaQA 66.7% / NQ 41.7% is the v5 + SiLU/Swish multi-reward, step 750 peak on a ~36-sample-per-dataset eval, historical, not canonical. Re-evaluating V5 at 500 samples dropped TriviaQA well below that peak — the later re-eval reads 44.0%, the earlier April-28 matrix 55.6% (the two 500-sample sets disagree, and only the ranking is treated as supported). This is precisely the eval-discipline finding of [Part IX](#) (small eval sets are unreliable. 500-sample minimum):

TriviaQA: small-eval peak vs. 500-sample re-eval			
Result	Checkpoint	Eval size	TriviaQA
Small-eval peak ¹	v5+SiLU/Swish @750	~36 / dataset	66.7%
500-sample re-eval ²	V5	500 / dataset	44.0%

Table A.2. TriviaQA accuracy per checkpoint; the 500-sample re-eval is the eval-discipline finding of Part IX. NQ figures and status are in the notes.

Table notes

1. **Small-eval peak (v5+SiLU/Swish @750):** NQ 41.7%. Historical — narrative, not the claim.
2. **500-sample re-eval (V5):** candidate; canonical pending eval-file lock (NQ@500); well below the small-eval peak, and the April-28 matrix reads 55.6% on the same dataset.

So the lesson the appendix carries forward is eval discipline and early stopping, not a 66.7% headline.

Entropy timeline (the spine of [Part IV](#)):

Approach	Entropy behavior	Accuracy behavior	Lesson
no entropy control	collapse	deceptive score	no entropy control fails
global bonus, too large	explosion	dead	global bonus too high
global bonus, smaller	slower explosion	temporary peak	smaller bonus not enough

Approach	Entropy behavior	Accuracy behavior	Lesson
Clip-Higher only	collapse delayed	unstable	Clip-Higher alone insufficient
Clip-Cov + gated reward	non-monotone stabilization	peaked then de-graded (small-eval)	Clip-Cov helps but needs eval discipline

Table A.3. The same v1–v5 saga read as an entropy timeline: behavior, accuracy, and the lesson each run taught.

Annotated index, summarized from our run logs. The canonical 500-per-dataset NQ@500 is 21.4%; the v5.1 turn-curriculum row is pending the raw eval files.

APPENDIX B

Annotated Logs, C9 MoE Stability (GSPO, FREEZE, JITTER, COMBO)

The MoE stability program.

The C9 program is the mixture-of-experts (MoE) RL work on Qwen3-30B-A3B. Its value is not a simple “GRPO fails, GSPO wins” story, vanilla GRPO showed scale-dependent robustness (the 128-expert MoE did not universally collapse under conservative settings, R.6). The lesson is that MoE routing stability is scale- and configuration-dependent, and GSPO / FREEZE / JITTER / COMBO change the tradeoff between entropy, validation, and internal drift. GSPO is the safer default, not proof that vanilla GRPO cannot work.

Cycle / arm	What it tested	Result
C9-GRPO-Collapse	BF16 baseline, no warmup	the routing-entropy decay (the problem)
C9-GRPO-FP8	FP8 <i>rollout</i> on MoE	block-size / TP / memory wall (not “viability”) — see below
C9-GSPO	sequence-level optimizer	productive narrowing: ~4× more accurate at <i>lower</i> entropy
C9-FREEZE	freeze router gradients (diagnostic)	isolates routing vs representation drift
C9-JITTER	<code>jitter_noise=0.01</code> on router inputs	architectural-fix arm (Switch-Transformers)
C9-COMBO	GSPO + JITTER stacked	constructive stack by matched-step validation

Table B.1. The C9 MoE-stability arms (GRPO collapse through GSPO+JITTER combo): what each tested and what it showed.

The corrected finding ([Ch 53](#), [Ch 117](#)). GSPO does not simply “prevent entropy decay.” The measured story is that GSPO makes the routing narrowing productive, under GSPO routing narrows from the early ~0.081–0.112 band ([Part XV](#)) to low entropy at deeper steps (e.g. ~0.04) while validation rises, roughly 4× the accuracy at lower entropy than vanilla GRPO. GSPO is therefore the safer default for MoE (its sequence-level ratio avoids the worst token-level pathologies), not proof that vanilla GRPO always breaks a MoE, vanilla GRPO showed scale-dependent robustness (a 128-expert MoE did not universally collapse under conservative settings). C9-COMBO (GSPO+JITTER) was the matched-step winner, with one caution: low entropy can be a false-positive, read it jointly with validation and reward variance ([Chapter 126](#)).

Do not use Appendix B to support the absolute claim that vanilla GRPO always breaks MoE.

On C9-GRPO-FP8: this hit an implementation wall, not a viability result, FP8 rollout at TP=8 failed (expert dim 768/8=96 not divisible by FP8 `block_n=128`. Only TP=1/2/3 valid), and TP=2 then OOMed (~15 GB/GPU + KV + BF16 sync buffer). See [Chapter 158](#).

Annotated index. The four-way VAL ladder (val@s10/s20/s30/s40 per arm) and the exact entropy traces are kept in our run logs rather than reproduced in full here.

APPENDIX C

Annotated Logs, SWE-bench SW1–SW6

The Qwen3-14B coding-agent saga. The worked postmortem of [Chapter 127](#).

The cross-domain validation of Clip-Cov, and the book’s central postmortem chain. Each version is one change. Do not attribute later settings to earlier versions.

Run	The one change	Result	Lesson
SW1	vanilla GRPO, binary test-pass reward, no Clip-Cov, no KL	entropy 0.42 → 0.03 in ~30 steps; ~4-token outputs	sparse binary SWE reward + full FT collapses fast
SW2	add Clip-Cov + KL	damped oscillation; the run survives	Clip-Cov is domain-agnostic, not search-specific
SW3	stronger Clip-Cov / KL / LR	entropy stable but reward declining	exposed a GRPO length bias
SW4	<code>seq_mean_token_mean</code> loss + continuous shaping rewards	first stable reward signal; overfits early	loss aggregation and continuous rewards matter
SW5	scale to ~280–300 instances, constant LR	collapse again at step 18–22	constant LR + falling entropy → late collapse
SW6	cosine LR, stronger Clip-Cov, lower LR, KL 0.01, max turns 10	stable to step 146 — but not learning	stability is necessary, not sufficient

Table C.1. The SWE-bench coding-agent saga SW1–SW6: each version’s one change, the failure it fixed, and the new failure it exposed.

Headline result carried forward ([Ch 127](#), [Ch 134](#)): stability was fixed at SW2, learning at SW4. SW6 plateaued at ~1.5B tokens, a capacity ceiling (SWE-Gym too hard for 14B), not a bug. LR for the 14B GRPO++ runs: 5e-7 to 1e-6.

Annotated index, summarized from our run logs.

APPENDIX D

Annotated Logs, GDPO G1–G4

The per-channel-normalization cautionary tale.

The GDPO (Group-Decomposed Policy Optimization) line tested per-channel / per-component advantage normalization. Its value to the book is negative: it is the evidence behind the warning ([Chapter 123](#), pitfall P22) that blindly normalizing every reward channel destroys the correctness-gradient hierarchy.

Run	Idea	Result	Diagnosis
G1	vanilla GDPO, 5 channels	failed	per-channel dilution
G2	strategy channels	crashed / unstable	channel conflict
G3	gated strategy channels	OOM ~step 213; no learning	implementation + no gain
G4	budgeted marginal utility	~step 250; zero learning	F1 signal diluted

Table D.1. GDPO variants (G1–G4): idea, result, and diagnosis. None beat the v5 scalar-gated Clip-Cov baseline.

Headline lesson carried forward ([Ch 123](#) P22): auxiliaries must be gated behind correctness, not inverse-variance-normalized alongside it. GDPO is why the book’s reward doctrine keeps correctness as the dominant, ungated signal.

Annotated index, summarized from our run logs.

APPENDIX E

Hydra Config Reference (Schematic)

A schematic reference, the key fields and their gotchas, not the fully-resolved per-run configs (the fully-resolved per-run configs, plus a `dump_resolved_config.sh + sha256sum` of the tokenizer/chat-template, are generated per run rather than reproduced here).

The book's experiments are configured through a Hydra overlay system: a base config, an algorithm overlay (GRPO / GSPO / GRPO++ / DAPO), a quantization overlay, and a tier overlay (1 / 4 / 8 GPU). The fields that matter:

```
# base
model_revision: <hf-id>@<commit>
algorithm: grpo | gspo | grpo_pp | dapop | opd # the optimizer
loss_agg_mode: token_mean | seq_mean_token_mean | sample_mean #
  ↳ the F3 length-bias lever
normalize_adv_by_std: true | false # advantage normalization
  ↳ (separate field)
kl: {placement: reward, estimator: kl, coef: 1.0e-3} # P29 (K1
  ↳ in reward)
clip: {low: 0.2, high: 0.28} # GRPO. GSPO uses ~3e-4
  ↳ sequence-level
quant: bf16 | fp8 | awq_int4 | gptq_int4 | qlora_nf4 | qerl_nvfp4
rollout: {engine: vllm, gpu_memory_utilization: 0.6,
  ↳ sleep_level: 1}

# tier overlays
1gpu: {gpus: 1, strategy: fsdp2_or_lora} # 0.6B/4B full-FT. 14B
  ↳ QLoRA. 32B QeRL
4gpu: {gpus: 4, layout: "2 teacher-vLLM + 2 student-FSDP"} # OPD
8gpu: {gpus: 8, strategy: fsdp2_full_shard | megatron_ep8} #
  ↳ flagship
```

Key reminders: the algorithm, the loss-aggregation mode (`loss_agg_mode`, the F3 length-bias lever), and advantage normalization are three separate fields, not one (a common conflation). Put the KL in the reward (K1), not the loss (P29). Set `loss_agg_mode` explicitly, since the cross-framework number mismatch (P30) traces to it. Use the GSPO sequence-level clip on MoE, never the GRPO 0.2 (F14). Start co-located rollout at `gpu_memory_utilization=0.6` or use Unsloth Standby for 0.95 (P16). Resolved per-run configs (Search-V5, C9-GSPO, SWE-SW6) are produced by the manifest dumper rather than reproduced here.

APPENDIX F

Algorithm Cheat Sheet

The one-page decision card.

Which optimizer?

Situation	Use	Why
Dense model, stable task	GRPO	simplest; works when entropy is well-behaved
Dense model, entropy collapsing	GRPO++ (Clip-Cov + Clip-Higher)	asymmetric clip lets good tokens grow (Part IV)
Mixture-of-experts	GSPO (safer default)	sequence-level ratio survives routing flips (F5/Ch 117)
Long task, base sometimes succeeds	GRPO++ + cosine LR + goal-gated shaping	only if real successes exist (the SW2→SW6 recipe, Appx C)
Long task, base almost never succeeds	SFT / OPD / an easier executable environment	outcome RL has no signal to amplify
Proxy shaping ↑ reward but not eval	stop and redesign the reward	you are training “looking busy,” not success

Table F.1. Which optimizer to reach for, by situation — and the reason behind each choice.

The non-negotiables: KL in the reward, not the loss (K1, P29). `loss_agg_mode` named explicitly (P30). Seeds ≥ 3 with the variance budget (Ch 87). The rollout-correctness preflight before any lossy acceleration (Ch 136). The diagnostic before launch: if the base solves the task ~0% or ~100%, RL has no signal (Ch 134), RL sharpens what the base can sometimes do, it does not expand. On reward shaping (Part VI): generic shaping on a sparse, hard task tends to reward activity rather than success, gate auxiliaries behind correctness, and if shaped reward rises while held-out eval does not, stop and redesign rather than push harder. When outcome reward is sparse and a stronger same-tokenizer teacher exists, prefer OPD over shaping.

APPENDIX G

Compute Reference

Machines, layouts, A100-40GB-vs-H100-80GB history.

The two machines.

Machine	GPUs
Node A	8× A100-40GB
Node B	8× H100-80GB

Table G.1. The two machines behind the book: the A100-40GB machine (8x A100-40GB) and the H100-80GB machine (8x H100-80GB).

Footprints (1× H100-80GB).

Model	BF16	FP8	INT4	Rollout tok/s (FP8)
Qwen3-14B (dense)	~28 GB	—	—	~97
Qwen3-32B (dense)	—	—	—	~46
Qwen3-30B-A3B (MoE)	~62 GB	~34 GB	~18 GB	~155

Table G.2. Single-H100 memory footprints by precision and rollout throughput for the three reference models.

Tiers: 1 GPU (pedagogy + small models. 14B via QLoRA. 32B via the NF4-LoRA proxy, QeRL/NVFP4 itself is external/planning only in this H100 stack, Blackwell-gated) · 4 GPU (OPD 2+2. Mid-size sharded) · 8 GPU (flagship: 14B full-FT, 30B MoE GSPO, OPD 4+4). Plan in wall-clock and relative compute. Budget rule: +20–30% for failed runs. Account for validation overhead (~72× a step on multi-turn) and FSDP→network-storage checkpoint I/O (20+ min). Full tables: [Chapters 133, 148–151, 163.](#)

APPENDIX H

Experiment Index (Main Rows. Full Q-program in Appendix R)

Every run cross-referenced to chapter and appendix.

Program	Runs	Chapters
Search-agent	v1–v5	Part IV; Appx A
SWE-bench coding agent	SW1–SW6	Ch 127; Appx C
GDPO line	G1–G4	Ch 123; Appx D
C9 MoE / GSPO	C9-Collapse, -FP8, -GSPO, -FREEZE, -JITTER, -COMBO	Ch 53; Appx B
Quantization (training-side)	Q1–Q4 (+Q5 closed)	Ch 132; Appx R
Inference / rollout	I-1...I-6	Part XIV; Appx P
Forensics (AWQ etc.)	A1, A3, E-A3.1–.6	Ch 141–143; Appx P
Internals re-mining	J4 §1.1–§2.4	Part VII; Appx Q
On-policy distillation	math + (code, pending)	Part VI

Table H.1. Every experiment program cross-referenced to its runs, chapters, and appendices.

The fully expanded per-run table, config hash, eval hash, tier, cost, best checkpoint, raw-log path, and canonical/historical/retracted status, for the Q-program rows is now in Appendix R, which carries the 55-cell grid, the multi-seed table, and the audit/retraction log. Remaining non-Q rows (search-agent, SWE, GDPO) are indexed to Appendices A–D.

APPENDIX I

The Complete Failure-Mode Catalog

The master F-table.

Evidence status (read every row through this): Lived = directly observed in this project's runs. Measured = confirmed by a specific eval/preflight. External = documented in framework issues/papers and confirmed in our stack. Provisional = ID or mechanism not yet final. F1–F17 are Lived/Measured except the framework-seam pitfalls (F7, F8), which are External-confirmed-here. The rollout-side family below (F18–F21) is Measured (forensic), with permanent IDs now assigned.

ID	Failure	Signals	Time	Root cause	First fix
F1	Entropy collapse	entropy → floor; identical gens	~30 steps	symmetric clip caps prob gains	Clip-Cov + Clip-Higher
F2	Advantage col- lapse	reward-std → 0	varies	base solves ~0% or ~100%	change data / cold-start
F3	Length bias	completions grow, accuracy flat	50+	length-biased loss aggregation	seq_mean_ token_mean
F4	Repetition col- lapse	n-gram loops	50+	reward ex- ploitable by repetition	repetition penalty; KL
F5	Ratio explosion (MoE)	ratio_max fat tails; clip_frac_ high→1	<50	routing changes across the up- date → token- level ratio breaks; risk is scale-dependent (expert count, LR, backend consistency) — <i>not</i> proof GRPO always breaks MoE	GSPO as safer default; lower LR

ID	Failure	Signals	Time	Root cause	First fix
F6	Reference drift	kl_to_ref ↑, eval ↓, reward ↑	100–300	KL too low / no refresh	$\beta \geq 1e-3$; joint reward
F7	Throughput collapse	rollout sawtooth; util ~5%	continuous	vLLM resync per micro-batch	sleep mode; packing
F8	VRAM leak	peak VRAM monotone ↑	200–500	Torch-2.6 FSDP2 bug	Torch 2.7; veRL ≥ 0.5
F9	Tokenizer mismatch (OPD)	KL flat; no improvement	step 1	teacher/student vocab differ	sha256 tokenizer; within-family
F10	Tool-token contamination	eval ↓; regurgitates context	50+	loss over tool/retrieval tokens	response mask
F11	Disk-full crash	crash, no warning	~80	save_total_limit unset	save_total_limit=2; object storage
F12	Chat-template drift	reward ↑, eval ↓; format-err eval-only	step 1	train/eval render differ	hash template; same path
F13	Late collapse	sudden collapse at peak	varies	constant LR at peak	cosine decay
F14	GSPO clip misconfig	NaN loss, grad-norm 1e6	step 1	GRPO 0.2 clip on GSPO	sequence-level clip $\sim 3e-4$
F15	OPD mode collapse	student $\equiv$ teacher; diversity $\rightarrow 0$	200+	reverse-KL too aggressive	temp 0.7; stop at plateau
F16	QLoRA rollout slowdown	5–10 $\times$ slower rollout	continuous	bnb 4-bit forward kernel	Unsloth kernels; BF16/LoRA
F17	AWQ-teacher logprob drift	unstable OPD reward; KL oscillates	~50	quant noise in teacher logits	GROUP_SIZE=128; FP8 teacher

Table I.1. The complete failure-mode catalog F1–F17: signals, onset time, root cause, and first fix for each.

Rollout-side lossy family (F18–F21) (permanent IDs, [Part XIV](#)):

ID	Failure	Signal	Verdict
F18	FP8-KV distribution corruption	KL 7.61; ~40% disagree; 0.97 $\times$	reject
F19	AWQ-rollout IS-ratio break	KL 6.86; 42.8%; one token 6.72 nats $\rightarrow$ 825 $\times$; 18.6% outside clip	reject unless pre-flight passes

ID	Failure	Signal	Verdict
F20	dense-AWQ NaN-at-L2 kernel cascade	NaN at layer 2 (AutoAWQ)	no clean H100 path
F21	MoE router-flip (quantized rollout)	31% top-1; 87% top-8; gate unquantized	reject for MoE rollout

Table I.2. The rollout-side lossy failure family F18–F21: signal and RL verdict for each.

Evidence/scope: F18–F21 are measured (forensic) in [Part XIV](#), F18 FP8-KV, F19/F20 dense-AWQ, F21 from the Qwen3-30B-A3B AWQ-rollout path specifically, not a universal-MoE claim.

Operational cousins (crash/mislead, not silent-corrupt): vLLM `n_samples>1 else T=0`. SGLang↔vLLM CUDA collision. Zombie VRAM after `pkill`. Overlay-vs-network-mount path trap. EAGLE-3 cross-version regression.

APPENDIX J

The Pain-Point Index, by Symptom

Framework and operational traps, by symptom. Issue IDs as reported in the pitfall table. (This in-print index gives each trap's symptom, first-check / fix command, debug-time-saved, and evidence status.)

veRL (P1–P9): P1 vLLM gibberish after weight resync (sleep-2. Pin in-place loader) · P3 Ray kills workers ~100 steps (host RAM leak) · P5 padding in loss when no EOS (response_mask, silent) · P8 sleep-mode one vLLM per process.

TRL / vLLM (P10–P18): P10 length-biased aggregation default (set `loss_agg_mode`) · P12 colocate+sleep → no learning (silent) · P13 format/accuracy relabel · P16 OOM at 0.8 vanishes at 0.6 (Unslloth Standby → 0.95) · P17 QLoRA loads 16-bit anyway.

Reward / multi-turn (P19–P25): P19 256-token garbage collapse (only-KL gradient) · P20 near-zero loss ≠ convergence (centered advantages. Unanimous groups) · P21 `None` reward → silent 0 (assert numeric) · P22 composite collapses to highest-variance component (correctness-gate, don't inverse-variance-normalize) · P23 training on tool-output tokens (turn-level mask) · P24 naïve dense per-turn rewards -14pp (calibrate) · P25 QwQ-32B template (patched revision).

OPD / KL / eval-drift (P26–P33): P26 cross-tokenizer distillation trains on wrong logprobs · P27 sampled-token OPD biased (+0–2% → top-K reverse-KL) · P29 K3-in-loss is a biased gradient estimator → put K1 in the reward · P30 cross-framework number mismatch (hash algorithm, `loss_agg_mode`, and advantage-normalization mode) · P31 BFCL version mismatch · P32 τ -bench domain ambiguity · P33 math extraction fails on LaTeX, use strict extraction first and report the extract rate. A regex fallback is a labeled diagnostic pass only, never the path to a canonical accuracy number (a loose fallback

hides truncation and grabs intermediate numbers, the OPD/GSM8K lesson).

Silent killers / reporting (P34–P37): P34 `TOKENIZERS_PARALLELISM` deadlock · P35 flash-attn ABI mismatch after restart · P36 sporadic NCCL hangs (all three: set env in the launch script) · P37 single-seed reporting bias (≥ 3 seeds as a minimum. Attach the [Ch 87](#) budget).

APPENDIX K

Internals Telemetry Methods

The probe protocols that make Parts V and VII reproducible.

The telemetry layer instruments a training run to “watch the model think,” at a measured ~9% overhead on the experiment budget. The toolkit:

Probe	What it measures	Reads out
NS5 forward-KL neutral probes	KL on held-out neutral prompts at matched steps	drift not attributable to the task
Residual drift ($\Delta\mu$ / $\Delta\max$)	per-layer change in residual-stream statistics	where representations move
Per-layer routing entropy	expert-selection entropy by layer (MoE)	routing collapse vs stability
Logit lens	intermediate-layer logit projections	when predictions form across depth
Attention forensics	head attention patterns, matched-step	which heads attend where, and how it shifts

Table K.1. The internals-telemetry probes: what each measures and what it reads out.

Protocol notes: compute KL in fp32 (the fp64-vs-fp32 sensitivity of [Chapter 129](#)). Compare at matched steps across arms. Budget the ~9% overhead into the run cost. The re-mining scripts (checkpoint census, log census, the small per-probe utilities) are the reusable implementation. They are what let the re-mining program (Appendix Q) answer questions without new training.

APPENDIX L

The Hyperparameter Sensitivity Reference

Hyperparameter	Safe range / value	Sensitivity notes
Learning rate (14B GRPO++)	5e-7 to 1e-6	1e-5 collapses small models (entropy at step 30)
Clip (GRPO)	low 0.2 / high 0.28 (Clip-Higher)	asymmetry is the point (Part IV)
Clip (GSPO)	~3e-4 sequence-level	the GRPO 0.2 → instant NaN (F14)
KL coefficient β	$\geq 1e-3$ minimum	too low → reference drift (F6); place in <i>reward</i> (P29)
Group size n	bounded by the KV ceiling	~239 @ 1K ctx → ~7 @ 32K on 1 × H100 (Ch 146)
Rollout temperature	0.7–1.0 (RL range)	spec-decode speedup is temperature-robust (Ch 137)
gpu_memory_utilization	0.6 co-located, 0.95 with Standby	the 0.6-vs-0.8 OOM cliff (P16)
Student temp (OPD)	0.7 + entropy bonus 1e-4	prevents mode collapse (F15)

Table L.1. Hyperparameter sensitivity reference: safe ranges and the failure each knob triggers when misset.

The meta-rule: the most sensitive knobs are the learning-rate schedule (constant-at-peak → F13) and the clip regime (algorithm-specific, F14). The most forgiving is the backend (FA vs FlashInfer, ~5–7%). Change one knob per run (Chapter 127).

APPENDIX M

The Decision Matrices

Data → algorithm (the band).

Base success on the task	RL verdict
~0%	too hard — cold-start (SFT) or change task
20–80%	RL helps — informative advantage
~100%	saturated — eval only

Table M.1. Base success rate on a task mapped to the RL verdict: too hard, RL helps, or saturated.

Stability ranking (most → least robust): GSPO (MoE) $\approx$ GRPO++ (dense) > DAPO > vanilla GRPO. Reward 2×2: correctness-only (plateaus) vs shaped (risk of hacking) × single-objective (stable) vs multi-objective (collapse risk → correctness-gate). Annealing: cosine LR decay for long runs (prevents F13). Hold constant only on short, stable tasks.

APPENDIX N

The Named Concepts

(Glossary of Citable Ideas)

The terms below are *this book's own* coined ideas and experiments, not external methods; external works it builds on are collected in the References appendix. Each is tagged by status: [Doctrine] repeated across programs · [Mechanism] supported by measurement · [Hypothesis] plausible, not proven · [Operational] engineering rule · [Coined] naming layer.

- The Correctness Contract [Doctrine], a rollout acceleration is “free” only if it passes the gates for its class: Gate A forward-KL $KL(target \parallel accel)$ over generated response tokens at matched seed, threshold 0.05. Gate B end-to-end step wall-clock. Gate C held-out eval (lossy methods only). Safe-by-construction / exact / lossy taxonomy (Ch 136).
- Re-Mine Before Re-Run [Doctrine], the cheapest experiment is the one already in your checkpoints and logs. A forensic pass over existing artifacts answers many questions for a fraction of a fresh run (Ch 163. The re-mine-vs-re-run exhibit, a forensic pass cost a small fraction of a fresh run).
- Fork-Token Credit [Mechanism], RL's gradient concentrates on the tokens where rollouts diverge (the forks): Spearman ~ 0.95 / 99.8% concentration on fork tokens, measured on the matched checkpoints of Part VII (exact eval conditions in Appendix Q. The Spearman against the raw probe output).
- Update-Style Preservation (LoRA-Preserves, Full-FT-Forgets) [Mechanism], at 14B, LoRA-RL and OPD preserve thinking-mode capability while full-FT GRPO destroys it (Q4 BF16 $-23.6pp$, FP8 $-10.8pp$, multi-seed). The update style, not the precision, decides what survives, and the loss is invisible in attention-space KL (Ch 132. Appx R.3).

- Sharpen-Not-Expand [Doctrine], under the recipes and tasks studied here, RL post-training sharpens capabilities the base model can already sometimes exhibit rather than adding new ones (Part VII), hence the 20–80% data band. Stated as an empirical pattern in this setting, not an eternal law.
- The Variance Budget, the seed-to-seed standard deviation that a claimed improvement must exceed to be real. A 2–3pp gain at one seed is indistinguishable from noise (Ch 87).
- Productive vs Wasted Entropy, low entropy is not automatically failure: it is productive sharpening when held-out eval and reward variance are still improving, and collapse only jointly with their stalling (Ch 126, 131).
- The Lossy-Saturation Hypothesis [Hypothesis], once a perturbation's $|\Delta\text{logit}|$ exceeds the top-1/top-2 gap (typically 1–3 logits), the argmax flips near-Bernoulli, so token disagreement saturates near ~40% regardless of the lossy source (Ch 144). Explicitly a hypothesis with a specified-but-unrun test (gap vs $|\Delta\text{logit}|$ histograms), not an established result.
- The Router-Flip Mechanism [Mechanism], on the Qwen3-30B-A3B-AWQ rollout path (64 math probes, 48 layers, 128 experts, top-8, gate left unquantized), ~31% of tokens reroute their top-1 expert (87% change the top-8 set), upstream-activation-driven (Ch 142). Scoped to this model/quant/path, rollout-side, not a universal MoE theorem.
- Update-Style-Decides-What-Survives, the policy-update style (token- vs sequence-level ratio, clip regime, KL placement) determines which behaviors survive optimization, more than the reward alone.

Plus the core mechanisms named earlier: Clip-Cov (covariance-aware clipping), GRPO++ (Clip-Cov + Clip-Higher + KL recipe), and the gated reward mechanism (correctness-gated auxiliaries).

APPENDIX O

Companion Papers

Four extractable papers. These are *this book's own* venue-shaped contributions, not external references (those are in the References appendix). The book is the integrated reference.

1. Search-agent reward & entropy dynamics [priority paper]. The SiLU/Swish-gate observation, Clip-Cov, and gradient-diversity findings from the v1–v5 saga (Appx A). Reports the canonical 500-sample results (TriviaQA 44.0%) and treats the historical small-eval peak (66.7% / 41.7% on ~36 samples) and the 36→500 correction as part of the contribution, not as a headline. Evidence state: Clip-Cov mechanism solid. Canonical numbers pending the NQ@500 raw eval.
2. What RL does to a model, a mechanistic account. Sharpen-not-expand, fork-token credit, productive entropy, update sparsity, the full-FT regression, the MoE residual-direction ceiling ([Part VII](#)). Absorbs the former MoE-internals paper).
3. Practical eval traps in RL post-training. The [Part IX](#) doctrine, methodology-is-the-result, the variance budget, the benchmark-version and extraction traps (P31–P33).
4. Lossy inference is unsafe for RL (proposed, systems-venue). The correctness contract, the FP8-KV and AWQ quantitative KL numbers, the router-flip mechanism, W8A8 as the safe recipe, and the saturation hypothesis ([Part XIV](#)). Evidence state: the KL numbers, router-flip, and W8A8 are measured. The saturation claim is a hypothesis. The MTP-vs-EAGLE comparison is directional (different base models), so claim “several measured findings plus one hypothesis,” not “4 proven novel findings,” until each is evidence-locked. Recommendation: stand-alone (different venue/audience. Priority claims time-sensitive).

APPENDIX P

Annotated Journal, The Inference Program

I-cycles, with the forensics master tables.

The inference program (Part XIV) ran I-cycles for negligible compute. The master tables:

Acceleration × RL verdict.

Method	Class	Gate A (KL)	Speed	Verdict
Spec-decode (MTP)	safe-by-construction	—	1.27×→1.18× across T	adopt
Prefix caching (APC)	exact	—	26% / 1.40× (long-prefix seq)	adopt at the seam
Disaggregation	exact	—	25–50%	adopt if 2 GPUs
W8A8	lossy	0.0055	—	safe
FP8-KV	lossy	7.61	0.97×	reject
AWQ-INT4	lossy	6.86	0.95×	reject

Table P.1. Each rollout acceleration by class, Gate-A KL, speedup, and adopt-or-reject verdict.

The correctness-contract schema (the gates behind the verdicts above):

```
Gate A - distribution preservation
  comparison: baseline rollout engine vs accelerated rollout
    ↪ engine
  prompt set: fixed. Decoding mode: matched (T, top-p, seed)
  KL direction: forward, KL(target || accel). Unit: per token
    ↪ (response tokens only)
  threshold: 0.05 nats/token
  secondary: token-disagreement %, max per-token logprob gap,
    ↪ ratio out-of-[0.8,1.2] fraction
```

```
Gate B - wall-clock
  metric: end-to-end RL step (not only tok/s). Batch size / n /
  ↪ context fixed
Gate C - downstream eval
  required for LOSSY methods only. Fixed held-out set.
  ↪ Same-script base control
```

MoE router flips (AWQ-MoE path): 31% top-1, 87% top-8, gate unquantized ([Ch 142](#)). AWQ-MoE preflight: mean-seq-KL 0.2864. Max per-token logprob gap 6.72 nats $\rightarrow$ 825 $\times$ ratio. 18.6% of ratios outside the clip band. MTP vs EAGLE-3: acceptance 0.601 vs 0.283 ($\approx 2.1\times$), directional, different base models/stacks, not a pure draft-head head-to-head. W8A8: safe under this tested preflight (KL 0.0055), not a universal guarantee. Verbatim internal I-cycles. Evidence boundaries in [Chapter 147](#).

APPENDIX Q

Annotated Journal, The Internals Re-Mining Program

The re-mining program, with the checkpoint census.

The internals re-mining program forensically mined existing checkpoints and logs rather than running new training, the empirical engine behind [Part VII](#) and the Re-Mine-Before-Re-Run principle.

The checkpoint census mapped ~14 TB of existing checkpoints across programs. The findings, by evidence tier:

Finding	Tier	Check-points	Probe	Main num-ber	Chapter	Status
Sharpen-not-expand	A	base + trained	pass@k	pass@1 ↑, pass@64 ↓	Part VII	canonical (with the in-setting caveat)
Fork-token credit	A/B	matched checkpoints	credit map	Spearman ~0.95	Part VII	verify against raw probe
Residual-direction ceil-ing (MoE)	A/B	FREEZE/JITTER/COMBO	residual drift	direction sign	Part V/VII	canonical if raw table present
W8A8 load time	B	inference run	load pro-filer	~6× (49 vs 8 s/shard)	Ch 129/143	operational

Table Q.1. The checkpoint-census findings by evidence tier: probe, main number, chapter, and status.

Method: a checkpoint census, then a log census, then small per-probe scripts applied the telemetry toolkit of Appendix K to artifacts already on disk, re-mining instead of re-running. Headline forensics-three ([Ch 129](#)): a NaN that was really a masking bug (symptom ≠ cause). The fp64-vs-fp32 KL estimator changing a gate decision. W8A8's ~6× load time.

APPENDIX R

The Q-Program Master Record

This record consolidates the Q-program’s completed grid, cross-evals, audits, and corrections.

The Q-program asked whether quantization (FP8, NF4, NVFP4) preserves accuracy under GRPO, across model scale (4B/14B/32B) and benchmark. The answer, in one line: RL fine-tuning leaves the model essentially intact (internals preserved), quantization is BF16-equivalent where it ran, and the RL transfer that does occur is scale-dependent and quantization-invariant.

R.1 The canonical correction: “Q1” is NF4, not NVFP4

The single most important entry (V2 Cycle 290). Q1 was planned as NVFP4 (the QeRL thesis: a 32B model in ~5 GB on one H100), but could not run on Hopper: vLLM’s NVFP4 kernel (`modelopt_fp4 / petit_nvfp4`) is compiled for `sm_100`, Blackwell only. On the H100 (`sm_90`) it dequantizes NVFP4 to BF16 at load, so the ~5 GB footprint becomes ~64 GB and the memory thesis evaporates. The run that landed under the `q1_nvfp4_*` directory therefore used NF4 (bitsandbytes) as the Hopper proxy. Read every “Q1” as the NF4-LoRA proxy. True NVFP4 is deferred, gated on Blackwell access.

R.2 The 55-cell master grid, NF4 vs BF16 at 32B (final, corrected)

The cleanest complete comparison (Q1 NF4 vs Q1 BF16, both 32B, LoRA $r=64$, 100 GRPO steps):

Benchmark	Q1 NF4	Q1 BF16	Δ (BF16–NF4)	Verdict
math_dapo (n=1000)	0.173	0.158	–1.5pp	parity (was a 0.470 anomaly — see R.6)
GSM8K (n=1319)	0.917	0.9212	+0.42pp	parity
MATH-500 (n=500)	0.770	0.786	+1.6pp	parity
AIME-24 greedy	0.100	0.100	0	identical
AIME-24 thinking (4 seeds)	0.525	0.500	–2.5pp	parity (NF4 numerically higher)

Table R.1. NF4 vs BF16 at 32B across five benchmarks: parity on every row.

NF4 = BF16 across all five benchmarks at 32B, which upholds and extends the QeRL parity claim (QeRL never ran thinking-mode evals). NF4 also showed lower seed-variance (std 0.029 vs BF16’s 0.085).

R.3 The chapter-defining finding: multi-seed AIME-24 thinking (4 seeds × n=30)

Variant	Tier	Method	Steps	Mean acc	std (n)	Δ vs base
Base 32B	32B	—	0	0.475	0.041 (n=4)	—
Q1 NF4	32B	LoRA NF4	100	0.525	0.029 (n=4)	+5.0pp
Q1 BF16	32B	LoRA BF16	100	0.500	0.085 (n=4)	+2.5pp
Base 14B	14B	—	0	0.458	0.032 (n=4)	—
Q2 OPD	14B	OPD distill	200	0.475	0.069 (n=4)	+1.7pp (noise)
Q3 NF4	14B	LoRA NF4	100	0.467	0.027 (n=4)	+0.9pp (noise)
Q4 BF16	14B	full-FT BF16	200	0.222	0.038 (n=3)	–23.6pp
Q4 FP8	14B	full-FT FP8	200	0.350	0.117 (n=2)	–10.8pp

Table R.2. Multi-seed AIME-24 thinking-mode accuracy: RL transfer is scale-dependent and full-FT GRPO regresses at 14B.

(Final Cycle 286 multi-seed values. These supersede the earlier single-seed Q4 snapshot. Q4 BF16 0.222/–23.6pp and Q4 FP8 0.350/–10.8pp are the multi-seed-confirmed numbers.)

RL thinking-transfer is scale-dependent, quantization-invariant, and update-style-sensitive: +2–5pp at 32B (NF4 $\approx$ BF16), flat at 14B for LoRA/OPD (+0.9/+1.7pp, noise), but full-FT GRPO at 14B regresses hard (Q4 BF16 –23.6pp, Q4 FP8 –10.8pp, both multi-seed). The named lesson: LoRA-RL preserves thinking-mode capability. Full-FT GRPO destroys it. Two corollaries: (1) the 14B’s +16.7pp on AIME greedy (R.5) is a greedy-decoding artifact, the same checkpoint regresses on thinking-mode (sharpen-not-expand). (2) the regression is invisible in attention-space, Q4 BF16’s per-head KL is only $\sim 2.1\times$ the no-regression Q3 (R.4) while thinking-accuracy halves, so attention-space KL is a poor predictor of thinking-mode capability loss. (Aside: Q1 NF4 step_50 $\approx$ step_100, 0.533 $\rightarrow$ 0.525, transfer fully realizes by step_50, so this recipe’s budget can halve.)

R.4 KL scaling and internals (per-head KL global_mean, full_probes_v1)

Ckpt	Size	Steps	Per-head KL	MMLU-mini	NIAH
Q3 NF4	14B	50 $\rightarrow$ 100	1.08e-4 $\rightarrow$ 1.04e-4	100%	100%
Q4 BF16	14B	200	2.22e-4	100%	100%
Q4 FP8	14B	200	2.19e-4	100%	100%
Q1 NF4	32B	50 $\rightarrow$ 100	8.50e-5 $\rightarrow$ 8.33e-5	100%	100%
Q1 BF16	32B	100	$\sim 8.5e-5$	100%	100%

Table R.3. Per-head KL drift by checkpoint: FP8 matches BF16 and the model stays intact in attention-space.

Three findings: (1) FP8 = BF16 in attention-space at 14B (BF16 2.22e-4, FP8 2.19e-4, a drop-in for training). (2) larger model drifts less, 32B is $\sim 22\%$ lower KL/step than 14B. (3) NF4 LoRA drift is non-monotonic (step_50 KL > step_100, drifts away, then partially returns. LoRA may saturate KL drift where full-FT does not). MMLU-mini and NIAH 100% throughout $\rightarrow$ RL leaves the model intact at the attention level.

R.5 The Q4 14B audit (the “base 4B” mislabel correction)

V2 Cycles 243–274 propagated a wrong “base 4B” assumption. Q4 was 14B-base. Corrected transfers vs the 14B base:

Benchmark	base 14B	Q4 BF16	Q4 FP8
GSM8K	0.9174	0.9378 (+2.04pp)	0.9265 (+0.91pp)
MATH-500	~0.778	0.792 (+1.4pp)	0.798 (+2.0pp)
AIME-24 greedy	0.100	0.200 (+10pp)	0.267 (+16.7pp)
AIME-24 thinking (4-seed)	0.458	0.222 (–23.6pp)	0.350 (–10.8pp)

Table R.4. The corrected Q4-14B transfers vs the 14B base: greedy gains but thinking-mode regression.

The “+16.7pp biggest-OOD” headline survives only on AIME greedy (deterministic, single-seed). On thinking-mode Q4 regresses (multi-seed, R.3). Q4 BF16-vs-FP8 paired (both 14B): GSM8K BF16 0.938 > FP8 0.927 (+1.13pp, rollout-noise not attention-drift). KL paired-equivalent.

R.6 C9-Q-1, vanilla GRPO on a 30B-A3B MoE (50 steps)

A datapoint that nuances [Part V](#)’s GSPO mandate: vanilla GRPO on this large MoE did NOT collapse in 50 steps (val 41.3% @ step 45). Transfer: +54pp on math_dapo ($\approx$ 0.595 vs 0.060 base), +0.5–2.7pp on GSM8K (shrinks with n. Within noise, needs $n \geq 500$ to confirm). Internals: 1 architecturally collapsed layer in the base. entropy rises with training (the opposite of the dense-collapse pattern). This is the evidence behind “GSPO is the safer default, not an absolute requirement” ([Chapter 156](#)).

R.7 Retractions and audits log

Item	Correction	Cycle
BF16 wins +6.7pp on thinking-AIME	RETRACTED — single-seed artifact; 4-seed $\Delta = -2.5\text{pp}$ (NF4 higher), not significant	280
Q1 BF16 +10pp transfer	partially retracted — single seed at top of distribution; multi-seed mean +2.5pp vs base 32B (0.475)	280
Q1 NF4 math_dapo = 0.470	corrected to 0.173 (0.470 was contaminated from C9-Q-1)	292
Q4 transfers “vs base 4B”	re-based to 14B (GSM8K/MATH-500 deltas shrank ~1–2pp)	274
Q1 = NVFP4	relabeled NF4 (Blackwell blocker, R.1)	290

Table R.5. The retractions and audits log: each corrected claim and the cycle that fixed it.

Lesson carried into the book’s method: never claim a multi-point gap on $n=30$ from a single seed (the 95% Wilson interval at $n=30$ is $\pm 18\text{pp}$). Re-read the source eval JSON rather than trusting cached recollection. The Q-program is data-locked and these corrections are reflected throughout Chapters 132–133, 151, and the appendices.

APPENDIX S

References

The algorithms, tools, models, and benchmarks this book builds on are named in prose throughout. This appendix collects the external works in full, grouped by the role each plays in the book. It is a reading list, not a claim of completeness; where a method is the book's own, it is marked as such in the text, not here.

RL and post-training algorithms

- Shao, Zhihong, et al. “DeepSeekMath: Pushing the Limits of Mathematical Reasoning in Open Language Models” (GRPO). arXiv:2402.03300, 2024.
- Zheng, Chujie, et al. “Group Sequence Policy Optimization” (GSPO). arXiv:2507.18071, 2025.
- Yu, Qiyang, et al. “DAPO: An Open-Source LLM Reinforcement Learning System at Scale” (also Clip-Higher). arXiv:2503.14476, 2025 (NeurIPS 2025).
- Liu, Zichen, et al. “Understanding R1-Zero-Like Training: A Critical Perspective” (Dr. GRPO). arXiv:2503.20783, 2025.
- Cui, Ganqu, et al. “The Entropy Mechanism of Reinforcement Learning for Reasoning Language Models” (Clip-Cov / KL-Cov). arXiv:2505.22617, 2025.
- Cui, Ganqu, et al. “Process Reinforcement through Implicit Rewards” (PRIME). arXiv:2502.01456, 2025.
- Schulman, John, et al. “Proximal Policy Optimization Algorithms” (PPO). arXiv:1707.06347, 2017.
- Rafailov, Rafael, et al. “Direct Preference Optimization: Your Language Model is Secretly a Reward Model” (DPO). NeurIPS 2023; arXiv:2305.18290.

- Ahmadian, Arash, et al. “Back to Basics: Revisiting REINFORCE-Style Optimization for Learning from Human Feedback in LLMs” (RLOO). ACL 2024; arXiv:2402.14740.
- Williams, Ronald J. “Simple Statistical Gradient-Following Algorithms for Connectionist Reinforcement Learning” (REINFORCE). *Machine Learning* 8(3–4), 1992.
- Ouyang, Long, et al. “Training Language Models to Follow Instructions with Human Feedback” (RLHF). NeurIPS 2022; arXiv:2203.02155.
- Jin, Bowen, et al. “Search-R1: Training LLMs to Reason and Leverage Search Engines with Reinforcement Learning.” arXiv:2503.09516, 2025.
- “S-GRPO: Early Exit via Reinforcement Learning in Reasoning Models.” arXiv:2505.07686, 2025.
- “DRA-GRPO: Exploring Diversity-Aware Reward Adjustment for R1-Zero-Like Training of Large Language Models.” arXiv:2505.09655, 2025.
- “MC-GRPO: Median-Centered Group Relative Policy Optimization for Small-Rollout Reinforcement Learning.” arXiv:2601.22582, 2026.
- Liu, Shih-Yang, et al. (NVIDIA). “GDPO: Group reward-Decoupled Normalization Policy Optimization for Multi-reward RL Optimization.” arXiv:2601.05242, 2026. (NVIDIA’s GDPO — normalizes each reward independently within the group; cited in the GDPO chapter to disambiguate from this book’s own Group-Decomposed Policy Optimization.)
- Tang, Yunhao, and Rémi Munos. “On a Few Pitfalls in KL Divergence Gradient Estimation for RL.” arXiv:2506.09477, 2025. (The biased-gradient property of the K3-in-loss KL estimator; cited at the K1-in-reward discussion.)

Training and serving infrastructure

- Sheng, Guangming, et al. “HybridFlow: A Flexible and Efficient RLHF Framework” (veRL). EuroSys 2025; arXiv:2409.19256.

- Kwon, Woosuk, et al. “Efficient Memory Management for Large Language Model Serving with PagedAttention” (vLLM). SOSP 2023; arXiv:2309.06180.
- Zheng, Lianmin, et al. “SGLang: Efficient Execution of Structured Language Model Programs.” NeurIPS 2024; arXiv:2312.07104.
- von Werra, Leandro, et al. “TRL: Transformer Reinforcement Learning.” GitHub, 2020 (software).
- Mei, Zhiyu, et al. “AReaL: A Large-Scale Asynchronous Reinforcement Learning System for Language Reasoning.” NeurIPS 2025; arXiv:2505.24298.
- Hu, Jian, et al. “OpenRLHF: An Easy-to-use, Scalable and High-performance RLHF Framework.” arXiv:2405.11143, 2024.
- Zhao, Yanli, et al. “PyTorch FSDP: Experiences on Scaling Fully Sharded Data Parallel.” VLDB 2023; arXiv:2304.11277. (FSDP2: see also the PyTorch documentation.)
- Rasley, Jeff, et al. “DeepSpeed: System Optimizations Enable Training Deep Learning Models with Over 100 Billion Parameters.” KDD 2020.
- Rajbhandari, Samyam, et al. “ZeRO: Memory Optimizations Toward Training Trillion Parameter Models.” SC20; arXiv:1910.02054.
- Shoeybi, Mohammad, et al. “Megatron-LM: Training Multi-Billion Parameter Language Models Using Model Parallelism.” arXiv:1909.08053, 2019.
- Moritz, Philipp, et al. “Ray: A Distributed Framework for Emerging AI Applications.” OSDI 2018; arXiv:1712.05889.
- Dao, Tri, et al. “FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness.” NeurIPS 2022; arXiv:2205.14135.
- Ye, Zihao, et al. “FlashInfer: Efficient and Customizable Attention Engine for LLM Inference Serving.” arXiv:2501.01005, 2025.
- Unsloth AI. “Unsloth.” GitHub, software.

- NVIDIA. “NeMo-RL: A Scalable Post-Training Library.” GitHub (NVIDIA-NeMo/RL), 2025 (software).
- Hugging Face. “Open-R1: A Fully Open Reproduction of DeepSeek-R1.” GitHub (huggingface/open-r1), 2025 (software).
- Pan, Jiayi, et al. “TinyZero: Minimal Reproduction of DeepSeek R1-Zero.” GitHub (Jiayi-Pan/TinyZero), 2025 (software).

Quantization and precision

- Lin, Ji, et al. “AWQ: Activation-aware Weight Quantization for LLM Compression and Acceleration.” MLSys 2024; arXiv:2306.00978.
- Frantar, Elias, et al. “GPTQ: Accurate Post-Training Quantization for Generative Pre-trained Transformers.” ICLR 2023; arXiv:2210.17323.
- Xiao, Guangxuan, et al. “SmoothQuant: Accurate and Efficient Post-Training Quantization for Large Language Models.” ICML 2023; arXiv:2211.10438.
- Dettmers, Tim, et al. “QLoRA: Efficient Finetuning of Quantized LLMs” (NF4). NeurIPS 2023; arXiv:2305.14314.
- Dettmers, Tim, et al. “LLM.int8(): 8-bit Matrix Multiplication for Transformers at Scale.” NeurIPS 2022; arXiv:2208.07339.
- Micikevicius, Paulius, et al. “FP8 Formats for Deep Learning.” arXiv:2209.05433, 2022.

Architecture components

- Qiu, Zihan, et al. “Gated Attention for Large Language Models.” NeurIPS 2025; arXiv:2505.06708.
- Yang, Songlin, Jan Kautz, and Ali Hatamizadeh. “Gated Delta Networks: Improving Mamba2 with Delta Rule.” ICLR 2025; arXiv:2412.06464.
- Dao, Tri, and Albert Gu. “Transformers are SSMs: Generalized Models and Efficient Algorithms Through Structured State Space Duality” (Mamba-2). ICML 2024; arXiv:2405.21060.

- DeepSeek-AI; Liu, Aixin, et al. “DeepSeek-V2” (Multi-head Latent Attention). arXiv:2405.04434, 2024.
- Su, Jianlin, et al. “RoFormer: Enhanced Transformer with Rotary Position Embedding” (RoPE). *Neurocomputing* 568, 2024; arXiv:2104.09864.
- Peng, Bowen, et al. “YaRN: Efficient Context Window Extension of Large Language Models.” ICLR 2024; arXiv:2309.00071.
- Shazeer, Noam. “GLU Variants Improve Transformer” (SwiGLU; cited as inspiration for the gating idea). arXiv:2002.05202, 2020.
- Ramachandran, Prajit, Barret Zoph, and Quoc V. Le. “Searching for Activation Functions” (SiLU / Swish). arXiv:1710.05941, 2017.
- Sun, Yutao, et al. “Retentive Network: A Successor to Transformer for Large Language Models” (RetNet). arXiv:2307.08621, 2023.
- Zhang, Biao, and Rico Sennrich. “Root Mean Square Layer Normalization” (RMSNorm). NeurIPS 2019; arXiv:1910.07467.
- Dai, Damai, et al. “DeepSeekMoE: Towards Ultimate Expert Specialization in Mixture-of-Experts Language Models.” ACL 2024; arXiv:2401.06066.

Base models

- Qwen Team. “Qwen3 Technical Report.” arXiv:2505.09388, 2025.
- Qwen Team. “Qwen3.5” (Qwen3.5-397B-A17B and the Qwen3.5 family, the book’s primary models in Part XV). Alibaba, 2026 (qwen.ai/blog; technical report forthcoming).
- DeepSeek-AI, et al. “DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning.” arXiv:2501.12948, 2025.
- DeepSeek-AI, et al. “DeepSeek-V3 Technical Report.” arXiv:2412.19437, 2024.
- Luo, Michael, et al. “DeepSWE: Training a Fully Open-Sourced, State-of-the-Art Coding Agent by Scaling RL.” Agentica & Together AI, 2025 (technical report).

- Touvron, Hugo, et al. “Llama 2: Open Foundation and Fine-Tuned Chat Models.” arXiv:2307.09288, 2023.
- Kimi Team. “Kimi k1.5: Scaling Reinforcement Learning with LLMs.” arXiv:2501.12599, 2025.
- Groeneveld, Dirk, et al. “OLMo: Accelerating the Science of Language Models.” ACL 2024; arXiv:2402.00838.

Benchmarks and datasets

- Cobbe, Karl, et al. “Training Verifiers to Solve Math Word Problems” (GSM8K). arXiv:2110.14168, 2021.
- Hendrycks, Dan, et al. “Measuring Mathematical Problem Solving with the MATH Dataset” (MATH). NeurIPS Datasets & Benchmarks 2021; arXiv:2103.03874. MATH-500 subset: Lightman, Hunter, et al. “Let’s Verify Step by Step.” arXiv:2305.20050, 2023.
- American Mathematics Competitions. AIME (2024, 2025).
- Joshi, Mandar, et al. “TriviaQA: A Large Scale Distantly Supervised Challenge Dataset for Reading Comprehension.” ACL 2017; arXiv:1705.03551.
- Kwiatkowski, Tom, et al. “Natural Questions: A Benchmark for Question Answering Research.” TACL 7, 2019.
- Yang, Zhilin, et al. “HotpotQA: A Dataset for Diverse, Explainable Multi-hop Question Answering.” EMNLP 2018; arXiv:1809.09600.
- Jimenez, Carlos E., et al. “SWE-bench: Can Language Models Resolve Real-World GitHub Issues?” ICLR 2024; arXiv:2310.06770.
- Pan, Jiayi, et al. “Training Software Engineering Agents and Verifiers with SWE-Gym.” ICML 2025; arXiv:2412.21139.
- Jain, Naman, et al. “R2E-Gym: Procedural Environments and Hybrid Verifiers for Scaling Open-Weights SWE Agents.” COLM 2025; arXiv:2504.07164.

- Johnson, Justin, et al. “CLEVR: A Diagnostic Dataset for Compositional Language and Elementary Visual Reasoning.” CVPR 2017; arXiv:1612.06890.
- Yu, Licheng, et al. “Modeling Context in Referring Expressions” (RefCOCO). ECCV 2016; arXiv:1608.00272.

Textbooks, courses, and neighboring guides

- Sutton, Richard S., and Andrew G. Barto. “Reinforcement Learning: An Introduction.” 2nd ed., MIT Press, 2018.
- Szepesvári, Csaba. “Algorithms for Reinforcement Learning.” Synthesis Lectures on AI and ML, Morgan & Claypool, 2010.
- OpenAI. “Spinning Up in Deep RL.” spinningup.openai.com, 2018 (educational resource).
- Lambert, Nathan. “Reinforcement Learning from Human Feedback” (the RLHF Book; rlhfbook.com). arXiv:2504.12501, 2025.
- Raschka, Sebastian. “Build a Reasoning Model (From Scratch).” Manning Early Access Program (MEAP), in progress.
- DeepLearning.AI. “Post-Training of LLMs” (short course). deeplearning.ai, 2025.
- DeepLearning.AI. “Reinforcement Fine-Tuning LLMs with GRPO” (short course). deeplearning.ai, 2025.
- Hugging Face. “The LLM Course: Building Reasoning Models (GRPO and TRL).” huggingface.co/learn, 2025.

Reproducibility and methodology

- Fanelli, Daniele. “Negative Results Are Disappearing from Most Disciplines and Countries.” *Scientometrics* 90(3), 2012.
- Rosenthal, Robert. “The File Drawer Problem and Tolerance for Null Results.” *Psychological Bulletin* 86(3), 1979.
- Henderson, Peter, et al. “Deep Reinforcement Learning That Matters.” AAAI 2018; arXiv:1709.06560.

- Agarwal, Rishabh, et al. “Deep Reinforcement Learning at the Edge of the Statistical Precipice” (rliable). NeurIPS 2021; arXiv:2108.13264.
- Godbole, Varun, George E. Dahl, Justin Gilmer, Christopher J. Shallue, and Zachary Nado. “Deep Learning Tuning Playbook.” GitHub (google-research/tuning_playbook), 2023.